

# Kafka Streams IN ACTION

Real-time apps and  
microservices with the  
Kafka Streaming API

Bill Bejeck



MANNING



**MEAP Edition**  
**Manning Early Access Program**  
**Kafka Streams in Action**  
**Real-time apps and microservices with the Kafka Streaming API**  
**Version 8**

Copyright 2018 Manning Publications

For more information on this and other Manning titles go to  
[www.manning.com](http://www.manning.com)

# welcome

---

After several years of subscribing to Manning's MEAPs, I'm excited now to welcome you to the MEAP for my own book, *Kafka Streams in Action*!

These days, you can't afford to ignore the hot topics of Big Data, streaming data, and distributed programming. We're at a point where stream processing has become an increasingly important factor for businesses looking to harness the power from data generated in real time. With 13 years in software development, I've spent the last six years working exclusively on the back end, leading ingest teams, and handling large volumes of data daily.

Kafka Streams makes it seamless to implement stream processing on the data flowing into Kafka. So, while Kafka is a de facto standard in the industry for feeding and exporting data, Kafka Streams represents a powerful new feature. In this book, I will teach you Kafka Streams, so you, too, can add stream processing to your toolkit.

This is a particular exciting time, as this is the last MEAP release for Kafka Streams in Action! All chapters and code examples have been updated for the new 1.0 API which came out in October of 2017.

In addition the book has been updated per reviewer comments, so what you read here should be very close the final published version.

While I've made every attempt to make sure the book is clear and accurate, feel free to join me on the *Kafka Streams in Action* [Author Forum](#) at manning.com, to ask questions, offer feedback, and participate in the conversation to shape this book.

—Bill Bejeck

# *brief contents*

---

## **PART 1: GETTING STARTED WITH KAFKA STREAMS**

- 1 Welcome to Kafka Streams*
- 2 Kafka Quickly*

## **PART 2: KAFKA STREAMS DEVELOPMENT**

- 3 Developing Kafka Streams*
- 4 Streams and State*
- 5 The KTable API*
- 6 The Processor API*

## **PART 3: ADMINISTERING KAFKA STREAMS**

- 7 Monitoring and Performance*
- 8 Testing a Kafka Streaming Application*

## **PART 4: ADVANCED CONCEPTS WITH KAFKA STREAMS**

- 9 Advanced Applications with Kafka Streams*

## **APPENDICES:**

- A Additional Configuration Information*
- B Exactly Once Semantics*

# *Part 1: Getting Started with Kafka Streams*



In part one we'll discuss the start of the "big data" era and how it started with the need for processing large amounts of data batches and eventually progressed to the need for processing the data as it becomes available, stream processing. We'll also discuss what Kafka Streams is and present a "mental model" of how it works without any code so you can just focus on the big picture.

Since Kafka Streams is a library that runs on Kafka, knowledge of Kafka, while not strictly required is very useful. So we'll also cover Kafka quickly to get you up to speed on how to work with Kafka.

# Welcome to Kafka Streams

## In this chapter

- The "Big Data" movement and how it changed the programming landscape
- The Need for Stream Processing
- How Stream Processing works
- Introducing Kafka Streams
- What Problems Kafka Streams solves

### **1.1 The "Big Data" movement and how it changed the programming landscape**

If you take a look around at the modern programming landscape, it has exploded with 'Big Data' frameworks and technologies. Sure client side development has undergone transformations of its own and mobile device applications have exploded as well. But no matter how big the mobile device market gets or how client side technologies evolve, there is one constant; we need to process more and more data every day. As the amount of data grows, the need to analyze and leverage the benefit from that data grows at the same rate.

But having the ability to process large quantities of data in bulk amounts (batch processing) is not always enough. Increasingly, organizations are finding there is a need to process data as it is available (stream processing). Kafka Streams, a cutting edge approach to stream processing, is a library that allows you to perform per-event processing of records. By per-event processing, I mean the ability to work on data as it arrives. Per-event processing means you process each single record as soon as it is available; there is no grouping data together in small batches, no 'micro-batching,' required.

**SIDE BAR****What is Micro Batching?**

When the need for processing data as it arrives became more and more apparent a new strategy developed, that of 'micro' batching. As the name implies micro-batching is nothing more than batch processing, but in smaller quantities of data. In some cases, by reducing the size of the 'batch' one can indeed process results quicker, but micro-batching is still batch processing, although at faster intervals, and does not give you real per-event processing

In this book, you will learn how to use Kafka Streams to solve your streaming application needs. From basic ETL work, complex stateful transformations to joining records we'll cover the components of Kafka Streams to enable you to solve these kinds of challenges in your streaming applications.

Before we dive into Kafka Streams, I want to take a few minutes to explore the history of big data processing. I feel its useful to understand the history of big data processing, because as we explore the problems and solutions we'll clearly see how the need for Kafka and then Kafka Streams evolved. In the next several sections I will elaborate on my views and experiences of how the 'Big Data' era got started and what led us to the Kafka Streams solution.

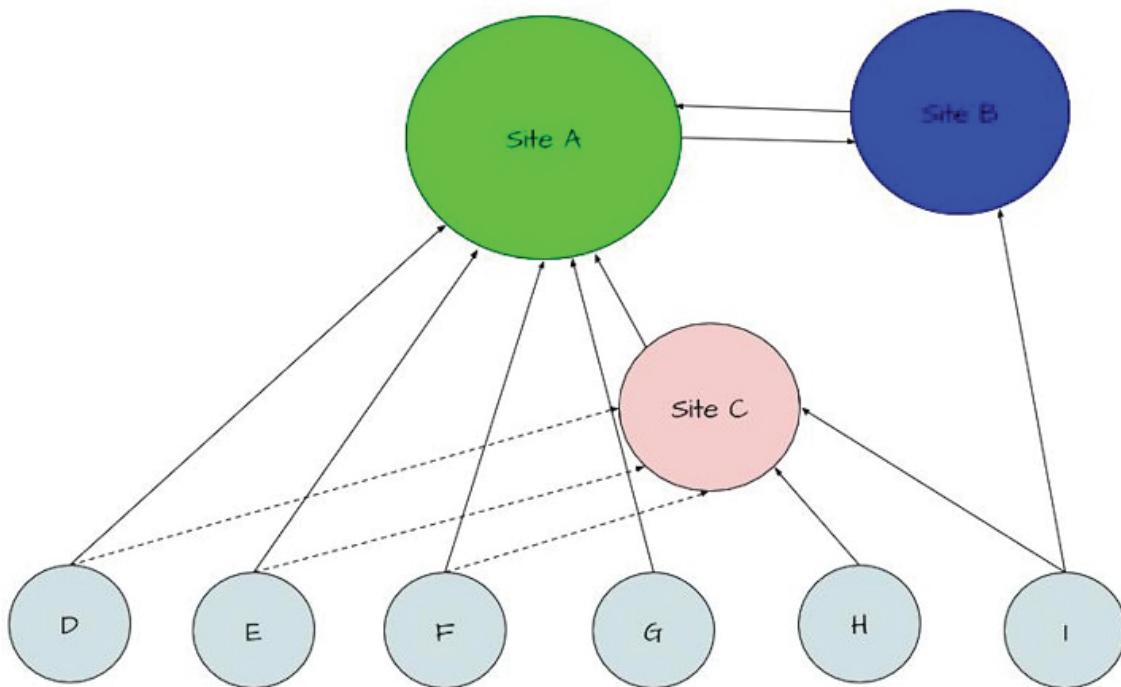
### **1.1.1 The Genesis of Big Data**

While one could argue that the internet started back in the late 1960's to early 1970's it was mainly in the hands of the U.S Department of Defense and university researchers at the time. It wasn't until the mid-1990's that the internet started to have a real impact on our daily lives. Since then, the connectivity provided by the web has given us unparalleled access to information and the ability to communicate instantly with anyone, anywhere in the world.

There was an unexpected byproduct of all this connectivity, the generation of massive amounts of data.

But in my mind, I'm going to say that the 'Big Data' era officially began in 1998, the year Sergey Brin and Larry Page formed Google. Brin and Page developed a new way of ranking web pages for searches, the PageRank algorithm. At a very high level, the PageRank algorithm works like this: an algorithm rates websites by counting the number and quality of links pointing to that site. The assumption is the more important or relevant a web page is, more sites will refer to it. Here is a graphical representation of the PageRank algorithm:

PageRank - a simple illustration - the largest circle represents the best site, the most relevant search result because it receives the most references.



**Figure 1.1 The PageRank algorithm in action. The circles represent websites, and the larger ones represent sites having a high number of links pointing to it from other sites.**

Let's take a quick moment to explain the mass of arrows and circles represented in the graphic.

- Site A is the most important because it has the most references pointing to it.
- Site B is somewhat important, although it does not have as many references, it does have a site pointing to it.
- Site C is less important than A or B. More references are pointing to site C than site B, but the quality of those references is less.
- All the rest of the sites at the bottom (small circles D through I) have no references pointing to them. This makes them the least valuable.

While the image is an oversimplification of the PageRank algorithm, it does serve to give you the big picture of how it works.

### 1.1.2 How to Implement PageRank

At the time this was a revolutionary approach because, until Google's approach, page rank searches on the web were more likely to use boolean logic to return results. If a website contained all or most of the terms you were looking for, that website was in the search results, regardless of actual quality of the content. To perform the PageRank algorithm on *all internet content* the traditional approaches to working with data weren't going to be effective for Google. The traditional methods of working with data at the time would take too long, so another approach was required.

For Google to survive and grow it needed to index all that content quickly (quickly being a relative term) and present quality results to the public, otherwise Google wouldn't be in business for long. What Google developed for processing all that data was another revolutionary approach, the MapReduce paradigm. Not only did MapReduce enable Google to do the work it needed to as a company; it inadvertently spawned an entire new industry in computing.

### 1.1.3 Important Concepts from Map Reduce

To be clear, map and reduce functions were not new concepts at the time. What was unique about Google's approach was to apply those simple concepts at a massive scale across many machines. At its heart, MapReduce has roots in functional programming. A "map" function takes some input and "maps" that input into something else, without changing the original value. Here's a simple example in Java 8 where a `LocalDate` object is "mapped" into a `String` message, while the original `LocalDate` object is left unmodified.

#### **Listing 1.1 Simple Map Function**

```
Function<LocalDate, String> addDate = (date) -> "The Day of the week is " + date.getDayOfWeek();
```

Though simple, this short example is sufficient for demonstrating what a map function does.

On the other hand, a "reduce" function takes N number of parameters and "reduces" them down into a singular, or at least smaller, value.

Taking a collection of numbers and adding them together is a good example. To perform a reduction on a collection of numbers you first provide an initial starting value, in this case, we'll use 0 (the identity value for addition).

The first step is adding the seed value to the first number in the list. You then add the result of that first addition to the second number in the list. The function repeats this process until reaching the last value, producing a single number. Here are the steps to reduce a `List<Integer>` containing the values 1,2,3:

#### **Listing 1.2 Reduce walkthrough**

0 + 1 = 1	1
1 + 2 = 3	2
3 + 3 = 6	3

- ➊ Adding seed value to the first number.
- ➋ Taking result from step one and adding to the second number in the list.
- ➌ Adding the sum of step 2 with the third and last number.

As you can see a "reduce" function collapses results together to form smaller results.

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

Similar to the map function, the original list of numbers is left unchanged.

Now, let's take a look at implementing such a simple reduce function using a Java 8 lambda:

### **Listing 1.3 Sample Reduce Function**

```
List<Integer> numbers = Arrays.asList(1, 2, 3);
int sum = numbers.reduce(0, (i, j) -> i + j);
```

Since the main topic of this book is not MapReduce, will stop our background discussion here. However, some the key concepts introduced by the MapReduce paradigm (later implemented in Hadoop, the original open source version based on Google's MapReduce white paper) we'll see come into play in Kafka Streams:

1. How to distribute data across a cluster to achieve scale in processing.
2. The use of key-value pairs and partitions to group distributed data together.
3. Instead of avoiding failure, you embrace failure by using replication.

Next, we are going to discuss these concepts in general terms, but pay attention, because as we go through the book, you'll see them coming up again.

#### **1.1.4 Distributing data across a cluster to achieve scale in processing**

If you have 5TB (5000GB) of data, that could be an overwhelming amount of data to work with on one machine at one time. But if you can split the data up to more machines to the point where each server is processing a manageable amount, your problem of a large amount of data is minimized. Table 1 illustrates this point very clearly:

**Table 1.1 How Splitting Up 5TB Helps Processing**

Number of Machines	Amount of Data Processed Per Server
10	500GB
100	50GB
1000	5GB
5000	1GB

As you can see from the table, you start out with an unwieldy amount data to process. But by just spreading the load out to more servers, the processing of data is no longer an issue. By looking at the last entry in the table, 1GB of data is something a laptop could trivially handle. So the first key concept to understand from map-reduce is by spreading the load across a cluster of machines, we can make what was originally an imposing problem concerning the size of data, into a manageable one.

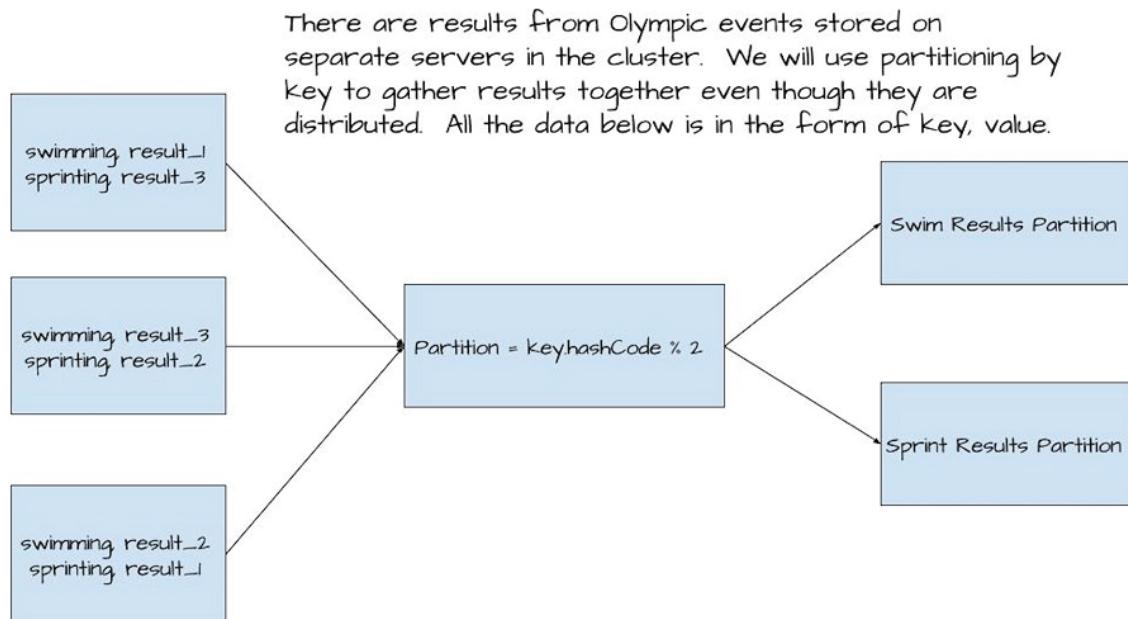
### 1.1.5 Using key-value pairs and partitions to group distributed data together

The key-value pair is a simple data structure with powerful implications. In the previous section, you saw how to process a massive amount of data; you spread that data over a cluster of machines. Distributing your data solves the processing problem, but now you have the problem of getting data back together during your processing. To address the issue of grouping distributed data you use the keys from the key-value pairs as a partition of the data. By using keys as a partition, grouping distributed data is now an easy problem to solve using the following formula:

#### **Listing 1.4 Determine the partition with a hashing function**

```
int partition = key.hashCode() % numberOfPartitions;
```

Here is a graphical depiction of how to apply a hashing function so you can group distributed data together on the same partition.



**Figure 1.2** Here we see grouping records by key on the same partition. Even though the records started out on separate servers, they arrive in the same location as a result of using partitions.

Partitioning is an important concept, and we will see detailed examples in later chapters.

### 1.1.6 Embracing failure by using replication.

Another key component of Google's MapReduce was the Google File System (GFS). In conjunction with Hadoop being the open source implementation of MapReduce HDFS (Hadoop File System) was the open source implementation of GFS.

At a very high level, GFS/HDFS split data up into blocks and distributed those blocks across the cluster. But the essential part of GFS/HDFS is the approach to server and

disk failure. Instead of trying to prevent failure, the framework embraces failure by using replication. The blocks of data are replicated across the cluster (by default the replication factor is 3).

By replicating data blocks on different servers, you no longer have to worry about disk failures or even complete server failures causing a halt in production. Replication of data is crucial in giving distributed applications fault tolerance, which is essential for a distributed application to be successful.

The key point of this section is the replication of data across a cluster gives you resilience in the face of the disk and machine failures. We'll see later on how partitions and replication work in Kafka Streams.

Now, let's take a look at how we evolved from batch processing to the need for stream processing (again based on my personal experiences and journey through programming).

### **1.1.7 Batch Processing is Not Enough**

Hadoop caught on with the computing world like wildfire. Now you can process vast amounts of data, have fault tolerance while using commodity hardware (cost savings). But Hadoop/MapReduce is a batch-oriented process. Being batch oriented means, you collect large amounts of data, process it, then store the output for later use.

Batch processing is a perfect fit for something like PageRank; you can't make determinations of what is a valuable resource across the entire internet by watching user clicks in real-time. But business came under ever increasing pressure to respond to important questions more quickly. Questions like:

- What is trending right now?
- How many invalid login attempts in the last 10 minutes?
- How is the recently released feature being utilized by the user base?

It was apparent that another solution was needed, and that solution is stream-processing.

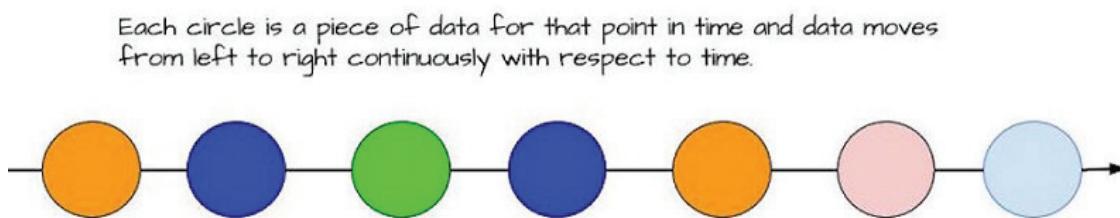
Every business, government or organization today has one thing in common, and that is the byproduct of conducting your business is the generation of data, lots, and lots of data. For example, there are estimates that a single flight of a commercial airplane can generate up to half a terabyte of data *per flight*. When one stops to consider the number of domestic flights every day in the United States alone, you can see this data volume amounts to a staggering amount of data to be analyzed.

For all that data to be of any value, people and or machines need to review it and provide assessments. All that data can be collected and used for making decisions to cut costs, improve the performance or any number of decision points.

Now that we have established the need for stream processing let's go on to define what it means.

## 1.2 Introducing Stream Processing

While there are varying definitions of stream processing, in this book we are going to define stream processing as working with data as it is arriving into your system. We can further refine the definition to say that stream processing is the ability to work with an infinite stream of data with continuous computation, as it flows, no need to collect or store the data to act on it. This simple diagram below represents a stream of data each circle on the line represents data at a point in time. Data is continuously flowing, as data in stream processing is unbounded, think of the number line we learned in school, you can stop to look at or work with a single number, but the line keeps going.



**Figure 1.3** This marble diagram is a simple representation of stream processing. Each circle represents some information or event occurring at a particular point in time. The number of events is unbounded and move continually from left to right just like a number line.

So who need to use stream processing? Well anyone who needs feedback very quickly from any observable event. Let's discuss some examples below.

### 1.2.1 When to use Stream Processing and when not to use it

Like any technical solution, stream-processing is not a "one size fits all" solution. So let's look at some scenarios where stream processing is a fit first.

The need to quickly respond or report on incoming data is a good use case for stream-processing.

A good illustration is credit card fraud. The owner may not notice a stolen credit card, but by reviewing purchases as they happen against established patterns (location, general spending habits), you may be able to detect a stolen credit card and alert the owner. Another good example would be intrusion detection into a network. While analyzing application log files after a breach has occurred may be helpful to prevent future attacks or where to improve security, the ability to monitor aberrant behavior in real-time is critical.

Imagine a large race such as the NYC Marathon. Almost all runners will have a chip on their shoe, and when runners go by sensors spread out across the course, you can use that information to track the position of runners. By using the sensor data, we can determine the leaders, spot potential cheating as well as detecting if a runner has potentially run into problems.

In terms of time sensitivity for data, the financial industry is probably at the top of the list. The ability to track market prices and direction in real time is essential for brokers or consumers to make effective decisions regarding when to sell or buy stocks or other

financial instruments.

On the other hand, stream-processing is not a solution for all problem domains, so let's take a quick look at where stream-processing might not be appropriate.

To effectively make forecasts of future behavior to predict the outcome of future events, one needs to use a large amount of data. Only by using vast quantities of data accumulated over time can one eliminate anomalies and start to identify patterns and trends of problem domain you are measuring.

Economic forecasting is a good example. Information is collected on many variables over an extended period in an attempt to make an accurate forecast, interest rates in the housing market for example. Another good example is assessing the effect of implementing a change in curriculum in schools. It's only after one or two testing cycles that school administrators can measure if the changes made are achieving their goals.

In this section, we have discussed areas where stream-processing is a good fit and some others where isn't appropriate.

Here are the key points to remember. If you need to report on or take action immediately as data arrives, stream-processing is a good candidate for you. On the contrary, if you need to perform in-depth analysis or are only compiling a large repository of data for later analysis a stream-processing approach may not be a fit.

Since this is a book about Kafka Streams, a stream-processing technology, our next step is to walk through a concrete example where we have a need to do stream-processing.

### **1.3 Handling a Purchase Transaction**

Our first step is to apply a general stream-processing approach to an example from the world of retail sales. After we have walked through the general solution, we are going to show how we can use Kafka Streams specifically to implement the stream processing application.

Jane Doe is on her way home from work and remembers she needs toothpaste and stops at a ZMart that is on the way home. She goes in to pick up the toothpaste and heads to the checkout to pay for her items. The cashier asks Jane if she is a member of the "ZClub" and she scans her membership card, so Jane's membership info is now part of the purchase transaction.

Jane has no cash (does anyone use cash anymore?) so when the total is rung up Jane hands the cashier her debit card. The cashier swipes the card and gives Jane the receipt, and she leaves. As Jane is walking out of the store, she checks her email, and there is a message from Zmart thanking her for her patronage with various coupons attached for discounts on Jane's next visit.

While the transaction above is a normal occurrence that a customer would not give a second thought about, a while back you recognized it for what it is; a wealth of information that can help ZMart run more efficiently and serve the customer base in a better fashion. Let's go back in time a little bit to see how you were able to make this situation a reality.

### 1.3.1 Weighing the Stream Processing Option

You are the lead developer for ZMart's streaming data team. ZMart is a big box retail store with several locations across the country. ZMart does a great business with total sales for any given year upwards of \$1B. You would like to start mining the data from your company's transactions to make the business more efficient. You know you have a tremendous amount of data to work with from ZMart's daily sales, so whatever technology you implement, it will need to be able to work fast and scale to handle this vast volume of data.

You decide to use stream processing because there are business decisions and opportunities that can be leveraged as each transaction occurs, no need to wait to gather data and make decisions hours later. You get together with management and your team to come up with four primary requirements that are necessary for the stream processing initiative to succeed.

So our next step is to go over each one quickly.

### 1.3.2 The Business Requirements for ZMart's Purchase Transaction Tracking Program

- First and foremost you value your relationship with your customers. With all the privacy concerns of today, your first goal is to protect their privacy, with protecting their credit cards numbers as the highest priority. So however we use the information from the transaction above, Jane's credit card information should never at risk of exposure.
- There is a new customer rewards program in place where customers earn bonus points for the amount of money they spend on certain items. The goal is to notify customers very quickly once they have received a reward, you want them back in the store! Again appropriate monitoring of activity is required here. Remember how Jane received an email immediately after leaving the store? That's the kind of exposure you want for the company.
- The company would like to refine further its advertising and sales strategy. The company wants to track purchases by region to figure out which items are more popular in certain parts of the country. The goal is to target sales and specials for specific items that are best sellers in a given area of the country.
- All purchase records need to be saved in an off-site storage center for historical and ad-hoc analysis.

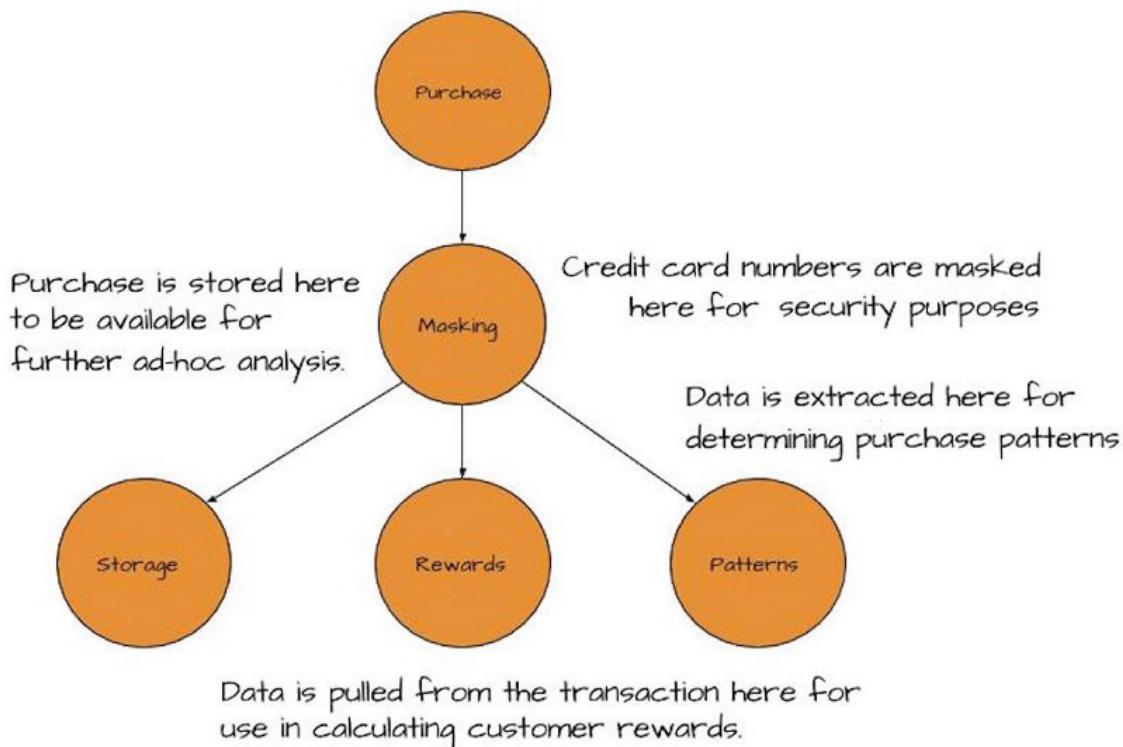
The requirements you have are straight forward enough on their own, but how would you go about implementing those various requirements against a single purchase transaction like the one outlined above?

### 1.3.3 Deconstructing the Requirements into a Graph

If we stop and look at the requirements outlined above, we can quickly recast them into a graph, namely a directed acyclic graph or DAG. The point where the customer completes the transaction at the register is the source node for the entire graph. Viewing the purchase transaction as the node of origin, then the list of requirements simply become child nodes of the main source node.

Here's the view of the business requirements recast as a DAG.

The point of purchase is the source or parent node for the entire graph



**Figure 1.4** Here we are recasting the business requirements for your streaming application as a directed acyclic graph (DAG). Each vertex represents a requirement, and the edges show the flow of data through the graph.

Next, we will show how to map a purchase transaction to our requirements graph.

## 1.4 Changing the Perspective of a Purchase Transaction

In this section, we will walk through every step of the purchase to show how we apply it to the requirements graph. This walkthrough is essentially how Kafka Streams works. For now, we will stick with generic graph flow, but in the next section, we will demonstrate how we will specifically apply Kafka Streams to this issue. Also, we will use very high-level language as we map the purchase to the requirements graph.

### 1.4.1 Source Node

The source node of the graph is where the application consumes the purchase transaction.

The point of purchase is the source or parent node for the entire graph



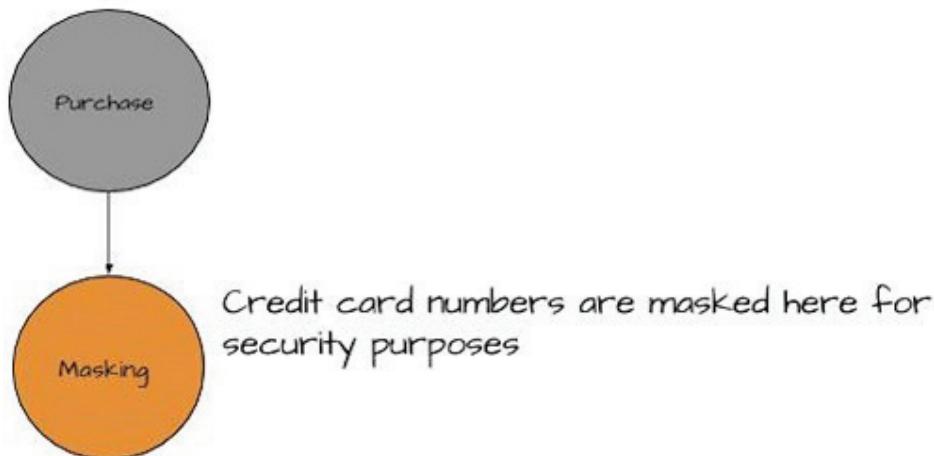
**Figure 1.5** The simple start for the sales transaction graph. This node is the source of raw sales transaction information that will flow through

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

the graph.

### 1.4.2 Credit Card Masking Node

The child node of the graph source is where the credit-card masking takes place. For the credit-card masking operation, we make a copy of the data then convert all the digits on the credit card number to an "x" except for the last four digits.



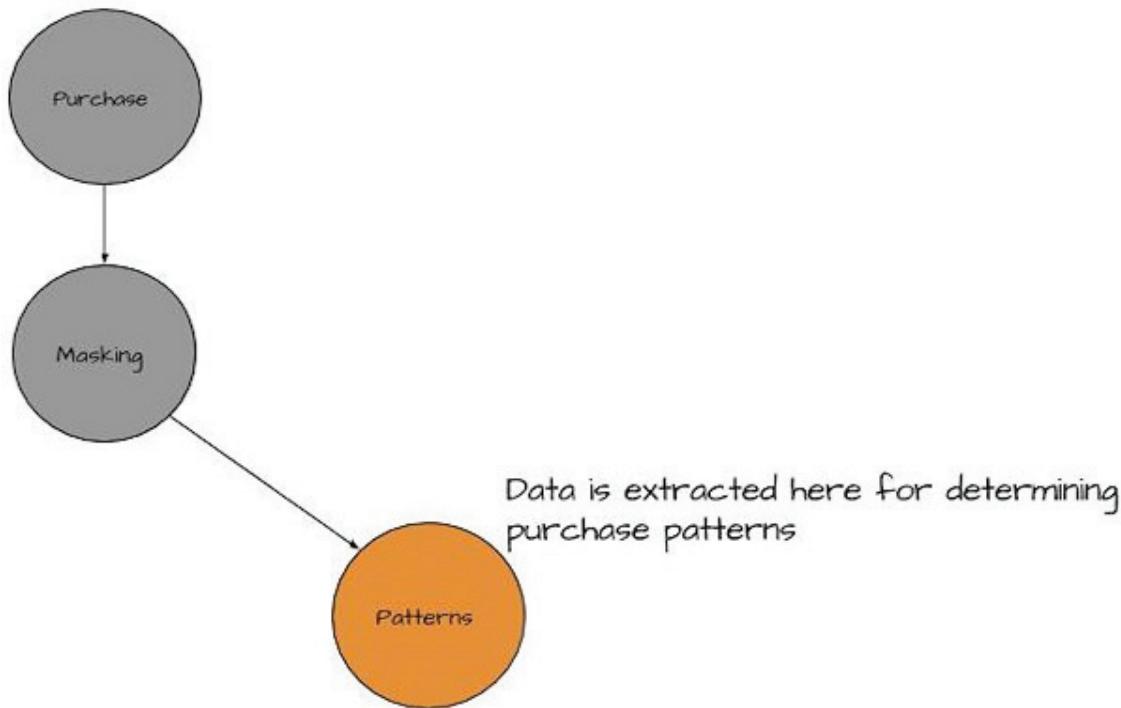
**Figure 1.6** The first vertex or node in the graph representing the business requirements. This node is responsible for masking credit card numbers and is the only node that receives the raw sales data from the source node effectively making this node the source for all other nodes connected to it.

This node is the only direct child of the source, all other nodes in the graph are children of the credit-card masking node, so data flowing through the rest of the graph will have the credit-card field converted to the 'xxxx-xxxx-xxxx-1122' format. The remaining nodes that are

- The Purchase patterns node
- The Rewards accumulator node
- The Storage node that writes the sales transaction to storage for ad-hoc analysis later.

### 1.4.3 Patterns Node

The purchase patterns node extracts the relevant information to establish where customers purchase products throughout the country. I

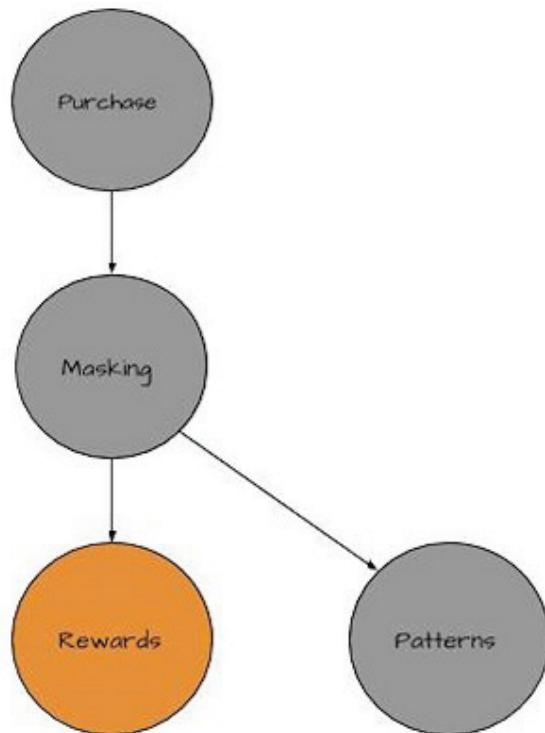


**Figure 1.7** Here we are adding the patterns node in the graph. It consumes purchase information from the masking node and converts it into a record showing when a customer purchased an item and the zip code where the customer completed the transaction.

Instead of making a copy of the data; the patterns node will retrieve the item, date and the zip code where the item was purchased and create a new object containing those fields. Now, let's move on to the next node in our graph.

#### 1.4.4 Rewards Node

The next child node in the process is the rewards accumulator. ZMart has a customer rewards program that gives customers points for purchases made in the store. This processor's role is to extract dollar amount spent and the client's id from the and create a new object containing those two fields.



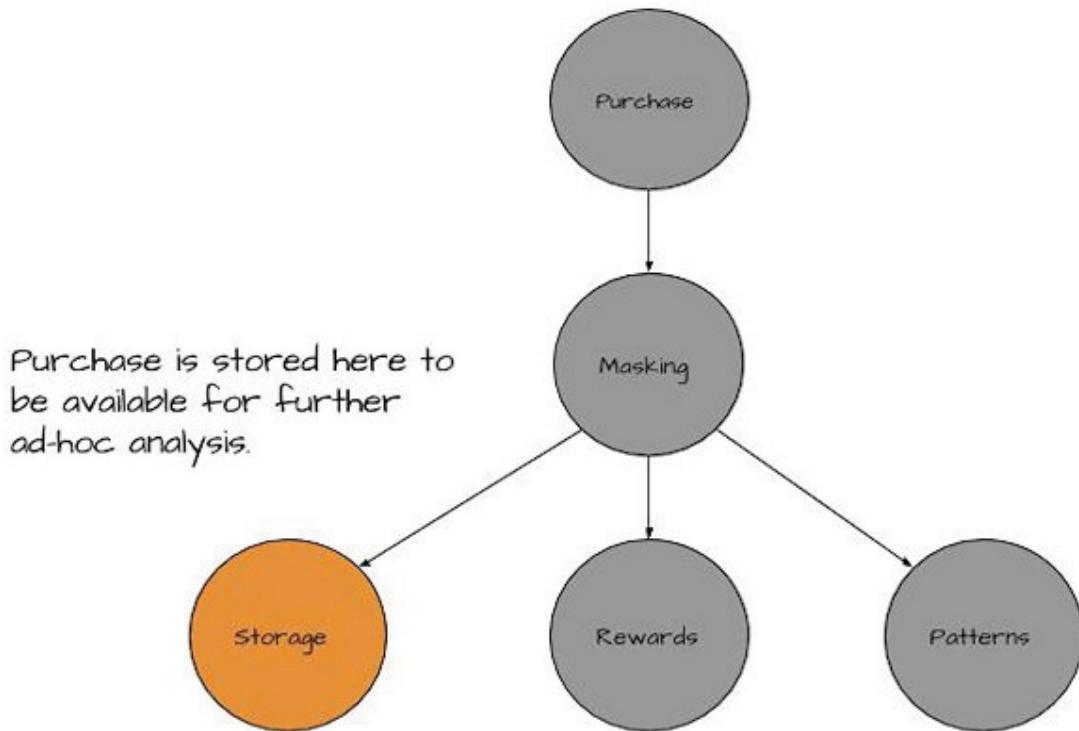
Data is pulled from the transaction here for use in calculating customer rewards.

**Figure 1.8** This is the rewards node in the graph and is responsible for consuming sales records from the masking node and converting them into records containing the total of the purchase and the customer id.

As with the patterns node, this is where we will stop with the rewards node in the graph, but when we are applying Kafka Streams, we will be adding an extra node as well.

#### 1.4.5 Storage Node

The final child node writes the entire purchase data out to storage for further ad-hoc analysis.



**Figure 1.9** Here is the last node in the graph and consumes records from the masking node as well. These records aren't converted into any other format but are stored in a NoSQL data store for ad-hoc analysis later.

We have now finished tracking the results of the example purchase transaction through the graph of your requirements. So the next step is, what will we use to convert this graph into live, functional streaming application? In the next section, we'll see exactly how we can use Kafka Streams to build a solution to meet our goals.

## 1.5 Applying Kafka Streams to the Purchase Processing Graph

Before we get into the technical details of implementing Kafka Streams to our issue, let's start with a 'plain-speak' definition.

Kafka Streams is a library that allows you to perform per-event processing of records. What do we mean by per-event processing? You have the ability to work on data as it arrives, no grouping data together in small batches, you process each single record *as soon as* it is available no 'micro-batching' required.

Most of the goals here you have time sensitive in that you want to be able to take action as soon as possible. Preferably you will be able to collect information as it occurs. Additionally, there are several locations across the country, so you need all transaction records to funnel down into a single, constant flow or "stream" of data for your analysis. For these reasons, Kafka Streams is a perfect fit. Kafka Streams will allow you to process records as they arrive and give you the low latency processing you require.

## 1.6 Kafka Streams as a Graph of Processing Nodes

In Kafka Streams, you define a topology of processing "nodes". One or more nodes will have Kafka topic(s) as its source, and you add additional processing nodes considered child nodes. Each child node itself can define other child nodes (we will use the terms processor and node interchangeably) or processors. Each processor node performs its assigned task then forwards the message to each of its child nodes. This process of performing work then forwarding data to any child nodes continues until every child node has executed its function.

Does this process sound familiar? It should as this is the exact transformation performed with your business requirements, we transformed them into a graph of processing nodes. Traversing a graph is how Kafka Streams works; it is a directed acyclic graph (DAG) or topology of processing nodes. You start with a source or parent node, and it has one or more children. Data always flows from the parent to the child nodes, never from child to parent.

Each child node, in turn, can define child nodes of their own and so on. Messages flow through the graph in a depth-first manner. This depth-first approach has significant implications. Each record (a key-value pair) is processed *in full* by the entire graph before accepting another one for processing. Having each record processed depth-first through the whole DAG eliminates the need to have back pressure built into Kafka Streams.

### SIDE BAR    What is Back Pressure?

While there are varying definitions of back pressure, here I define backpressure as the need to restrict the flow of data by buffering or some blocking mechanism. Back pressure is necessary when a "source" is producing data faster than a "sink" can receive and process the data.

By being able to connect or chain together multiple processors in this manner, you can quickly build up complex processing logic, while at the same time having each component itself remain relatively straightforward and easy to keep the details straight. It is this composition of processors where the power and complexity come into play.

### SIDE BAR    What is a Topology?

A topology is a way you arrange parts of an entire system and have connected them with each other. When we say Kafka Streams has a topology, we are referring to transforming data by running through one or more processors.

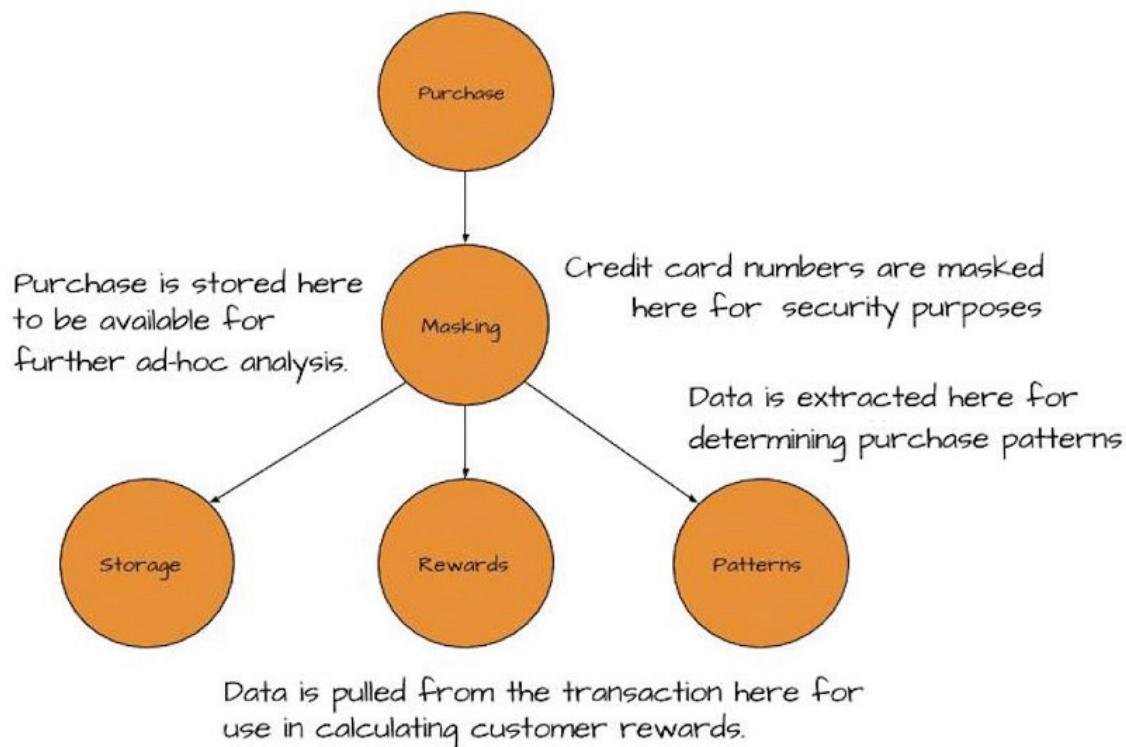
**NOTE****Kafka Streams and Kafka**

As you might have guessed from the name, Kafka Streams runs on top of Kafka. In this introductory chapter knowledge of Kafka itself won't be strictly necessary, as we will focus more how Kafka Streams works conceptually. While there may be a few Kafka specific terms mentioned, for the most part, we will be concentrating on the stream processing aspects of Kafka Streams. But, knowledge of Kafka is essential to work effectively with Kafka Streams. So for readers who are new to Kafka or are unfamiliar with Kafka altogether, we will teach you what you need to know about Kafka in chapter 2.

## 1.7 Applying Kafka Streams to the Purchase Transaction Flow

Now let's walk through building a processing graph again, but this time we will create a Kafka Streams program. To refresh our memory, here is the requirements graph that we just constructed:

The point of purchase is the source or parent node for the entire graph



**Figure 1.10** Here's the view of our business requirements again cast a graph. Remember vertexes are processing nodes that handle data and the edges show the flow of data.

While we will be building a Kafka Streams program as we build this graph, this will still be a relatively high-level approach, and some details will be left out. We will go into more detail later in the book when we take a look at the actual code.

Once our Kafka Streams program starts to consume records, we convert the raw messages into `Transaction` objects. Specifically here is the pieces information that will make up the transaction object:

1. Your ZMart customer id (you scanned your member card).
2. The total dollar amount spent.
3. The item(s) you purchased.
4. The zip code of the store where the purchase took place.
5. The date and time of the transaction.
6. Your debit/credit card number.

### **1.7.1 Defining the Source**

The first step in any Kafka Streams program is to establish a source for the stream. The source could be any of the following:

1. A single topic
2. Multiple topics in a comma separated list
3. A regex that can match one or more topics

**NOTE**
**Terminology**

While my goal is to keep the description on at a high level here, I may mention some Kafka specific terms. Not to worry we will cover these terms in Chapter 2.

In this case, it will be the "transaction" topic. A Kafka Streams program runs one or more `StreamThread` (the number of Stream threads are user defined) instances. Each `StreamThread` has an embedded `Consumer` and `Producer` that handles all reading from and writing to Kafka. In addition to specifying source topic(s), you also provide `Serdes` for the keys and values. A `Serde` instance contains the serializer and deserializer needed to convert objects to byte arrays/byte arrays to objects, respectively.

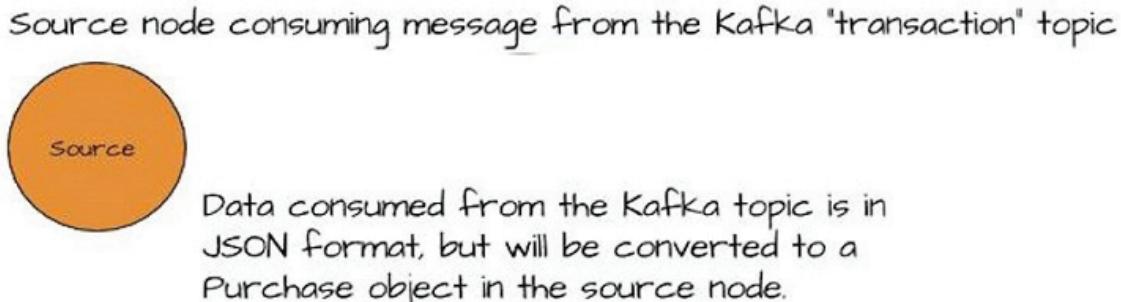
**SIDE BAR**
**What is a Serdes?**

Serdes is a class containing a Serializer/Deserializer pair for a given object.

In this case, we will work with the `Purchase` object, as it defines methods that will aid in the processing work. You provide key serdes, value serdes and source topic(s) to a `KStreamBuilder` which returns a `KStream` instance. You use the `KStream` object throughout the rest of the program. The value deserializer automatically converts the incoming byte arrays into to `Purchase` objects as messages continue to flow into the stream.

It is important to note that to Kafka itself, the Kafka Streams program appears as any other combination of consumers and producers. There could be any other number of

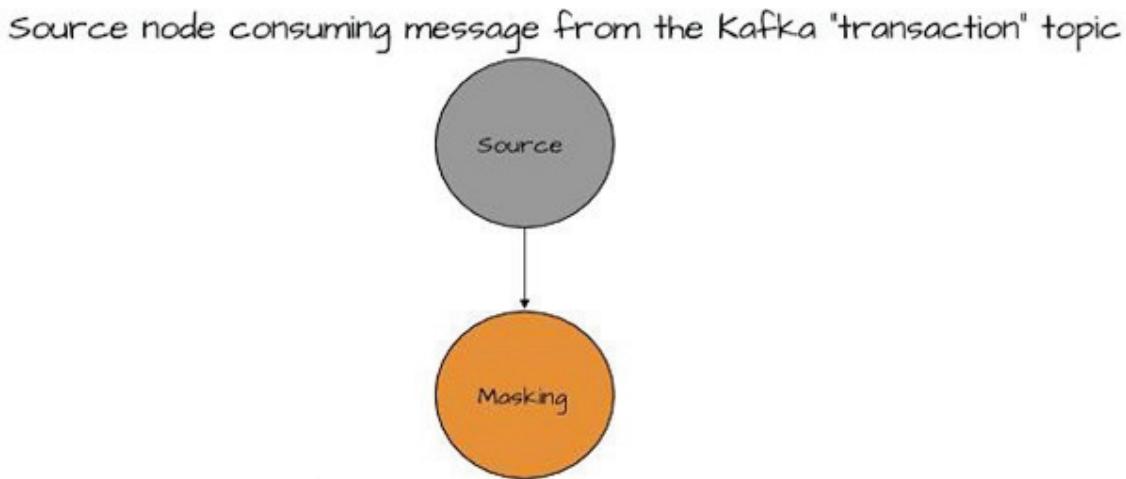
applications reading from the same topic in conjunction with your streaming program.



**Figure 1.11** Here is the source node, but instead using the vague term "source" we'll call this what it is - a Kafka topic.

### 1.7.2 The First Processor - Masking Credit Card Numbers

Now you have your source defined, and we can start creating processors that will do the work on the data. The first goal from above is to mask the credit-card numbers recorded in the incoming purchase records. The first processor will be used to convert credit-card numbers from 1234-5678-9123-2233 to xxxx-xxxx-xxxx-2233. The `KStream.mapValues` method performs the masking. The `KStream.mapValues` method returns a new `KStream` instance that changes the values as specified by the given `ValueMapper` as records flow through the stream. This particular `KStream` instance will be the parent processor for any other processors you define. Our new parent processor provides any downstream processors with `Purchase` objects with the credit-card number already masked.



Child node of the "Source" node. A copy of the `Purchase` object is made at this point and the credit-card number is converted from  
1234-5678-9123-5678 to xxxx-xxxx-xxxx-5678

**Figure 1.12** This is the masking processor and is a child of the main source node. It receives all the raw sales transactions and emits new records with the credit card number masked. All other processors receive their records from this processor.

## CREATING PROCESSOR TOPOLOGIES

A few words on creating processor topologies. Each time you create a new `KStream` instance by using a transformation method, you are in essence building a new processor. That new processor is connected to the other processors already created. By composing processors, we can use Kafka Streams to create complex data flows elegantly.

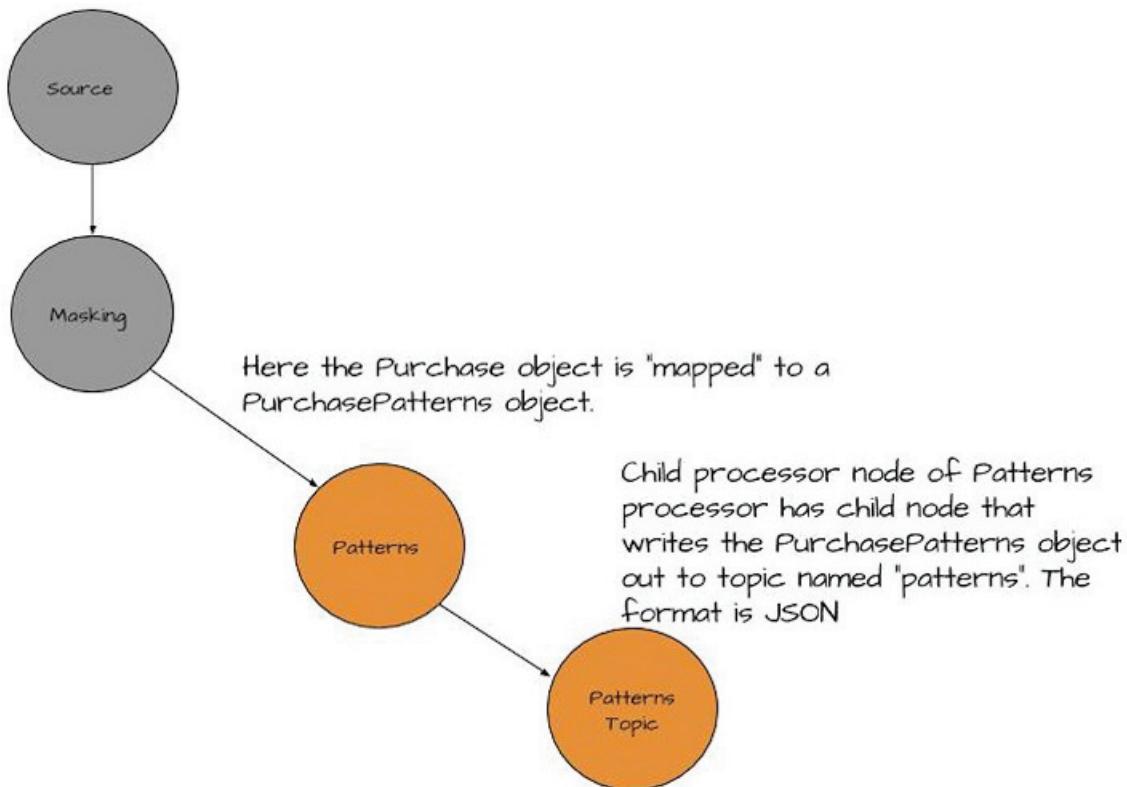
It is important to note that by calling a method that returns a new `KStream` instance does not cause the original one to stop consuming messages. When using a transforming method, a new processor is created and added to the existing processor topology. The updated topology is then used as a parameter to create the new `KStream` instance. The new `KStream` object starts receiving consumed messages from the point of its creation.

It is very likely that you will build new `KStream` instances to perform additional transformations while retaining the original stream to continue consuming incoming messages for an entirely different purpose. You will see an example of this when we define the second and third processors.

While it is possible to have a `ValueMapper` convert the incoming value to an entirely new type, in this case, it will just return an updated copy of the `Purchase` object. Using a "mapper" to update an object is a pattern we will see over and over in KStreams. Each transforming method returns a new instance of a `KStream` that will apply some function to all incoming data at that point.

Now you should have a clear image of how we can build up our processor "pipeline" to transform and output data. Next, we move on to creating the next two processors using the first processor built as a parent processor.

### 1.7.3 The Second Processor - Purchase Patterns



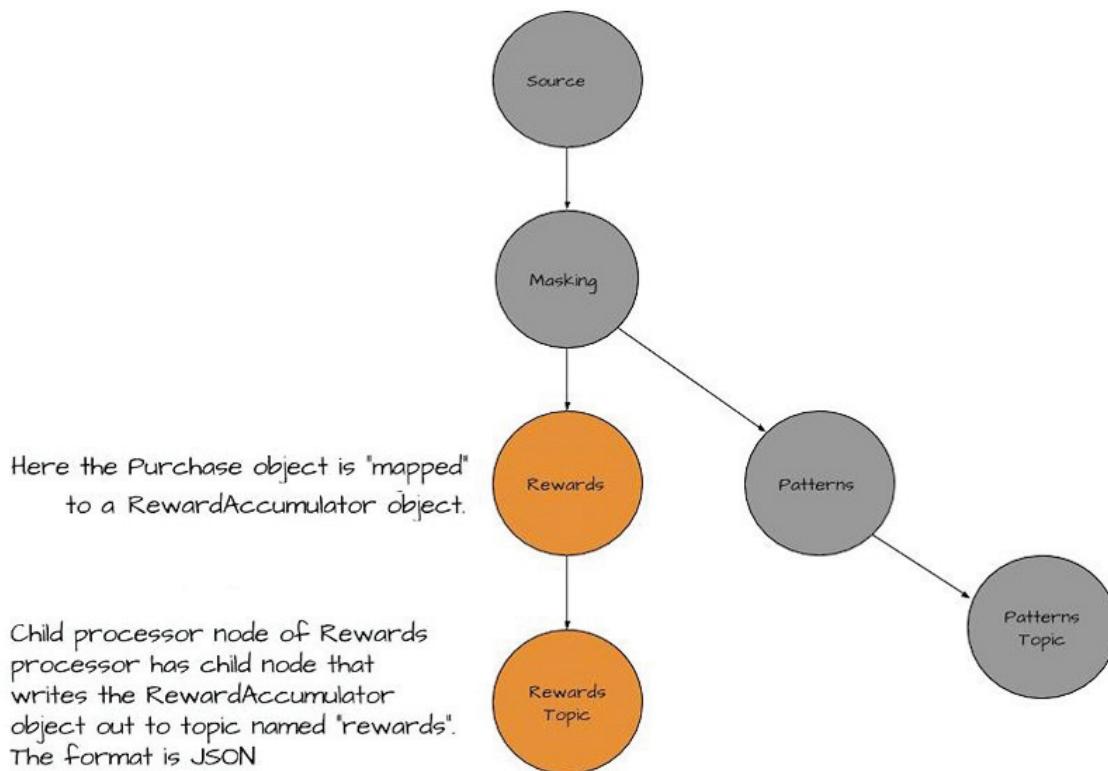
**Figure 1.13** Here is the purchase patterns processor. It's responsible for taking Purchase objects and converting them into PurchasePattern objects containing the item purchases and the zip code of the area of the country where the transaction took place. We've added a new processor here that takes records from the patterns processor and writes them out to a Kafka topic.

Next in the list of processors to create is one that can capture information necessary for determining purchase patterns in different regions of the country. To do this, we will add a child processing node to the first processor/KStream we created. Let's call this first KStream instance purchaseKStream from here on out. The purchaseKStream produces Purchase objects with the credit-card number masked.

The purchase patterns processor receives a Purchase object from its parent node and "maps" the object of a new type, a PurchasePattern object. The mapping process extracts the actual item purchased (toothpaste) and the zip code it was bought in and uses that information to create the PurchasePattern object. We will go over exactly how this mapping process occurs in chapter 3.

Next, the purchase patterns processor adds a child processor node that receives the newly mapped object and writes it out to a Kafka topic named "patterns". The PurchasePattern object is converted to some form of transferable data when written to the topic. Other applications can then consume this information and use it to determine the inventory levels as well as purchasing trends in a given area.

### 1.7.4 The Third Processor - Customer Rewards

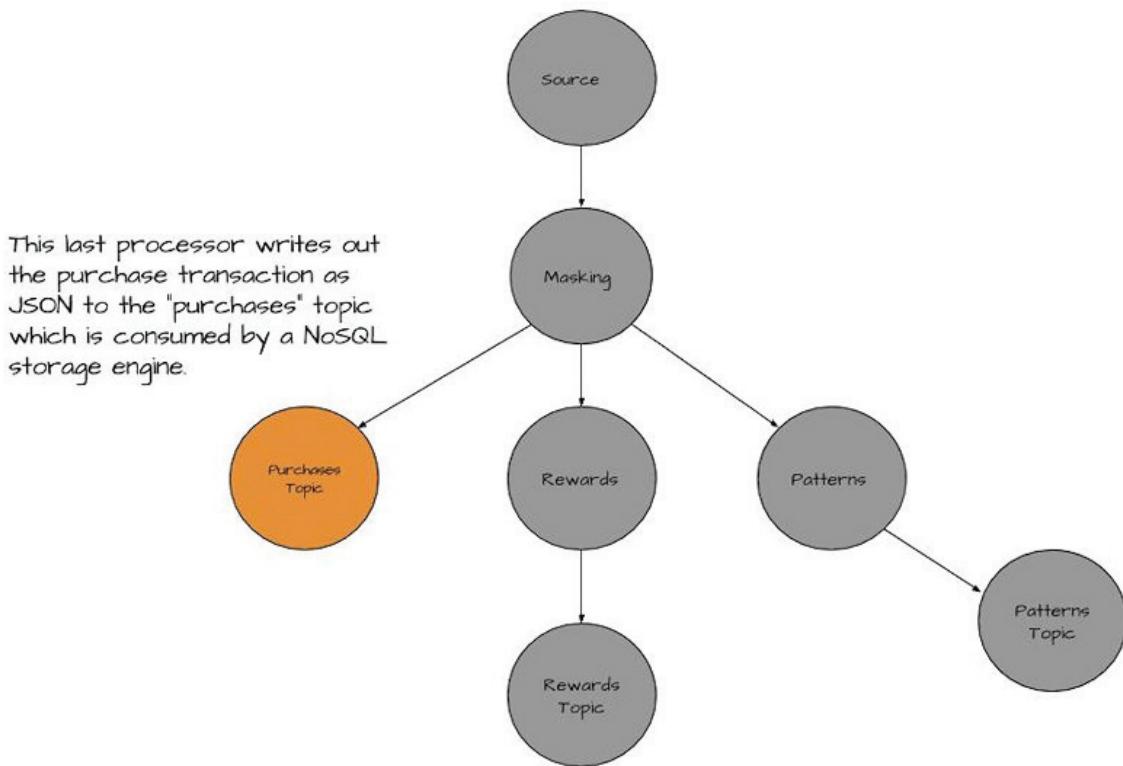


**Figure 1.14 Here is the customer rewards processor responsible for transforming Purchase objects into a RewardAccumulator object containing the customer id, date and total dollar amount of the transaction. We've also added another child processor here to write the Rewards objects to another Kafka topic.**

The third processor created will extract the required information for the customer rewards program. This processor is also a child node of the purchaseKStream processor. The rewards processor, similar to the purchase patterns processor, receives the Purchase object and "maps" it into another type the RewardAccumulator object.

The reward processor also adds a child processing node to write the RewardAccumulator object out to a Kafka topic, "rewards". It's from consuming records from the "rewards" topic that another application would determine any rewards for ZMart customers, the email that Jane Doe received from the purchase scenario for example.

### 1.7.5 Fourth Processor - Writing Purchase Records



**Figure 1.15** This is the final processor in the topology. This processor is responsible for writing out the entire Purchase object out to another Kafka topic. The consumer for this topic will store the results in a NoSQL store such as MongoDB

The last processor and third child node of the masking processor node write the purchase record out to a topic "purchases". The purchases topic will be used to feed a NoSQL storage application that will consume the records as they come in. These records will be used for ad-hoc analysis later. This record contains all the information from the purchase example above.

We have taken the first processor that masked the credit-card number and have used it to feed three other processors; 2 that further refined or transformed the data and one that wrote the masked results out to a topic for further use by other consumers. You can see now that by using Kafka Streams, we can build up a powerful processing "graph" of connected nodes to perform data integration tasks in a stream processing manner.

## 1.8 Summary

We have covered a lot of information in this chapter. Let's do a quick review of the key points you should remember:

- Kafka Streams is a graph of processing nodes that when combined provide powerful and complex stream processing.
- How batch processing is very powerful but is not enough to satisfy real-time needs for working with data.
- How distributing data, key-value pairs, partitioning and data replication are critical for

distributed applications.

To understand Kafka Streams, you really should know some Kafka. If you don't we'll cover the essentials so you can get started with Kafka Streams:

- Installing Kafka and sending a message.
- A deep dive into Kafka's architecture, what is a distributed log.
- What are Topics and how are they used in Kafka
- Understanding how producers and consumers work and how to write them effectively.

However, if you are already comfortable with Kafka, feel free to go straight to Chapter 3 where we'll build a Kafka Streams application to tackle the concrete example discussed in this chapter.

# Kafka Quickly

## In this chapter

- Kafka Architecture
- Sending Messages with Producers
- Reading Messages with Consumers
- Installing and Running Kafka

While this is a book about Kafka Streams, it's impossible to have a conversation about Kafka Streams without discussing Kafka itself. After all Kafka Streams is a library that runs on Kafka. Kafka Streams is designed very well so strictly speaking it is possible to get up and running with minimal to no Kafka experience. However, your progress and ability to fine tune are going to be limited. Having a good fundamental knowledge of Kafka is essential to get the most out of Kafka Streams.

## 2.1 Who Should Read This Chapter

This chapter is for developers whom are interested in getting started with Kafka Streams but have little no experience with Kafka itself. If you have a good level of experience with Kafka, feel free to skip this chapter and proceed directly to Chapter 3.

### NOTE

#### Kafka Is A Big Topic

Kafka is too large a topic to cover in its *entirety* in one chapter. We'll cover enough information to give you a good understanding how Kafka works and a few of the core configuration settings you'll need to know. For in depth coverage of Kafka, look into *Kafka In Action* by Dylan Scott [www.manning.com/books/kafka-in-action](http://www.manning.com/books/kafka-in-action)

## 2.2 The Data Problem

Organizations today are swimming in data. Internet companies, financial businesses, and large retailers are better positioned now more than ever to leverage the available data to serve their customers better and find more efficient ways of conducting business (we are going to take a positive outlook on this situation and assume only good intentions when looking at customer data).

Let's make a list of the various requirements we'd like to have in our data management solution.

1. What we need is a way to send data to a central storage quickly.
2. Because machines fail all the time, we also need the ability to have our data replicated, so those inevitable failures don't cause downtime and data loss.
3. Finally, we need the potential to scale to any number of consumers of data without having to keep track of different applications. We need to make the data available to anyone in an organization, but not have to keep track of who has and has not viewed the data.

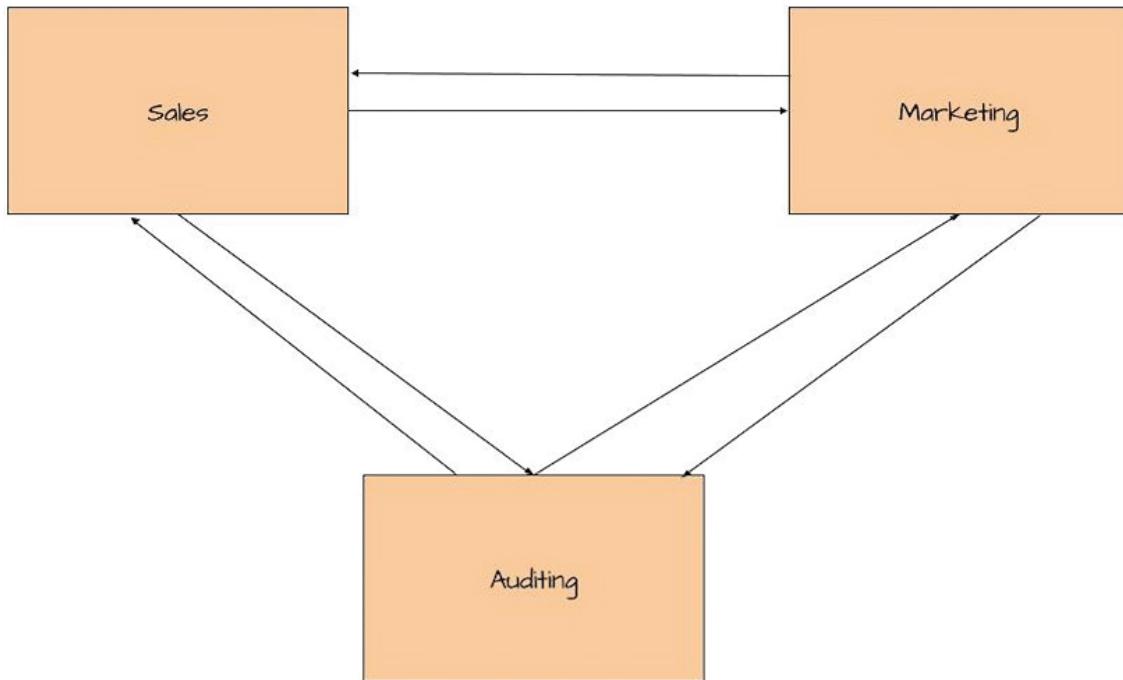
## 2.3 Using Kafka to Handle Data

In chapter 1 we were introduced to the large retail company ZMart. At that point ZMart wanted a streaming platform to leverage their sales data to offer better customer service and improve sales overall. But if we went back six months before that point, we would have seen ZMart looking to get a handle on its data situation. ZMart had a custom solution that worked well initially but is now becoming unmanageable for reasons we will soon see.

### 2.3.1 ZMarts Original Data Platform

Originally ZMart was a small company and had retail sales data flowing into its system by separate applications. The initial approach worked fine initially, but over time it was evident that a new approach would be needed. Data from sales in one department is not an isolated concern for that department alone. Several areas of the company are interested, and each area has a different take on what's important and how they want to the data structured. Here's how ZMart's data platform looked initially:

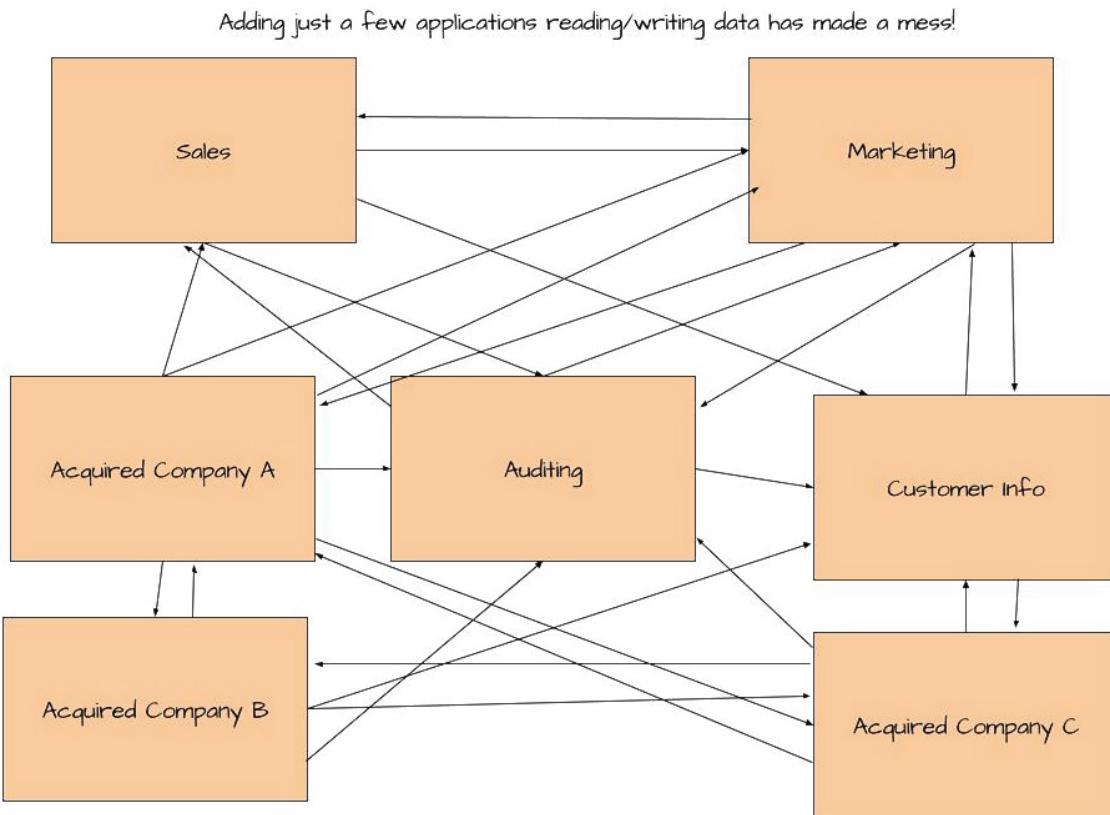
Initially, only a small number of applications consuming/writing data  
not an ideal architecture, but manageable at this point.



**Figure 2.1 This is the original data architecture for ZMart. In the beginning, it was simple enough to have information flowing to and from each source of information.**

Over time ZMart has continued to grow by acquiring other companies and expanding their offerings in existing stores. With each addition, the "web" of connected applications continued to become more complicated. What initially started out with a handful of applications communicating with each other, has turned into a veritable pile of spaghetti trying to make all these existing applications aware of each other to share information.

As you can see from the graphic below even with just three applications, the amount connections needed is cumbersome and confusing.

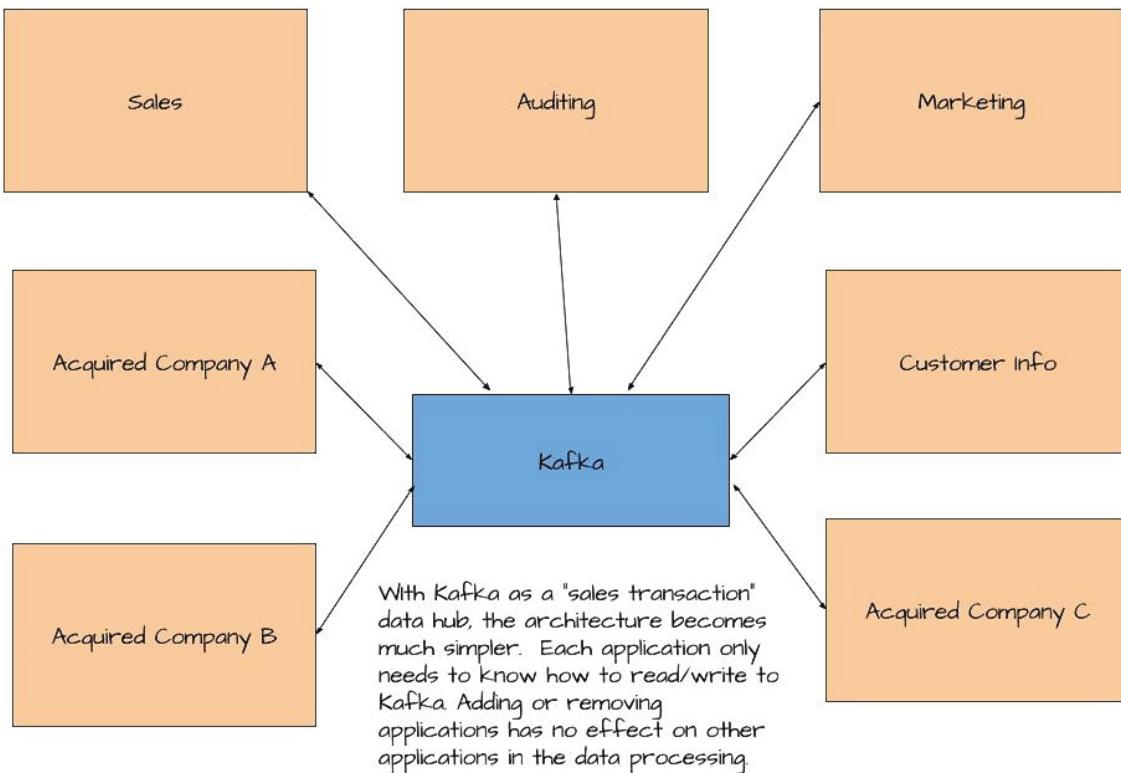


**Figure 2.2 With more machines added over time, the level of complexity trying to connect all these information sources became somewhat of a mess.**

You should be able to see how adding any additional applications will make this data architecture unmanageable over time.

### 2.3.2 Sales Transaction Data Hub

What's needed is one intake process to hold all transaction data, a transaction data hub. Additionally, this transaction data hub should be stateless, in that it accepts transaction data and stores it in such a fashion that any consuming application can pull the information it needs and it's up to the consuming application to keep track of what it has seen. Here's an image showing how you could envision Kafka as the sales transaction data hub.



**Figure 2.3** Here we can see how using Kafka as a sales transaction hub simplifies the ZMart data architecture significantly. Now each machine does not need to know about every other source of information. All that is needed is how to read from and write to Kafka.

The transaction data hub will only maintain how long it's been holding any transaction data and when that data should be rolled off or deleted.

### 2.3.3 Enter Kafka

In case you haven't guessed it yet, we have the perfect use-case here for Kafka. Kafka is a fault tolerant, robust publish/subscribe system. A single Kafka node is called a broker and other than development on your local machine, you will have more than one Kafka servers that make up a cluster. Kafka stores messages are written by Producers in topics. Consumers subscribe to topics and contact Kafka to see if there are messages available in those subscribed topics(s).

The preceding section is an overview of Kafka from 50,000 feet now let's take a deeper look in the following sections.

## 2.4 Kafka Architecture

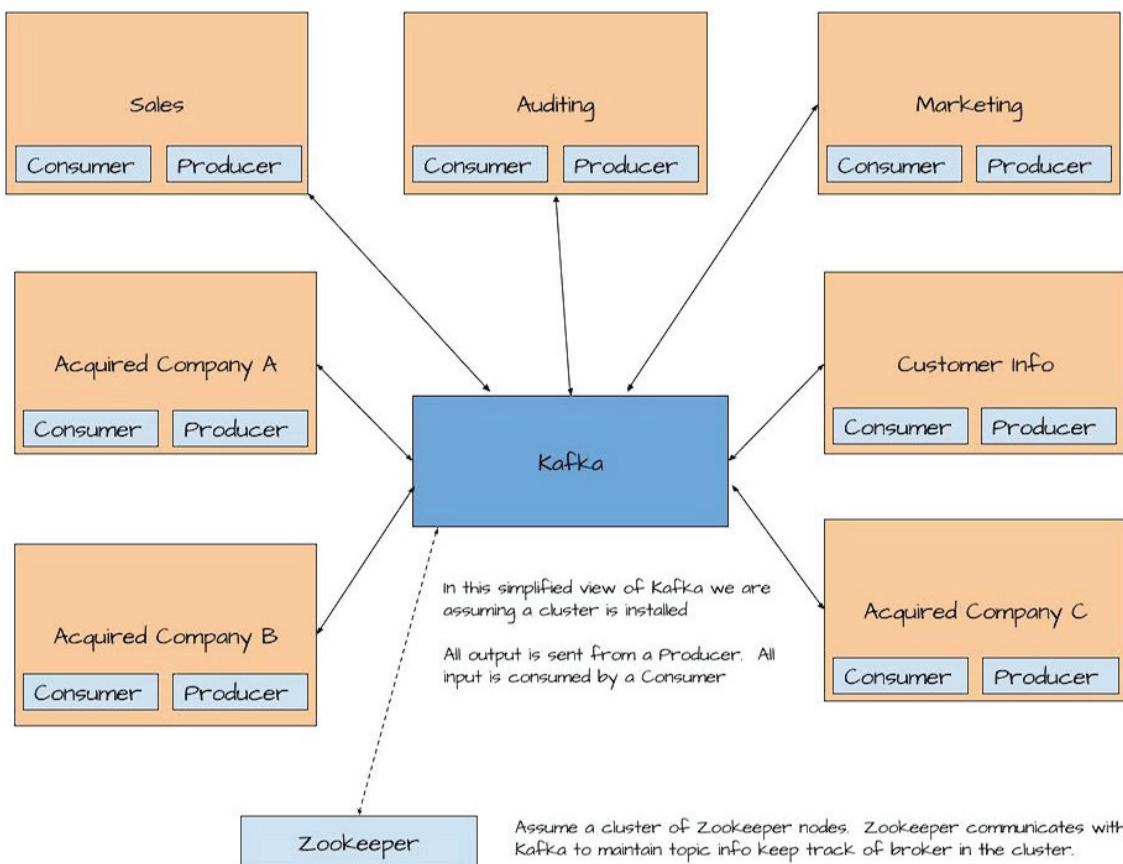
In this section, we will take a much deeper look into how Kafka works and key parts of its architecture.

**NOTE****For the Impatient**

In the next several sections we will be discussing the architecture of Kafka in depth. If you are interested in kicking the tires on Kafka sooner rather than later feel free to skip ahead to the section titled "Installing Kafka." After going through the install process come back to this point to continue learning about Kafka.

### 2.4.1 Kafka Is A Message Broker

Earlier we stated that Kafka is a publish/subscribe system, but to be more precise Kafka acts a message broker. A broker acts as an intermediary bringing two parties together (that don't necessarily know each other) for a mutually beneficial exchange or deal. In figure 4 below we see the evolution of the ZMart data infrastructure from what we covered in sections 2 and 3. In this image, we've added the producer(s) and consumer(s) to show how the individual parts communicate with Kafka, but not directly with each other.



**Figure 2.4 Kafka is a message broker. We can see here how producers send messages to Kafka and those messages are stored and made available to consumers via subscriptions to topics.**

Kafka stores messages in topics and retrieves messages from topics. There is no connection at all between the producers and the consumers of the messages. Additionally,

Kafka does not keep any state regarding the producers and consumers; it acts solely as a message "clearing house."

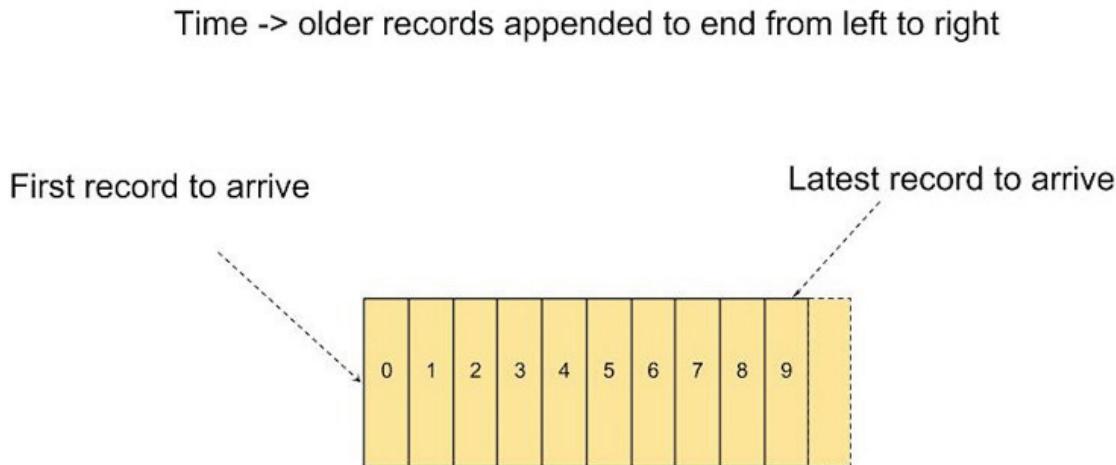
The underlying technology of a Kafka topic is a "log". We'll learn what a log is in the next section, but for now, we can say that it is a file where Kafka appends incoming records to the end. To help manage the load of messages coming into a topic Kafka uses partitions. We covered partitions in Chapter 1 and if you recall one use of partitions is to bring data located on different machines together on the same server. We'll discuss partitions in detail in the section after we discuss logs.

### 2.4.2 Kafka Is A Log

The mechanism underlying Kafka is the notion of a log. Now, most software engineers have a familiarity with logs regarding tracking what an application is doing. Having performance issues or errors in your application? The first place to check is the application logs. But, that is not the definition of a log we are using here. The definition of a log in the context of Kafka (or any other distributed system) is "It is an append-only, totally ordered sequence of records ordered by time"<sup>1</sup>. Here's an image of what a log looks like:

---

Footnote 1 Jay Kreps [log-what-every-software-engineer-should-know-about-real-time-datas-unifying](#)



**Figure 2.5 A log is a file where incoming records are always appended. A newly arrived record is placed immediately after the last record received. This process gives time ordering of records contained in the file.**

An application appends records to the end of the log as they arrive. Records have an implied ordering by time since the earliest records are to the left and the last record to arrive is at the end to the right, even though there might not be a timestamp associated with each record. Logs are a very simple data abstraction with very powerful implications. If you have records in complete order with respect to time, making conflict resolution or determining which update to apply to different machines becomes straight forward, the latest record wins.

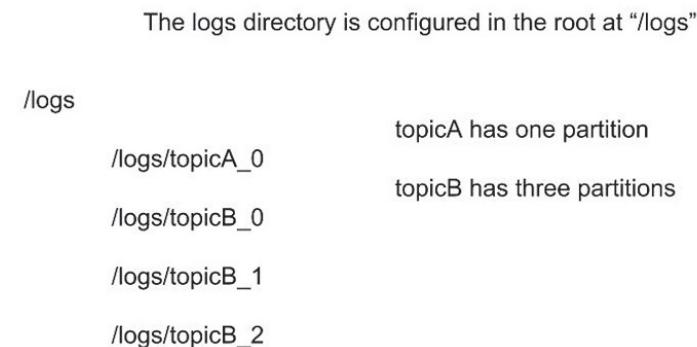
Now to take the next step, topics in Kafka are simply logs that are segregated by the

topic name. We could almost think of topics as "labeled" logs.

If the log is replicated among a cluster of machines, if a single machine goes down, it becomes easier to bring that server back up, just replay the log file. The ability to recover from failure is precisely the role of a distributed commit log. Hopefully, now the concept of a "distributed commit log" is a little more clear at this point. We have only scratched the surface of a very deep topic when it comes to distributed applications and data consistency, but what we have covered so far should be enough to get an understanding of what is going on under the covers with Kafka. When we start to discuss how Kafka works, all the information presented here will be very apparent.

### 2.4.3 How logs work in Kafka

When installing Kafka, one of the configuration settings is the "logs" directory. This configuration parameter specifies where Kafka stores messages in the "log" that just discussed. Each topic maps to a directory under the specified log directory location. If the topic is partitioned, then there are as many directories as partitions with a number appended to the end (we cover partitions in detail in the next section). Inside each directory is the log file where incoming messages get added to the end. Once the log files reach a certain size (either number of records or size on disk), or when a configured time difference between message timestamps is met, the log file is "rolled", and Kafka appends incoming messages to a new log. Here's a quick graphical view:



**Figure 2.6 The logs directory is base storage for messages. Each directory under /logs represents a topic-partition. The structure of the directory starts with the name of the topic followed by an underscore which is followed by a number representing the partition number.**

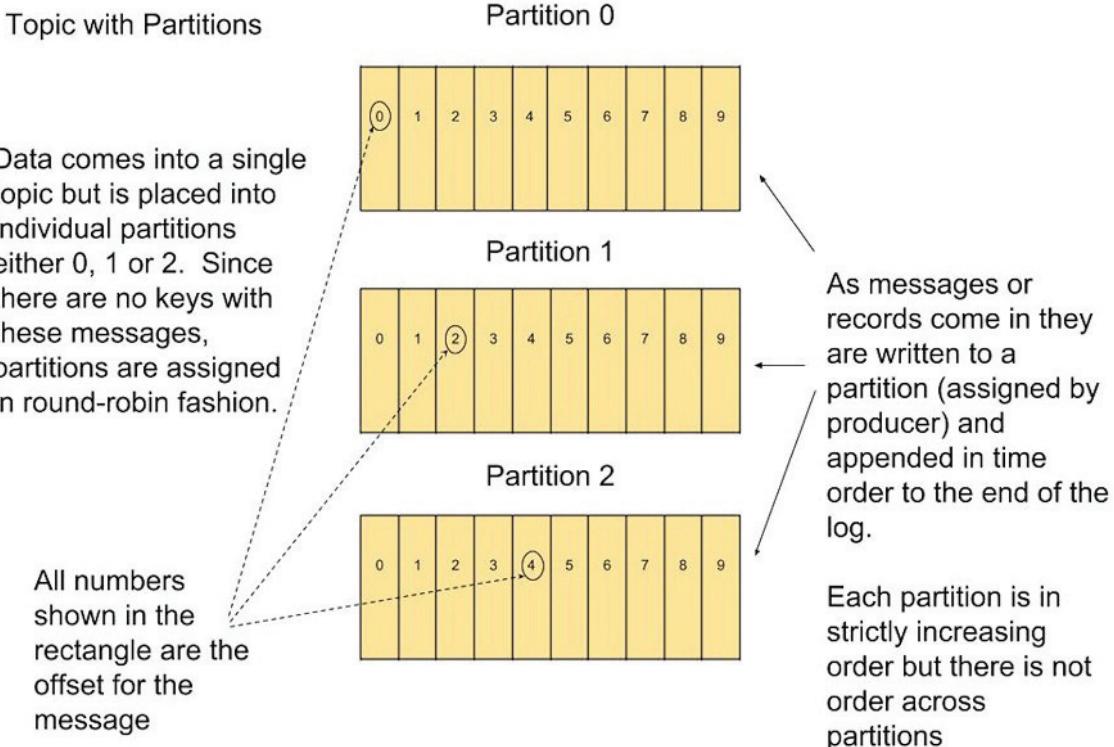
At this point, we can see that a log and topic are highly connected concepts. We could say that a topic is a log or represents a log and the topic name gives us a good handle on which log we need to store the messages sent to Kafka via producers.

Now that we have covered the concept of a log let's discuss another fundamental concept in Kafka, partitions.

### 2.4.4 Kafka and Partitions

Partitions are a critical part of Kafka's design as they are essential for performance and guarantee data with the same keys will be sent to the same consumer and in order. In figure 7, we dissect a topic to give us a clear idea of the role partitions play.

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.



**Figure 2.7 Kafka uses partitions to achieve high throughput and spread the amount of messages sent to an individual topic across several machines in the cluster.**

Partitioning a topic is essentially splitting the data forwarded to a topic across parallel streams. Partitioning is key to how Kafka can achieve its tremendous throughput. While we discussed that a topic is a distributed log, each partition is a log unto itself and follows the same rules. Kafka appends each incoming message to the end of the log, and all messages are strict time ordered. Each message has an offset number assigned to it. Order of messages across partitions is not guaranteed, but the order of each partition is guaranteed.

Partitioning serves another purpose aside from increasing throughput. Partitioning allows for spreading topic messages across several machines, so as to not restrict the capacity of a given topic to the available disk space on one server. Now that we've seen the role partitions play in helping throughput, next we'll take a look at another critical part partitions play; ensuring messages with the same keys end up together.

#### 2.4.5 Partitions Group Data by Key

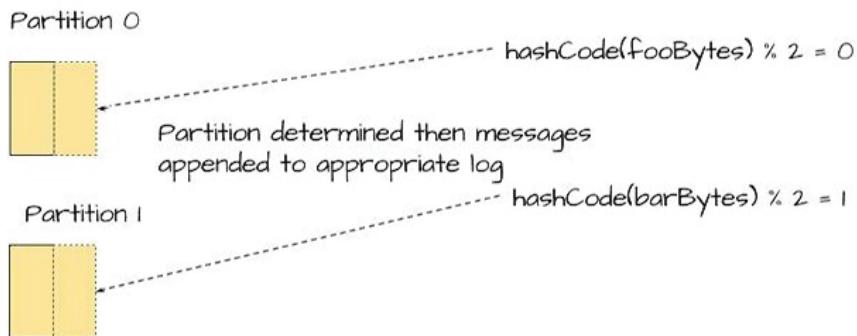
Kafka works with data in key-value pairs. If the keys are null, the Kafka Producer will write records to partitions in a round-robin fashion. Here's a graphical representation of how round-robin partition assignment operates:

Incoming messages:

```
{foo, message data}
{bar, message data}
```

Keys of the message are used to determine the partition the message should go to. In this case obviously the keys are not null.

The bytes of the key are used to calculate the hash



**Figure 2.8** Here we can see that "foo" is sent to partition 0 and "bar" gets sent to partition 1. You obtain the partition by hashing the bytes of the key modulus the number of partitions.

However, if the data has populated the keys Kafka uses the following formula (in pseudo code) to determine which partition to send the key-value pair to:

#### **Listing 2.1 Determining Partitions**

```
HashCode.(key) % number of partitions
```

By using a deterministic approach to finding the partition to use, records with the same key will *always* be sent to the same partition and in order. We should note the `DefaultPartitioner` uses this approach and should you need a different strategy to select partitions; you have the ability to provide your custom partitioner.

#### **2.4.6 Writing a Custom Partitioner**

Why would you want to write a custom partitioner? While there are several possibilities, we will provide one simple use case here; that is the use of composite keys.

For example, let's say you have purchase data flowing into Kafka and the keys contain two values: a customer id and a transaction date. But you need to group values by the customer-id, so taking the hash of the customer-id and the purchase date will not work.

In this case, we will need to write a custom partitioner that "knows" the correct part of our composite key to use for determining which partition to use.

Here the composite key looks like this:

## Listing 2.2 Composite Key found in src/main/java/bbejeck/model/PurchaseKey.java

```
public class PurchaseKey {

    private String customerId;
    private Date transactionDate;

    public PurchaseKey(String customerId, Date transactionDate) {
        this.customerId = customerId;
        this.transactionDate = transactionDate;
    }

    public String getCustomerId() {
        return customerId;
    }

    public Date getTransactionDate() {
        return transactionDate;
    }
}
```

But when it comes to partitioning, we need to ensure that all customer transactions go to the same partition. By using the key in its entirety, this won't happen, though. Since purchases are across many dates, including the date makes the key non-unique placing the transactions across random partitions. We need to ensure we send transactions with the same customer id to the same partition. The only way to do that only uses the customer id when determining the partition.

Here we present an example of a custom partitioner that does what's required. The `PurchaseKeyPartitioner` extracts the customer-id from the key for determining the partition to use.

## Listing 2.3 Custom Partitioner Code found in src/main/java/bbejeck/chapter\_2/partitioner/PurchaseKeyPartitioner.java

```
public class PurchaseKeyPartitioner extends DefaultPartitioner {

    @Override
    public int partition(String topic, Object key, byte[] keyBytes, Object value,
                        byte[] valueBytes, Cluster cluster) {

        if (key != null) {
            PurchaseKey purchaseKey = (PurchaseKey) key;
            key = purchaseKey.getCustomerId(); 1

            keyBytes = ((String) key).getBytes(); 2
        }

        return super.partition(topic, key, keyBytes, value, valueBytes, cluster); 3
    }
}
```

- ➊ If the key isn't null extract customer id
- ➋ Set the key bytes to the new value
- ➌ Return the partition with the updated key, delegating to the superclass.

With our custom partitioner here we have chosen to extend the `DefaultPartitioner`. While we could have implemented the `Partitioner` interface directly, there is existing logic in the `DefaultPartitioner` that we want to leverage. Keep in mind when creating a custom partitioner you are not limited to using only the key. Using the value alone or in combination with the key is valid as well.

**NOTE****Partitioner Interface**

The Kafka API provides a Partitioner interface should you need to write a custom partitioner. We won't be covering writing your partitioner from scratch, but the principals would be the same as presented in the example.

We've just seen how to construct a custom partitioner; next is how to wire the partitioner up with Kafka.

#### **2.4.7 Specifying a Custom Partitioner**

Now that we have written a custom partitioner we need to tell Kafka we want to use it over the default partitioner. Although we haven't covered producers yet, you specify a different partitioner when configuring the Kafka Producer.

##### **Listing 2.4 Partitioner Config Setting**

```
partitioner.class=bbjeck_2.partition.PurchaseKeyPartitioner
```

We'll go over producer configuration in detail when we cover using Kafka producers, but as you may have surmised, by setting a partitioner per producer instance we are free to choose to use any partitioner class per producer.

**WARNING****Choosing Keys**

You must exercise some caution when choosing keys to use and when using parts of a key/value to partition on. Make sure that the choice of the key to use has a fair distribution across all of your data. Otherwise you'll end up with a "data skew" problem as most of your data is centered on just a few of your partitions.

#### **2.4.8 Determining the Correct Number of Partitions**

The number of partitions to use when creating a topic is part art and part science. One of the key considerations for the number of partitions is the amount of data flowing into a given topic. More data implies more partitions for higher throughput. But as with anything in life, there are tradeoffs for everything.

Increasing the number of partitions increases the number of TCP connections and open file handles. Additionally, depending on how long it takes to process an incoming

record in a consumer will also determine throughput. If you have heavyweight processing in your consumer, adding more partitions may help some, but ultimately the slower processing will hinder performance.<sup>2</sup>.

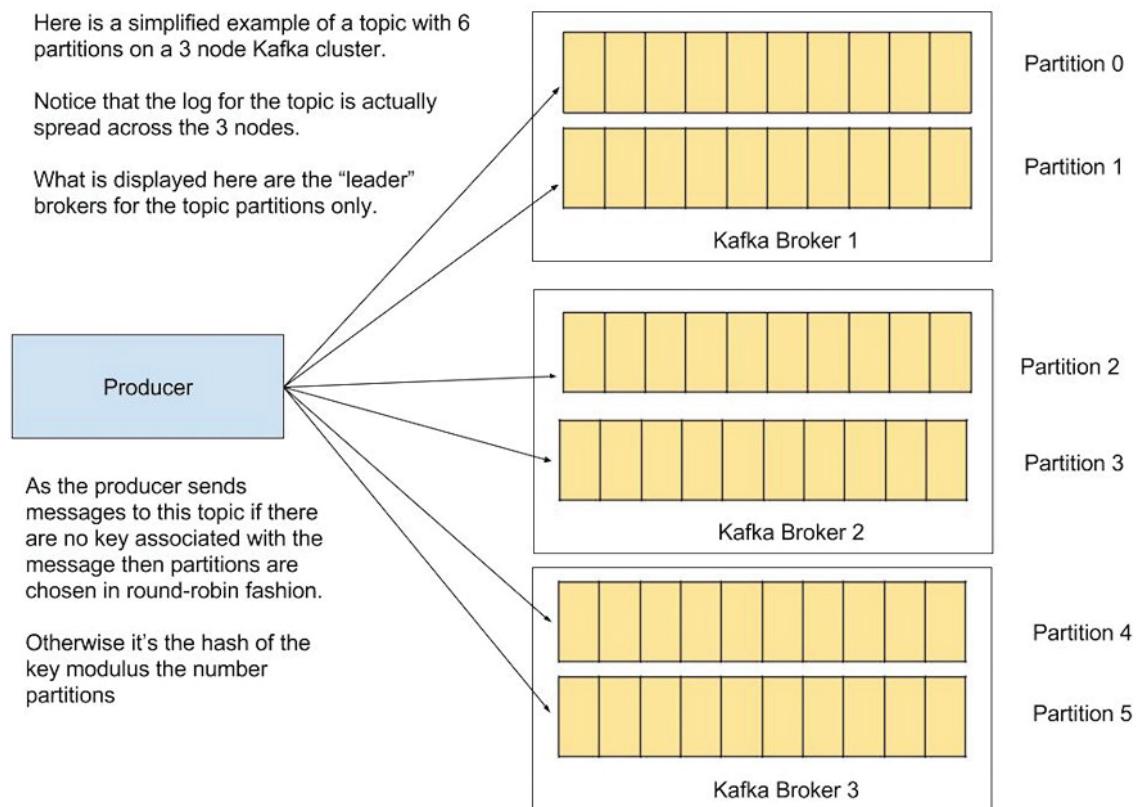
Footnote 2 Jun Rao

[www.confluent.io/blog/how-to-choose-the-number-of-topics-partitions-in-a-kafka-cluster/](https://www.confluent.io/blog/how-to-choose-the-number-of-topics-partitions-in-a-kafka-cluster/)

### 2.4.9 The Distributed Log

Up to this point, we have discussed the concept of a log and partitioned topics. Let's take a minute to explain those two concepts together to help us understand a distributed log.

So far, all of our discussions on Kafka have centered on logs and topics on one Kafka server or broker. But typically a Kafka production cluster environment includes several machines. We have intentionally kept the discussion centered around one node as it is easier to understand the concepts introduced when considering one node. But in practice, we will always be working with a cluster of machines in Kafka. The following image gives you a pictorial representation of message distribution to partitions across a cluster.



**Figure 2.9 A producer writes messages to partitions of a topic. If there is no key associated with the message the producer chooses one in a round-robin fashion. Otherwise the hash of the key modulo the number of partitions is used.**

When a topic is partitioned Kafka does not allocate those partitions on one machine, Kafka spreads them across other machines in the cluster. As Kafka appends records to the log, Kafka is distributing those records across several machines by partition. If you take another look at figure 9, you can view this process in action. Let's walk through a

quick example using the image above as a guide. For this example, we will assume one topic and null keys, so the producer assigns partitions in a round-robin manner.

The producer sends its first message to partition-0 on Kafka-Broker-1, the second message to partition-1 on Kafka-Broker-1, and the third message to partition-2 on Kafka-Broker-2. When the producer sends its sixth message, it goes to partition-5 on Kafka-Broker-3, and the next message starts over going to partition-0 on Kafka-Broker-1. Message distribution continues in this manner spreading message traffic across all nodes in the Kafka cluster.

While that might sound like your data is at risk, Kafka offers data redundancy. Data is replicated to one or more machines in the cluster as you write to one broker in Kafka (we'll cover replication in an upcoming section).

#### **2.4.10 Zookeeper - Leaders, Followers, and Replication**

Up to this point, we have discussed the role topics in Kafka and how and why topics are partitioned. We discussed how partitions are not co-located on one machine but are spread out on brokers throughout the cluster. Now we have come to another crucial juncture in Kafka's architecture and how Kafka provides data availability in the face of machine failures.

Kafka has the notion of leader and follower brokers. In Kafka, for each topic-partition, one broker is "chosen" as the *leader* for the other brokers (the *followers*). One of the chief duties of the leader is to assign *replication* of topic-partitions to the following brokers. Just as Kafka allocates partitions for a topic across the cluster, Kafka also places replication of the partitions across machines.

But before we go into details on how leaders, followers and replication work we need to discuss the technology Kafka uses to achieve this.

#### **2.4.11 Apache Zookeeper**

For the complete Kafka newbie, you may be asking yourself "Why are we talking about Apache Zookeeper in a Kafka book?". Apache Zookeeper is integral to Kafka's architecture, and it's Zookeeper that enables Kafka to have leader brokers and to do such things as tracking the replication of topics.

ZooKeeper is a centralized service for maintaining configuration information, naming providing distributed synchronization, and providing group services. All of these kinds of services are used in some form or another by distributed applications.

-- [zookeeper.apache.org](http://zookeeper.apache.org)

Given that Kafka is a distributed application, it should start becoming clearer how ZooKeeper is involved in Kafka's architecture. (For our discussion here we will only consider Kafka installations where there are 2 or more Kafka servers installed.) In a Kafka cluster, one of the brokers is "elected" as the "controller."

We covered partitions in the previous section and discussed how Kafka spreads partitions across different machines in the cluster. Topic partitions have a leader and

follower(s) (the level of replication determines the degree of replication). When producing messages, Kafka sends the record to the broker that is the *leader* for the record's partition.

### 2.4.12 Electing A Leader

Kafka uses Zookeeper for the election process for the controller broker. Discussing consensus algorithms are way beyond the scope of this book, so we'll take the 50,000-foot view by stating Zookeeper elects a broker out of the cluster to be the controller (hopefully with less contention than 2016).

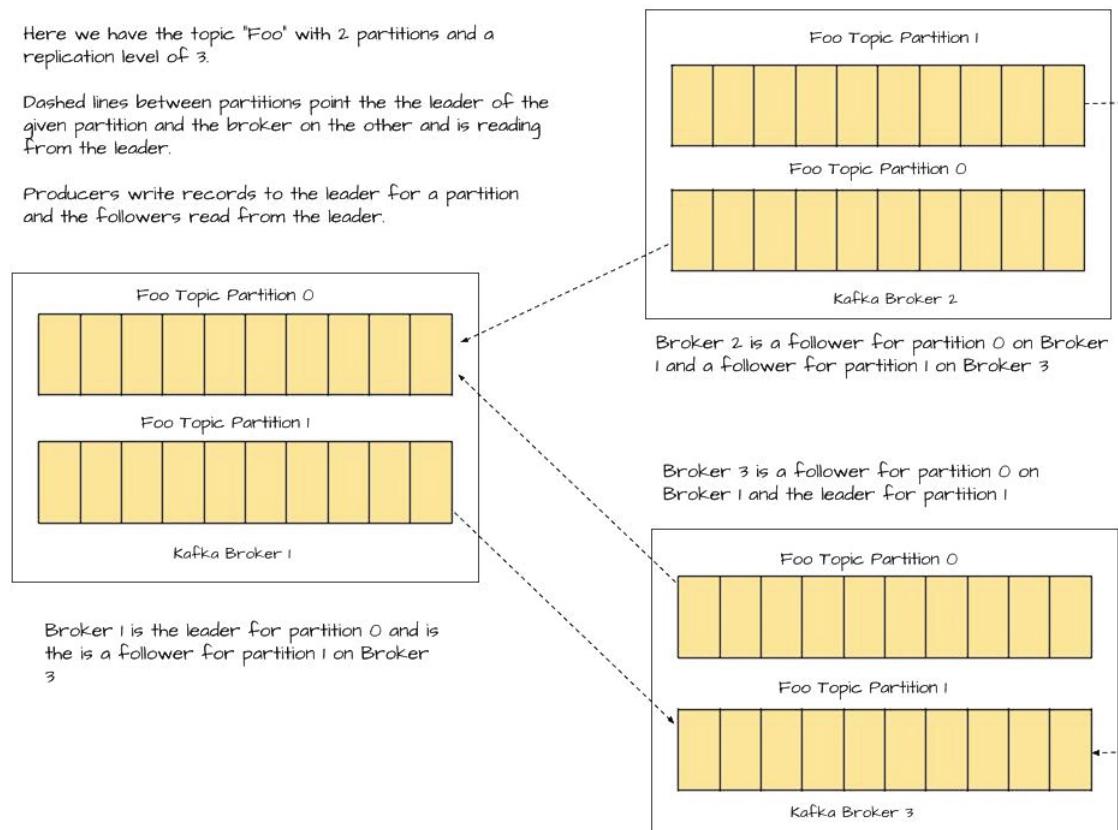
If the controlling broker fails or becomes unavailable for any reason, Zookeeper elects a new controller from a set of brokers that are considered to be caught up with the leader (an in-sync replica or ISR). The brokers that make up this set are dynamic, and zookeeper only recognizes brokers included in this set for election as leader.<sup>3</sup>

---

Footnote 3 [kafka.apache.org/documentation/#design\\_replicatedlog](https://kafka.apache.org/documentation/#design_replicatedlog)

### 2.4.13 Replication

Kafka replicates records between brokers to ensure data availability should a broker in the cluster fail. You can set the level of replication individually for each topic (as we saw in our publishing and consuming example above) or for all topics in the cluster. The following image (figure 11) demonstrates the replication flow between brokers.



**Figure 2.10** Brokers 1 and 3 in this image are a leader for one topic-partition and a follower for another, while Broker 2 is a follower only. Replication is the process of the following broker copying data from the lead broker. The dashed lines represent where the

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

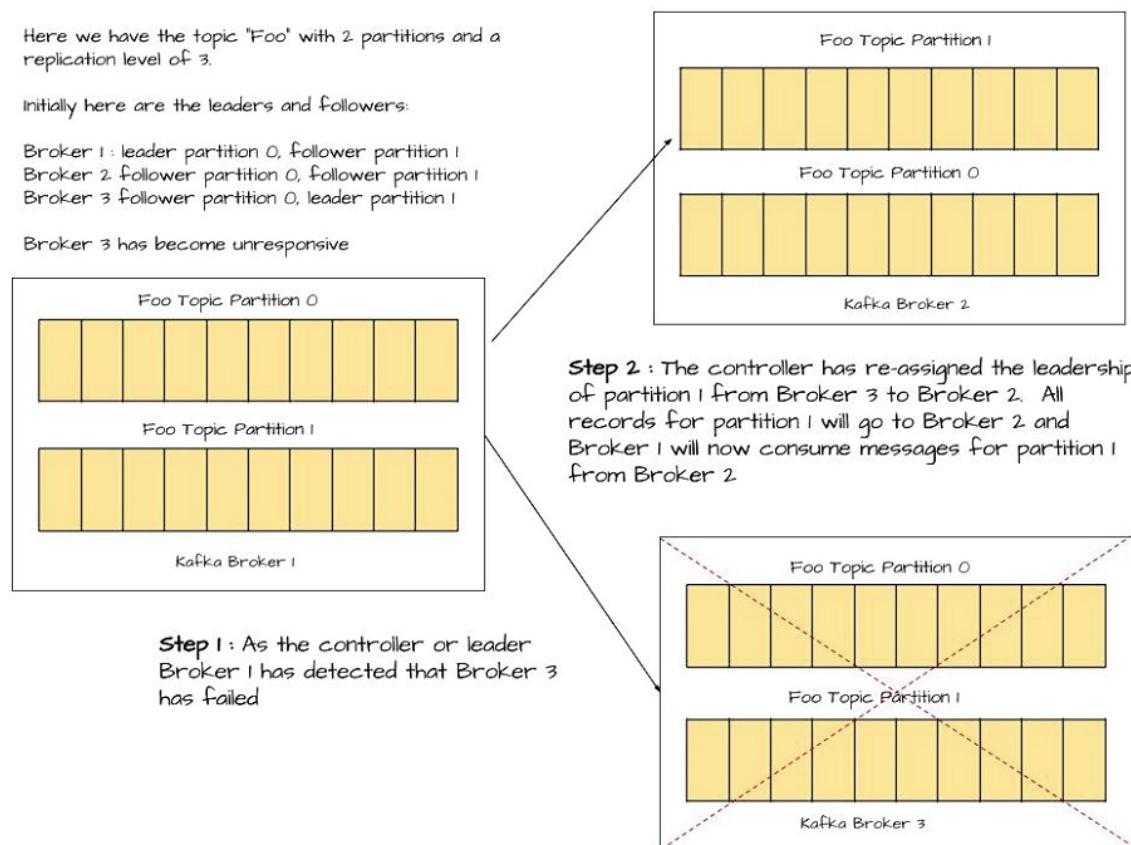
**following broker will copy data from.**

The Kafka replication process is straight forward. Brokers that "follow" a topic-partition consume messages from the topic-parition leader and append the records to their log. As we discussed in the section "Electing A Leader," "following" brokers that are caught up with the leader broker are considered an in-sync replica (ISR). ISR brokers are available to be elected leader should the current leader fail or become unavailable<sup>4</sup>.

Footnote 4 [kafka.apache.org/documentation/#replication](https://kafka.apache.org/documentation/#replication)

#### 2.4.14 Controller Responsibilities

The controller broker is responsible for setting up leader/follower relationships for all partitions of a topic. If a Kafka node dies or is unresponsive (to Zookeeper heartbeats) all of its assigned partitions (leading and following) are reassigned by the controller broker. Figure 10 below gives us a brief example of a controller broker in action:



**Figure 2.11 The controller broker is responsible for making assignments to other brokers to be the lead broker for some topics/partitions and a follower for other topics/partitions. When a broker becomes unavailable, the controller broker will re-assign the failed broker's assignments to other brokers in the cluster.**

In the image above we see the demonstration of a simple failure scenario. In step one the controller broker detects that broker 3 is not available. In step two, the controller broker re-assigns leadership of the partitions on broker-3 to broker-2.

Zookeeper is also involved in the following aspects of Kafka operations:

- Cluster Membership - Joining the cluster and maintaining membership in the cluster. If a broker becomes unavailable, Zookeeper removes the broker from cluster membership.
- Topic configuration - what are the topics in the cluster, which broker is the leader for a topic, what are the number of partitions for a topic and what are the specific configuration overrides for a given topic.
- Access Control - who can read and write from particular topics.

5

---

Footnote 5 This section derived information from answers given by Gwen Shapira on [www.quora.com/What-is-the-actual-role-of-ZooKeeper-in-Kafka/answer/Gwen-Shapira](http://www.quora.com/What-is-the-actual-role-of-ZooKeeper-in-Kafka/answer/Gwen-Shapira)

In this section, we covered why Kafka has a dependency on Apache Zookeeper. It's important for you to come away with the idea that its Zookeeper that enables Kafka to have a leader broker with followers. The head broker has the critical role of assigning topic-partitions for replication to the followers as well as reassignments in the case of when a member broker fails.

#### **2.4.15 Log Management**

While we have covered appending messages, we haven't talked about how to manage the logs as they continue to grow. The amount of space on spinning disks in a cluster is a finite resource, so it is important that Kafka removes messages over time.

When it comes to handling the removal of old data in Kafka, there are two approaches, the "traditional" log deletion approach and compaction.

#### **2.4.16 Deleting Logs**

The log deletion strategy is a two-phased approach. The first phase is rolling the log into segments. The second phase is deleting the oldest segments.

To manage the increasing size of the logs, Kafka "rolls" them into segments. The timing of log rolling is based on timestamps embedded in the messages. Kafka "rolls" a log when a new message arrives and its timestamp is greater than the timestamp of the first message in the log plus the `log.roll.ms` configuration. At that point, the log is rolled and a new segment is created serving as the active log. The previous active segment is still used to retrieve messages for consumers. Log rolling is a configuration setting you can specify when standing up a Kafka broker<sup>6</sup>. There are two options for log rolling:

---

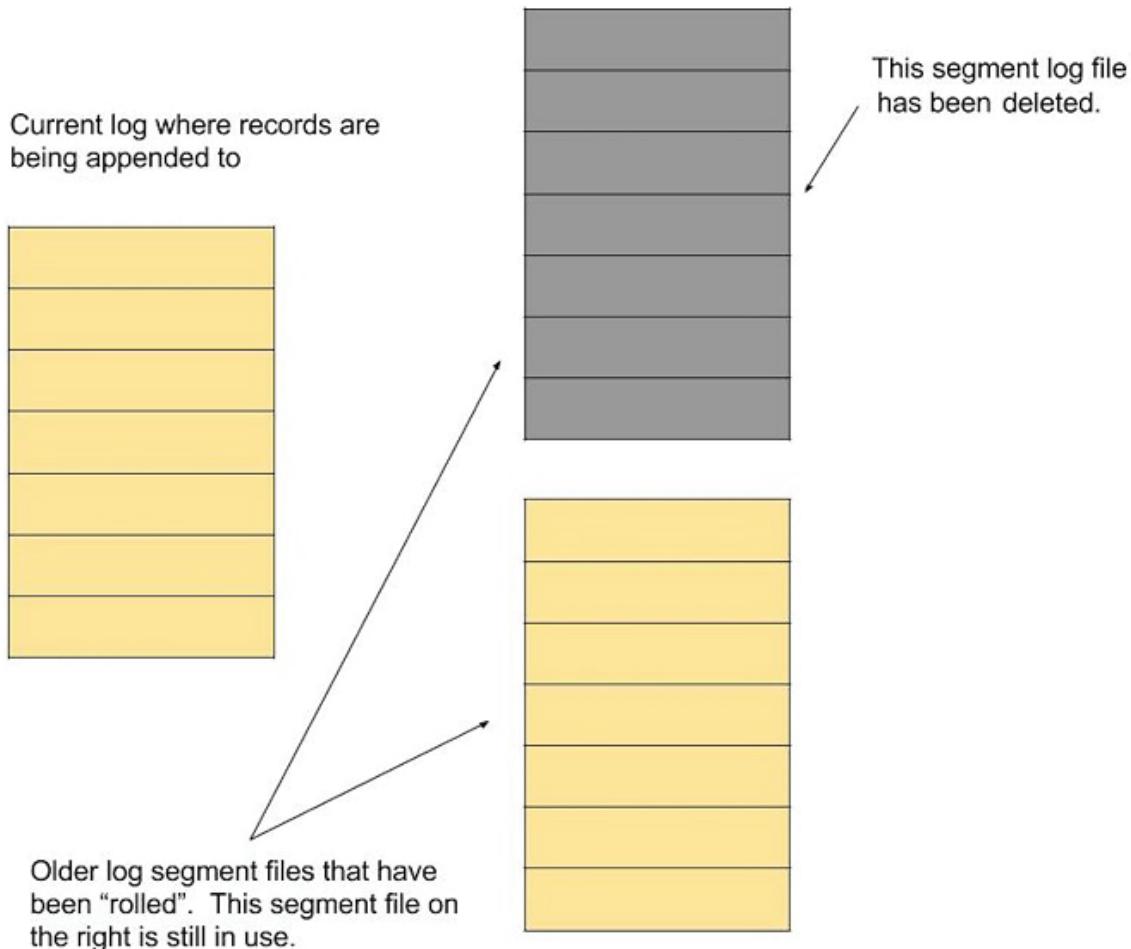
Footnote 6 [kafka.apache.org/documentation/#brokerconfigs](http://kafka.apache.org/documentation/#brokerconfigs)

#### **Listing 2.5 Log Rolling Configuration**

```
log.roll.ms - primary configuration, but no default value
log.roll.hours - secondary configuration, only used if log.roll.ms is not set
```

Over time the number segments will continue to grow, and older segments will need to be deleted to make room for incoming data. To handle the deletion you can specify how long to retain the segments. In figure 12 you see the process of "log rolling," and the

gray section represents a recently deleted segment.



**Figure 2.12** Here on the left side of the image we have the current log segments. The gray portion accounts for a deleted log segment, and the one below represents a recently rolled one still in use.

Like the log rolling removal of segments is based off timestamps in the message and not just wallclock time or when the file was last modified. Log segment deletion uses the largest timestamp in the log. We will list three settings in order of priority meaning that configurations earlier in the list are used first if present.

#### **Listing 2.6 Log Segment Retention**

```
log.retention.ms - how long to keep a log file in milliseconds.  
log.retention.minutes - how long to keep a log file only in minutes.  
log.retention.hours - log file retention in hours.
```

I presented these settings with the implicit assumption of very high volume topics, wherein a given period you are guaranteed to reach this maximum file size. Another configuration setting `log.retention.bytes` could be specified with a higher rolling

time threshold to keep additional I/O operations down. Finally, to guard against the case of a significant spike in volume when there are relatively large roll settings, there is the `log.segment.bytes` setting which governs how large an individual log segment can be.

We've covered deletion of logs in this section. While this approach works well for non-keyed records or records that stand alone, if you have keyed data and expected updates there is another method that will suit your needs better.

#### 2.4.17 Compacted Logs

Consider the case where you have keyed data, where you will receive updates for that key over time, meaning a new record with the same key, is an update of the previous value. Think of a simple case of stock ticker symbol as the key and the price per share as the value. If you are using that information to display some dashboard application in the case of a crash or restart, you will need to be able to start back up with the latest data for each key.<sup>7</sup>

---

Footnote 7 [kafka.apache.org/documentation/#compaction](https://kafka.apache.org/documentation/#compaction)

But under the deletion policy, a segment could get removed in between the time the next update comes, and the application crashes or has a restart. You will not have all the records you received for a given key on startup. It would be better to be able to maintain the last known value for a given key, treating the new record with the same key as we would if it were an update to a database table.

Updating records by key is the behavior that compacted topics (logs) deliver. Instead of a coarse grained approach of deleting entire segments by time or size, compaction is more fine grained and deletes old records *per key* in a log. At a very high level, a log cleaner (a pool of threads) runs in the background recopying log segment files, removing records where there is an occurrence later in the log with the same key. Here in figure 13, we can see how log compaction works leaving the most recent message for a given key.

Before Compaction			After Compaction		
offset	key	value	offset	key	value
10	foo	A			
11	bar	B			
12	baz	C			
13	foo	D	11	bar	B
14	baz	E	15	boo	F
15	boo	F	16	foo	G
16	foo	G	17	baz	H
17	baz	H			

**Figure 2.13** On the left is a log before compaction. You'll notice duplicate keys with different values which are essentially updates for the given key. On the right is the log after compaction. Notice the latest value for each key is retained but the log is smaller in size.

Using this approach guarantees we have last seen record for a given key is in the log. You can specify log retention per topic, so it's entirely possible to have some topics that use time-based retention and other topics using compaction.

By default the log cleaner is enabled and to utilize compaction for a topic, you will need to set this property `log.cleanup.policy=compact` when creating the topic in question. We should note at this point that compaction is used in Kafka Streams when using State Stores but you will not be creating those logs/topics yourself; the framework handles that task. But it's important to have an understanding of how compaction works at any rate.

**NOTE**
**How are Records Removed from a Compacted Topic?**

With a `cleanup.policy` of `compact` you might ask how to remove a record from the log. With a compacted topic the deletion semantics are to provide a `null` value for a given key, setting a tombstone marker. Any key with a `null` value ensures that any prior record with the same key is removed and the tombstone marker itself is removed after a period of time.

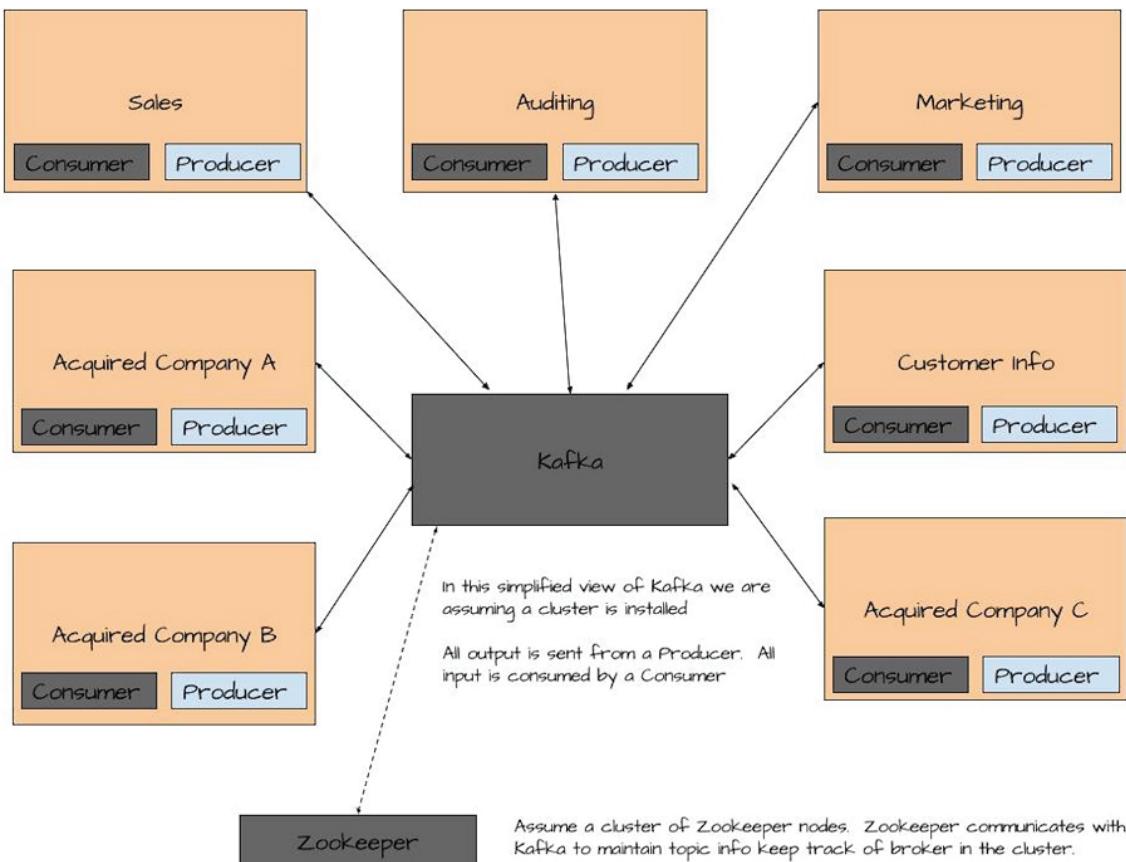
Log compaction is a broad subject which could take almost an entire chapter by itself. For more information, the reader is encouraged to view the [compaction documentation](#).

The key takeaway from this section is this: if you have independent, stand-alone events/messages use log deletion otherwise if you have updates to events/messages you'll want to use log compaction.

We've spent a good deal of time covering how Kafka handles data internally; now it's time to move outside of Kafka and discuss how we send messages to Kafka with producers and read messages from Kafka with consumers.

## **2.5 Sending Messages with Producers**

Going back to ZMart's need for a centralized sales transaction data hub, we now have Kafka in place. Now let's get specific on how we send purchase transactions into Kafka. In Kafka, the `Producer` is the client used for sending messages. Below we revisit the view of ZMart's data architecture and have highlighted the producers to emphasize where they fit into the data flow.



**Figure 2.14** Pictured here are the Producers used to send messages to Kafka . Keep in mind that each producer has no idea about which consumer or when the consumers read the messages. Some parts of the diagram are shaded to draw your attention to the producers.

Although ZMart has a lot of sales transactions for simplicity, we are going to consider the purchase of a single item a book costing \$10.99. When the customer completes the sales transaction, the information is converted into a key-value pair and sent to Kafka via a Producer.

The key is the customer id "123447777" and the value is in JSON format "`{"item": "book", "price": 10.99}`" (we've escaped the double quotes so we can represent the JSON as a String literal in Java). With the data in this format we use a Producer to send the data to our Kafka cluster:

#### Listing 2.7 Producer Example found in src/main/java/bbejeck.chapter\_2/producer/SimpleProducer.java

```

Properties properties = new Properties();
properties.put("bootstrap.servers", "localhost:9092");
properties.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
properties.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");
properties.put("acks", "1");
properties.put("retries", "3");
properties.put("compression.type", "snappy");

properties.put("partitioner.class", PurchaseKeyPartitioner.class.getName());

PurchaseKey key = new PurchaseKey("1234568", new Date());
    
```

1

```

try(Producer<PurchaseKey, String> producer = new KafkaProducer<>(properties)) {
    ProducerRecord<PurchaseKey, String> record = new ProducerRecord<>("transactions", key,
        "{\"item\":\"book\", \"price\":10.99}");
    ②
    ③
}

Callback callback = (metadata, exception) -> {
    if (exception != null) {
        System.out.println("Encountered exception " + exception);
    }
};

Future<RecordMetadata> sendFuture = producer.send(record, callback);
} ④
⑤
}

```

- ① Properties for configuring a producer.
- ② Creating the KafkaProducer.
- ③ Instantiating the ProducerRecord.
- ④ Building a Callback.
- ⑤ Sending the record and setting the returned Future to a variable.

Kafka producers are thread-safe. All sends to Kafka are asynchronous as `Producer.send` returns immediately once the producer places the record in an internal buffer. The buffer is used to send records in batches. Depending on your configuration, there could be some blocking if one attempts to send a message while producers' buffer is full.

The `Producer.send` method depicted here takes a `Callback` instance. Once the lead broker acknowledges the record, the producer fires `Callback.onComplete` method. Only one of the arguments will be non-null. In this case, we are only concerned with printing out the stacktrace in the event of error. The returned `Future` yields a `RecordMetadata` object once the server acknowledges the record.

#### NOTE

#### What's a Future?

In the previous example, we discussed how the `Producer.send` method returns a `Future` object rather casually. A `Future` object represents the result of an asynchronous operation. More importantly, a `Future` gives you the option lazily retrieve asynchronous results vs. waiting for completion. For more information look here - [docs.oracle.com/javase/..//concurrent/Future](https://docs.oracle.com/javase/..//concurrent/Future)

When creating the `KafkaProducer` instance above, we passed a `java.util.Properties` parameter containing the configuration for the producer. While the configuration of a `KafkaProducer` is not complicated, there key properties to consider when setting up the producer. It's in these settings where we would specify a custom partitioner like the one demonstrated before.

There are too many properties to cover entirely here, so we will include the ones

presented from the example above. You are encouraged to review the [full producer configurations](#).

### 2.5.1 Producer Properties

#### **Listing 2.8 Bootstrap Servers**

```
bootstrap.servers is a comma separated list of host:port values.
```

Eventually, the producer will use all brokers in the cluster, this initial list for initially connecting to the cluster.

#### **Listing 2.9 Serialization**

```
key.serializer and value.serializer instruct Kafka how to convert the keys and values into byte arrays
```

Internally Kafka uses byte arrays for keys and values. So we need to provide Kafka the correct serializers to convert objects to byte arrays before sending across the wire.

#### **Listing 2.10 Acks**

```
acks are the minimum number of acknowledgments from a broker that the producer will wait for before considering a record send completed.
```

Valid values for "acks" are all, 0 or 1. With a value of "all", the producer will wait for a broker to receive confirmation that all followers have committed the record. When set to 1 the broker writes the record to its log but does not wait for any followers to acknowledge committing the record. Using a 0 means the producer will not wait for any acknowledgment, mostly fire and forget.

#### **Listing 2.11 Retries**

```
retries if sending a batch results in a failure, the number of times to attempt to re-send
```

Retries specify how many times to attempt to send a record when met with a failure. If record order is important, you should consider setting `max.in.flight.requests.per.connection` to 1 to prevent the scenario of a second batch being sent successfully then a failed record sent after a re-try.

#### **Listing 2.12 Compression Type**

```
compression.type is used to determine what compression algorithm to apply if any
```

If set, the "compression.type" instructs the producer to compress a batch before sending. Note that it's the entire batch that is compressed and not individual records.

### **Listing 2.13 Partitioner Class**

```
partitioner.class name of class implementing Partitioner interface
```

The `partitioner.class` included for completeness and related to our custom partitioner discussion in the section on partitions in Kafka.

### **2.5.2 Specifying Partitions and Timestamps**

When creating a `ProducerRecord`, we have the option of specifying a partition and or a timestamp. Instantiating the `ProducerRecord` in the above example is just one of 4 overloaded constructors. We could have used other constructors that allow for setting a partition and timestamp or just a partition.

### **Listing 2.14 ProducerRecord Constructors**

```
ProducerRecord(String topic, Integer partition, String key, String value)
ProducerRecord(String topic, Integer partition, Long timestamp, String key, String value)
```

### **2.5.3 Specifying a Partition**

In a previous section, we discussed the importance of partitions in Kafka. We also discussed how the `DefaultPartitioner` works and how we can supply our custom partitioner. So why would you explicitly set the partition? While there are a variety of business reasons why you would do so, here's one example.

You have keyed data coming in, but it does n't matter that the records go to the same partition as the consumers have logic to handle any values that the key might contain. Additionally, the distribution of the keys might not be even, so you want to ensure all partitions receive roughly the same amount of data. Here's a rough implementation of how you could do this:

### **Listing 2.15 Manually Setting the Partition**

```
AtomicInteger partitionIndex = new AtomicInteger(0); ①
int currentPartition = Math.abs(partitionIndex.getAndIncrement()) % numberPartitions;
ProducerRecord<String, String> record = new ProducerRecord<>("topic", currentPartition,
    "key", "value"); ②
```

- ① Creating an `AtomicInteger` instance variable.
- ② Getting current partition and using as a parameter.

**NOTE****What's an AtomicInteger?**

An `AtomicInteger` belongs to the `java.util.concurrent.atomic` package which comprises classes supporting lock-free thread-safe operations on single variables. For more information can be found here - [docs.oracle.com/javase/..//package-summary](https://docs.oracle.com/javase/..//package-summary)

### **2.5.4 Timestamps in Kafka**

Kafka version 0.10 added timestamps to records. You set the timestamp when creating a `ProducerRecord` via the overloaded constructor call `ProducerRecord(String topic, Integer partition, Long timestamp, K key, V value)`. If you don't set one, then the Producer sets a timestamp (current wallclock time) before sending the record to the Kafka broker.

Accompanying the addition of timestamps is the broker configuration setting, `log.message.timestamp.type`. This configuration can be set to either `CreateTime` or `LogAppendTime` (`CreateTime` is the default). Like many other broker settings, the value configured on the broker applies to all topics as a default value. But when creating a topic, you can specify a different value applying only to that topic.

If you specify `LogAppendTime` and the topic does not override the `log.message.timestamp.type` config, the broker will overwrite the timestamp to the current time when it appends the record to the log. Otherwise, the timestamp from `ProducerRecord` is used. Why would you choose one setting over another? Using the setting `LogAppendTime` is considered "processing time" and a setting of `CreateTime` is considered "event time". So it depends on your business requirements. You'll need to decide does it matter more to you when Kafka processed the record, or when the actual event occurred.

We will see the important role timestamps have regarding controlling data flow in Kafka Streams in later chapters. Now we will discuss the other side of the equation, consuming messages.

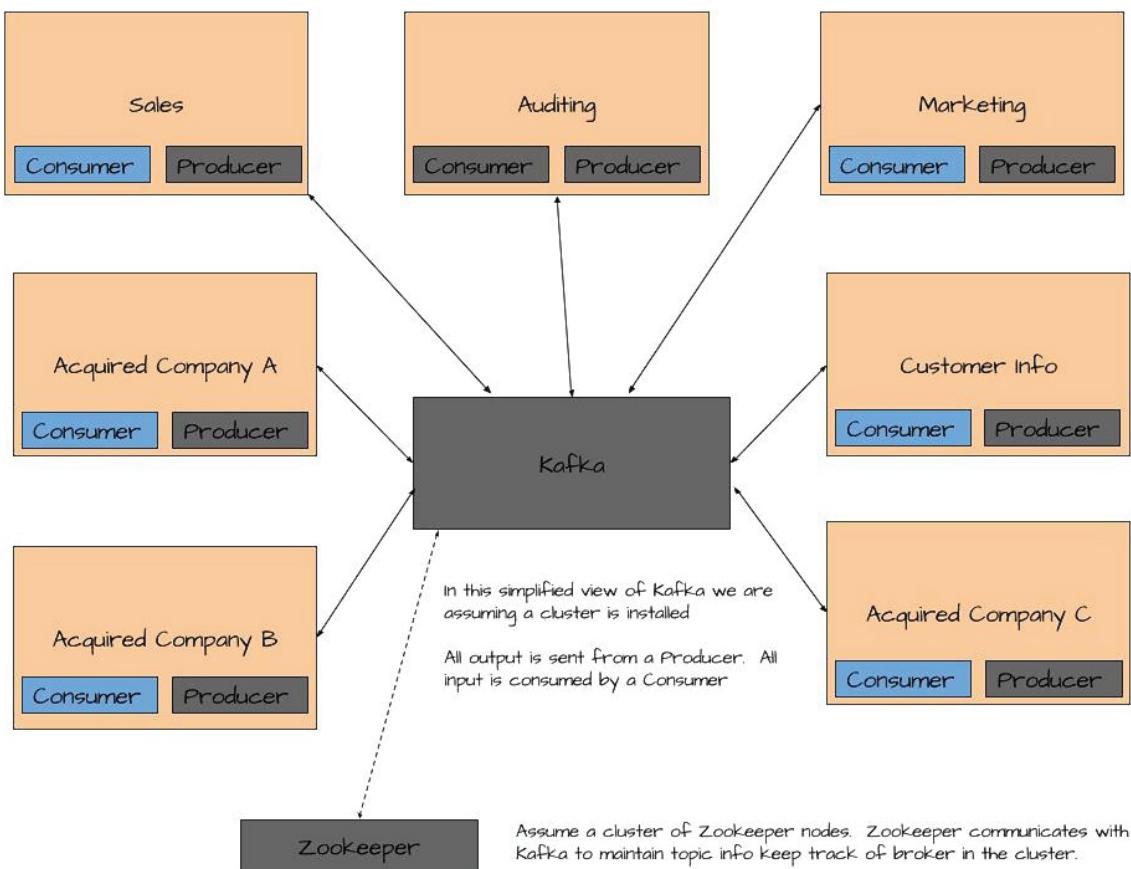
### **2.6 Reading Messages with Consumers**

Now let's move to the "other side" of producing, consuming messages. Now we in an internal application at ZMart and we are building a prototype application to show the latest sales statistics. Again even though ZMart has a large sales volume, we are just going to consider consuming the message we sent in the `Producer` example above. Since this prototype is in its earliest stages, all we are going to do at this point is consume the message and print out the information to the console.

The `KafkaConsumer` is the client used to consume messages from Kafka. The `KafkaConsumer` class is straight forward to use, but there are a few operational considerations to take into account. Since Kafka Streams requires Kafka version 0.10.2 or higher, we will only discuss the new consumer that was part of the Kafka 0.9 release.

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

Below we have updated our ZMart architecture diagram to highlight the where the consumer plays a role in the data flow.



**Figure 2.15** Here are the consumers that read messages from Kafka. Just as producers have no knowledge of the consumer, consumers read messages from Kafka with no knowledge of who produced the message.

## 2.6.1 Managing Offsets

While the `KafkaProducer` is essentially stateless, the `KafkaConsumer` manages some state by periodically committing the offsets of messages consumed from Kafka. Offsets uniquely identify a message and represent its starting position in the log. Consumers will need to commit the offsets of messages they have received periodically. Committing an offset has two implications for a consumer:

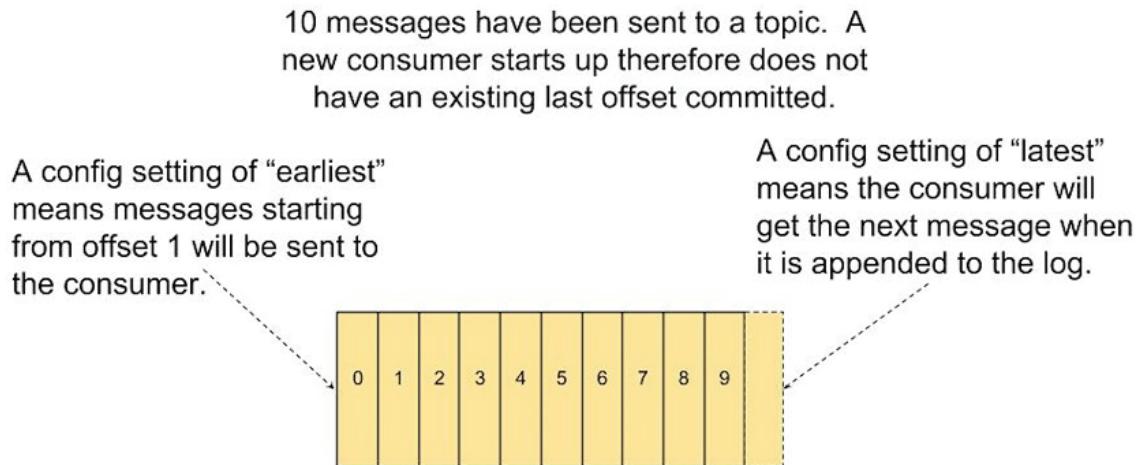
1. Committing implies the consumer has fully processed the message
2. It also represents the starting point for that consumer in the case of failure or a re-start.

If you have a new consumer instance or some failure has occurred, and the last committed offset is not available, where the consumer will start from depends on your configuration.

- `auto.offset.reset="earliest"` retrieves messages starting at the earliest available offset, implying any messages that have not yet been removed by the log management process.

- `auto.offset.reset="latest"` starts retrieving messages from the latest offset, essentially only consuming messages from the point of joining the cluster.
- `auto.offset.reset="none"` no reset strategy specified, the broker throws an exception back to the consumer.

In the image below you can see the impact of choosing an auto offset reset setting. By selecting "earliest," you receive messages starting at offset one. If you choose "latest", you'll get a message starting at offset 11.



**Figure 2.16** Here is the graphical representation of "earliest" vs. "latest". As you can see, a setting of "earliest" will give you everything you have missed (any message not yet deleted), and "latest" means you will wait for the next available message to arrive.

Next, we need to discuss the options we have for committing offsets, either automatically or manually.

### 2.6.2 Automatic Offset Commits

Automatic offset commits are enabled by default and is represented by the `enable.auto.commit` property. There is a companion configuration option `auto.commit.interval.ms` which specifies how often the consumer commits offsets (default value is 5 seconds). You should take care when adjusting this value, too small a value increases network traffic and to large a value could lead to large amounts of repeated data in the case of a failure before a commit occurs.

### 2.6.3 Manual Offset Commits

Within the category of manually committing offsets, there are two types: synchronous and asynchronous.

#### Listing 2.16 Synchronous Commits

```
consumer.commitSync()
consumer.commitSync(Map<TopicPartition, OffsetAndMetadata>)
```

The no-arg `commitSync()` method blocks until all offsets returned from the last retrieval (poll) succeeds. This call applies to all subscribed topics and partitions. The

other version that takes a `Map<TopicPartition, OffsetAndMetadata>` parameter commits only the offsets, partitions, and topics as specified in the map.

### ASYNCHRONOUS COMMITS

There are analogous `consumer.commitAsync()` methods that as the name implies are completely asynchronous and return immediately. There is an option with 2 of the `consumer.commitAsync` methods to provide an `OffsetCommitCallback` object, called when the commit had concluded either successfully or with an error. Providing a callback instance allows for asynchronous processing/error handling.

The advantage of using manual commits is that it gives you direct control over when a record is considered processed.

#### 2.6.4 Creating the Consumer

We create a `Consumer` in a similar way to creating the producer. We supply configurations in the form a Java `java.util.Properties` object and get back a `KafkaConsumer` instance. The consumer instance then subscribes to topic from a supplied list of topic names or a regular expression. Typically we run the `Consumer` in a loop where we poll for a period specified in milliseconds.

A `ConsumerRecords<K, V>` object is the result of the polling. `ConsumerRecords` implements the `Iterable` interface and the object returned from each call to `next()` returns a `ConsumerRecord` object, containing metadata about the message in addition to the actual key and value.

After we have exhausted all of the `ConsumerRecord` objects returned from the last call to `poll` we return to the top of the loop, polling again for the specified period. In practice, consumers are expected to run indefinitely in this manner unless an error occurs or the application needed to be shut down and restarted (this is where committed offsets come into play as on reboot the consumer will pick up where it left off).

#### 2.6.5 Consumers and Partitions

The previous section described how one `Consumer` works. In practice we'll probably want multiple `Consumer` instances, one for each partition of a topic. While it is possible to have one `Consumer` reading from multiple partitions, it's not an uncommon practice to have a thread pool with the number of threads equal to the number of partitions. Each thread in the pool runs a `Consumer` assigned to one partition. This `Consumer` per partition pattern maximizes throughput, but keep in mind if you spread your consumers across multiple applications or machines the *total* thread count across all instances should not exceed the total number of partitions in the topic. Any threads in excess over the total partition count will be idle.

If a `Consumer` fails, the lead broker assigns its partitions to another active consumer.

**NOTE** In this example, we show subscribing a `Consumer` one topic. But this is for demonstration purposes only; you can subscribe a `Consumer` to an arbitrary number of topics.

The lead partition leader assigns topic-partitions to all available consumers with the same `group.id`. The `group.id` is a configuration setting that identifies the consumer as belonging to a "consumer group," that way consumers don't need to reside on the same machine. It is probably preferred to have your consumers spread out across a few machines. In the case of a failure of one machine, the lead broker assigns topic-partitions to consumers on good machines.

### 2.6.6 Rebalancing

What we have described in the previous section is termed "rebalancing." Topic-partition assignments to a `Consumer` are not static, but very dynamic in nature. As you add consumers to the group, some of the current topic-partition assignments are taken from active consumers and given to the new consumers. This reassignment process continues until every partition has been assigned to a consumer reading data. After that equilibrium point, any additional consumers will remain idle. And as we discussed just a moment ago, consumers leaving the group for whatever reason have their topic-partition assignments reassigned to other consumers.

### 2.6.7 Finer Grained Consumer Assignment

In the previous section "Consumers and Partitions" we described the situation of using a thread pool and subscribing multiple consumers (belonging to the same consumer group) to the same topic(s). Although Kafka will balance the load of topic-partitions across all consumers, the assignment of the topic and which partition is not deterministic. We don't know what topic-partition coupling(s) each consumer will receive.

The `KafkaConsumer` has methods that allow for subscribing to a particular topic and partition only.

#### Listing 2.17 Subscribing to Individual Topic and Partition

```
TopicPartition fooTopicPartition_0 = new TopicPartition("foo", 0);
TopicPartition barTopicPartition_0 = new TopicPartition("bar", 0);

consumer.assign(Arrays.asList(fooTopicPartition_0, barTopicPartition_0));
```

There are tradeoffs to consider with using manual topic-partition assignment:

- Any failures will not result in topic-partitions to be re-assigned even for consumers with the same group id. Any needed changes in assignment require another call to `consumer.assign`.
- The group specified by the consumer is still used for committing. However, since each consumer will be acting independently, it is a good idea to give each consumer a unique group id when doing the manual topic-partition assignment.

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

## 2.6.8 Consumer Example

Here is the Consumer code for our ZMart prototype that consumes transactions and prints them to the console:

### Listing 2.18 ThreadedConsumerExample found in src/main/java/bbejeck.chapter\_2/consumer/ThreadedConsumerExample.java

```

public void startConsuming() {
    executorService = Executors.newFixedThreadPool(numberPartitions);
    Properties properties = getConsumerProps();

    for (int i = 0; i < numberPartitions; i++) {
        ①
        Runnable consumerThread = getConsumerThread(properties);
        executorService.submit(consumerThread);
    }
}

private Runnable getConsumerThread(Properties properties) {
    return () -> {
        Consumer<String, String> consumer = null;
        try {
            consumer = new KafkaConsumer<>(properties);
            consumer.subscribe(Collections.singletonList("test-topic"));
            while (!doneConsuming) {
                ②
                ConsumerRecords<String, String> records = consumer.poll(5000);
                for (ConsumerRecord<String, String> record : records) {
                    String message = String.format("Consumed: key = %s value = %s
                        with offset = %d partition = %d",
                        record.key(), record.value(), record.offset(), record.partition());
                    ③
                    System.out.println(message);
                }
                ④
            } catch (Exception e) {
                e.printStackTrace();
            } finally {
                ⑤
                if (consumer != null) {
                    consumer.close();
                }
            }
        };
    }
}

```

- ① Building consumer thread.
- ② Subscribing to the topic.
- ③ Polling for 5 seconds.
- ④ Printing formatted message.
- ⑤ Closing the consumer regardless, will leak resources otherwise.

This example has left out other sections of the class for clarity; this example is not meant to stand on its own. You'll find the full example in the accompanying source code with the book.

## 2.7 Installing and Running Kafka

You will be using the 1.0.0 version of Kafka (the most recent as of the writing of this book). As Kafka is a Scala project, each release comes in two versions; one for Scala 2.11 and another for Scala 2.12. We will be using the 2.12 Scala version of Kafka. Although you can download the release, I have included the binary distribution of Kafka that will work with Kafka Streams as demonstrated and described in this book.

To install Kafka extract the tgz file found in the source code repo in the "libs" folder somewhere on your machine.

**NOTE**
**Kafka Dependencies**

The binary distribution of Kafka includes Apache Zookeeper, so no extra installation work is required.

### 2.7.1 Kafka Local Configuration

To run locally on your machine, Kafka requires minimal configuration if you choose to accept the default values. By default Kafka will use port 9092 and zookeeper will use port 2181. Assuming you have no applications already using those ports, you are all set. Kafka will write its logs to /tmp/Kafka-logs and zookeeper uses /tmp/zookeeper for its log storage. Depending on your machine you may need to change permission/ownership of those directories or modify the location where you want to write the logs. To change the Kafka logs directory cd into <kafka-install-dir>/config and open the server.properties file. Find the "log.dirs" entry and change the value to what you would rather use. In the same directory open the zookeeper.properties file and change the "dataDir" entry. Later on configuring Kafka will be covered in detail, but this is all the configuration you will do for now. Keep in mind these "logs" are the actual data used by Kafka and zookeeper to their work, and not the application level logs that track the application's behavior. The application logs are found in the <kafka-install-dir>/logs directory.

### 2.7.2 Running Kafka

Kafka is very simple to get started. Since Zookeeper is essential for the Kafka cluster to function properly (zookeeper determines the lead broker, holds topic information, performs health checks on cluster members, etc.), you will need to start zookeeper before starting Kafka.

**NOTE**
**Kafka Directories**

From now on all directories referenced will assume you are working in directory <kafka-install-dir>/ If you are using a Windows machine the directory is <kafka-install-dir>/bin/windows

## STARTING ZOOKEEPER

To start zookeeper open a command prompt and enter the following command:

```
bin/zookeeper-server-start.sh config/zookeeper.properties
```

you will see a lot of information run on the screen, and it should end up looking something like this:

```

bin          bash          bash          bash          java          bash          bash          ...
odoo@odoo:~/kafka_2.11.0.1.0-SNAPSHOT$ ./zookeeper-server-start.sh ./config/zookeeper.properties
SLF4I: Class path contains multiple SLF4J bindings.
SLF4I: Found binding in [/var/file:/usr/local/kafka_2.11-0.10.1.0-SNAPSHOT/lib/slf4j-log4j12-1.7.18.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4I: Found binding in [/var/file:/usr/local/kafka_2.11-0.10.1.0-SNAPSHOT/lib/slf4j-log4j12-1.7.21.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4I: See http://www.slf4j.org/codes.html#multiple_bindings for explanation.
SLF4I: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
[2016-08-13 09:46:12,863] INFO Read configuration from: ./config/zookeeper.properties [org.apache.zookeeper.server.ZooKeeperQuorumPeerConfig]
[2016-08-13 09:46:12,863] INFO Using configuration from: ./config/zookeeper.properties [org.apache.zookeeper.server.ZooKeeperQuorumPeerConfig]
[2016-08-13 09:46:12,863] INFO Outgoing purge-interval set to 6 [org.apache.zookeeper.server.ZooKeeperQuorumPeerConfig]
[2016-08-13 09:46:12,863] INFO Purge task is not scheduled. [org.apache.zookeeper.server.ZooKeeperQuorumPeerConfig]
[2016-08-13 09:46:12,863] INFO Configuration or log dir not defined in config, running in stand-alone mode [org.apache.zookeeper.server.ZooKeeperQuorumPeerConfig]
[2016-08-13 09:46:12,863] INFO Reading configuration from file: ./config/zookeeper.properties [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,883] INFO Starting server [org.apache.zookeeper.server.ZooKeeperServer$Init]
[2016-08-13 09:46:12,893] INFO Server environment:java.version=1.8.0_66 [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:javacmd=/usr/bin/java [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:java.vendor=Oracle Corporation [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:java.vendor.version=Java(TM) SE Runtime Environment [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:java.home=/Library/Java/JavaVirtualMachines/jdk1.8.0_66.jdk/Contents/Home [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:java.vm.name=Java HotSpot(TM) 64-Bit Server VM [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:java.vm.version=25.6-b10 [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:java.vm.vendor=Oracle Corporation [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:java.vm.specification.name=Java Platform API Specification [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:java.vm.specification.vendor=Oracle Corporation [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:java.vm.specification.version=1.8 [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:os.name=Mac OS X [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:os.arch=x86_64 [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:os.version=10.11.5 [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:user.name=odoo [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:user.home=/Users/odoo [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:user.dir=/Users/odoo [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Binding to port 0.0.0.0.0.0.0.0.2181 [org.apache.zookeeper.server.KafkaServer$ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:java.io.tmpdir=/var/folders/lv/49_qrs38sdqghbwy4z3800np/T/ [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:javacompiler=<NA> [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:os.name=Mac OS X [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:os.arch=x86_64 [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:os.version=10.11.5 [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:user.name=odoo [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:user.home=/Users/odoo [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO Server environment:user.dir=/Users/odoo [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO tickTime set to 3000 [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO maxSessionTimeout set to -1 [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO minSessionTimeout set to -1 [org.apache.zookeeper.server.ZooKeeperServer]
[2016-08-13 09:46:12,893] INFO binding to port 0.0.0.0.0.0.0.2181 [org.apache.zookeeper.server.KafkaServer$ZooKeeperServer]
[2016-08-13 09:46:12,933] INFO binding to port 0.0.0.0.0.0.0.2181 [org.apache.zookeeper.server.KafkaServer$ZooKeeperServer]

```

**Figure 2.17 Output visible on the console when starting Zookeeper starts up.**

## STARTING KAFKA

To start Kafka open another command prompt and type this command:

```
bin/Kafka-server-start.sh config/server.properties
```

again you will see text scroll by on the screen, and when Kafka has fully started you should see something similar to this:

The log output from the Kafka startup command is a long series of command-line arguments and environmental variables. It starts with 'bash' and ends with 'java'. The log contains many references to Kafka's internal classes and configurations.

**Figure 2.18 Out put from Kafka when starting up.**

#### NOTE

#### Reverse the Order on Shutdown

We mentioned why Zookeeper is essential for Kafka to run. Therefore it is important to reverse the order when shutting down; that is, stop Kafka first then stop Zookeeper. To stop Kafka, you can enter control-c from the terminal Kafka is running in or run `kafka-server-stop.sh` from another terminal. The same goes for zookeeper except for the script to shutdown Zookeeper is `zookeeper-server-stop.sh`

### 2.7.3 Sending Your First Message

Now that you have Kafka up and running, it is time to use Kafka what is meant to do, sending and receiving messages. Before you send a message, you will need to define a topic for a producer to send a message.

#### YOUR FIRST TOPIC

Creating a topic in Kafka is very simple; it is just a matter of running a script with some configuration parameters. Although the configuration is an easy task, the settings you provide have broad performance implications. By default Kafka is configured to auto-create topics, meaning if you attempt to send or read a non-existent topic, the Kafka broker will create one for you (using default configurations in the `server.properties` file). However, this is rarely a good practice to rely on the broker to create topics, even in development. The issue is the first produce/consume attempt will fail, as it takes time for the metadata about the topic's existence to propagate. So be sure always proactively to create your topics.

#### CREATING A TOPIC

To create a topic, you will run the `Kafka-topics.sh` script. Open a terminal window and run this command:

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

```
bin/kafka-topics.sh --create --topic first-topic
                   --replication-factor 1
                   --partitions 1
                   --zookeeper localhost:2181
```

When the script executes, you should see something similar to this in your terminal:

```
bin — bash — 255x5
bash                                bash                                bash                                java
oddball:bin bbejeck$ ./kafka-topics.sh --create --topic first-topic --replication-factor 1 --partitions 1 --zookeeper localhost:2181
Created topic "first-topic".
oddball:bin bbejeck$
```

**Figure 2.19** Here are the results from creating a topic. It's important to create your topics ahead of time so you can supply your topic specific configurations. Otherwise, auto created topics use default configurations or the configuration in the `server.properties` file.

Most of the configuration flags in the above command are obvious, but let us quickly review two of them, replication-factor and partitions. The replication-factor determines how many copies of the message the lead broker distributes in the cluster. In this case, there will be no copies made; just the original message will reside in Kafka. A replication factor of 1 is fine for a quick demo or prototype, but in practice, you almost always want a replication factor of 2-3 to provide data availability in the face of machine failures. Partitions are the number of partitions that the topic will use. Again just one partition is fine here, but depending on the load, you will certainly want more partitions. Determining the correct number of partitions is not an exact science and strategies, and pros/cons for determining the number of partitions will be covered later in this chapter.

## SENDING A MESSAGE

Sending a message in Kafka requires writing a `Producer` client, but fortunately, Kafka comes with a handy script called `kafka-console-producer` that allows you to send a message from a terminal window. You use the console producer in this example, but we cover the `Producer` client in some detail in section 5 of this book. To send your first message run the following command:

```
bin/kafka-console-producer.sh --topic first-topic --broker-list localhost:9092
```

```
bin — bash — 255x5
bash                                bash                                bash
oddball:bin bbejeck$ ./kafka-console-producer.sh --topic first-topic --broker-list localhost:9092
the quick brown fox jumped over the lazy dog
```

**Figure 2.20** The Console Producer is a great tool to quickly test your configuration and to ensure the end to end functionality.

There are several options for configuring the console producer, but for now, you will only use the required ones, the topic to send the message to and a list of Kafka brokers to connect to (in this case just the one on your local machine). Starting a console producer is

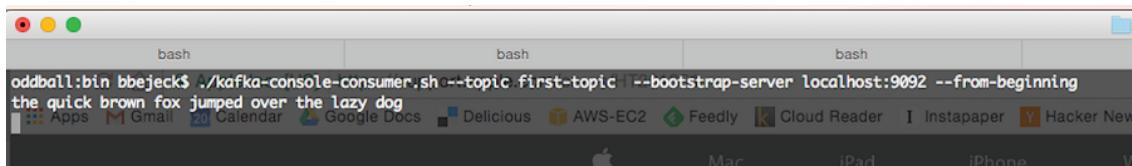
a "blocking script," so after executing the above command, you just need to enter some text and hit enter. You can send as many messages as you like, but for demo purposes here you will type a single message "the quick brown fox jumped over the lazy dog", hit the enter key and then enter control-c to exit the producer.

## READING A MESSAGE

Kafka also provides a console consumer for reading messages from the command line. Both the console producer and consumer are very handy tools for troubleshooting or doing quick prototyping. Running the console consumer works in similar fashion, once started it will keep reading messages from a particular topic until the script is stopped by you (again control-c). To launch the console-consumer run this command:

```
bin/kafka-console-consumer.sh --topic first-topic
    --bootstrap-server localhost:9092
    --from-beginning
```

After starting the console consumer, you should see this in your terminal:



**Figure 2.21 The Console Consumer is a super handy tool to quickly get a feel if data is flowing and messages contain the expected information.**

The "--from-beginning" parameter specifies you will receive any message not deleted from that topic. Since the console consumer will not have any committed offsets, without the "--from-beginning" setting, you will only get messages sent after the console consumer has started.

You have just completed a whirlwind install of Kafka and producing/consuming your first message. Now it's time go back where you left off in this chapter to learn the details of how Kafka works!

## 2.8 Conclusion

We have covered a lot of ground in this chapter. As I stated before, since this book is about Kafka Streams this chapter is not intended to be a thorough treatment of Kafka itself, but a solid introduction to give those readers new to Kafka an understanding of how Kafka works. However, there are key points that we learned in this chapter that you need to take with you on your journey through the rest of this book:

- Kafka is message broker that receives messages and stores them in a way that makes it easy and fast to respond to consumer requests. Messages are never pushed out to consumers, and message retention in Kafka is entirely independent of when and how often messages are consumed.
- Kafka uses partitions for achieving high throughput and provide a means for grouping messages with the same key in order.

- Producers are used for sending messages to Kafka.
- Null keys mean round robin partition assignment otherwise the producer uses the hash of the key modulo the number of partitions for partition assignment.
- Consumers are what you use to read messages from Kafka.
- Consumers that are part of the same consumer group are given topic-partition allocations in an attempt to distribute messages evenly.

In the next chapter, we start looking at Kafka Streams with a concrete example from the world of retail sales. While Kafka Streams will handle the creation of all Consumer and Producer instances, you should be able to see the concepts we introduced here come into play.

# P2

## *Part 2: Kafka Streams Development*

In this part, we'll build on the previous section and take the metal model you have and put into action by developing a basic first Kafka Streams application. With your feet wet from that first basic application, you'll walk through the significant API points of Kafka Streams so you can learn how to build an application that will suit any of your needs. You'll first learn about the need for providing state to a streaming application, and how you use the state for performing joins, much like the joins you would perform when running SQL queries. Then you'll move on to a new abstraction from Kafka Streams, the KTable API. While you've spent the first parts of this section on the "high-level" DSL, you'll wrap up this section discussing the lower-level Processor API, where you'll learn how to make Kafka Streams do pretty much anything you need it to do.

# *Developing Kafka Streams*

## In this chapter

- Hello World Kafka Streams
- Introducing the KStreams API
- The ZMart Kafka Stream application in depth
- Splitting an incoming stream into multiple streams
- Advanced KStreams API methods

In the first chapter, learned about the Kafka Streams library. We talked about building a topology of processing nodes, or a graph used to transform data as it is streaming into Kafka. In this chapter, we will show you how to create this processing topology by using the KStreams API. The KStreams API is important to learn because it is the core tool with which you build Kafka Streams applications. By learning the KStream API you will gain the knowledge necessary to assemble Kafka Streams programs, but more importantly, you will gain a deeper understanding of how the components work together and how they can be used to achieve your stream processing goals.

### **3.1 KStreams API**

The KStreams API is a high-level DSL that enables you to build Kafka Streams applications quickly. The API is very well thought out, and there are methods to handle most stream processing needs out of the box, so creating a sophisticated processing program can be achieved without much effort. At the heart of the API is the `kstream` object which is used to represent the streaming key-value pair records.

Most of the methods in the KStreams API return a reference to a KStream object, allowing for a "fluent interface" style of programming. Additionally, a good percentage of the KStream methods accept types comprised of single method interfaces allowing for using Java 8 lambda expressions. Taking the previous two factors into account and you can see the simplicity and ease with which you can build a Kafka Streams program.

**NOTE****What's a Fluent Interface?**

Back in 2005 Martin Fowler and Eric Evans developed the concept of the fluent interface ([martinfowler.com/bliki/FluentInterface.html](http://martinfowler.com/bliki/FluentInterface.html)). A fluent interface is one where the return value of a method call is the same instance which originally called the method. This approach is very useful when constructing objects with several parameters, for example, `Person.builder().firstName("Beth").withLastName("Smith").withOccupation("CEO")`. In Kafka Streams there is one small but important difference; the returned KStream object is a new instance, not the same instance that made the original method call.

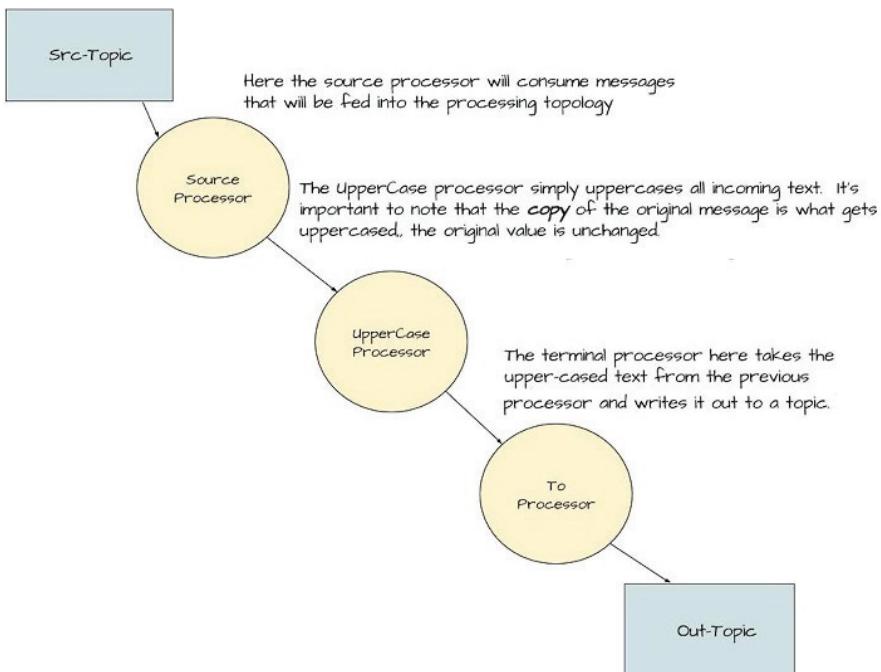
There is also a lower-level API, the Processor API that is not as succinct as the KStreams API but allows for more control. We will cover the Processor API in a later chapter. With the introduction out of the way, let's dive into the requisite hello world program for Kafka Streams.

### **3.2 Kafka Streams Hello World**

For our first example, we will deviate from the problem outlined in chapter 1 to a more simple use case. I did this so you can quickly get off the ground and see how Kafka Streams works. We will get back to the problem statement from chapter 1 later for more realistic, concrete examples.

Your first program will be a toy application that takes incoming messages and turns them into upper-case characters, effectively yelling at anyone who reads the message. The title of this application is called the "Yelling Application."

Before diving into the code let's take a look at the processing topology you will assemble for the first application. We will follow the same pattern from chapter 1, where we built up a processing graph topology, where each node in the graph has a particular function. The main difference will be in the simplicity of the graph. The following image shows the topology we will build.



**Figure 3.1 Graph or Topology the our first Kafka Streams application, the 'Yelling' App**

As you can see from the image, we are building a very simple processing graph. So simple in fact it resembles a linked list of nodes more than the typical tree-like structure of a graph. But there is enough here to give us strong clues as to what expect in the code. There will be a source node, a processor node transforming incoming text to all uppercase, and finally, a processor writing results out to a topic.

While this is a trivial example, the code shown here is representative of what you will see in our Kafka Streams programs. For most of our examples you will see a familiar structure:

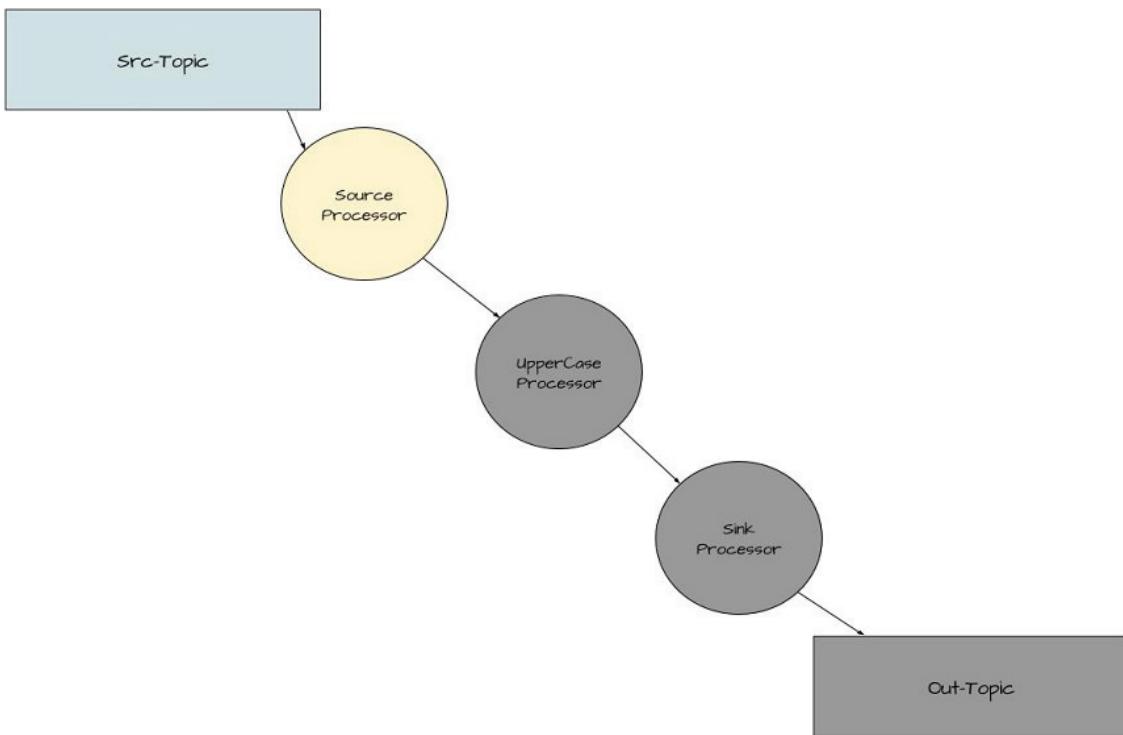
1. Define the configuration items
2. Create any Serdes instances either custom or pre-defined.
3. Build the processor topology.
4. Create and start the KStreams.

When we get into more advanced examples, the principal difference will be in the complexity of processor topology. With that in mind let's take a walk through building our first application.

### 3.2.1 Creating the Topology for the Yelling App

The first step to creating any Kafka Streams application is to create a source node. The source node is responsible for consuming records from a topic which will flow through the application.

We'll present the graph with the relevant point highlighted, so it will be easier to follow how the code relates to the processing topology. Here in this image, we've added the source node indicated by the highlighted portion of the graph.



**Figure 3.2 Creating the Source Node of the 'Yelling' App**

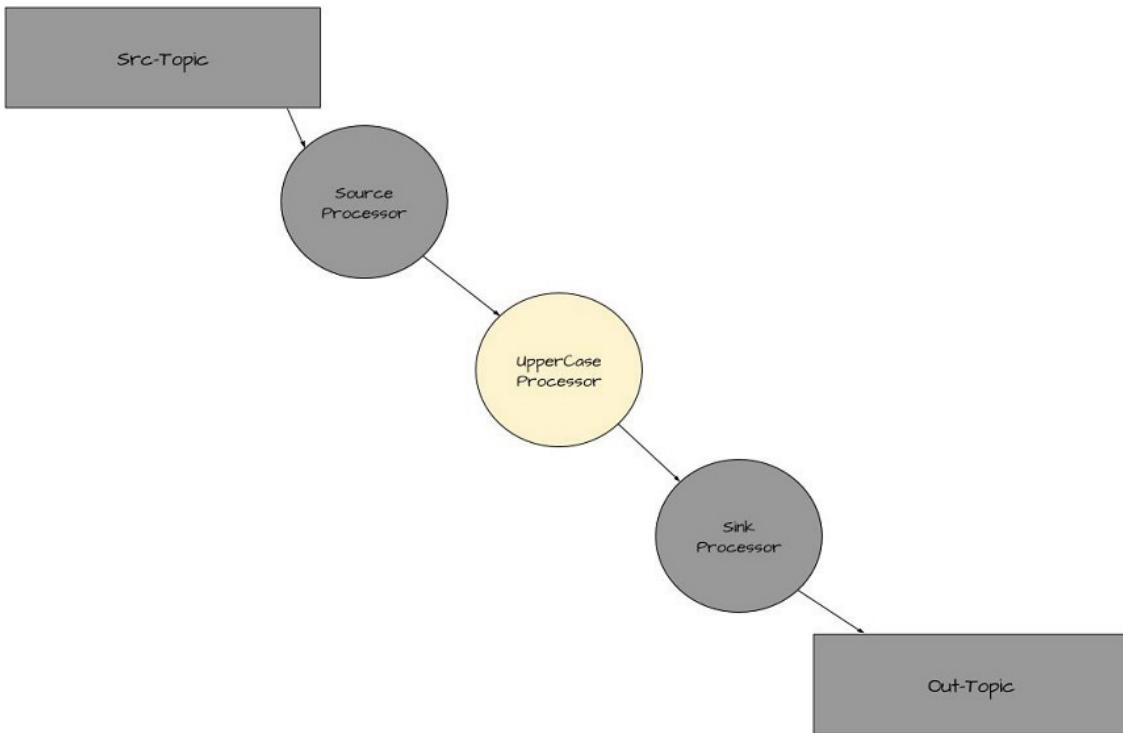
The code we use behind the graph image follows:

### Listing 3.1 Defining the Source for the Stream

```
KStream<String, String> simpleFirstStream = builder.stream("src-topic",
    Consumed.with(stringSerde, stringSerde));
```

With this line of code, you are creating the source or parent node of the graph. The source node of a topology consumes records from a given topic or topics. In this case, it's aptly named "src-topic." The KStream instance `simpleFirstStream` is set to consume messages written to the "src-topic" topic. In addition to the topic name parameter, you also provide serde objects (via a `Consumed` instance) for deserializing the records from Kafka. You'll use the `Consumed` class for any optional parameters whenever you create a source node in Kafka Streams.

We have a source node for our application, but now we need to attach a processing node to make use of the data. Our next step is to add a processing node as shown in the following image:



**Figure 3.3 Adding the UpperCase Processor to the 'Yelling' App**

The code used to attach the processor (a child node of the source node) looks like this:

### Listing 3.2 Mapping Incoming Text to Uppercase

```
KStream<String, String> upperCasedStream = simpleFirstStream.mapValues(String::toUpperCase);
```

With this line, you are creating another KStream instance that is a child node of the parent node. By calling the `KStream.mapValues` function you are creating a new processing node whose inputs are the results of going through the `mapValues` call. It's important to remember you should not modify the *original* value in the `ValueMapper` provided to `mapValues`. The `upperCasedStream` instance receives transformed copies of the initial value from the `simpleFirstStream.mapValues` call. In this case, it's upper-case text.

The `mapValues` method takes an instance of the `ValueMapper<V, V1>` interface. The `ValueMapper` interface defines only one method `ValueMapper.apply`, making it an ideal candidate for using a Java 8 lambda expression. And this is what we've done here with `String::toUpperCase`, which is a method reference, an even shorter form of a Java 8 lambda expression.

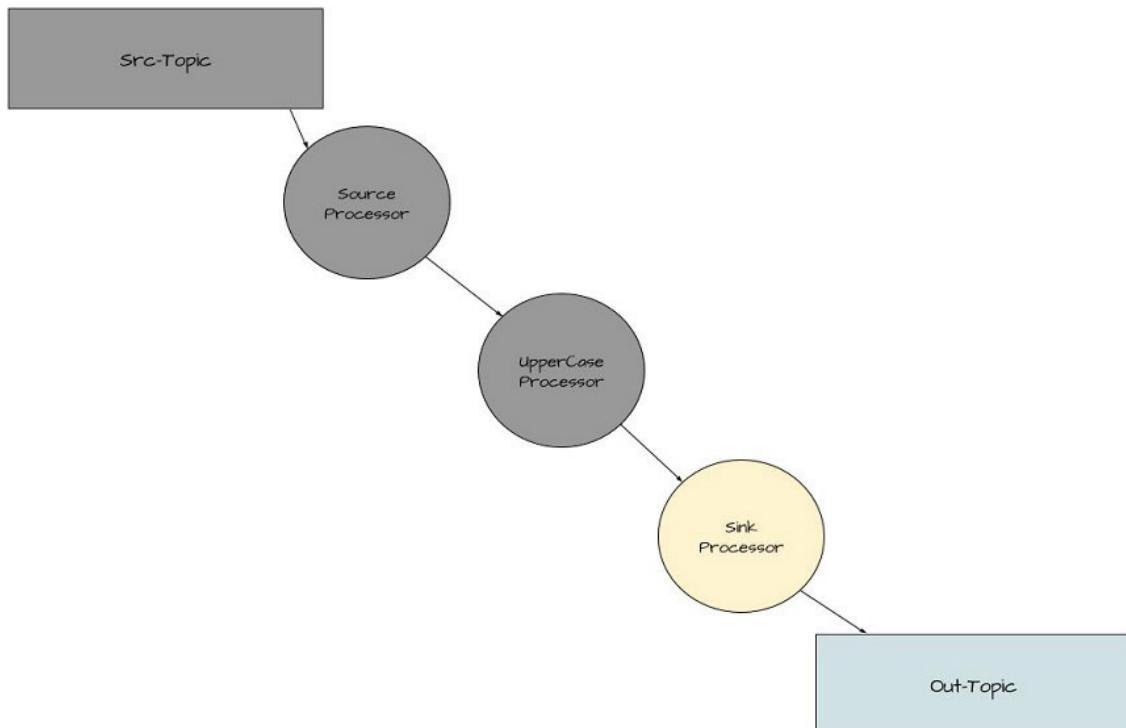
#### NOTE

While there many Java 8 tutorials available for lambda expressions and method references good starting points are [lambdaexpressions](#) and [methodreferences](#) respectively.

You could have easily used the form of `s s.toUpperCase()`, but since `toUpperCase` is an instance method on the `String` class, we can use a method reference.

Using lambda expressions vs. concrete implementations is a pattern we will see over and over with the KStreams API. Since most of the methods expect types that are single method interfaces, you can easily use Java 8 lambdas.

So far our Kafka Streams application is consuming records and then transforming them to upper-case. The final step for us to take is adding a processor/node that writes the results out to a topic. The following image shows us where we are in the construction of the topology:



**Figure 3.4 Adding Processor for Writing 'Yelling' App Results**

### **Listing 3.3 Creating a Sink Node**

```
upperCasedStream.to("out-topic", Produced.with(stringSerde, stringSerde));
```

Now we are adding the last processor in the graph. By calling the `KStream.to` method you are creating a terminal node in the topology. To be more precise, you are creating a sink node where the `upperCasedStream` processor writes the key-value pairs to a topic named "out-topic." Any KStream method that has a return type of void is a terminal node and will end stream processing at that point, but it may not necessarily be a sink node.

Again we see the pattern of providing serde instances needed for serialization/deserialization when reading from or writing to a Kafka topic. But in this case you are using a `Produced` instance which provides any optional parameters when

creating a sink node in Kafka Streams.

**NOTE** One note about serde usage in Kafka Streams. You don't always have to provide serdes to either the `Consumed` or `Produced` objects. In those cases, the application will use the serializer/deserializer listed in the configuration instead. Additionally, you can specify serde for either the key or value only with the `Consumed` and `Produced` classes.

In our example above we used three lines to build the topology.

#### **Listing 3.4 Building Processor Topology in Stages**

```
KStream<String, String> simpleFirstStream = builder.stream("src-topic",
    Consumed.with(stringSerde, stringSerde));
KStream<String, String> upperCasedStream = simpleFirstStream.mapValues(String::toUpperCase);
upperCasedStream.to("out-topic", Produced.with(stringSerde, stringSerde));
```

Having each step on an individual line was used to demonstrate each stage of the building process. However, all methods in the KStream API that don't create terminal nodes return a new KStream instance. Returning a `KStream` object allows for using the "fluent interface" style of programming referenced earlier. To demonstrate this idea, here's another way you could have constructed the yelling application topology:

#### **Listing 3.5 Building Processor Topology Using Fluent Style**

```
builder.stream("src-topic", Consumed.with(stringSerde,
    stringSerde)).mapValues(String::toUpperCase).to("out-topic",
Produced.with(stringSerde, stringSerde));
```

We just shortened our program from three lines to one without losing any clarity or purpose. From this point forward all of our examples will be written using the "fluent" interface style unless doing so will make the clarity of the program suffer.

We've built our first Kafka Streams topology. But we've glossed over the important steps of configuration and Serde creation which we'll address now.

### **3.2.2 Configuration**

While Kafka Streams is highly configurable with several properties you can adjust for your specific needs; our first example uses two configuration settings. Both of these are required as there are no default values provided. Attempting to start a Kafka Streams program without these two properties defined results in throwing a `ConfigException`.

#### **Listing 3.6 Setting Configuration Items**

```
props.put(StreamsConfig.APPLICATION_ID_CONFIG, "yelling_app_id");
props.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
```

- `StreamsConfig.APPLICATION_ID_CONFIG` : This property identifies your Kafka
- ©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

Streams application and must be a unique value for the entire cluster. The application id config also serves as a default value for the client-id prefix and group-id parameters if you don't set either value. The client-id prefix is the user-defined value used to identify clients connecting to Kafka uniquely. The group-id is used to manage the membership of a group of consumers all reading from the same topic and assures that many consumers can effectively read subscribed topics in the group.

- StreamsConfig.BOOTSTRAP\_SERVERS\_CONFIG : The bootstrap servers configuration can be a single hostname: port setting or multiple values comma separated. The value of this setting should be very apparent as it points the Kafka Streams application to the Kafka cluster.

We will cover several more configuration items as we explore more examples throughout the book.

### 3.2.3 Serde Creation

#### **Listing 3.7 Instantiating a String Serde**

```
Serde<String> stringSerde = Serdes.String();
```

This line from point 3 above is where you create the `Serde` instance required for serialization/deserialization of the key/value pair of every message and store it in a variable for use later on. The `Serdes` class provides default implementations for the following types:

- Strings
- Byte arrays
- Long
- Integer
- Double

Implementations of the `Serde` interface are extremely useful as they contain the serializer and deserializer, which eliminates the repetition for having to specify four parameters (key serializer, value serializer, key deserializer and value deserializer) every time you need to provide a `Serde` in a `KStream` method. In our next example, we will show how to create your `Serde` implementation to handle serialization/deserialization of more complex types.

Now let's take a look at the whole program you just put together.

#### **Listing 3.8 Hello World - The Yelling Application found in src/main/java/bbejeck/chapter\_3/KafkaStreamsYellingApp.java**

```
public class KafkaStreamsYellingApp {
    public static void main(String[] args) {
        Properties props = new Properties();
        props.put(StreamsConfig.APPLICATION_ID_CONFIG, "yelling_app_id");
```

1

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

```

props.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");

StreamsConfig streamingConfig = new StreamsConfig(props); ❷

Serde<String> stringSerde = Serdes.String(); ❸

StreamsBuilder builder = new StreamsBuilder(); ❹

KStream<String, String> simpleFirstStream = builder.stream("src-topic", ❺
    Consumed.with(stringSerde, stringSerde));

KStream<String, String> upperCasedStream = ❻
    simpleFirstStream.mapValues(String::toUpperCase);

upperCasedStream.to("out-topic", Produced.with(stringSerde, stringSerde)); ❼

KafkaStreams kafkaStreams = new KafkaStreams(builder.build(), streamsConfig);

kafkaStreams.start(); ❽
Thread.sleep(35000);
LOG.info("Shutting down the Yelling APP now");
kafkaStreams.close();

}
}

```

- ❶ Properties for configuring the Kafka Streams program,
- ❷ Creating the StreamsConfig with the given properties.
- ❸ Creating the Serdes used to serialize/deserialize keys and values.
- ❹ Creating the StreamsBuilder instance used to construct our processor topology.
- ❺ Creating the actual stream with a source topic to read from (Parent node in the graph).
- ❻ A processor using a Java 8 method handle (first child node in the graph).
- ❼ Writing the transformed output to another topic (sink node in the graph).
- ❽ Kicking off the Kafka Stream thread(s).

In this section, you have constructed your first Kafka Streams application. Let's quickly review the steps taken as it's a general pattern we'll see in most of our Kafka Streams applications. Specifically, you took the following steps:

1. Created a StreamsConfig instance.
2. Then you created a Serde object.
3. You constructed a processing topology.
4. After putting all the pieces into place, you started the Kafka Streams program.

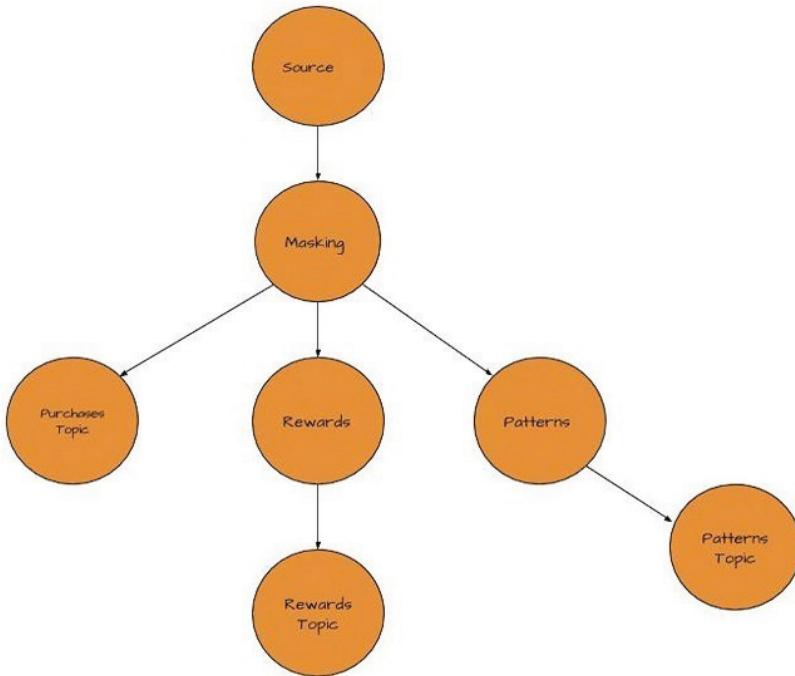
Aside from the general construction of a Kafka Streams application, a key takeaway here is to leverage lambda expressions whenever possible to make your programs more concise.

Next, we'll move on to a more complex example that will allow us to explore more of

the Kafka Streams API. The example will be new, but the scenario is one you are already familiar with, the solution to the ZMart data processing goals.

### 3.3 Working with Customer Data

In chapter one, we discussed the new customer data processing requirements you have to help ZMart do business more efficiently. We demonstrated how to build a topology of processors that will work on purchase records as they come streaming in from transactions in ZMart stores. Here we present the completed graph again for review:



**Figure 3.5 Initial Complete Topology for ZMart Kafka Streams program**

While we're at it, let's briefly go over the requirements for our streaming program which also serves as a good description of what our program will do:

- All records need to have credit-card numbers protected, in this case masking the first 12 digits
- Extract the item(s) purchased and the zip code for determining purchase patterns in the company. This data will be written out to a topic.
- Capture the customer ZMart member number and amount spent and written to a topic. Consumers of the topic will use the data to determine rewards.
- Write the entire transaction out to topic to be consumed by a storage engine for ad-hoc analysis.

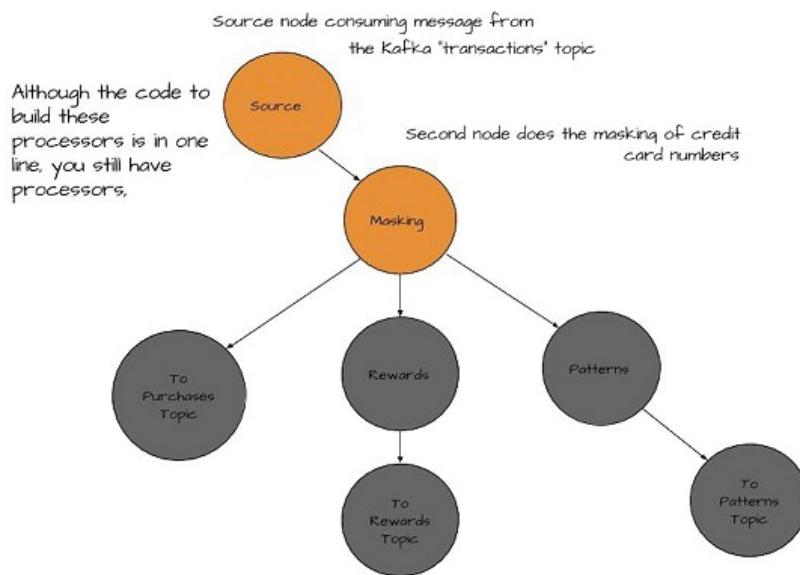
As in our first example, we will see the fluent interface approach combined with Java 8 lambdas when building the application. As this application is more involved than our first example, the use of these two techniques is pervasive. While sometimes it's clear the return type of a method call is a `KStream` object other times it may not be so clear. A key point to remember is that the majority of the methods in the KStream API return *new* KStream instances.

Now let's get to work building the streaming application that will satisfy your business requirements.

### 3.3.1 Constructing the Processors

Now let's dive into building the processing topology. To help us make the connection between the code and the processing topology graph we've created in the previous chapter, we will highlight the part of the graph that we are currently working on.

#### BUILDING THE SOURCE NODE



**Figure 3.6** The source processor consumes from a Kafka topic, but it feeds the masking processor exclusively, actually making it the source for the rest of the topology.

#### Listing 3.9 Building the Source Node and First Processor

```
KStream<String, Purchase> purchaseKStream = streamsBuilder.stream("transactions",
    Consumed.with(stringSerde, purchaseSerde))
    .mapValues(p -> Purchase.builder(p).maskCreditCard().build());
```

Here we are building the source node and first processor of the topology by chaining two calls of the KStream API together. While it should be fairly obvious what the role of the origin node is, the first processor in the topology is responsible for masking credit-card numbers to protect customer privacy.

You create the source node with a call to the `StreamsBuilder.stream` method using a default String Serde, a custom Serde for `Purchase` objects and the name of the topic that is the source of messages for the stream. In this case, we are only specifying one topic, but we could have provided a comma separated list of names or a regular expression to match topic names instead.

Here we chose to provide the serdes with a `Consumed` instance, but we could have left out the `Consumed` instance and only provide the topic name and relied on the default serdes provided via configuration parameters.

The next immediate call is to the `kstream.mapValues` method, taking a `ValueMapper<V, V1>` instance as a parameter. Value mappers take a single parameter of one type (a `Purchase` object in this case) and 'map' that object a to a new value, possibly of another type. In the example here the `KStreams.mapValues` returns an object of the same type (`Purchase`), but with a masked credit-card number.

You should take note that when using the `kstream.mapValues` method the original key is unchanged and is not factored into mapping a new value. If we had wanted to generate a new key-value pair or include the key in producing a new value, we would use the `kstream.map` method that takes a `KeyValueMapper<K, V, KeyValue<K1, V1>>` instance.

### HINTS OF FUNCTIONAL PROGRAMMING

An important concept to keep in mind with the 'map' or 'mapValues' functions is that they are expected to operate without side-effects. Meaning the functions do not modify the object or value presented as a parameter.

It's at this point we can get a sense of functional programming aspects in the KStreams API. While functional programming is a deep topic and a full discussion is beyond the scope of this book, we can briefly discuss two central principles of functional programming here.

The first would be to avoid state modification. If an object requires a change or update, we pass the object to a function, and a copy or entirely new instance is made containing the desired changes or updates. In the example here the lambda passed to `KStream.mapValues` is used to update the `Purchase` object having a masked credit-card number, the credit card field on the original `Purchase` object is left unchanged.

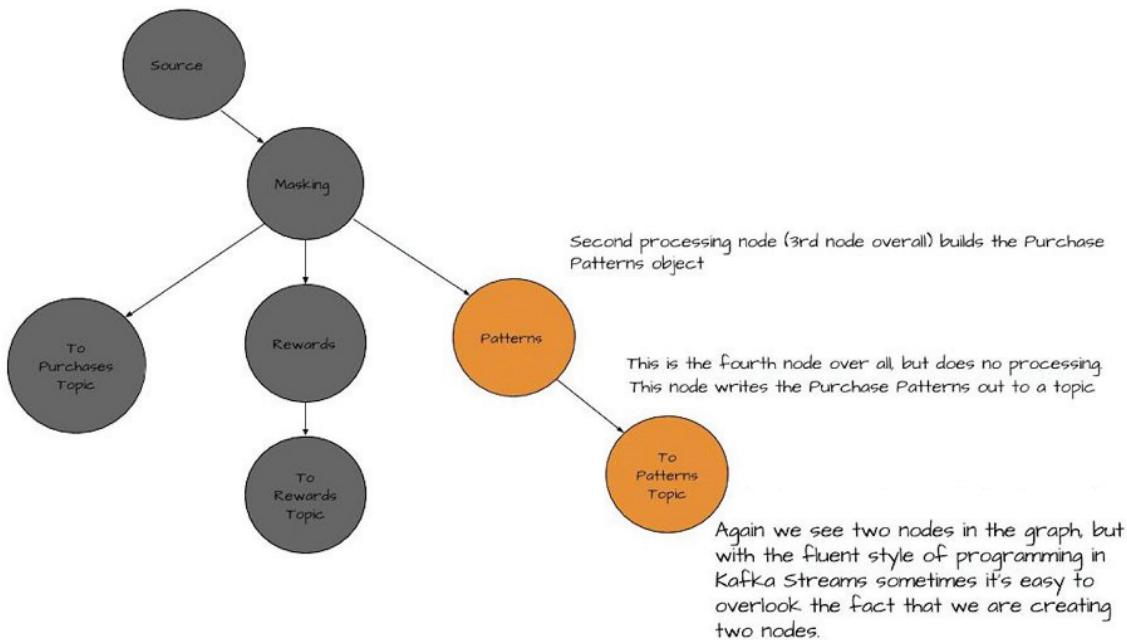
The second principle is building complex operations by composing several smaller single purpose functions together. The composition of functions is a pattern we will frequently see working with the KStream API.

#### SIDE BAR    What is Functional Programming

While there may be other definitions of functional programming, we will define it here in the following manner. Functional programming is a programming approach where functions are first class objects themselves. Furthermore, functions are expected to avoid providing side-effects, that is the modification of state or mutable objects.

### BUILDING THE SECOND PROCESSOR

Now we are going to go over building the second processor, listing number 3 from our example. The second processor is responsible for extracting pattern data sent to a topic used by ZMart to determine purchase patterns in regions of the country.



**Figure 3.7** Here we are adding the second processor, which build purchase pattern information. Additionally, we are adding a terminal node which writes the PurchasePattern object out to Kafka.

#### **Listing 3.10 Code for the Second Processor and a Terminal node to write to Kafka**

```
KStream<String, PurchasePattern> patternKStream =
purchaseKStream.mapValues(purchase -> PurchasePattern.builder(purchase).build());

patternKStream.to("patterns", Produced.with(stringSerde,purchasePatternSerde));
```

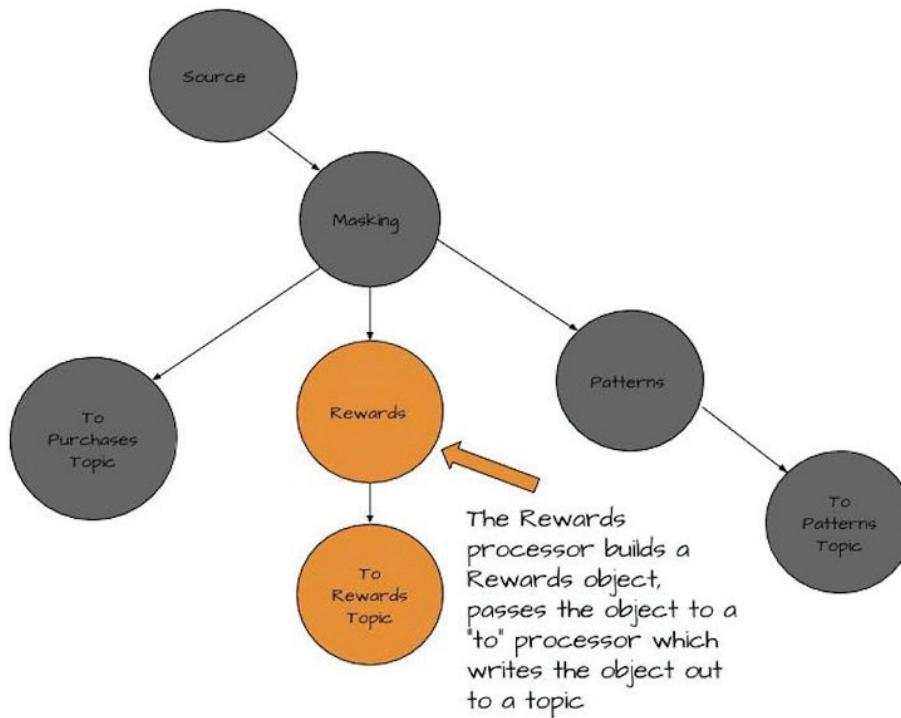
Here we see the `purchaseKStream` processor using the familiar `mapValues` call to create a new `KStream` instance. Here this new `KStream` will start to receive `PurchasePattern` objects created as a result of the `mapValues` call.

One thing to notice here is that we declare a variable to hold the reference of the new `KStream` instance as you'll use it to print the results of the stream to the console with a `print` call, which is very useful during development and for debugging. Here the purchase patterns processor forwards the records it receives to a child node of its own, defined by the method call `KStream.to` and using the previously built serdes and the name of the topic.

The `KStream.to` method is a mirror image of the `KStream.source` method, instead of setting a source for the topology to read from, the `KStream.to` method defines a sink node which is used to realize the data contained in a particular `KStream` instance to a Kafka topic. The `KStream.to` method also provides overloaded versions in which you can leave out the `Produced` parameter and use the default serdes defined in the configuration. One of the optional parameters you can set with the `Produced` class is a `StreamPartitioner` that we will discuss next.

## BUILDING THE THIRD PROCESSOR

The third processor in the topology is the customer rewards accumulator node used that will be used to track purchases made by members of ZMart's preferred customer club. More specifically, the reward accumulator sends data to a topic consumed by applications at ZMart HQ to determine rewards as customers complete purchases.



**Figure 3.8** At this point we are building the third processor which creates the RewardAccumulator object from the purchase data. We also are adding a terminal node to write the results out to Kafka. Each terminal node in the graph writes to separate topic.

### Listing 3.11 Code for the Third Processor and Terminal node which writes to Kafka

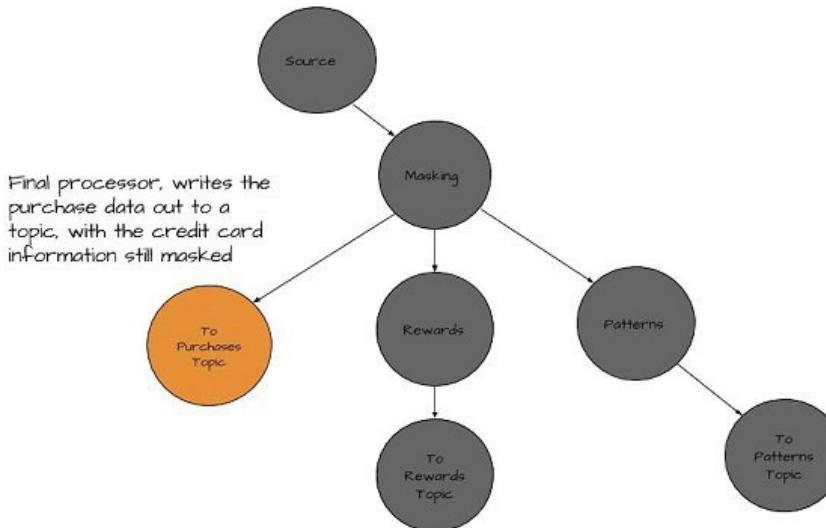
```
KStream<String, RewardAccumulator> rewardsKStream =
purchaseKStream.mapValues(purchase -> RewardAccumulator.builder(purchase).build());
rewardsKStream.to("rewards", Produced.with(stringSerde,rewardAccumulatorSerde));
```

At this point in the topology, you are building the rewards accumulator processor using what should be by now a familiar pattern: creating a new `KStream` instance that maps the raw purchase data contained in the value to a new object type. We also attach a sink node to the reward accumulator so the results of the rewards `KStream` can be realized to a topic and used for determining customer reward levels.

## LAST PROCESSOR

Finally, we take the first `KStream` we created, `purchaseKStream`, and attach a sink node to write out the raw purchase records (credit-cards masked of course) out to a topic called "purchases". The purchases topic will be used to feed into a NoSql store such as Cassandra or Presto<sup>8</sup> to perform ad-hoc analysis.

Footnote 8 [prestodb.io/](https://prestodb.io/)



**Figure 3.9** This is the last node we're adding to our topology (at least at this point). This node writes out the entire purchase transaction out to a topic whose consumer is a NoSQL data store such as MongoDB. The following graphic depicts where we are at this point in our application.

### Listing 3.12 Final Processor

```
purchaseKStream.to("purchases", Produced.with(stringSerde, purchaseSerde));
```

Now that we've built our application piece by piece let's look at the entire application. You'll quickly notice it's more complicated than our "hello world" example.

### Listing 3.13 ZMart Customer Purchase KStreams Program src/main/java/bbejeck/chapter\_3/ZMartKafkaStreamsApp.java

```

public class ZMartKafkaStreamsApp {

    public static void main(String[] args) {
        // some details left out for clarity

        StreamsConfig streamsConfig = new StreamsConfig(getProperties());

        JsonSerializer<Purchase> purchaseJsonSerializer = new JsonSerializer<>();
        JsonDeserializer<Purchase> purchaseJsonDeserializer = new JsonDeserializer<>(Purchase.class);
        Serde<Purchase> purchaseSerde = Serdes.serdeFrom(purchaseJsonSerializer,
            purchaseJsonDeserializer);
        //Other Serdes left out for clarity

        Serde<String> stringSerde = Serdes.String();

        StreamsBuilder streamsBuilder = new StreamsBuilder();

        KStream<String, Purchase> purchaseKStream = streamsBuilder.stream("transactions",
            Consumed.with(stringSerde, purchaseSerde))
            .mapValues(p -> Purchase.builder(p).maskCreditCard().build()); ②

        KStream<String, PurchasePattern> patternKStream = purchaseKStream.mapValues(purchase ->
            PurchasePattern.builder(purchase).build()); ③
    }
}
  
```

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

```

    patternKStream.to("patterns", Produced.with(stringSerde,purchasePatternSerde));

    KStream<String, RewardAccumulator> rewardsKStream = purchaseKStream.mapValues(purchase ->
        RewardAccumulator.builder(purchase).build());
    rewardsKStream.to("rewards", Produced.with(stringSerde,rewardAccumulatorSerde)); ④

    purchaseKStream.to("purchases", Produced.with(stringSerde,purchaseSerde)); ⑤

    KafkaStreams kafkaStreams = new KafkaStreams(streamsBuilder.build(),streamsConfig);
    kafkaStreams.start();

}

```

- ① Creating Serde(s), the data format is in JSON.
- ② Building the source and first processor
- ③ Building the PurchasePatterns processor
- ④ Building the RewardAccumulator processor
- ⑤ Building the storage sink, the topic used by the storage consumer.

**NOTE** We have left out some details here for clarity. Throughout the book code examples aren't necessarily meant to stand on their own. For the full example, the reader is left to look at the accompanying source code.

As you can see this is a little more involved than our original introductory application, but it has a similar flow. Specifically, you still performed the following steps:

- A StreamsConfig instance.
- You build one or more Serde instances.
- The processing topology is constructed.
- Assemble all the component then start the KafkaStreams program.

While building the application we've mentioned using a Serde while not taking any time to explain why or how we create them. Let's take some time now to discuss the role of the Serde in a Kafka Streams application.

### 3.3.2 Creating a Custom Serde

Kafka transfers data in byte array format. Since the data format is JSON, we'll need to tell Kafka how to convert an object first into JSON then into a byte array when sending data to a topic. Conversely, we'll need to specify how to convert consumed byte arrays into JSON then into the object type our processors will use. This conversion of data to/from different formats is why we need to have a Serde. While there are some Serdes provided out of the box (String, Long, Integer, etc.) we'll need to create custom Serdes for our objects.

In our first example, we only needed a serializer/deserializer for Strings, and there is an implementation provided by the `Serdes.String()` factory method. In this example,

however, we need to create custom Serde instances, as the type of the objects is arbitrary. We will walk through building the Serde for the `Purchase` class. We won't cover making the other Serde instances as they follow the same pattern, just with different types.

Building a Serde requires implementations of the `Deserializer<T>` and `Serializer<T>` interfaces. Here are the implementations of the `Serializer` and `Deserializer` that we will use throughout the examples. Also, we have chosen to use the Gson library from Google to convert objects to/from JSON.

#### **Listing 3.14 The Generic Serializer found in src/main/java/bbejeck/util/serializer/JsonSerializer.java**

```
public class JsonSerializer<T> implements Serializer<T> {

    private Gson gson = new Gson(); 1

    @Override
    public void configure(Map<String, ?> map, boolean b) {
    }

    @Override
    public byte[] serialize(String topic, T t) { 2
        return gson.toJson(t).getBytes(Charset.forName("UTF-8"));
    }

    @Override
    public void close() {
    }
}
```

- ➊ Creating the Gson object
- ➋ Serializing an object to bytes

For serialization, we will first convert an object to JSON, then get the bytes from the string. To handle the conversions from/to JSON we are using Gson<sup>9</sup>.

---

Footnote 9 [github.com/google/gson](https://github.com/google/gson)

For our deserializing process we take different steps; create a new String from a byte array then use Gson to convert the JSON string into a Java object.

#### **Listing 3.15 The Generic Deserializer found in src/main/java/bbejeck/util/serializer/JsonDeserializer.java**

```
public class JsonDeserializer<T> implements Deserializer<T> {

    private Gson gson = new Gson(); 1
    private Class<T> deserializedClass; 2

    public JsonDeserializer(Class<T> deserializedClass) {
        this.deserializedClass = deserializedClass;
    }
}
```

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

```

public JsonDeserializer() {
}

@Override
@SuppressWarnings("unchecked")
public void configure(Map<String, ?> map, boolean b) {
    if(deserializedClass == null) {
        deserializedClass = (Class<T>) map.get("serializedClass");
    }
}

@Override
public T deserialize(String s, byte[] bytes) {
    if(bytes == null){
        return null;
    }

    return gson.fromJson(new String(bytes),deserializedClass); ③
}

@Override
public void close() {
}
}

```

- ① Creating the Gson object.
- ② Instance variable of Class to deserialize.
- ③ Deserializing bytes to an instance of expected Class.

Now going back to the lines at listing 1 :

### **Listing 3.16 Creating a Serde**

```

JsonDeserializer<Purchase> purchaseJsonDeserializer =
    new JsonDeserializer<>(Purchase.class); ①
JsonSerializer<Purchase> purchaseJsonSerializer = new JsonSerializer<>();
Serde<Purchase> purchaseSerde =
    Serdes.serdeFrom(purchaseJsonSerializer,purchaseJsonDeserializer); ②③

```

- ① Creating the Deserializer for the Purchase class
- ② Creating the Serializer for the Purchase class
- ③ Creating the Serde for Purchase objects

As you can see a `serde` object is very useful as it serves as a container for a serializer and deserializer of a given object.

We've covered a lot of ground so far in developing a Kafka Streams application. While we still have much more to cover, let's pause for a moment and talk about the development process itself and one way we can make life easier for us while developing a Kafka Streams application.

### 3.4 Interactive Development

So we have covered in depth the graph that has been built to process purchase records from ZMart in a streaming fashion. The purchase data is now flowing through the topology and so far so good as ZMart prepares to start taking advantage of your streaming application.

But there is one challenge here that may not be obvious at first. You have four processors, and 3 of those write out to individual topics. While it is certainly possible to have a console consumer running to view results as you develop, it'd be good to have a more convenient solution, like the ability to watch data flowing through the topology in the console:

```
[patterns]: null, PurchasePattern(zipCode='21842', item='beer', date='Thu Feb 18 12:07:10 EST 2016')
[rewards]: null, RewardAccumulator(customerName='Doe, Andrew', purchaseTotal=19.5598)
[purchases]: null, Purchase(firstName='Andrew', lastName='Doe', creditCardNumber='xxxx-xxxx-xxxx-3020', itemPurchased='beer', quantity=4, price=4.6377, purchaseDate='Thu Feb 18 12:07:10 EST 2016', zipCode='21842')
[patterns]: null, PurchasePattern(zipCode='10005', item='eggs', date='Thu Feb 11 22:03:37 EST 2016')
[rewards]: null, RewardAccumulator(customerName='Grange, Eric', purchaseTotal=26.8865)
[purchases]: null, Purchase(firstName='Eric', lastName='Grange', creditCardNumber='xxxx-xxxx-xxxx-3020', itemPurchased='eggs', quantity=2, price=10.4043, purchaseDate='Thu Feb 11 22:03:37 EST 2016', zipCode='10005')
[patterns]: null, PurchasePattern(zipCode='20852', item='batteries', date='Tue Feb 16 14:17:45 EST 2016')
[rewards]: null, RewardAccumulator(customerName='Boggins, Andrew', purchaseTotal=5.8758)
[purchases]: null, Purchase(firstName='Andrew', lastName='Boggins', creditCardNumber='xxxx-xxxx-xxxx-3020', itemPurchased='batteries', quantity=1, price=5.8758, purchaseDate='Tue Feb 16 14:17:45 EST 2016', zipCode='20852')
[patterns]: null, PurchasePattern(zipCode='20852', item='shampoo', date='Sat Feb 13 04:58:42 EST 2016')
[rewards]: null, RewardAccumulator(customerName='Loxley, Eric', purchaseTotal=10.2052)
[purchases]: null, Purchase(firstName='Eric', lastName='Loxley', creditCardNumber='xxxx-xxxx-xxxx-1938', itemPurchased='shampoo', quantity=2, price=5.3567, purchaseDate='Sat Feb 13 04:58:42 EST 2016', zipCode='20852')
[patterns]: null, PurchasePattern(zipCode='19971', item='diapers', date='Mon Feb 15 21:23:01 EST 2016')
[rewards]: null, RewardAccumulator(customerName='Black, Andrew', purchaseTotal=11.7635)
[purchases]: null, Purchase(firstName='Andrew', lastName='Black', creditCardNumber='xxxx-xxxx-xxxx-1938', itemPurchased='diapers', quantity=1, price=11.7635, purchaseDate='Mon Feb 15 21:23:01 EST 2016', zipCode='19971')
[patterns]: null, PurchasePattern(zipCode='20852', item='eggs', date='Tue Feb 18 17:31:14 EST 2016')
[rewards]: null, RewardAccumulator(customerName='Grange, Eric', purchaseTotal=6.4234)
[purchases]: null, Purchase(firstName='Eric', lastName='Grange', creditCardNumber='xxxx-xxxx-xxxx-1938', itemPurchased='eggs', quantity=1, price=6.4234, purchaseDate='Tue Feb 18 17:31:14 EST 2016', zipCode='20852')
[patterns]: null, PurchasePattern(zipCode='21842', item='diapers', date='Fri Feb 19 10:23:28 EST 2016')
[rewards]: null, RewardAccumulator(customerName='Grange, Eric', purchaseTotal=10.1552)
[purchases]: null, PurchasePattern(zipCode='20852', item='batteries', date='Thu Feb 18 21:18:06 EST 2016')
[rewards]: null, RewardAccumulator(customerName='Boggins, Steve', purchaseTotal=29.1552)
[purchases]: null, Purchase(firstName='Steve', lastName='Boggins', creditCardNumber='xxxx-xxxx-xxxx-1938', itemPurchased='batteries', quantity=4, price=7.2888, purchaseDate='Thu Feb 18 21:18:06 EST 2016', zipCode='20852')
[patterns]: null, PurchasePattern(zipCode='21842', item='doughnuts', date='Sat Feb 13 21:28:45 EST 2016')
[rewards]: null, RewardAccumulator(customerName='Smith, Bob', purchaseTotal=13.3516)
[purchases]: null, Purchase(firstName='Bob', lastName='Smith', creditCardNumber='xxxx-xxxx-xxxx-1938', itemPurchased='doughnuts', quantity=2, price=6.6758, purchaseDate='Sat Feb 13 21:28:45 EST 2016', zipCode='21842')
[patterns]: null, PurchasePattern(zipCode='20852', item='beer', date='Fri Feb 12 19:55:06 EST 2016')
[rewards]: null, RewardAccumulator(customerName='Boggins, Eric', purchaseTotal=14.2085)
[purchases]: null, Purchase(firstName='Eric', lastName='Boggins', creditCardNumber='xxxx-xxxx-xxxx-3058', itemPurchased='beer', quantity=2, price=14.1433, purchaseDate='Fri Feb 12 19:55:06 EST 2016', zipCode='20852')
[patterns]: null, PurchasePattern(zipCode='10005', item='beer', date='Fri Feb 12 02:26:32 EST 2016')
[rewards]: null, RewardAccumulator(customerName='Doe, Sarah', purchaseTotal=13.8466)
[purchases]: null, Purchase(firstName='Sarah', lastName='Doe', creditCardNumber='xxxx-xxxx-xxxx-8783', itemPurchased='beer', quantity=1, price=13.8466, purchaseDate='Fri Feb 12 02:26:32 EST 2016', zipCode='10005')
```

**Figure 3.10 A great tool while you are developing is the capacity to print data output from each node to the console. To enable printing to the console just replace any of the to methods with a call to print.**

There is a method on the `KStream` interface that can be very useful during development, the `KStream.print` method. The `KStream.print` method takes an instance of the `Printed<K, V>` class. `Printed` provides two static methods allowing you print to `stdout`, `Printed.toSysOut()` or to write results to a file: `Printed.toFile(filePath)`.

Additionally you can label your printed results by chaining the `withLabel` method allowing you to print an initial "header" with the records, which is very useful when dealing with results from different processors.

It's important that your objects provide a meaningful `toString` implementation for useful results of printing your stream either to the console or a file.

Finally, if you don't want to use `toString` or want to customize how Kafka Streams prints records there is the `Printed.withKeyValueMapper` method which takes a `KeyValueMapper` instance so you can format your records in any way you want. The same caveat of not modifying the original records I mentioned earlier applies here as well.

In the book, I focus on printing the records to the console for all examples. Below are examples of using `KStream.print` extracted from the example above:

### Listing 3.17 Using KStream Print

```

patternKStream.print(Printed.<String, PurchasePattern>toSysOut().withLabel("patterns"));

rewardsKStream.print(Printed.<String, RewardAccumulator>toSysOut().withLabel("rewards"));

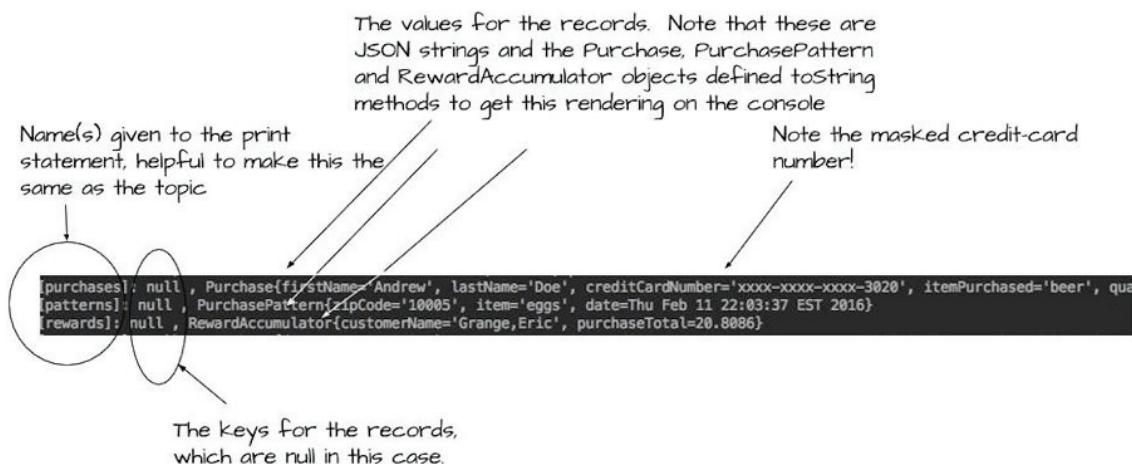
purchaseKStream.print(Printed.<String, Purchase>toSysOut().withLabel("purchases"));

```

①  
②  
③

- ① Setting up to print the PurchasePattern transformation to the console.
- ② Setting up to print the RewardAccumulator transformation to the console.
- ③ Printing the purchase data to the console.

Let's take a quick second to review the output we see on the screen and how it can help us during development.



**Figure 3.11 This a detailed view of the data on the screen. You'll quickly see if your processors are working correctly with printing to the console enabled.**

When developing your Kafka Streams application, you can replace the `to` method calls with `print` method calls. Now you can run the Kafka Streams application directly from your IDE as you make changes, stop and start the application and confirm the output is what you expect. While this is no substitution for unit and integration tests, viewing streaming results directly as you develop is a great tool.

One downside of using the `print` method is that it creates a terminal node, meaning you can't embed it in a chain of processors, but need to have a separate statement.

However, there is the `KStream.peek` method which takes a `ForeachAction` instance as a parameter and returns a new `KStream` instance. The `ForeachAction` interface has one method `apply`, which has a return type of `void`, so nothing from `KStream.peek` is forwarded downstream, making it ideal for operations like printing and allowing you to embed it in a chain of processors without the need for a separate `print` statement. We'll see the `KStream.peek` method used in this manner in other examples in the book.

## 3.5 Next Steps

At this point, you have your Kafka Streams purchase analysis program running well. Other applications have been developed to consume the messages written to the "purchase-patterns", "rewards" and "purchases" topics and the results for ZMart have been good. But alas, no good deed goes unpunished, and now that the ZMart executives can see what your streaming program can provide, a slew of new requirements come your way.

### 3.5.1 New Requirements

You have new requirements for each of the three categories of results you are producing so far. The good news is you still use the same source data. You are being asked to refine, and some cases further break down the data you are providing. The new requirements may apply to current topics or may require you to create entirely new topics.

- Purchases under a certain dollar amount need to be filtered out. Upper management is not that interested in the small item purchases for general daily articles.
- ZMart has expanded and has bought an electronics store chain and a popular coffee house chain. All purchases from these new stores will be flowing through the streaming application you have setup. You have been asked to send the purchases from these new subsidiaries to their topics.
- The NoSql solution you have chosen stores items in key-value format. Although Kafka uses key-value pairs also, the records coming into your Kafka cluster don't have any keys defined. So you'll need to generate a key for each record before the topology forwards it to the "purchases" topic.
- Unfortunately there seems to be the possibility of fraud in one of the stores. You have been asked to filter any purchases with a particular employee id and if possible, due to the sensitivity of the situation, write the results to a relational database instead of a topic.

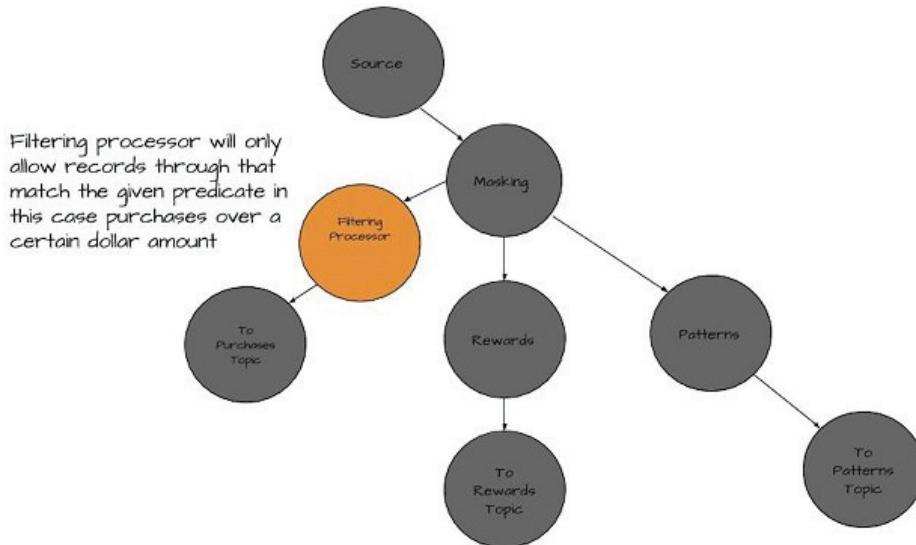
While you are confident that more requirements will be coming your way, you start to work on the current set of new requirements now. You start to look through the KStream API and are relieved to see that there are several methods already defined that will make fulfilling these new demands easy.

#### NOTE

From this point forward all code examples will be pared down to the essential parts of the discussion. Unless there is something new to introduce, you can assume the configuration and set up code remain the same. These truncated examples aren't meant to stand alone but are trimmed down to maximize clarity. The full code listing for this example can be found in  
src/main/java/bbejeck/chapter\_3/ZMartKafkaStreamsAdvancedReqsApp.java

## FILTERING PURCHASES

You decide to tackle the requirements in order, so you start with filtering out purchases that don't hit the minimum threshold. To remove low dollar purchases, you will need to insert a "filter" processing node between the `kstream` instance and the sink node. We are updating our processor topology graph in the following manner:



**Figure 3.12** Here we are placing a processor between the masking processor and the terminal node that writes to Kafka. This filtering processor will drop purchases under a given dollar amount.

### Listing 3.18 Filtering on KStream

```
KStream<Long, Purchase> filteredKStream = purchaseKStream.filter((key, purchase) ->
    purchase.getPrice() > 5.00).selectKey(purchaseDateAsKey); ①
```

- 1 Filtering purchases less than five dollars and selecting a key of the purchases date as long value.

Here you decide to use the `KStream.filter` method which takes a `Predicate<K, V>` instance as a parameter. Although we are chaining method calls together here, we are creating a new processing node in the topology. The `Predicate` interface has one method defined, `test` which takes two parameters, the key, and the value, although at this point we only need to use the value. Again we see how we can use a Java 8 lambda in place of a concrete type defined in the KStreams API.

Additionally, you want to use the purchase timestamp as a key so you use the `selectKey` processor which uses the `KeyValueMapper` defined above to extract the purchase date as a long value for use as the key, I'll go into detail about selecting the key in an upcoming section.

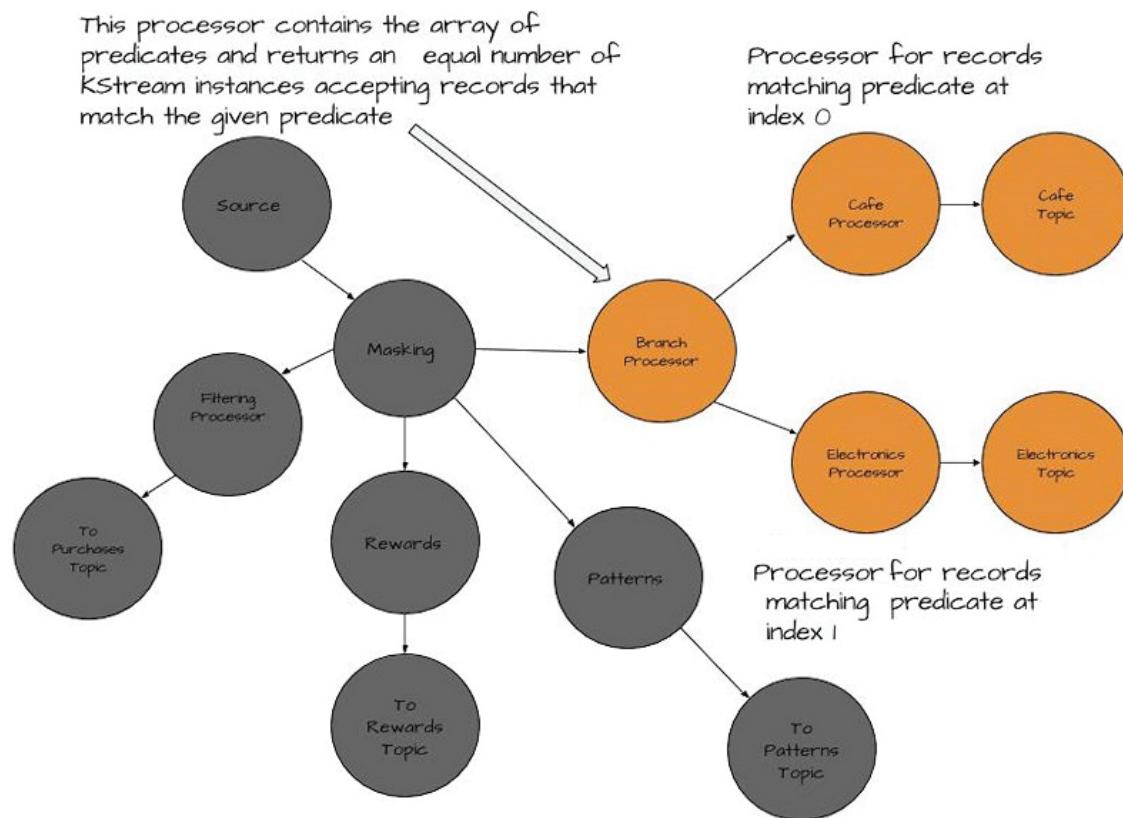
## SIDE BAR What is a Predicate

If you are familiar with functional programming, then you should feel right at home with the Predicate interface. If the term 'predicate' is new to you, it is nothing more than a given statement such as 'x < 100' and an object either matches the predicate statement or not.

There is a mirror image function `KStream.filterNot` that performs the same functionality but in reverse. Only records that don't match the given predicate are processed further in the topology.

## SPLITTING/BRANCHING THE STREAM

Now when it comes to splitting the stream of purchases into separate streams that can write to different topics has you scratching your head. Then you happen to find the `KStream.branch` method, perfect! The `KStream.branch` method takes an arbitrary number of `Predicate` instances and returns an array of `KStream` instances. The size of the returned array matches the number of predicates supplied in the call to `KStream.branch`. Where with the previous change we were modifying an existing "leaf" on the processing topology, with this requirement to branch the stream we will be creating brand new "leaf" nodes to our graph of processing nodes or topology.



**Figure 3.13 The branch processor is splitting the stream at this point into two streams; one stream consists of purchases from the cafe the other stream contains purchases from the electronics store.**

As records from the original stream flow through the branch processor, each record gets matched against the supplied predicates in the same order as provided. Matching is 'short-circuit,' the processor assigns records to stream on the first match; no attempts are made to match additional predicates.

The branch processor drops records if they do not match any of the given predicates. The order of the streams in the returned array match the order of the predicates provided to the branch method. You start to think that a separate topic for each department may not be the only approach, but you stick with what you have come up with for now and mark this task down as one requirement you will revisit later.

### **Listing 3.19 Splitting The Stream**

```
Predicate<String, Purchase> isCoffee = (key, purchase) ->
    purchase.getDepartment().equalsIgnoreCase("coffee");
Predicate<String, Purchase> isElectronics = (key, purchase) ->
    purchase.getDepartment().equalsIgnoreCase("electronics");

int coffee = 0;
int electronics = 1;

KStream<String, Purchase>[] kstreamByDept =
    purchaseKStream.branch(isCoffee, isElectronics);

kstreamByDept[coffee].to("coffee", Produced.with(stringSerde, purchaseSerde));
kstreamByDept[electronics].to("electronics", Produced.with(stringSerde, purchaseSerde));
```

- ➊ Creating the predicates, as Java 8 lambdas.
- ➋ Labeling the expected indices of the returned array.
- ➌ Calling branch to split the original stream into two streams.
- ➍ Writing the results of each stream out to a topic.

#### **WARNING Topic Creation**

In the above examples, we are sending records to several different topics. Although Kafka can be configured to automatically create topics when attempting to produce/consume for the first time from non-existent topics, it's not a good idea to rely on this mechanism. If you rely on auto-creating topics, topic configuration uses default values from the `server.config` properties file, which may or may not be the settings you need. You should always think about what topics you will need, the level of partitions and replication factor ahead of time and create them before running your Kafka Streams application.

In this example, we define the predicates ahead of time, as passing four lambda expression parameters would be a little unwieldy to view. For the same reason, we have labeled the indices of the returned array to maximize our readability here.

## SIDE BAR      Splitting vs Partitioning Streams

While "splitting" and "partitioning" may seem like similar ideas, regarding Kafka/Kafka Streams they are unrelated. Splitting a stream with the `KStream.branch` method is creating one or more streams that could ultimately send records to another topic. Partitioning is how Kafka distributes messages for one topic across servers, and aside from configuration tuning, it is the principal means of achieving high throughput in Kafka.

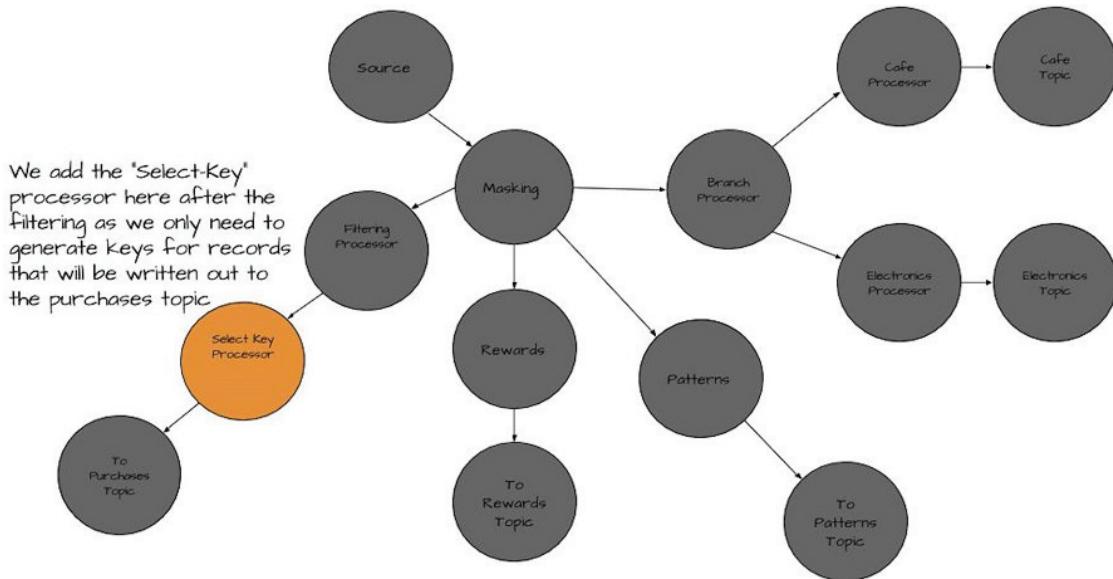
This example demonstrates the power and flexibility of Kafka Streams. You have been able to take the original stream of purchase transactions and split them into four streams, with very few lines of code. Also, you are starting to build up a more complex processing topology all while reusing the same source processor.

So far so good, you have been able to meet two of the three new requirements with ease. Now it's time to go to the last additional requirement, generating a key for the purchase record to be stored.

### GENERATING A KEY

Kafka messages are in key-value pairs, so all records flowing through a Kafka Streams application are key-value pairs as well. However, no requirements are stating that keys can't be null. In practice, if there is no need, having a null key will reduce the overall amount of data that travels the network. All of the records flowing into the ZMart Kafka Streams application have null keys.

That has been fine, up until the point you realize you NoSQL storage solution stores data in key-value format. So you need a way to create a key from the `Purchase` data before it gets written out to the "purchases" topic. While you certainly could use `KStream.map` to generate a key and return a new key-value pair (where only the key would be 'new'), there is the more succinct `KStream.selectKey` method that returns a new KStream instance that produces records with a new key (possibly a different type) and the same value. This change to the processor topology is similar to the filtering in that we will add a processing node between the "filter" and "to" processor.



**Figure 3.14** The NoSQL data store we are using will use the purchase date as a key for the data it stores. Here we're adding a selectKey processor to extract the purchase date to use as a key right before we write the data to Kafka.

### Listing 3.20 Generating A New Key

```

KeyValueMapper<String, Purchase, Long> purchaseDateAsKey = (key, purchase) ->
    purchase.getPurchaseDate().getTime(); 1

KStream<Long, Purchase> filteredKStream = purchaseKStream.filter((key, purchase) ->
    purchase.getPrice() > 5.00).selectKey(purchaseDateAsKey); 2

filteredKStream.print(Printed.<Long, Purchase>toSysOut().withLabel("purchases")); 3
filteredKStream.to("purchases", Produced.with(Serdes.Long(), purchaseSerde)); 4

```

- ① The KeyValueMapper extracting the purchase date and converted to a long
- ② Filtering out purchases and selecting the key in one statement
- ③ Printing results to the console
- ④ Materializing results to a Kafka topic.

To create the new key, we are taking the purchase date and converting it to a long. Although we could have passed a lambda expression, we have assigned it to a variable to help with readability. Also, note that we needed to change the serde type used in the KStream.to method, as we have changed the type of the key.

The previous example is a very simple display of mapping to a new key. Later in another example, we will show how to select keys to enable joining separate streams. Also, all of our examples up until this point have been stateless, but there are several options available to us to do stateful transformations as well, which will see a little later on.

### 3.5.2 Writing Records Outside of Kafka

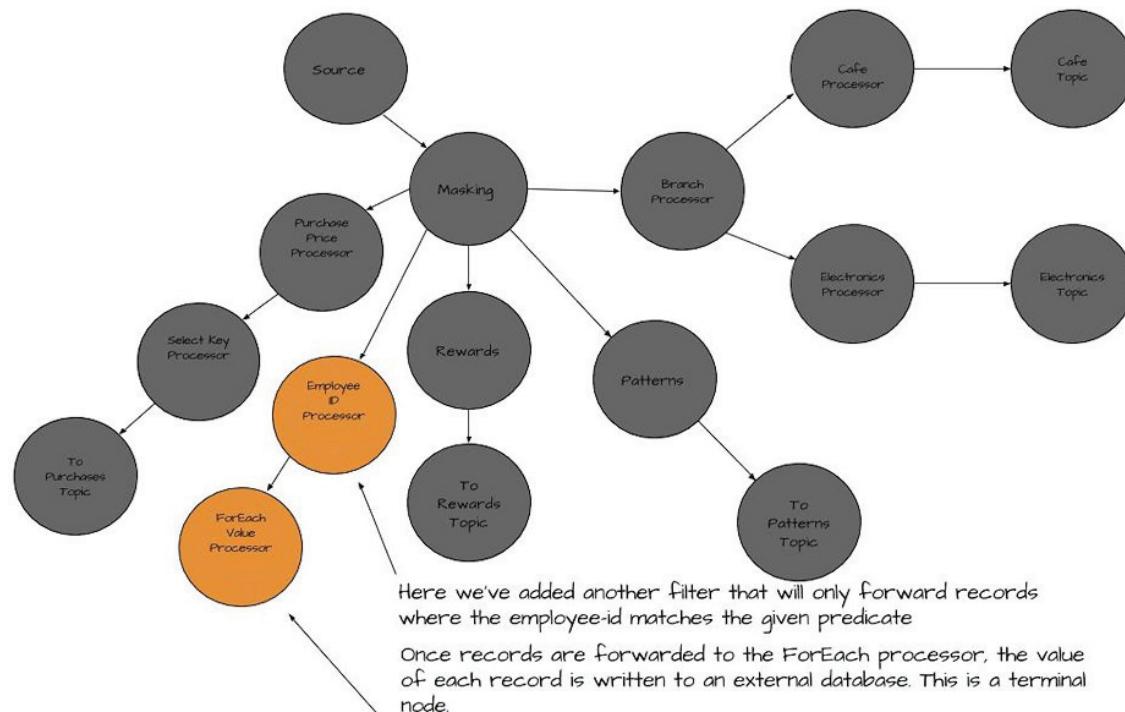
The security department at ZMartYou has approached you. Apparently, in one of the stores, there is a suspicion of fraud. There have been reports that one of the managers of the store is entering invalid discount codes for purchases. Security is not sure of what is going on, but they are asking for your help.

The security folks don't want this information to go into a topic. You talk to them about securing Kafka, and access controls and how you can lock down access to a topic, but the security folks are standing firm, these records need to go into a relational database where they have full control. You sense this is a fight you can't win, so you relent and resolve to get this task done as requested.

#### FOREACH ACTIONS

The first thing you need to do is create a new `KStream` that has results filtered down to a single employee id. Even though you have a large amount of data flow through your topology, this will reduce the volume to a tiny amount.

Here we are going to use `KStream.filter` with the predicate looking to match a specific employee id. This filter is completely separate from the previous filter and is attached to the source `KStream` instance. Although it is entirely possible to chain filters, this is not the case here; we want full access to data in the stream for this filter. As you go through the API, there is a method that catches your eye, the `KStream.foreach` method. The method predictable takes a `ForeachAction<K, V>` instance, and is a terminal node. The `KStream.foreach` is a simple processor and as the name implies `KStream.foreach` uses the provided `ForeachAction` instance to perform an action on each record it receives.



**Figure 3.15 To write purchases involving a given employee outside of our Kafka Streams application we'll first add a filter processor to extract purchases by employee id then use a foreach operator to write each record out to an external relational database.**

### **Listing 3.21 Foreach Operations**

```

ForeachAction<String, Purchase> purchaseForeachAction = (key, purchase) ->
    SecurityDBService.saveRecord(purchase.getPurchaseDate(), purchase.getEmployeeId(),
        purchase.getItemPurchased());

purchaseKStream.filter((key, purchase) ->
    purchase.getEmployeeId().equals("1234567")).foreach(purchaseForeachAction);

```

Here we created the `ForeachAction` using a Java 8 lambda (again!) and have stored it in a variable, although it's an extra line of code, the clarity gained by doing so more than makes up for it.

On the next line, we have created another `KStream` instance that will send the filtered results to the `ForeachAction` defined directly above it.

We should mention here, that `KStream.foreach` is stateless and if you need state to perform some action per record you will use the `KStream.process` method. The `KStream.process` method is discussed in the next chapter when we cover adding state to a Kafka Streams application.

If you step back and take a look at what you have accomplished at this point, it's pretty impressive considering the lines of code written. Don't get too comfortable, though, because the upper management team at ZMart has taken notice of your productivity and are excited about the results, but more changes and refinements to the purchases streaming analysis program are coming.

## **3.6 Summary**

Well, we have certainly covered a lot of ground in this chapter. You learned about the `KStream` API, and we were able to see how quickly we could develop an application that not only met your original goals but also allowed you to adapt to ever evolving changes and requirements quickly.

Specifically, we learned:

- How to use the `KStream.mapValues` function to map incoming record values to a new value, possibly of a different type. You also learned that these mapping changes should not modify the original object. There is another method `KStream.map` that performs the same action but can be used to map the key and the value to something new.
- A predicate is a statement that accepts an object as a parameter and returns true or false if that object matches a given condition. We were able to use predicates in the `filter` function to remove records from being forwarded in the topology for that don't match the given predicate.
- The `KStream.branch` method uses predicates to split records into new streams when a record matches a given predicate. The processor assigns a record to a stream on the first match and dropping unmatched records.
- How to select or modify a key using the `KStream.selectKey` method.

In the next chapter, we start to look at state, the required properties to use state with a steaming application, and why need to add state at all. Then we'll show how we add state to a `KStream` application first by using stateful versions of `KStream` methods we've seen in this chapter (`KStream.mapValues`). For a more advanced example, we'll show how we can perform joins between two different streams of purchases to help ZMart improve customer service.

# *Streams and State*

In this chapter:

- Applying Stateful Operations to Kafka Streams.
- Using State Stores for Lookups and Remembering Previously Seen Data.
- Joining Streams for Added Insight.
- How Time and Timestamps Drive Kafka Streams.

In the last chapter, we dove head first into the KStreams API and built a processing topology to handle streaming requirements from purchases made at ZMart locations. While we built a non-trivial processing topology, it was one-dimensional in that all transformations and operations were stateless. We considered each transaction in isolation without any regard to other events occurring at the same time or within certain time boundaries either before or after the transaction. Also, we only dealt with individual streams, ignoring any possibility of gaining additional insight by joining streams together.

Now we want to extract the maximum amount information from our Kafka Streams application. To get this level of maximum information, we are going to need to use state. State is nothing more than the ability to recall information we've seen before and use it to connect to current information. We can utilize state in different ways. We'll see one example when we explore the stateful operations, such as accumulation of values, provided by the KStreams API.

Another example of state we'll discuss is the joining of streams. Joining streams is closely related to the joins performed in database operations, joining records from the employee and department tables to generate a meaningful report on who staffs which departments in a company.

We'll also define what our state needs to look like and the requirements around using state when we discuss StateStores in Kafka Streams. Finally, we'll weigh the importance

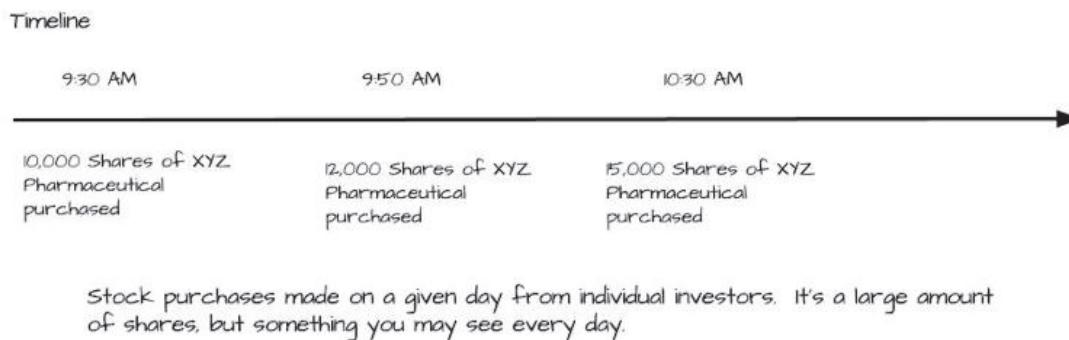
of timestamps and how they direct decisions for working with stateful operations from ensuring we only work with events occurring within a given timeframe to helping us work with data arriving out of order.

## 4.1 Thinking of Events

When it comes to event processing, sometimes an event by itself requires no further information or context. Other times we can understand an event in a literal sense, but without some added context we might miss the actual meaning of what is occurring or think of the situation in a whole new light given some additional information.

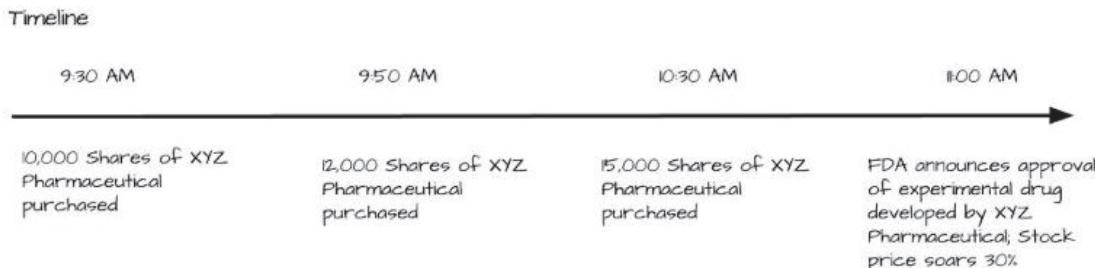
As an example of an event that does not require additional information, consider the attempted use of stolen credit card. The transaction is canceled immediately once the stolen card's attempted use is detected. We don't need any additional information to be sure of the decision.

But in some case, a singular event in and of itself doesn't give you enough information to make impactful decisions. Consider a series of stock purchases by three individual investors within a short period. On the face of it, there isn't anything inherently odd about those transactions that would give us pause. Investors buying shares of the same stock is something that happens every day on Wall Street.



**Figure 4.1 A stock transaction without any extra information doesn't look like anything out of the ordinary.**

Now let us add some context; within a short period after the individual stock purchases, XYZ Pharmaceutical announced government approval for a new drug highly anticipated by the public, which sent the stock price to historic highs. Additionally, those three investors had close ties to XYZ Pharmaceutical. Now the transactions in question can be view in a whole new light.



Well the timing of the purchases and information release raises questions:

- Were these investors leaked information ahead of time?
- Is this actually one investor with inside information trying to cover their tracks?

**Figure 4.2** However when we add some additional context about the timing of the stock purchase, and we'll see it in an entirely new light.

### 4.1.1 Streams Need State

What the fictional scenario example from above shows us is something that most of us already know instinctively. Sometimes it's easy to reason about what's going on, but most of the time we need some context to make good decisions. When it comes to stream processing, we call that added context "state."

At first glance, the notion of state and stream processing may seem to be at odds with each other. Stream processing infers a constant flow of discrete events that don't have much to do with each other and need to be dealt with as they occur.

The notion of state might conjure images of a static resource such as a database table.

In actuality we could view these two as one in the same, it's just that the rate of change in a stream is potentially much faster and frequent than a database table <sup>10</sup>.

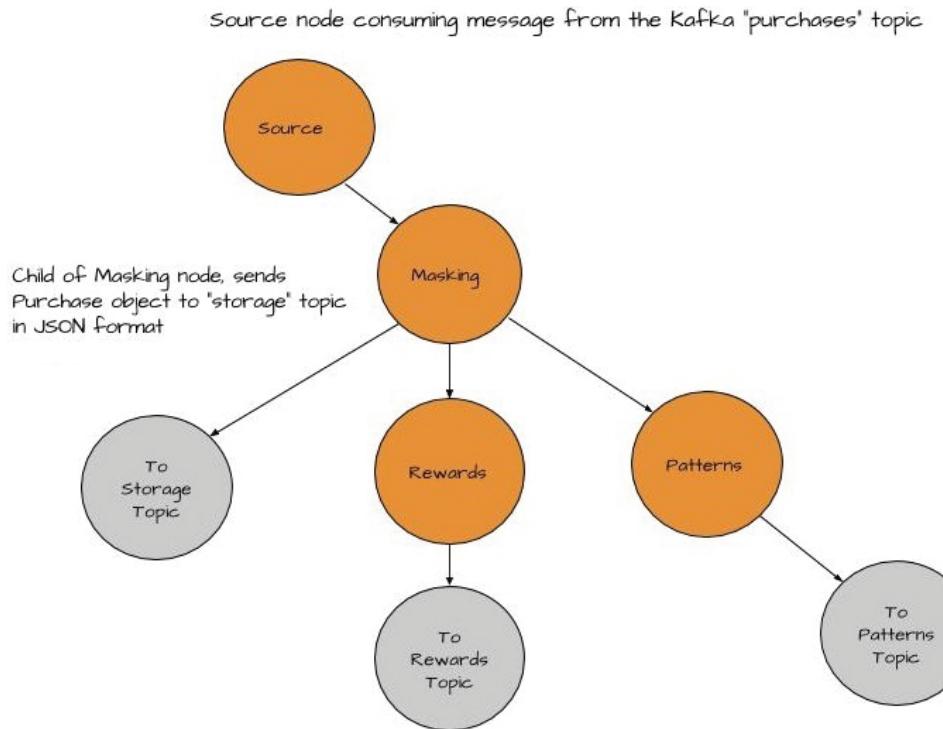
---

Footnote 10 [www.oreilly.com/ideas/why-local-state-is-a-fundamental-primitive-in-stream-processing](http://www.oreilly.com/ideas/why-local-state-is-a-fundamental-primitive-in-stream-processing)

However, you don't *always* need state to work with streaming data. In some cases, you may truly have discrete events/records that carry enough information to be valuable on their own. But more often than not our incoming stream of data either needs enrichment from sort of store, using information from events that arrived before or joining related events together from different streams.

### 4.2 Applying Stateful Operations to Kafka Streams

In this section, we are going to demonstrate how we can add a stateful operation to an existing stateless one to improve the information collected by our application. We are going to modify the original topology we built in Chapter 3 pictured below to refresh your memory.



**Figure 4.3 Here's another look at the original topology you built in Chapter three.**

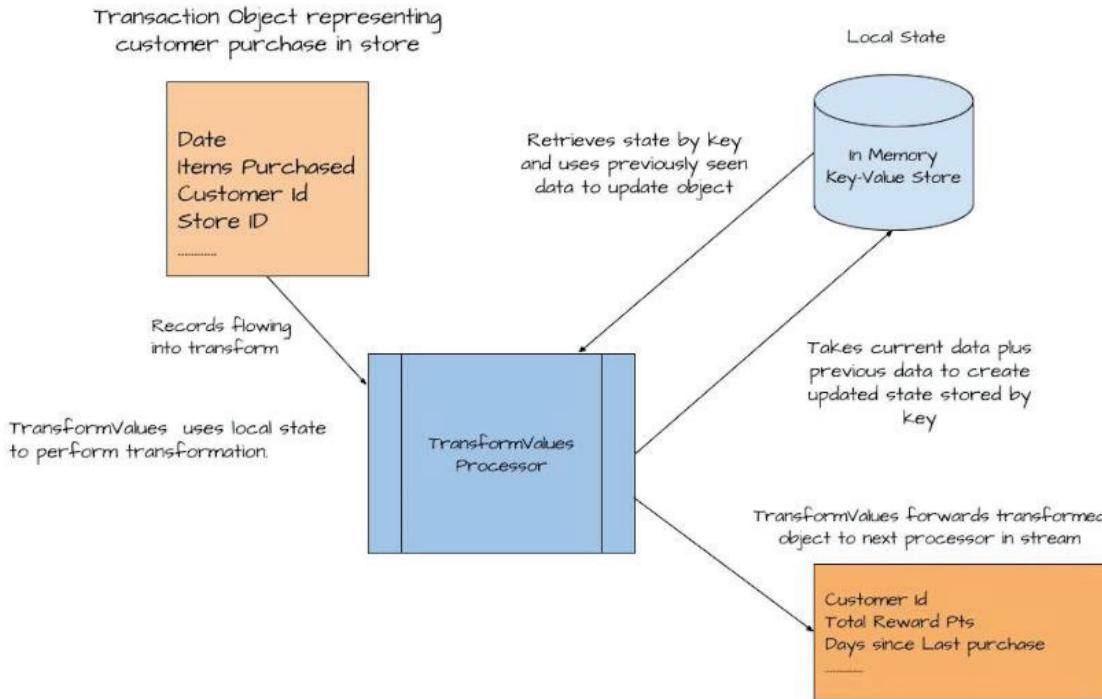
In the topology picture above, we produced a stream of purchase transaction events. One of the processing nodes in the topology, named "Rewards," calculated reward points for customers based on the amount of the sale. But in that processor, we were just calculating the total number of points for a single transaction and forwarding the results.

If we added some state to the processor, we could keep track of the number of customer reward points, and the consuming application at ZMart would just need to check the total and send out a reward if needed.

Now that we have discussed how and why we need state to work with Kafka Streams (or any other streaming application) let's go over some concrete examples of how we can apply stateful operations; starting with transforming the stateless rewards processor into a stateful processor using `TransformValues`.

#### 4.2.1 Transform Values Processor

The most basic of the stateful functions is `KStream.transformValues`. The image below is a visual guide to how the `KStream.transformValues` method operates.



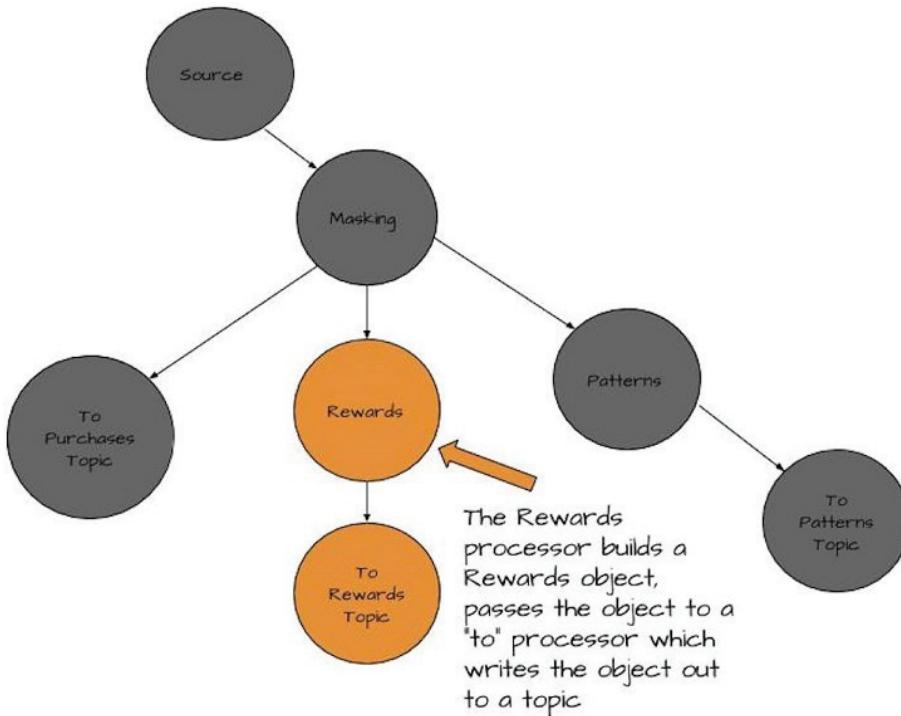
**Figure 4.4** The `transformValues` processor uses information stored in local state to update incoming records. In this case, the customer id is the key used to retrieve and store the state for a given record.

This method is semantically the same as `KStream.mapValues` with a few exceptions. One difference is `transformValues` has access to a `StateStore` instance to accomplish its given task. The other difference is the ability to schedule operations to occur at regular intervals via a `punctuate` method. The `punctuate` method will be discussed in detail when we cover the Processor API.

Next, we are going to re-work the original stateless rewards process to a stateful one, keeping track of the total bonus points achieved so far and the amount of time between purchases to provide more information to downstream consumers of your application.

#### 4.2.2 Stateful Customer Rewards

In chapter 3 we built a streaming topology for ZMart, and one of the processors in the topology extracts information for customers belonging to ZMart's rewards program. Here's the correspondent node in the topology:



**Figure 4.5** This is the original rewards processor you added to the topology in Chapter three. This processor is going to get an upgrade and go from making "stateless" transformations to "stateful" ones.

Initially, the rewards processor used the `KStream.mapValues` method to map the incoming `Purchase` object into a `RewardAccumulator` object. The `RewardAccumulator` object originally consisted of just two fields, the `customerId` and the purchase total for the transaction.

#### **Listing 4.1 RewardAccumulator Object**

```
public class RewardAccumulator {

    private String customerId; ①
    private double purchaseTotal; ②
    private int currentRewardPoints; ③

    //details left out for clarity
}
```

- ① The customer id.
- ② Total Dollar amount of purchase.
- ③ The current number of reward points.

Now the requirements have changed some, and now there are points associated with the ZMart rewards program. Where before there was an application reading from the "rewards" topic calculating customer achievements, there is a desire from management to have the point system maintained and calculated by the streaming application instead.

Additionally, you need to capture the amount of time between the current and the last purchase the customer made.

When the application reads records from the "rewards" topic, the consuming application will only need to check if the total points are above the given threshold to distribute an award. To meet this new goal, you add the `totalRewardPoints` and `daysFromLastPurchase` fields to the `RewardAccumulator` object and use the local state to keep track of accumulated points and the last date of purchase.

As a result here is the refactored `RewardAccumulator` code needed to support the changes.

#### **Listing 4.2 Refactored RewardAccumulator Object found in src/main/java/bbejeck/model/RewardAccumulator.java**

```
public class RewardAccumulator {

    private String customerId;
    private double purchaseTotal;
    private int currentRewardPoints;
    private int daysFromLastPurchase;
    private long totalRewardPoints; ①

    //details left out for clarity
}
```

- ① Field added for tracking total points

The updated rules for the purchase program are simple. The customer earns a point per dollar and transaction totals are rounded down to the nearest dollar. The overall structure of the topology won't change, but the rewards processing node will change from using the `KStream.mapValues` method to the `KStream.transformValues`. Semantically these two methods operate the same way in that we still map the `Purchase` object into a `RewardAccumulator` object. The difference lies in the ability to use local state to perform the transformation. Specifically, there are two main two steps we'll take:

1. Initialize the Value Transformer.
2. Map the `Purchase` object to a `RewardAccumulator` using state.

The `KStream.transformValues` takes a `ValueTransformerSupplier<V, R>` object which supplies an instance of the `ValueTransformer<V, R>` interface.

Our implementation of the `ValueTransformer` is `PurchaseRewardTransformer<Purchase, RewardAccumulator>`. For the sake of clarity, we won't reproduce the entire class here in the text. Instead, we'll walk through the important methods to cover the list of performed steps from above. Also note that these code snippets aren't meant to stand alone, and some details will be left out for clarity. Now let's move on with our first step, initializing the processor.

### 4.2.3 Initialize the Value Transformer

Your first step is to set up/create any instance variables in the transformer `init` method:

#### Listing 4.3 Init Method

```
private KeyValueStore<String, Integer> stateStore; ①

private final String storeName;
private ProcessorContext context;

public void init(ProcessorContext context) {
    this.context = context;
    stateStore = (KeyValueStore) this.context.getStateStore(storeName); ②
} ③
```

- ① Instance variables.
- ② Setting local reference to ProcessorContext.
- ③ Retrieving the StateStore instance by `storeName` variable, `storeName` set in constructor.

In the `init` method we retrieve the state store created when building the processing topology. Inside our transformer class, we cast to a `KeyValueStore` type as inside our transformer we are not concerned with the implementation (more on state store implementation types in the next section) at this point, just that we can retrieve values by key. There are other methods (`punctuate`, `close`) not listed here that belong to the `ValueTransformer` interface. We discuss The `punctuate` and `close` methods when covering the Processor API in Chapter six.

### 4.2.4 Map the Purchase object to a RewardAccumulator using state

Now that we have initialized our processor we move on to the main part of the action transforming a `Purchase` object using state. There a few simple steps we take performing the transformation:

1. Check for points accumulated so far by customer id.
2. Sum the points for the current transaction and present total
3. Set the reward points on the `RewardAccumulator` to the new total amount.
4. Save the new total points by customer id in the local state store.

#### Listing 4.4 Transforming the Purchase using State

```
public RewardAccumulator transform(Purchase value) {
    RewardAccumulator rewardAccumulator = RewardAccumulator.builder(value).build(); ①

    Integer accumulatedSoFar = stateStore.get(rewardAccumulator.getCustomerId()); ②

    if (accumulatedSoFar != null) {
        rewardAccumulator.addRewardPoints(accumulatedSoFar); ③
    }
    stateStore.put(rewardAccumulator.getCustomerId(),
```

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

```

    rewardAccumulator.getTotalRewardPoints()));

    return rewardAccumulator;
}

```

4

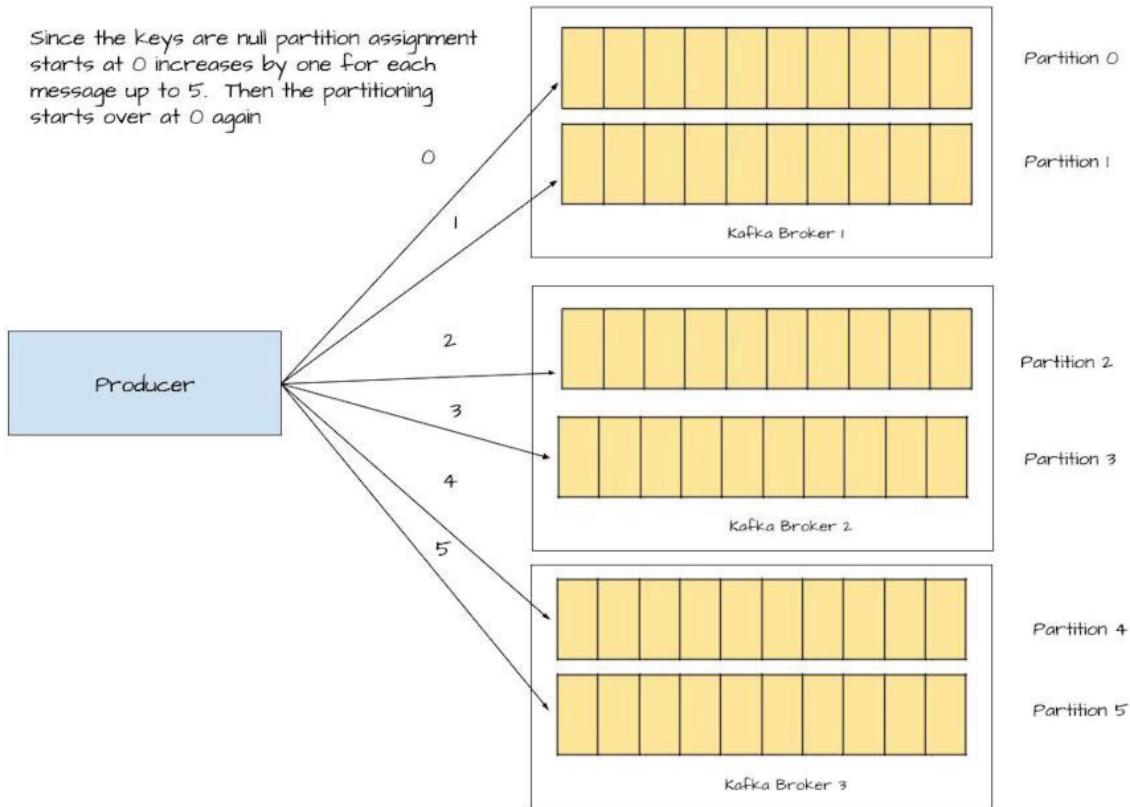
5

- 1 Building the RewardAccumulator object from the Purchase.
- 2 Retrieving latest count by customer id.
- 3 If an accumulated number exists, add it to the current total.
- 4 Store the new total points in the StateStore.
- 5 Return the new accumulated rewards points.

In the `transform` method we first do the mapping of a `Purchase` object into the `RewardAccumulator` and is the same operation used in the `mapValues` method. On the next few lines are where state gets involved in the transformation process. We do a lookup by key (customer id) and adding any points accumulated so far to the points from the current purchase. Then we place the new total in the state store until another it's needed again.

All that's left is to update our rewards processor. But before we update the processor, we have overlooked the fact that we need access to all sales by customer id. Gathering information per sale for a given customer implies all transactions for that customer are on the same partition.

However, because our transactions come into the application without a key, the producer assigns the transaction to partitions in a round-robin fashion. We covered round-robin partition assignment in Chapter 2, but it's worth reviewing the image here again:



**Figure 4.6** Here's another look at how a Kafka Producer distributes records evenly (round-robin) when the keys are null.

We'll have an issue here (unless you are using topics with one partition). Because of not having a populated key, round-robin assignment means all transactions for a given customer won't land on the same partition. Having customer transactions with the same id placed on the same partition is important because we look up records by id in the state store.

Otherwise, we'll have customers with the same id across different partitions, meaning lookups for the same customer in multiple state stores. The way for us to solve this problem is to repartition the data by customer id. Let's take a look at how we'll repartition our data in the next section.

#### NOTE

#### State Stores are Assigned Per Task

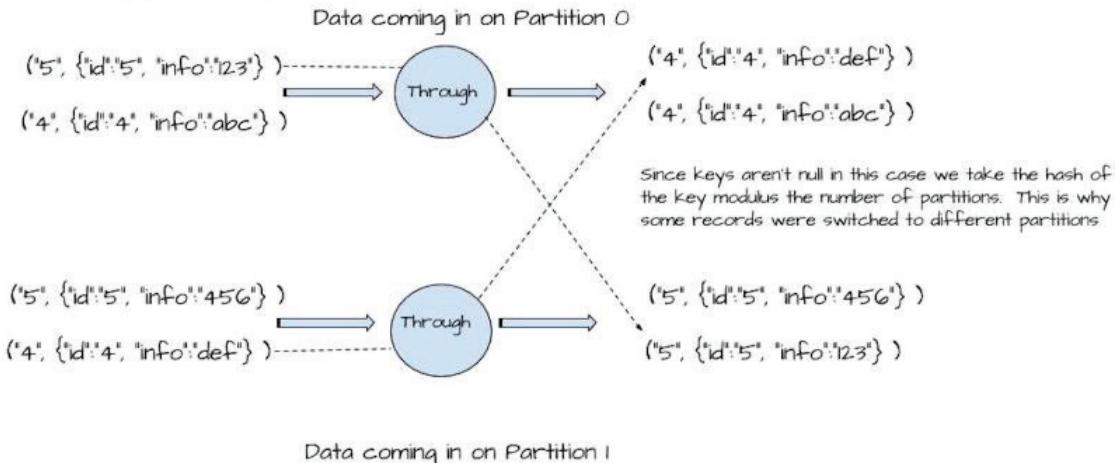
The statement above could be interpreted to mean that each partition has its own state store, but that is not the case. Partitions are assigned to a `StreamTask` and each `StreamTask` has its own state store.

## REPARTITIONING THE DATA

To repartition the data we write our records out to an intermediate topic then immediately read the records back into our topology. The following image gives us a good idea of how repartitioning works:

The 'through' processor writes records out to a new topic and returns a new KStream instance using the data written to that new topic as its source.

Since we didn't specify a StreamPartitioner we are using the DefaultPartitioner of the producer to write records out to the correct partitions. By taking the hash of the key modulus the number of partitions, records with newly generated keys are written out to a new a new topic.



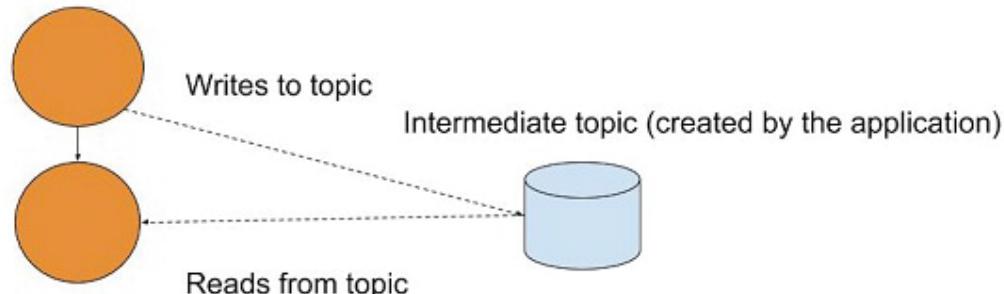
**Figure 4.7 Visual depiction on how repartitioning works**

Now we've discussed what repartitioning is and why we need to do it, let's demonstrate how. Fortunately, it's a trivial operation accomplished by using the `KStream.through` method. Here's an image depicting how the `KStream.through` method works:

The `KStream.through` method writes out to an intermediate topic then returns a new `KStream` instance that uses the intermediate topic as its source.

Original `KStream` node making the "through" call

Seamlessly repartitioning the data and keeping the flow going through the application.



`KStream` instance returned that immediately starts to consume from the intermediate topic.

**Figure 4.8 Writing out to intermediate topic then reading from it in a new `KStream` instance**

The `KStream.through` method creates an intermediate topic, and the *current*

`KStream` instance will start writing records to that topic. A *new* `KStream` instance is returned from the `through` method call, using the same intermediate topic for its source. This way we can seamlessly repartition our data and continue the data flow our Kafka Streams application. The code for the `KStream.through` method looks like this:

### Using `KStream.through` Method

```
RewardsStreamPartitioner streamPartitioner = new RewardsStreamPartitioner(); ①
KStream<String, Purchase> transByCustomerStream = purchaseKStream.through("customer_transactions",
Produced.with(stringSerde, purchaseSerde, streamPartitioner)); ②
```

- ① Instantiating our concrete `StreamPartitioner` instance.
- ② Creating a new `KStream` with `KStream.through` method. The `KStream.through` method takes two parameters: the topic name and a `Produced` instance which provides the key serde, value serde and the stream partitioner.

We could have taken the approach of generating new keys. But since we don't need keys beyond this point and the value contains the customer id, we've decided to use the value to determine the partition.

Since we've mentioned using a `StreamPartitioner` and it's essential for repartitioning to work, let's take a quick look at how we'd create one.

## USING A STREAMPARTITIONER

A `StreamPartitioner` instance is used to return a number representing the partition assignment. The partition assignment is calculated by taking the hash of an object modulo the number of partitions. Now partition assignments are determined by the value of the customer-id, and we know that data for a given customer is in the same state store.

Now let's look at our `StreamPartitioner` implementation:

### Listing 4.5 `RewardsStreamPartitioner` found in src/main/java/bbejeck/chapter\_4/partitioner/RewardsStreamPartitioner.java

```
public class RewardsStreamPartitioner implements StreamPartitioner<String, Purchase> {

    @Override
    public Integer partition(String key, Purchase value, int numPartitions) {
        return value.getCustomerId().hashCode() % numPartitions; ①
    }
}
```

- ① Determine partition by customer id.

Notice that we have not generated a new key and we have used a property of the value to determine the correct partition.

The key point to take away from this quick "detour" is when using state to update and modify data by key it's important for those records to be on the same partition.

**CAUTION** **Use repartitioning with care.**

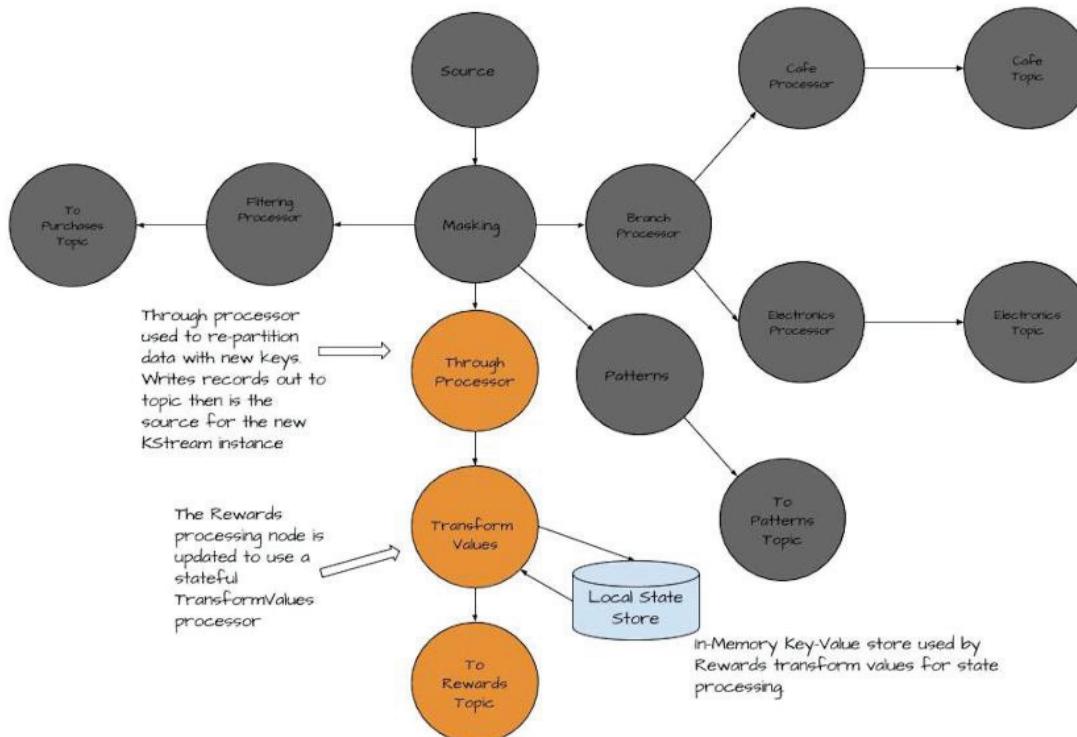
Don't mistake the ease of our repartitioning demonstration as a recommended best practice you can be cavalier with. While sometimes repartitioning is necessary it comes at the cost of duplication of your data. So my advice here is to use map/transform operations on values whenever possible and to use repartitioning logic sparingly.

Now let's pick up where we left off and go back to making changes in the rewards processor node to support stateful transformation.

#### 4.2.5 Updating The Rewards Processor

Up to this point, we've created a new processing node that writes purchase objects out to a topic using the customer id as the key. This new topic also becomes the new source for our soon to be updated rewards processor.

We did this to ensure all purchases for a given customer are written to the same partition, allowing us to have consistent results. We've updated our processing topology below by inserting the new processor between the credit card masking node (source for all purchase transactions) and the rewards processor. Now we'll move on to updating the rewards processor in the topology to use our new stateful processor. We'll call our new processor "Transform Values" as indicated below:



**Figure 4.9** Here we've added the through processor ensuring we send purchases to partitions by customer id, which allows the transformValues processor to make the right updates using local state.

Now we use our new `KStream` instance (created above using the `KStream.through` method) to update our processor and use the stateful transform approach with the following line of code:

#### **Listing 4.6 Changing Rewards Processor to Use Stateful Transformation**

```
KStream<String, RewardAccumulator> statefulRewardAccumulator =
    transByCustomerStream.transformValues(() ->
        new PurchaseRewardTransformer(rewardsStateStoreName), rewardsStateStoreName);
statefulRewardAccumulator.to("rewards", Produced.with(stringSerde, rewardAccumulatorSerde));
```

1

2

- ① Using a stateful transformation
- ② Writing the results out to a topic

The `KStream.transformValues` method takes a `ValueTransformerSupplier<V, R>` instance, which is provided via a Java 8 lambda expression.

In this section, we added stateful processing to a stateless node. By adding state to the processor, ZMart takes action sooner after a customer makes a reward qualifying purchase. We've seen how to use a state store and the benefits using one provides.

But we've glossed over important details that we'll need to understand how state can impact our applications entirely. With this in mind, the next section will specify which type of state store to use, what requirements we'll need to make state efficient and how we add the stores to a Kafka Streams program.

### **4.3 Using State Stores for Lookups and Seen Data**

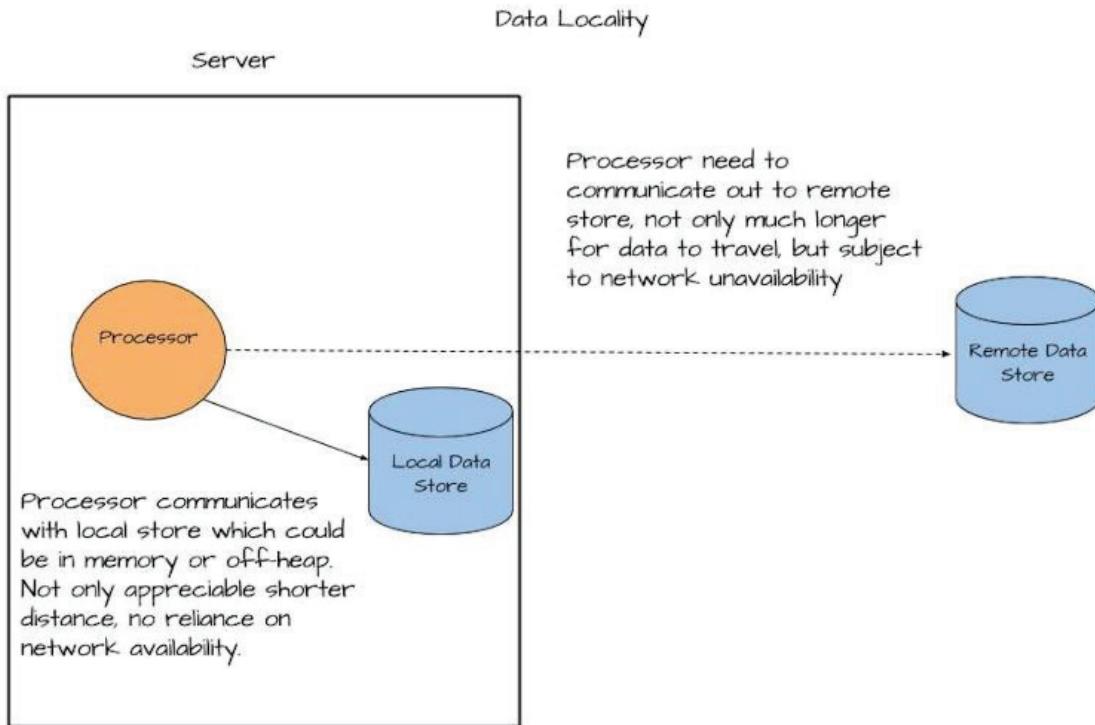
In this section, we are going to learn the essentials of using state stores in Kafka Streams and the key factors of using state in streaming applications in general. We'll do this so we can make useful choices when using state in our Kafka Streams applications.

So far we've discussed the need for using state with streams and have seen an example of one of the more basic stateful operations available in Kafka Streams. Before we describe the idea of state stores in Kafka Streams, let's briefly describe two important attributes of state that we'll need

1. Data locality
2. Failure Recovery

#### **4.3.1 Data Locality**

Our first key point is data locality. Data locality is critical for performance. While key lookups are typically very fast, at scale the latency introduced by using remote storage becomes a bottleneck. In the image below we show the principal behind data locality.



**Figure 4.10 Data locality is necessary for stream processing. First of all traveling across the network to retrieve data incurs a cost that at high data volume is noticeable, secondly it provides some isolation due to the fact if a process goes down other processes aren't affected.**

The dashed line represents a network call to retrieve data from a remote database hosted on another server. The solid line depicts a call to an in-memory data store co-located on the same server. While the image above is not drawn to scale, it's pretty clear that making a call to get data locally is more efficient than making a call across a network to a remote database.

The key point here is not the amount of latency per-record retrieval by itself, which may be minimal. The important factor here is you will potentially process millions to billions of records through a streaming application. When multiplied by a factor that large even a small amount of network latency has a huge impact.

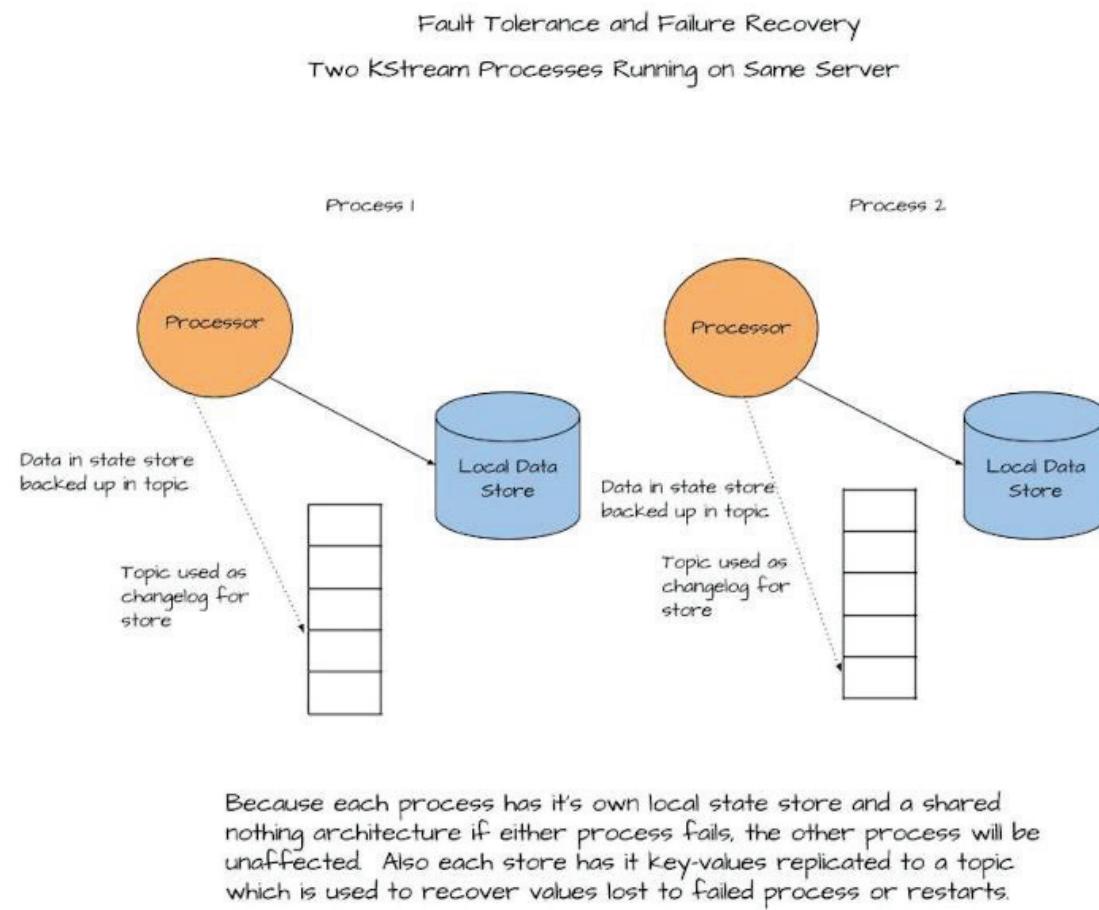
Data locality also means that the store is local to each processing node and there is no sharing across processes or threads. By not sharing a store, if a particular process fails, it shouldn't have an impact on the other stream processing processes/threads.

The key point here is while streaming applications will sometimes require state, it needs to be local to where the processing occurs, and each server/node in the application should have an individual data store.

### 4.3.2 Failure Recovery/Fault Tolerance

Application failure is inevitable, especially when we are talking about distributed applications. We need to shift our focus from preventing failure to how to recover quickly from failure or even restarts. The following image depicts the principle of data locality and fault tolerance. Each processor has its local data store, and a (changelog) topic is used to back up the state store.

While backing a state store with a topic may seem expensive, there are a couple of mitigating factors at play. First a `KafkaProducer` sends records in batches. Secondly, by default records are cached. It's only on cache flush when Kafka Streams writes records to the store, hence only the latest record for any given key is persisted. We discuss this caching mechanism with state stores in more detail in Chapter 5, section 2.



**Figure 4.11 The ability to recover from failure is important for stream processing applications. Kafka Streams persists data from the local in-memory stores to an internal topic so when resuming operations either from a failure or a re-start the data is re-populated.**

The state stores provided by Kafka Streams meet both the locality and fault-tolerance requirements. State stores are local to the defined processors and don't share access across processes or threads. State stores use topics as a backup and quick recovery, as shown in the image above.

Over the past two sections, we've covered the requirements we need for using state with a streaming application. The next step is to show how we enable the usage of state in our Kafka Streams applications.

### 4.3.3 Using State Stores in Kafka Streams

Adding a state store is a matter of creating a `StoreSupplier` instance with one of the static factory methods on the `Stores` class. In conjunction with the `StoreSupplier` there are two additional classes for customizing the state store, the `Materialized` and `StoreBuilder` classes. Which one you'll use depends on how you are adding the store to the topology.

If you are using the high-level DSL, typically you'll use the `Materialized` class and when working with the lower-level Processor API you'll use the `StoreBuilder`. Even though our current example is using the high-level DSL, we're adding a state store to a `Transformer` so we'll use the `StoreBuilder` for state store customization.

#### **Listing 4.7 Adding A StateStore**

```
String rewardsStateStoreName = "rewardsPointsStore";
KeyValueBytesStoreSupplier storeSupplier =
    Stores.inMemoryKeyValueStore(rewardsStateStoreName); 1

StoreBuilder<KeyValueStore<String, Integer>> storeBuilder =
    Stores.keyValueStoreBuilder(storeSupplier,
        Serdes.String(),
        Serdes.Integer()); 2

builder.addStateStore(storeBuilder); 3
```

- ① Creating the StateStore supplier.
- ② Creating the StoreBuilder and specifying the key and value types.
- ③ Adding State Store to the topology.

In listing one you create a `StoreSupplier` that provides an in-memory key value store. In listing two you combine provide the `StoreSupplier` as a parameter to create a `StoreBuilder` additionally specifying `String` keys and `Integer` values. Finally in listing three you add the `StateStore` to the topology by providing the `StoreBuilder` to the `StreamsBuilder`.

Here we created an in-memory key-value store with `String` keys and `Integer` values and have added the store to our application with the `StreamsBuilder.addStateStore` method. As a result, we now have the ability to use state in our processors by using a name for the state store.

We have shown an example of building a in-memory state store, but we have options at our disposal to create different types of `StateStore` instances, let's take a look at those options.

### 4.3.4 Additional Key Value Store Suppliers

In addition to the `Stores.inMemoryKeyValueStore` here are the other static factory methods for producing store suppliers:

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

1. Stores.persistentKeyValueStore
2. Stores.lruMap
3. Stores.persistentWindowStore
4. Stores.persistentSessionStore

It's worth noting that all persistent StateStore instances provide local storage using RocksDB <sup>11</sup>. Before we move on from state stores, I'd like to cover two additional important aspects of Kafka Streams state stores, how fault tolerance is provided with changelog topics and configuring changelog topics.

---

Footnote 11 [rocksdb.org/](https://rocksdb.org/)

### 4.3.5 StateStore Fault Tolerance

All the StateStoreSupplier types have logging enabled as a default. Logging in this context means there is a Kafka topic used as a changelog to backup the values in the store and provide fault tolerance.

For example, if you were to lose a machine running Kafka Streams, once you recovered your server and restarted your Kafka Streams application, the state stores for that instance would get restored to their original contents (the last commit before crashing).

This logging can be disabled when using the Stores factory with the method disableLogging(). However, one should give serious consideration before disabling logging; as this makes state stores lose fault tolerance and the ability to recover in the face of losing state on a machine.

### 4.3.6 Configuring Changelog Topics

The changelogs for state stores are configurable via the withLoggingEnabled(Map<String, String> config) method. You can use any configuration parameters available for topics in the map. Configuration of changelogs of state stores is an important consideration when building a Kafka Streams application. However, keep in mind you never need to create changelog topics, Kafka Streams handles changelog topic creation for you.

**NOTE**

**State Store Changelogs are Compacted Topics**

State Store changelogs are compacted topics. We discussed compacted topics back in Chapter 2. If you recall the delete semantics require a null value for a key, so if you want to remove a record from a state store permanently, you'll need to do a `put(key, null)` operation.

With Kafka topics the default for data retention setting for a log segment is one week and the size is unlimited. Depending on the volume of data, this might be acceptable, but

there is a good chance you'll want to adjust those settings. Additionally, the default cleanup policy is delete. First let's take a look at how you can configure your changelog topic to have a retention size of 10G and a retention time of 2 days :

#### **Listing 4.8 Setting Changelog properties**

```
Map<String, String> changeLogConfigs = new HashMap<>();
changeLogConfigs.put("retention.ms", "172800000");
changeLogConfigs.put("retention.bytes", "10000000000");

// to use with a StoreBuilder
storeBuilder.withLoggingEnabled(changeLogConfigs);

// to use with Materialized
Materialized.as(Stores.inMemoryKeyValueStore("foo") .withLoggingEnabled(changeLogConfigs));
```

In chapter 2 we discussed compacted topics offered by Kafka. To refresh your memory compacted topics use a different approach to cleaning a topic, instead of deleting log segments by size or time, log segments are "compacted" by keeping only the latest record for each key, older records with the same key are deleted. By default Kafka Streams creates changelog topics with a delete policy of "compact".

But if you have a changelog topic with a lot of unique keys, compaction might not be enough as the size of the log segment will keep growing. In that case the solution is simple, you specify a cleanup policy of delete and compact as demonstrated here:

#### **Listing 4.9 Setting Cleanup Policy**

```
Map<String, String> changeLogConfigs = new HashMap<>();
changeLogConfigs.put("retention.ms", "172800000");
changeLogConfigs.put("retention.bytes", "10000000000");
changeLogConfigs.put("cleanup.policy", "compact,delete");
```

Now your changelog topic will be kept at a reasonable size even with unique keys. This has been a brief section on topic configuration, but we'll discuss more about changelog topics and internal topic configuration in the configuration appendix.

We have just learned about the in-memory and persistent state stores Kafka Streams provides and how we can include them in our streaming applications. We also learned the importance of data locality and fault tolerance when using state in a streaming application. Now that we have covered basic stateful operations and state stores let's move on to joining streams.

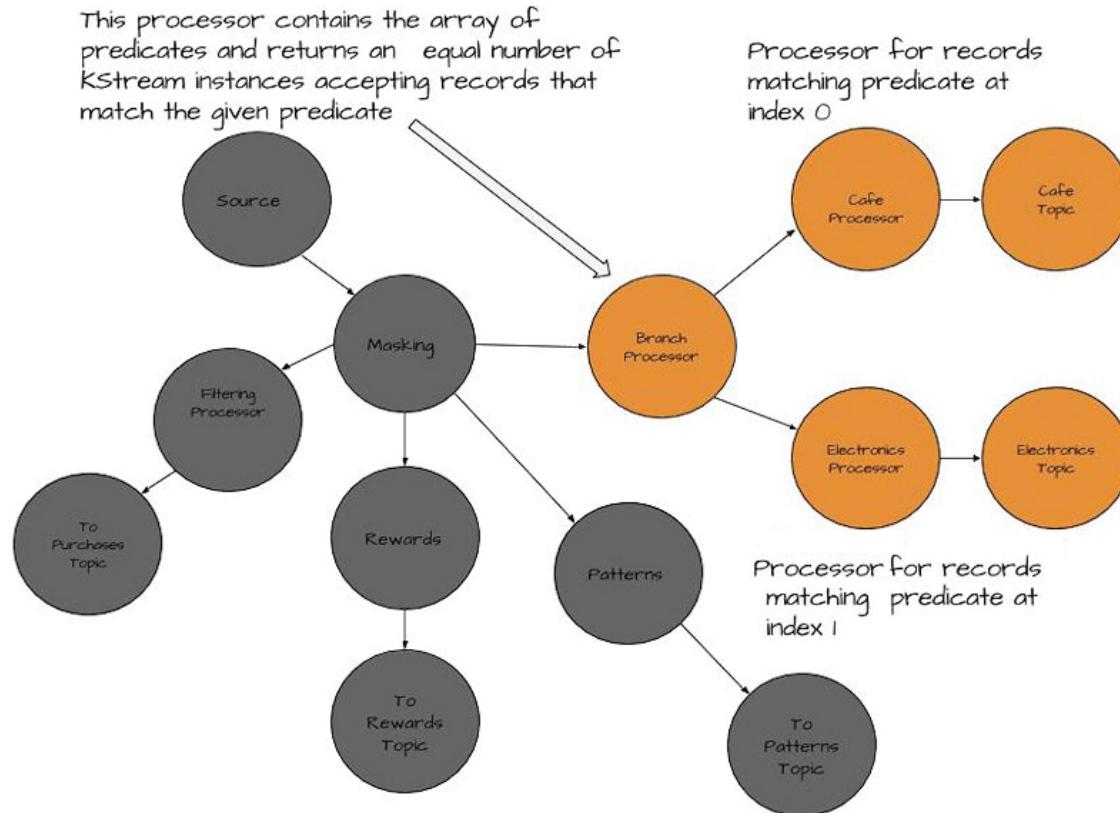
### **4.4 Joining Streams for Added Insight**

Remember earlier in the chapter we discussed the concept of streams needing state when events in the stream do not stand alone. In this section, we are going to learn how we can take different events from two streams with the same key and combine them to form a new event. The reason we are going to do this is sometimes the state or context we need is another stream.

The best way for us to learn about joining streams is a concrete example. For this

particular example, we will return to our ZMart scenario. If you recall, ZMart opened a new line of stores that carried electronics and related merchandise (CD's, DVD's and smartphones, etc.).

Trying a new approach, ZMart has partnered with a national coffee house and has embedded a cafe in each store. If you recall, in Chapter 3 you were asked to branch out the purchase transactions in those stores into two distinct streams. Here's how the topology looked for this requirement:

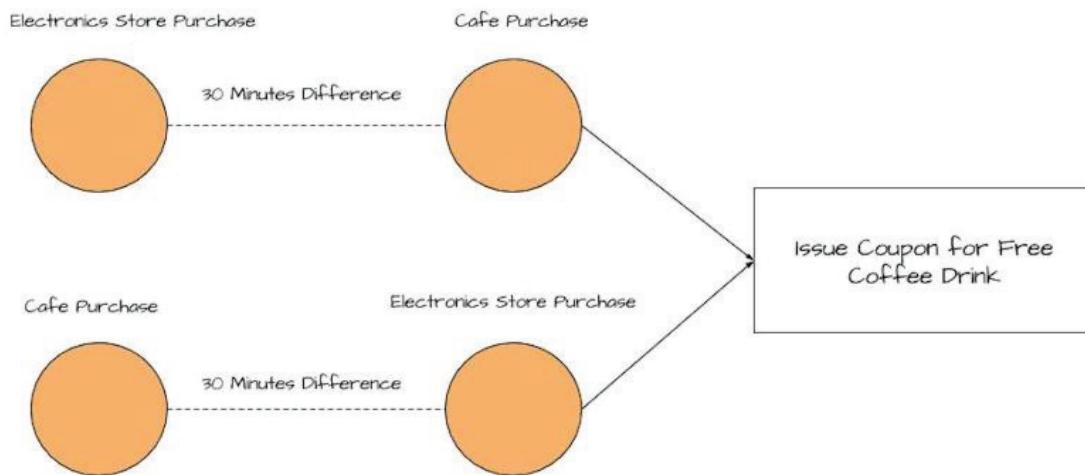


**Figure 4.12** Here's another look at the branch processor and where it stands in the overall topology.

Initially, this approach of embedding cafe has been a big success for ZMart. Naturally, ZMart would like to see this trend continue, so they have decided to start a new program. They want to keep traffic on the electronics store high (hoping that increased traffic leads to additional purchases) by offering coupons to the cafe.

ZMart wants to give coupons to customers that have bought coffee and made a purchase in the electronics store almost immediately after the second transaction. ZMart intends to see if they can generate some sort Pavlovian response to buying in their new store.

### Determining Free Coffee Coupons

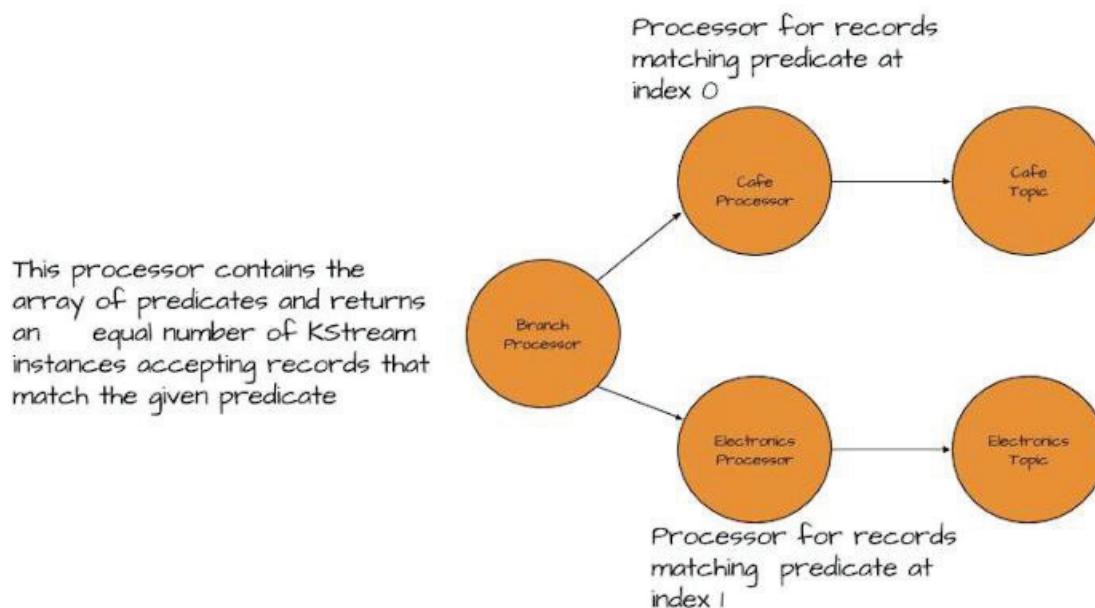


**Figure 4.13** Purchase records with timestamps within 20 minutes of each other are joined by customer id and used to issue a reward to the customer, a free coffee drink in this case.

To make the determination to issue a coupon, we will need to join the sales from the cafe against the sales in the electronics store. Joining streams is relatively straight forward as far as the code you will need to write. Let's get starting by setting up the data we'll need to process for doing joins.

#### 4.4.1 Data Setup

First off, let's take another view of the relevant portion of the topology responsible for branching the streams:



**Figure 4.14** In order to perform a join, we need more than one stream. The branch processor takes care of this for us by creating two streams; one stream contains cafe purchases and the other electronics purchases.

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

In addition, let's review the code used to implement the branching requirement:

#### **Listing 4.10 Branching Into Two Streams found in src/main/java/bbejeck/chapter\_3/ZMartKafkaStreamsAdvancedReqsApp.java**

```
Predicate<String, Purchase> coffeePurchase = (key, purchase) ->
    purchase.getDepartment().equalsIgnoreCase("coffee");
Predicate<String, Purchase> electronicPurchase = (key, purchase) ->
    purchase.getDepartment().equalsIgnoreCase("electronics"); ①

final int COFFEE_PURCHASE = 0; ②
final int ELECTRONICS_PURCHASE = 1;

KStream<String, Purchase>[] branchedTransactions =
    transactionStream.branch(coffeePurchase, electronicPurchase); ③
```

- ① Defining the predicates for matching records.
- ② Using labeled integers for clarity when accessing the corresponding array.
- ③ Creating the branched stream.

A quick review of the code above shows how to perform branching; you use predicates to match incoming records into an array of `KStream` instances. The order of matching is the position in the array of `KStream` objects. The branching process drops any record not matching any predicate.

Although we have our two streams to join, we still have an additional step to perform. Remember purchase records come into our Kafka Streams application with no keys. As a result, we'll need to add another processor to generate a key containing the customer-id. We need to have populated keys because it's what we use to join records together.

#### **4.4.2 Generating Keys of Customer ID To Perform Joins**

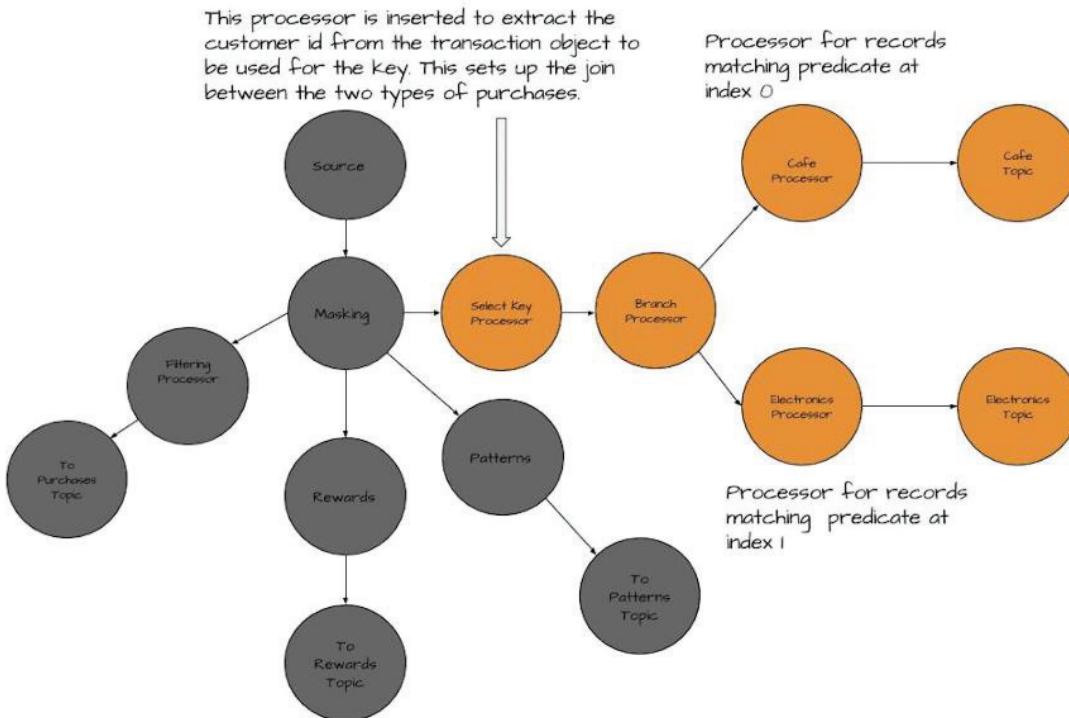
To generate a key, we'll select the customer-id from the purchase data in the stream. With this in mind we'll need to update the original `KStream` instance (`transactionStream`) and create another processing node between it and the branch node:

#### **Listing 4.11 Generating New Keys found in src/main/java/bbejeck/chapter\_4/KafkaStreamsJoinsApp.java**

```
KStream<String, Purchase>[] branchesStream = transactionStream.selectKey((k, v) ->
    v.getCustomerId()).branch(coffeePurchase, electronicPurchase); ①
```

- ① Inserting the "selectKey" processing node.

As a result of the code implemented above here's an updated view of the processing topology:



**Figure 4.15 We need to re-map the key-value purchase records into ones where the key contains the customer id. Fortunately, we can extract the customer id from the Purchase object.**

We've seen before that changing the key may require us to repartition the data. That is true in this example as well, so why didn't we see a repartitioning step?

In Kafka Streams whenever invoking a method that could result in generating a new key (`selectKey`, `map` or `transform`) an internal boolean flag is set to `true` indicating the new `KStream` instance requires repartitioning. With this boolean flag set if you perform a `join`, `reduce` or `aggregation` operation the repartitioning is handled for you automatically.

In this example you performed a `selectKey` operation on the `transactionStream`, meaning the resulting `KStream` is flagged for repartitioning. Additionally, you immediately performed a branching operation meaning each `KStream` resulting from the `branch` call is flagged for repartitioning as well.

#### NOTE

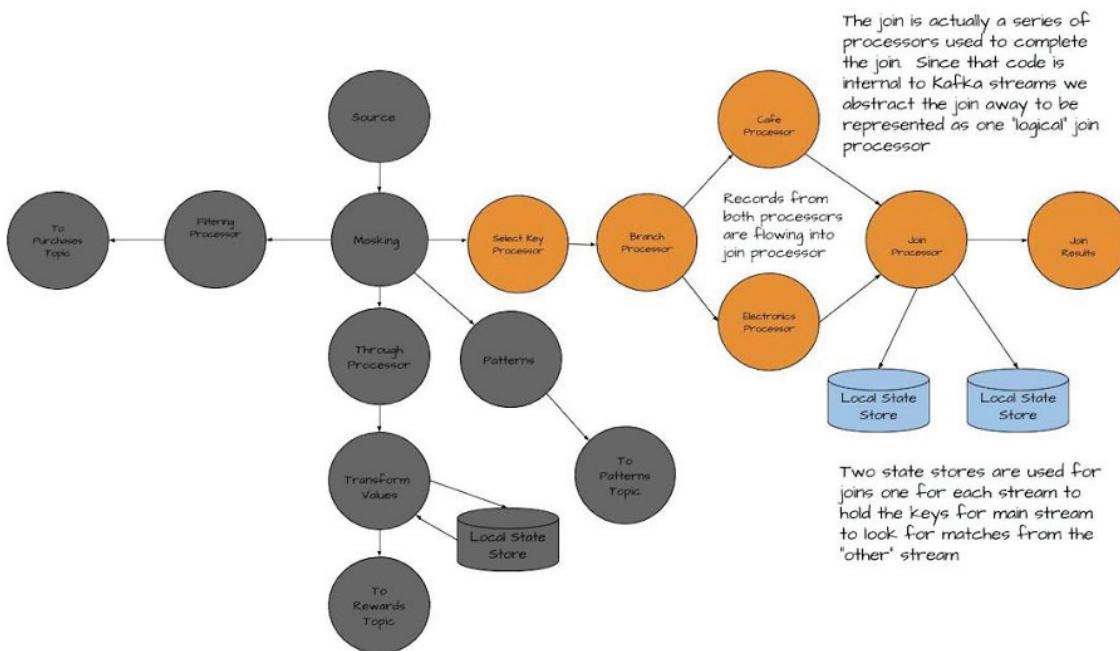
#### Finer Grained Control of Partition Assignment

In the current example we wanted to repartition by the key only. But there may be instances where you either don't want to use the key or you want to use some combination of the key and value. In these cases we can use the `StreamPartitioner<K, V>` interface as we saw in the example in section 4.2.4.

Now that we have two separate streams with populated keys we are ready for our next step: joining the streams by key.

#### 4.4.3 Constructing the Join

The next step is to perform the actual join. We'll take our two branched streams and join them by using the `kStream.join` method. Our topology will look like this:



**Figure 4.16** Here's how the updated topology looks now. Both the cafe and electronics processors forward their records to the join processor. The join processor uses two state stores to search for matches on a corresponding record from the "other" stream.

#### JOINING PURCHASE RECORDS

First off, to create the joined record, we need to create an instance of a `ValueJoiner<V1, V2, R>`. The `ValueJoiner` takes two objects, which may or may not be of the same type, and returns a single object, possibly of a different type. In this case, our `ValueJoiner` takes two `Purchase` objects and returns a `CorrelatedPurchase` object. Let's take a look at the code:

**Listing 4.12 ValueJoiner Implementation found in  
src/main/java/bbejeck/chapter\_4/joiner/PurchaseJoiner.java**

```
public class PurchaseJoiner implements ValueJoiner<Purchase, Purchase, CorrelatedPurchase> {

    @Override
    public CorrelatedPurchase apply(Purchase purchase, Purchase purchase2) {
        CorrelatedPurchase.Builder builder = CorrelatedPurchase.newBuilder();
        1
        Date purchaseDate = purchase != null ? purchase.getPurchaseDate() : null;
        Double price = purchase != null ? purchase.getPrice() : 0.0;
        2
        String itemPurchased = purchase != null ? purchase.getItemPurchased() : null;
        Date otherPurchaseDate = otherPurchase != null ? otherPurchase.getPurchaseDate() :
            null;
        Double otherPrice = otherPurchase != null ? otherPurchase.getPrice() : 0.0;
        String otherItemPurchased = otherPurchase != null ? otherPurchase.getItemPurchased() :
            null;
        3
    }
}
```

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

```

List<String> purchasedItems = new ArrayList<>();

if (itemPurchased != null) {
    purchasedItems.add(itemPurchased);
}

if (otherItemPurchased != null) {
    purchasedItems.add(otherItemPurchased);
}

String customerId = purchase != null ? purchase.getCustomerId() : null;
String otherCustomerId = otherPurchase != null ? otherPurchase.getCustomerId() : null;

builder.withCustomerId(customerId != null ? customerId : otherCustomerId)
    .withFirstPurchaseDate(purchaseDate)
    .withSecondPurchaseDate(otherPurchaseDate)
    .withItemsPurchased(purchasedItems)
    .withTotalAmount(price + otherPrice);

④
return builder.build();
}
}

```

- ① Instantiating the builder.
- ② Handling for null Purchase in the case of an outer join.
- ③ Handling for null Purchase in the case of a left outer join.
- ④ Returning the new CorrelatedPurchase object.

To create our `CorrelatedPurchase` object, we extract some information from each `Purchase` object. Because of the number of items we need to construct our new object, we use the builder pattern which makes the code clearer and eliminates any errors due to misplaced parameters. Additionally, the `PurchaseJoiner` checks for null values with both of the provided `Purchase` objects, hence it can be used for inner, outer and left-outer joins. I'll discuss the different join options in an upcoming section. Next, we'll move on to implementing the join between streams.

## IMPLEMENTING THE JOIN

Now you have seen how to merge records resulting from the join between the streams, let's move on to calling the actual `KStream.join` method.

### **Listing 4.13 Using the Join Method found in src/main/java/bbejeck/chapter\_4/KafkaStreamsJoinsApp.java**

```

① KStream<String, Purchase> coffeeStream = branchesStream[COFFEE_PURCHASE];
② KStream<String, Purchase> electronicsStream = branchesStream[ELECTRONICS_PURCHASE];
ValueJoiner<Purchase, Purchase, CorrelatedPurchase> purchaseJoiner = new PurchaseJoiner();
JoinWindows twentyMinuteWindow = JoinWindows.of(60 * 1000 * 20);

③ KStream<String, CorrelatedPurchase> joinedKStream = coffeeStream.join(electronicsStream,
    purchaseJoiner,
    twentyMinuteWindow,
    Joined.with(stringSerde,
    purchaseSerde,

```

```
purchaseSerde));
```

4

```
joinedKStream.print("joinedStream");
```

5

- ➊ Extracting the branched streams
- ➋ The ValueJoiner instance used to perform the join
- ➌ Calling the join method triggering automatic repartitioning of coffeeStream and electronicsStream
- ➍ Constructing the Join.
- ➎ Printing the join results to the console

Since we are supplying six parameters to the `kStream.join` method, let's quickly review what they are:

1. `electronicsStream` - the stream of electronic purchases to join with
2. `purchaseJoiner` - an implementation of the '`ValueJoiner<V1, V2, R>`' interface. The `ValueJoiner` accepts two values (not necessarily of the same type). The `ValueJoiner.apply` method performs the implementation specific logic and returns a (possibly new) object of type `R` (maybe a whole new type). In our example, the `PurchaseJoiner` will add some relevant information from both `Purchase` objects and return a `CorrelatedPurchase` object.
3. `twentyMinuteWindow` - A `JoinWindows` instance. The `JoinWindows.of` method specifies a maximum time difference between the two values to be included in the join. In this case, the timestamps are within 20 minutes of each other.
4. The final parameter is a `Joined` instance that provides optional parameters for performing joins. In this case it's the key and the value serde for the 'calling' stream and the value serde for the secondary stream. We only have one key serde because when joining records, keys must be of the same type.

#### NOTE

#### Serdess Required For Joins

Serdess are required for joins because join participants are materialized in windowed state stores. We only provide one serde for the key because both sides of the join must have a key of the same type.

Although we specified the purchases needed to be within 20 minutes of each other, there is no order implied. As long as the timestamps are within 20 minutes of each other, the join will occur.

There are two additional `JoinWindows` methods available where you can specify the order of events.

- `JoinWindows.after` -  
`streamA.join(streamB, ..., twentyMinuteWindow.after(5000), ...)` Here you specify the timestamp of the `streamB` record is at most 5 seconds *after* the timestamp of the `streamA` record. The starting time boundary of the window is unchanged.
- `JoinWindows.before` -  
`streamA.join(streamB, ..., twentyMinuteWindow.before(5000), ...)` Here you specify the timestamp of the `streamB` record is at most 5 seconds *before* the timestamp

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

of the `streamA` record. The ending time boundary of the window is unchanged.

With both the `before` and `after` methods above the time difference is always expressed in milliseconds.

#### **NOTE**

#### **Determining the timestamps for the join**

In this example we are relying on the actual timestamps of the transaction, not any timestamps set by Kafka. In order to use the timestamps embedded in the transaction we're specifying a custom timestamp extractor by setting `StreamsConfig.DEFAULT_TIMESTAMP_EXTRACTOR_CLASS_CONFIG` to use `TransactionTimestampExtractor.class`

The time windows used for the join are an example of "sliding" windows, and we'll cover windowing operations in detail in the next chapter.

Now we have constructed our joined stream; purchases made within 20 minutes of a coffee purchase will result in a coupon for a free drink on their next visit to ZMart.

Before we go any further, I would like to take a minute to explain an important requirement for joining data, co-partitioning.

### **CO-PARTITIONING**

In order to perform a join in KafkaStreams, you need to ensure all join participants are co-partitioned, meaning having the same number of partitions and keyed by the same type.

#### **NOTE**

`GlobalKTable` instances do not require repartitioning when involved in a join

As a result, when calling the `join` method above both `KStream` instances are checked to see if a repartition is required. Remember in section 4.2 you used the `selectKey` method on the `transactionStream` and immediately branched on the returned `KStream`.

As a result both the `coffeeStream` and `electronicsStream` require repartitioning. The repartitioning is handled automatically for you and is required because we need to ensure identical keys are written to the same partition. Additionally, when you start your Kafka Streams application topics involved in a join are checked for the same number of partitions and if any mismatches are found a `TopologyBuilderException` is thrown. However, it is the developer's responsibility to ensure the keys involved in a join are of the same type.

Co-partitioning also requires all Kafka Producers use the same partitioning class when writing to Kafka Streams source topics. Likewise you need to use the same

`StreamPartitioner` for any operations writing the same `KStream.to` operator. If you stick with the default partitioning strategies, then you don't need to worry about partitioning strategies.

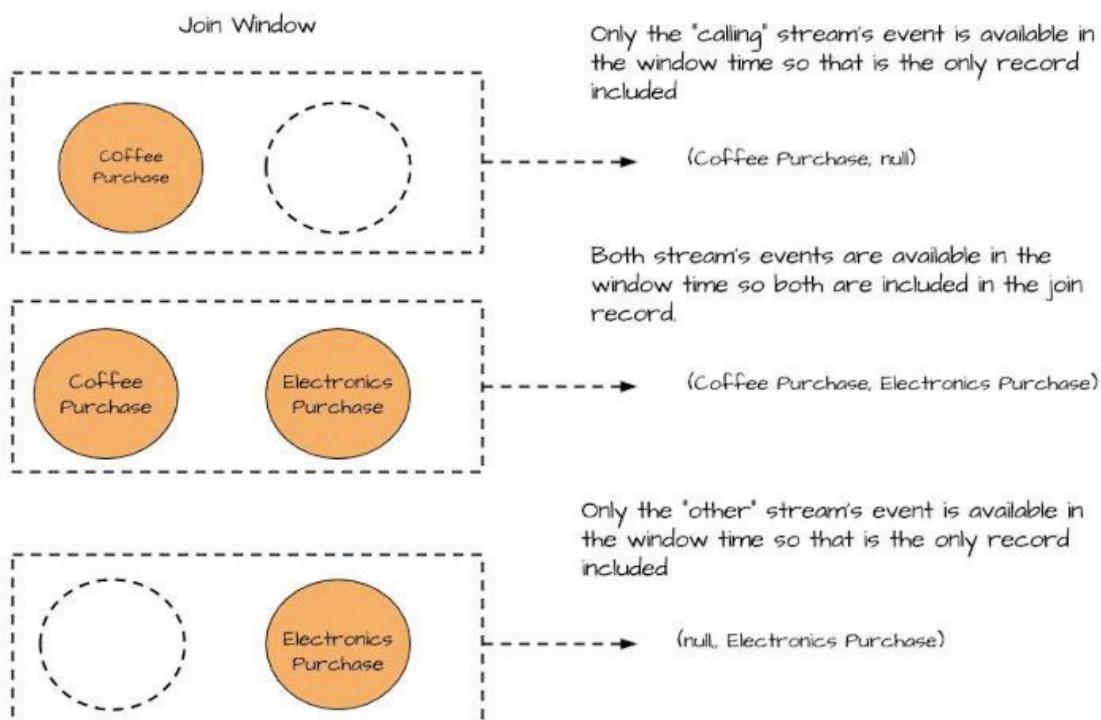
Next, let's continue on with joins by looking at other options available to you.

#### 4.4.4 Other Join Options

The example join presented here is a case of an 'inner join'. With an inner join, if either record is not present the join doesn't occur, and we don't emit a `CorrelatedPurchase` object. However, you have other options that don't require both records. These other join options are necessary if you have use cases where you need information even if the desired record for joining is not available.

#### OUTER JOINS

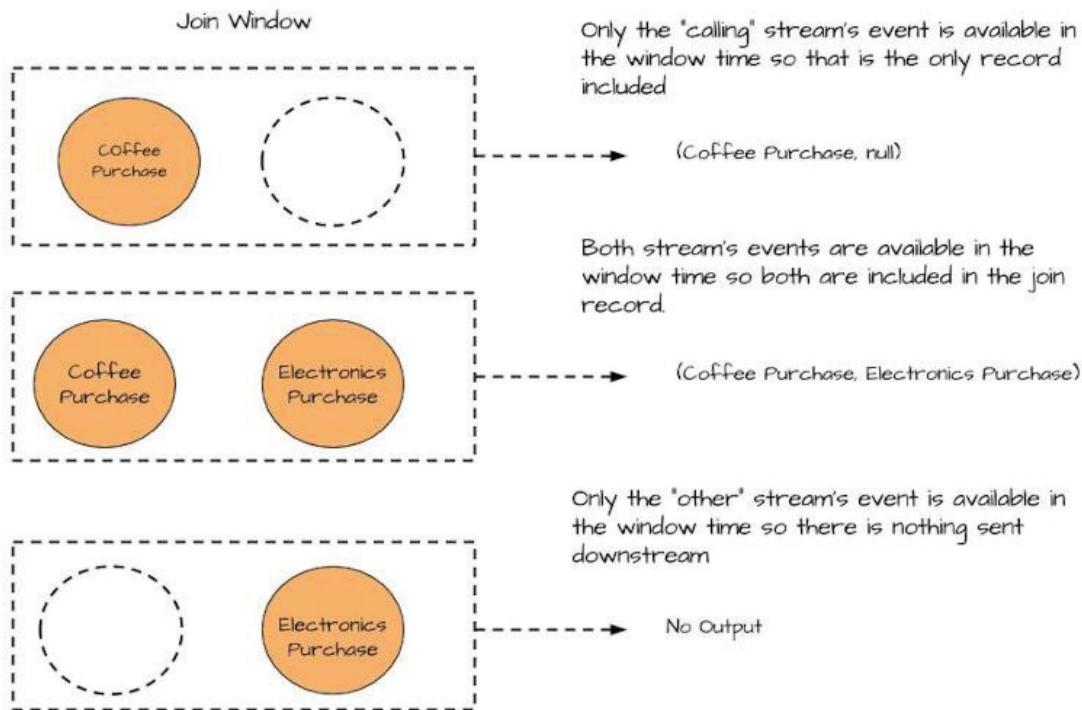
Outer-joins always output a record, but the forwarded join record may not include both of the events specified by the join. If either side of the join is not present once the join window expires, we send the record that is available downstream. Of course if both events are present within the window then the issued record contains both events. If you wanted to use an outer join in the example above you would do so like this `coffeeStream.outerJoin(electronicsStream..)`. The following image demonstrates the three possible outcomes of the outer join:



**Figure 4.17 Three outcomes are possible with outer joins, calling stream event only, both events, "other" stream event only.**

## LEFT OUTER JOIN

The records sent downstream from a left-outer join are similar to an outer-join with one exception. When the only event available in the join window is from the "other" stream, there is no output at all. If you wanted to use a left-outer join in the example above you would do so like this `coffeeStream.leftJoin(electronicsStream..)`. The graphic below shows the outcomes of the left-outer join



**Figure 4.18 Three outcomes are possible with the left outer join, but if only the "other" stream's record is available there is no output**

In this section, we have covered joining streams, but I have casually mentioned a concept that deserves a more detailed discussion. I'm referring to timestamps and the impact they have on your Kafka Streams application.

In our join example, we specified a max time difference between events of 20 minutes. In this case, it's the difference between purchases, but if you review the code for our example, how we set or extract these timestamps is not specified. With this in mind, let's move on to our next section.

## 4.5 Timestamps In Kafka Streams

In Chapter 2, section 2.5.4 we discussed timestamps in Kafka records. In this section, we'll discuss the use of timestamps specific to their role in Kafka Streams.

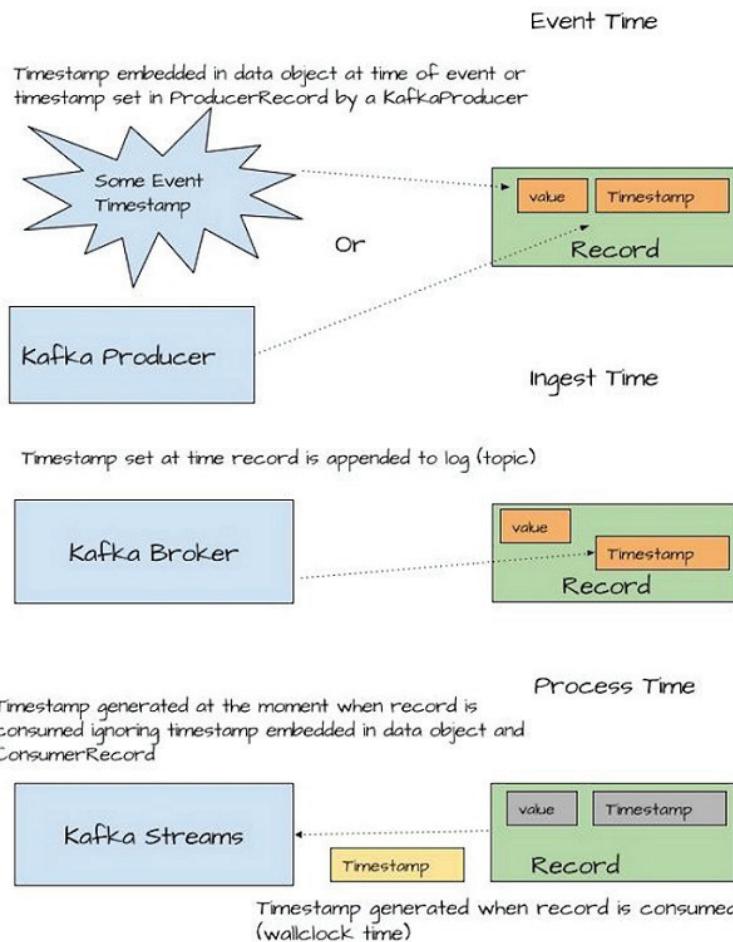
We need to cover timestamps because they come into play in key areas of Kafka Streams functionality such as:

1. The joining of streams.
2. Making an update to a changelog (KTable API).

3. Helps decide when the `Processor.punctuate` method is triggered (Processor API).

I know we haven't covered the KTable or Processor API at this point yet, but that's ok. We don't need them to understand this section.

In stream processing, in general, we can place timestamps into three categories summarized by this image:



**Figure 4.19 There are three categories of timestamps in Kafka Streams. Event time is a timestamp embedded in the message metadata by the producer when sending the message. Ingest time is the timestamp set by the broker when appending the message to the log (topic) and process time is the timestamp when Kafka Streams processes the record.**

Let's describe each category of the timestamp in a little more detail:

1. Event Time - timestamp when the event occurred. Usually, a timestamp embedded in the object used to represent the event. For our purposes, we will consider the timestamp set when the `ProducerRecord` is created as event-time as well.
2. Ingestion Time - ingestion time is a timestamp set when the data first enters into the data processing pipeline. Here we can consider the timestamp set by the Kafka broker (assuming configuration setting of `LogAppendTime`) at ingestion time.
3. Processing Time - processing time is a timestamp set when the data or event record is first started to flow through a processing pipeline.

**NOTE****Implications of time zones and timestamps**

So far we've had an implicit assumption that clients and brokers are located in the same timezone. But that might not always be the case. When using timestamps it's a safe bet to normalize the times to use the UTC timezone and thus eliminate any confusion if using broker set timestamps with a client in another timezone.

We'll see in this section how the Kafka Streams API supports all three types of processing timestamps.

We are going to consider three cases of timestamp processing semantics:

1. An embedded timestamp in the actual event/message object (event-time semantics).
2. Using the timestamp set in the record metadata when creating the `ProducerRecord` (event-time semantics).
3. Using the current timestamp (current local time) when the Kafka Streams application ingests the record (processing-time semantics).

For event-time semantics using the timestamp placed in the metadata by the `ProducerRecord` is sufficient. However, there may be cases when you have different needs. For example:

- You are sending messages to Kafka with events that have timestamps recorded in the message object itself. Additionally, there is some lag time in when these event objects are made available to the Kafka producer, so you want to consider only the embedded timestamp.
- You want to consider the time when your Kafka Streams application consumes records vs. the timestamps of the record itself.

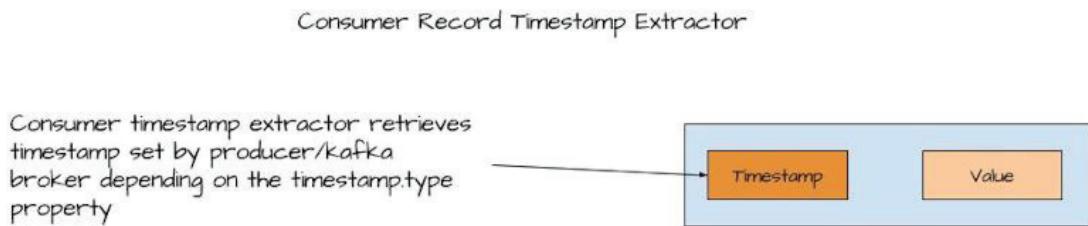
To enable different processing semantics, Kafka Stream provides a `TimestampExtractor` interface, with one abstract and four concrete implementations provided.

If you need to work with timestamps embedded in the record value, you'll need to create a custom `TimestampExtractor` implementation.

Let's briefly describe the included implementations and provide an example of implementing a custom `TimestampExtractor`.

#### **4.5.1 Provided `TimestampExtractor` Implementations**

Almost all of the provided `TimestampExtractor` implementations work with timestamps set by the producer or broker in the message metadata, thus providing event-time processing semantics. The figure below demonstrates the pulling the timestamp from of the `ConsumerRecord` object:



**Figure 4.20** Here we are using the timestamp found in the `ConsumerRecord` object. Either the producer or broker set this timestamp depending on your configuration.

Although we are assuming the setting of `CreateTime` for the timestamp, bear in mind if you were to use the `LogAppendTime` this would return the timestamp value when the Kafka broker appended the record to the log.

The `ExtractRecordMetadataTimestamp` is an abstract class providing the core functionality of extracting the metadata timestamp from the `ConsumerRecord`. Most of the concrete implementations extend this class.

There is an abstract method `ExtractRecordMetadataTimestamp#onInvalidTimestamp` implementors override to handle invalid timestamps (timestamp is less than 0).

Here's a list of the classes extending the `ExtractRecordMetadataTimestamp` class:

- `FailOnInvalidTimestamp` throws an exception in the case of an invalid timestamp.
- `LogAndSkipOnInvalidTimestamp` returns the invalid timestamp and logs a warning message that the record will be discarded due to the invalid timestamp.
- `UsePreviousTimeOnInvalidTimestamp` in the case of an invalid timestamp, the last valid extracted timestamp is returned.

We've covered the event-time timestamps extractors, but there is one more provided timestamp extractor that we'll cover in the next section.

#### 4.5.2 `WallclockTimestampExtractor`

The `WallclockTimestampExtractor` provides process/ingest-time semantics and doesn't extract any timestamps. Instead it returns the time in milliseconds from calling the `System.currentTimeMillis()` method.

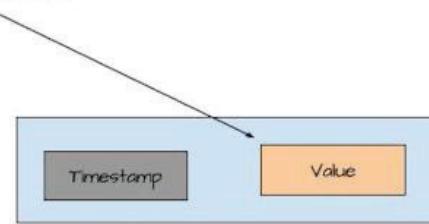
At this point we've covered working with the provided timestamp extractors, next we'll cover how to create a custom version.

#### 4.5.3 `Custom TimestampExtractor`

For working with timestamps (or calculating one) contained in the value object from the `ConsumerRecord`, you'll need a custom extractor implementing the `TimestampExtractor` interface. The following image depicts using the timestamp embedded in the value object vs. one set by Kafka (either producer or broker).

### Custom Timestamp Extractor

Custom timestamp extractor will pull out an embedded timestamp in the value object or perform some calculation to retrieve a timestamp. Kafka will allow null values so it's a good idea to have a fallback value (`System.currentTimeMillis()`).



**Figure 4.21** A custom `TimestampExtractor` provides a timestamp from the value contained in the `ConsumerRecord`. This timestamp could be an existing value or one calculated from properties contained in the value object.

Here's an example of a `TimestampExtractor` implementation, also used in our join example above.

**Listing 4.14 Custom `TimestampExtractor` found in  
`src/main/java/bbejeck/chapter_4/timestamp_extractor/TransactionTimestampExtractor.java`**

```
public class TransactionTimestampExtractor implements TimestampExtractor {

    @Override
    public long extract(ConsumerRecord<Object, Object> record, long previousTimestamp) {
        Purchase purchaseTransaction = (Purchase) record.value(); ①
        return purchaseTransaction.getPurchaseDate().getTime(); ②
    }
}
```

- ① Retrieving the `Purchase` object from the key-value pair sent to Kafka.
- ② Returning the timestamp recorded at the point of sale.

In our example, we chose to use a custom `TimestampExtractor` as we want to use the timestamps of the actual purchase time. This approach will allow us to join our records even with delays in delivery or out of order arrivals.

**WARNING Clever `TimestampExtractor` Implementations**

Take care when creating a custom `TimestampExtractor` to not get too 'clever'. Log retention / log rolling is now timestamp based and the timestamp provided by the extractor may become the message timestamp used by changelogs and output topics downstream.

#### 4.5.4 Specifying a `TimestampExtractor`

Now that we've discussed how timestamp extractors work, let's cover how we tell our application which one to use. We have two choices for specifying timestamp extractors.

The first choice is to set a "global" timestamp extractor specified in the properties when setting up your Kafka Streams application. If no property is set the default setting is `FailOnInvalidTimestamp.class`.

For example let's show how we'd configure the `TransactionTimestampExtractor` via properties when setting up the application.

### **Listing 4.15 Configuring A TimestampExtractor**

```
props.put(StreamsConfig.DEFAULT_TIMESTAMP_EXTRACTOR_CLASS_CONFIG,
         TransactionTimestampExtractor.class);
```

The second option is to provide a `TimestampExtractor` instance via a `Consumed` object, the advantage of which is to have a `TimestampExtractor` per input source, where before you would need handle records from different topics in one `TimestampExtractor` instance.

We've come to the end of the discussion on timestamp usage. In the coming chapters, we'll run into situations where the difference between timestamps drives some action, flushing the cache of a `KTable` for instance. While I don't expect you remember all three types of timestamp extractors, the key point here is timestamps are an important part of how Kafka/Kafka Streams function.

## **4.6 Conclusion**

We've covered a lot of ground in this chapter. This section could be considered the first half of the core concepts of Kafka Streams. Although we touched on several topics here are the key takeaways:

- Stream processing needs state. Sometimes events can stand on their own, but most of the time we need additional info to make good decisions.
- Kafka Streams provides useful abstractions for stateful transformations including joins.
- StateStores in Kafka Streams provide the type of state required for stream processing: data locality and fault tolerance.
- Timestamps control the flow of data in Kafka Streams. The choice of timestamp source(s) needs careful consideration.

Next chapter we'll continue exploring state in streams with more significant operations like aggregations and grouping. We'll also explore the `KTable` API. Where the `KStream` API concerns itself with individual discrete records, a `KTable` is an implementation of a changelog where records with the same key are considered updates. We'll also discuss joins between `KStream` and `KTable` instances. Finally, we'll explore one of the most exciting developments in Kafka Streams, Queryable State. Queryable state allows us to directly observe the state of our stream, without having to materialize the information by reading data from a topic in an external application.

# The KTable API

In this chapter:

- Relationship Between Streams and Tables and how Streams are Tables.
- Updates to records and the KTable abstraction.
- Aggregations and Windowing Operations
- Joining KStreams and KTables
- Global KTables
- Queryable State Stores

So far we have covered the KStream API and how to add state to our Kafka Streams application. In this chapter, we're going to look deeper into adding state. Along the way, we'll get introduced to new abstraction, the KTable.

When covering the KStream API, we are talking about individual events or an event stream. If we go back to the original ZMart example, when Jane Doe made a purchase, we considered the purchase as an individual event. We didn't keep track of how often or how many purchases Jane carried out before.

If you think about the purchase event stream with respect to a database, it can be considered a series of inserts. Since each record is new and unrelated to any other records, we can continually insert them into a table.

Now let's add a primary key (`customer-id`) to each purchase event. Going back to Jane Doe, we now have a series of *related* purchase events or updates. Since we are using a primary key each purchase is now update with respect to Jane's purchase activity. Treating an event stream as inserts and events with keys as updates is what we'll use to define the relationship between streams and tables.

Now let's discuss what we're going to cover in this chapter. First, we'll cover more in-depth the relationship between streams and tables. This relationship is important as it helps your understanding how the KTable operates.

Second, you'll learn about the KTable. The KTable API is necessary because it's designed to work with updates to records. We need the ability to work with updates for

operations like aggregations and counts. We touched on updates in Chapter 4 when we introduced stateful transformations. In section 4.2.5 of Chapter 4, you updated the Rewards Processor to keep track of customer purchases.

Third, we'll get into windowing operations. Windowing is the process of bucketing data within a given period. For example, how many purchases have been over the past hour, updated every ten minutes. Windowing allows us to gather data in chunks vs. having an unbounded collection.

**NOTE**
**What's Windowing and Bucketing?**

Windowing and bucketing are somewhat synonymous terms. Both operate by placing information into smaller chunks or categories. While windowing implies categorizing by time, the result of either operation is the same.

Our final concept in this chapter is queryable state stores. Queryable state stores are an exciting feature of Kafka Streams. I say exciting because they allow you to run direct queries against state stores. In other words, you can view stateful data without having to consume it from a Kafka topic or read it from a database.

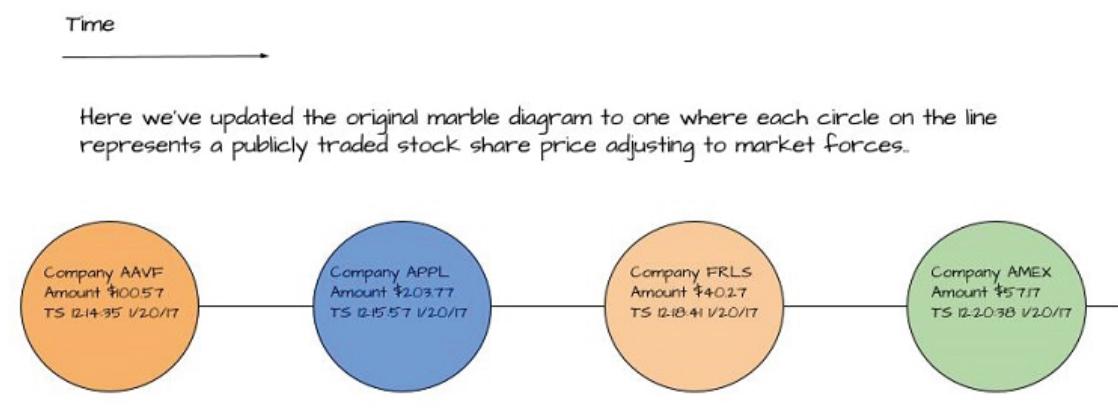
With these points in mind let's move on to our first topic.

## 5.1 The Relationship Between Streams and Tables

If you remember in Chapter one, I defined a stream as an infinite sequence of events. But stating 'infinite' sequence of events' is generic, so let's narrow it down with a specific example.

### 5.1.1 The Record Stream

Let's say we want to view a series of stock price updates, so we'll recast our generic diagram to look like this:



Imagine here you are observing stock ticker displaying updated share prices in real-time

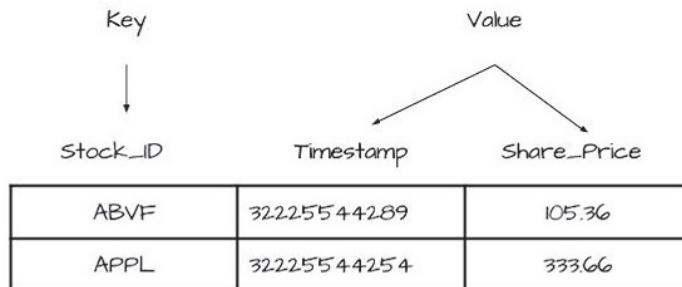
**Figure 5.1** Here we revised our simple marble diagram to be an unbounded stream of stock quotes.

We can see from this picture that each stock price quote is a discrete event and they

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

aren't related to each other. Even if the same company accounts for many price quotes, we are only looking at them one at a time.

This view of events is how the KStream works. The KStream is a stream of records. Now we can show how this concept ties into database tables by viewing this image of a simple stock quote table:



Rows from table above can be recast as key-value pairs. For example we've converted the first row from above to the key-value pair below

```
{key:{stockid:1235588}, value:{ts:32225544289, price:105.36}}
```

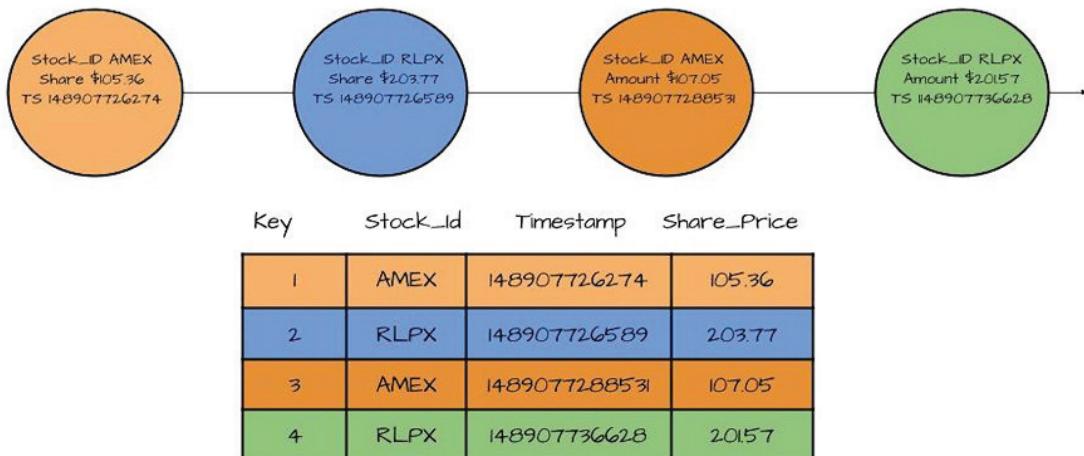
**Figure 5.2 A simple database table is representing stock prices for companies. There is a key column, and other columns are containing the value. But it is fair to say we have a key-value pair if we lump the other columns into a "value" container.**

**NOTE**

**Simplifying Assumptions**

To keep our discussion straight forward we are going to assume a key to be a singular value.

Now let's take another look at the record stream. Since each record stands on its own, the stream represents inserts into a table. In the following image we've combined the two concepts to illustrate this point:



Here we're showing the relationship between events and inserts into a database. Even though it's stock prices for two companies, it counts as 4 events total because we are only considering each item on the stream as a singular event.

As result each event is an insert and we increment the key by one for each insert into the table.

So with that in mind, each event is a new, independent record or insert into a database table.

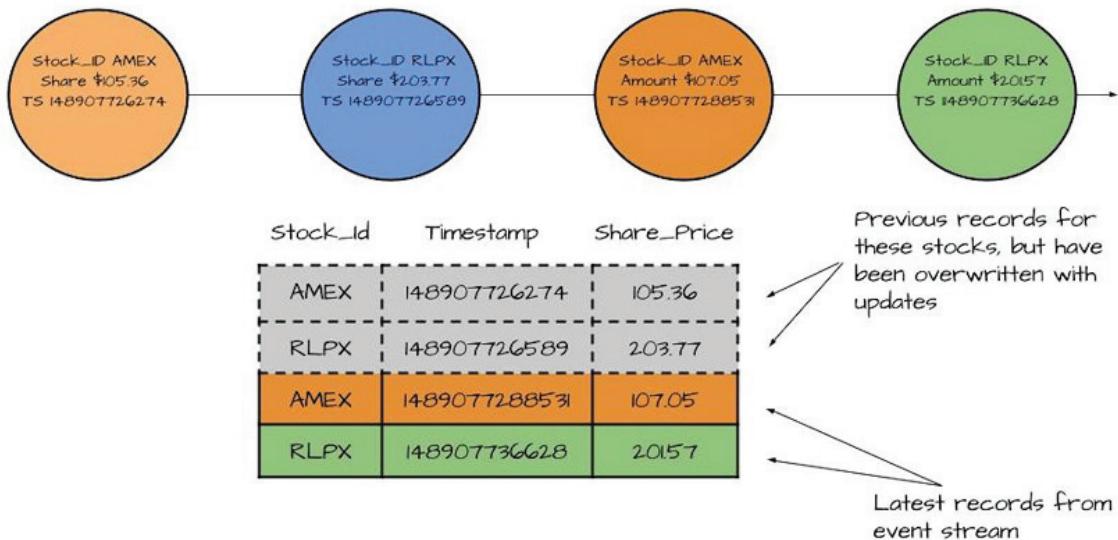
**Figure 5.3 Here we can see how a stream of individual events compares to inserts in a database table. We can flip this around and look at how we could stream each row of the table.**

What's important here is that we can view a stream of events in the same light as inserts into a table, which helps give us a deeper understanding of using streams for working with events.

The next step in our comparison of streams to tables is for us to consider the case when the events in the stream are related to one another.

### 5.1.2 Updates To Records or The Change Log

Let's take the same stream of customer transactions, but now we want to track activity over time. If we add a key of customer id, the purchase events can be related to each other and we now have an update stream as opposed to an event stream. If we consider our stream of events as a log, we can consider this stream of updates as a change log. The image below demonstrates this concept:



Here's how can view a change log stream. Now by taking the Stock ID into account as a primary key the subsequent events with the same key are updates.

In this case we only have 2 records, one per company. Although more records can arrive for the same company as share prices change, the records won't accumulate.

**Figure 5.4 Each record here does not stand on its own. With a record stream, we would have a total count of four events. In the case of updates or a change log, we only have two. Each incoming record overwrites the previous one with the same key.**

Here we can see the relationship between a stream of updates and a database table. Both the log and change log represent incoming records appended to the end of a file. In a log, we need to see all the records. But with the change log, we only want to keep the latest record for any given key.

#### SIDE BAR

#### Log vs. Change Log

With both the log and change log, as records come in they are always appended at the end of the file. The distinction between the two is in a log you want to read *all* records, but in change log, you only want the *latest* record for each key.

Now we need a way to trim down the log while maintaining those latest arriving records (per key). The good news is we have just such an approach. In chapter two we discussed the concept of log compaction. Let's take another look at the image for review:

Before Compaction

offset	key	value
10	foo	A
11	bar	B
12	baz	C
13	foo	D
14	baz	E
15	boo	F
16	foo	G
17	baz	H

After Compaction

offset	key	value
11	bar	B
15	boo	F
16	foo	G
17	baz	H

**Figure 5.5 On the left is a log before compaction. You'll notice duplicate keys with different values which are updates. On the right is the log after compaction. Notice we keep the latest value for each key, but the log is smaller in size.**

We can see the impact of the compacted log in this image. Since we only care about the latest values, we can remove older key-value pairs <sup>12</sup>.

---

Footnote 12 This section derived information from Jay Kreps in  
[www.confluent.io/blog/introducing-kafka-streams-stream-processing-made-simple/](http://www.confluent.io/blog/introducing-kafka-streams-stream-processing-made-simple/) and  
[log-what-every-software-engineer-should-know-about-real-time-datas-unifying](http://log-what-every-software-engineer-should-know-about-real-time-datas-unifying)

We're already familiar with event streams from working with the KStream. For a change log or stream of updates, we'll use an abstraction known as the KTable. Now that we've established the relationship between streams and tables, our next step is to compare an event stream to an update stream.

### 5.1.3 Event Streams vs Update Streams

In this section, we'll use the KStream and the KTable to drive our comparison of event stream vs. update stream. We'll do this by running a simple stock ticker application. The stock ticker application writes the current share price for three (fictitious!) companies. We'll produce three iterations of stock quotes for a total of 9 records.

A KStream and KTable instance reads the records and writes them to the console via the `print` method.

The following image shows the results of running the application:

Here we've done a simple stock-ticker for 3 fictitious publicly traded company. We use the data generator to produce 3 simulated updates for the stocks. You'll notice the KStream printed all records as they were received. However, the KTable only printed the last batch of records, as these were the latest "updates" for the given stock symbol.



**Figure 5.6 KTable vs KStream printing results from messages with the same keys**

As you can see from the results, the KStream printed all nine records out. We expect the KStream to behave this way as it views each record individually. In contrast, the KTable only printed three records, since the KTable views records as updates to previous ones.

#### NOTE

#### About Primary Keys

The example in the previous section demonstrated how a KTable works with updates. I made an implicit assumption that I'll make explicit here. When working with a KTable, your records must have populated keys in the key-value pairs. Having a key is essential for the KTable to work. As you can't update a record in a database table without having the key, the same rule applies here as well.

From the KTable point of view, it didn't receive nine individual records. The KTable received three original records and two rounds of updates, only printing the last round of updates. Notice the KTable records are the same as the last three records published by the KStream. We'll discuss the mechanisms of how the KTable emitted only the updates in the next section.

Here is our program printing stock ticker results to the console:

#### Listing 5.1 KTable and KStreams Printing to the Console found in src/main/java/bbejeck/chapter\_5/KStreamVsKTableExample.java

```
KTable<String, StockTickerData> stockTickerTable =
```

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

```

builder.table(STOCK_TICKER_TABLE_TOPIC);
KStream<String, StockTickerData> stockTickerStream =
    builder.stream(STOCK_TICKER_STREAM_TOPIC);

stockTickerTable.toStream().print(Printed.<String,
    StockTickerData>.toSysOut().withLabel("Stocks-KTable"));
stockTickerStream.print(Printed.<String,
    StockTickerData>.toSysOut().withLabel( "Stocks-KStream"));

```

- 1
- 2
- 3
- 4

- 1 Creating the KTable instance.
- 2 Creating the KStream instance.
- 3 KTable printing results to the console.
- 4 KStream printing results to the console.

**NOTE****Using Default Serdes**

I should point out here that in creating the KTable and KStream we did not specify any `Serdes` to use. The same is true with both calls to the `print` method. We were able to do this because we registered our default `Serdes` in the configuration like so:

```

props.put(StreamsConfig.DEFAULT_KEY_SERDE_CLASS_CONFIG,
    Serdes.String().getClass().getName());
props.put(StreamsConfig.DEFAULT_VALUE_SERDE_CLASS_CONFIG,
    StreamsSerdes.StockTickerSerde().getClass().getName());

```

If we used different types then we'd need to provide `Serdes` in any of the overloaded methods reading or writing records.

The takeaway here is records in a stream with the same key are updates, not new records by themselves. A stream of updates is the main concept behind the KTable.

Now we've seen the KTable in action, let's discuss the mechanisms behind the KTable functionality.

## 5.2 Record Updates and KTable Configuration

To explain how the KTable functions we should ask the following questions: where are records stored and how does the KTable make the determination to emit records? As we get into aggregation and reducing operations the answers to these questions are necessary.

For example, when performing an aggregation, you'll want to see updated counts. But you probably don't want an update each time the count increments by one.

To answer the first question let's take a look at the line creating the KTable:

```
builder.table(STOCK_TICKER_TABLE_TOPIC);
```

With this simple statement, the `streamsBuilder` creates a KTable instance and

simultaneously under the covers creates a `StateStore` for tracking the state of stream, thus creating an update stream. The `StateStore` created with this approach has an internal name and will not be available for interactive queries.

There is an overloaded version of `StreamsBuilder#table` that accepts a `Materialized` instance allowing you to customize the type of store and provide a name making it available for querying. We'll discuss interactive queries later in this chapter.

So we have our answer to the first question, the `KTable` uses the local state integrated with Kafka Streams for storage (We covered state stores in chapter 4, sections 4.3.3 through 4.3.10).

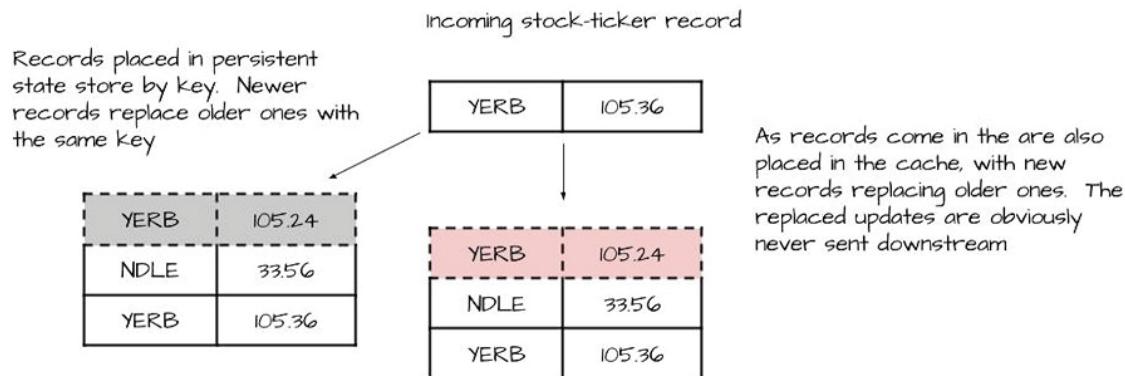
Now let's move on to what determines when the `KTable` emits updates to downstream processors. To answer this question, we need to consider a few factors:

1. The number of records flowing into the application. Higher data rates would tend to increase the rate of emitting updated records.
2. How many distinct keys you have in your data. Again the greater the number of distinct keys may lead to more updates sent downstream.
3. The configuration parameters "cache.max.bytes.buffering" and "commit.interval.ms"

From this list, we'll only cover what we can control, configuration parameters. First, let's take a look at the `cache.max.bytes.buffering` configuration.

### 5.2.1 Setting Cache Buffering Size

The `KTable` cache serves to deduplicate updates to records with the same key. This deduplication allows for receiving only the most recent update vs. receiving all updates, reducing the amount of data processed. The following image gives a visual description of the cache operation.



**Figure 5.7 KTable caching deduplicates updates to records with the same key, preventing a flood of consecutive updates to child nodes of the KTable in the topology.**

As we can see in the picture above, with caching enabled all record updates do not get forwarded downstream. The cache keeps only the latest record for any given key.

**NOTE****What does downstream mean?**

A Kafka Streams application is a topology or graph of connected nodes (processors). Any given node may have one or more child nodes. Unless it's terminal processor, once a processor is completed working with a record it forwards the record "downstream" to its respective child nodes.

Since the KTable represents a change log of events in a stream, we expect to work with the only latest update at any given point. Using the cache enforces this behavior. If we want to process all records in the stream, then we would use an event stream, the KStream, covered earlier.

All other factors being equal, a larger cache will reduce the amount of updates emitted. Additionally, caching reduces the amount of data written to disk by persistent stores (RocksDB) and if logging is enabled, the amount of records sent to the changelog topic for any store. Cache size is controlled by the "cache.max.bytes.buffering" setting and specifies the amount of memory allocated for the record cache.

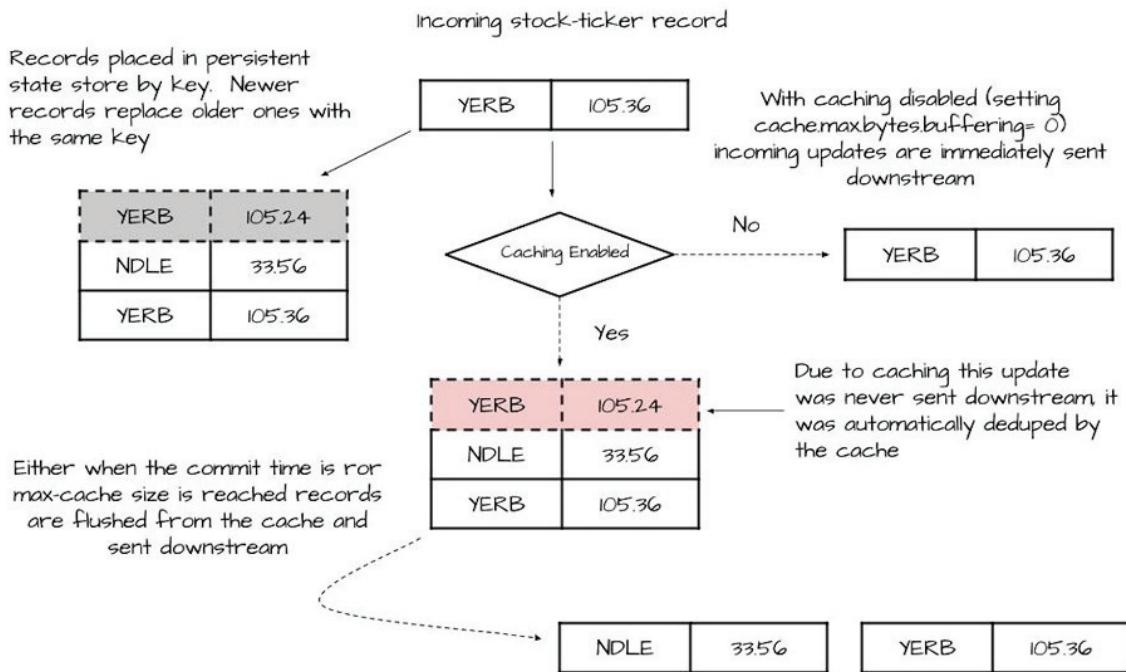
The amount of memory specified is divided evenly across the number of stream threads. (The number of stream threads is specified by the StreamsConfig.NUM\_STREAM\_THREADS\_CONFIG setting with 1 being the default).

**WARNING****Disabling the Cache**

To turn off caching one can set the "cache.max.bytes.buffering" setting to "0". But this setting will result in sending every KTable update downstream effectively turning our change log stream into an event stream. Also no caching means more IO as persistent stores will now write each update to disk vs. only the latest.

### **5.2.2 Setting the Commit Interval**

The other setting is the "commit.interval.ms" parameter. The commit interval determines how often (in milliseconds) to save the state of a processor. When saving the state of the processor (committing), forces a cache flush, sending the latest updated, deduplicated records downstream.



**Figure 5.8 Full caching workflow. If caching enabled, records are de-duped and sent downstream on cache-flush or commit**

In the full caching workflow above, we see two forces at play when it comes to sending records downstream. Either a commit or the cache reaching its max size, will send records downstream. Conversely, by disabling the cache, all records are sent downstream, including duplicate keys. Generally speaking, it's best to have caching enabled when using a KTable.

So you can see we need to strike a balance between cache size and commit time. A large cache with a small commit time will still result in frequent updates. But, a longer commit interval could lead to less updates (depending on memory settings) as cache evictions occur to free up space. There are no hard rules here, only trial and error will determine what works best for you. It's best to start with the default values of thirty seconds (commit time) and ten MB (cache size).

The key thing to remember is the rate of updated records sent from a KTable is configurable.

Next, let's take a look at how we can use the KTable in our applications.

### 5.3 Aggregations and Windowing Operations

In this section, we move on to cover some of the most potent parts of Kafka Streams. We can consider this section the "capstone" of your Kafka Streams experience. So far we've covered:

1. How to set up a processing topology
2. How to use state in our streaming application.
3. Performing joins between streams.
4. The difference between an event stream (KStream) and an update stream (KTable).

In the examples to follow we are going to tie all these elements together. Additionally, we are going to introduce Windowing, another powerful tool in streaming applications. In the following section, we'll tackle our first example, a straightforward aggregation.

### 5.3.1 Aggregating Share Volume by Industry

All our examples to this point have been in the retail sales arena. We are going to shift gears and draw our examples from the financial industry in this chapter.

We are going to take on the role of being a day trader for the next set of examples. In the first one, we'll track the share volume of companies across a list of selected industries. In particular, we are only interested in the top five (by share volume) in each industry.

Aggregation and grouping are necessary tools to use when working with streaming data. Many times reviewing single records as they arrive is not enough. To gain any information/insight you'll need grouping and combining of some sort.

To do this aggregation, we'll need a few steps to set the data up in the correct format. At a high level here are the steps we'll take:

1. Create a source from a topic publishing raw stock-trade information. We'll need to map a `StockTransaction` object into a `ShareVolume` object. The reason for performing this mapping step is simple. The `StockTransaction` object contains metadata about the trade, but we only want the volume of shares involved in the trade.
2. Then we'll group `ShareVolume` by its ticker symbol. Once grouped by symbol, we can reduce to a rolling total of share volume. We should note here calling `KStream.groupBy` returns a `KGroupedStream` instance. Then calling `KGroupedStream.reduce` is what gets us to our `KTable` instance.

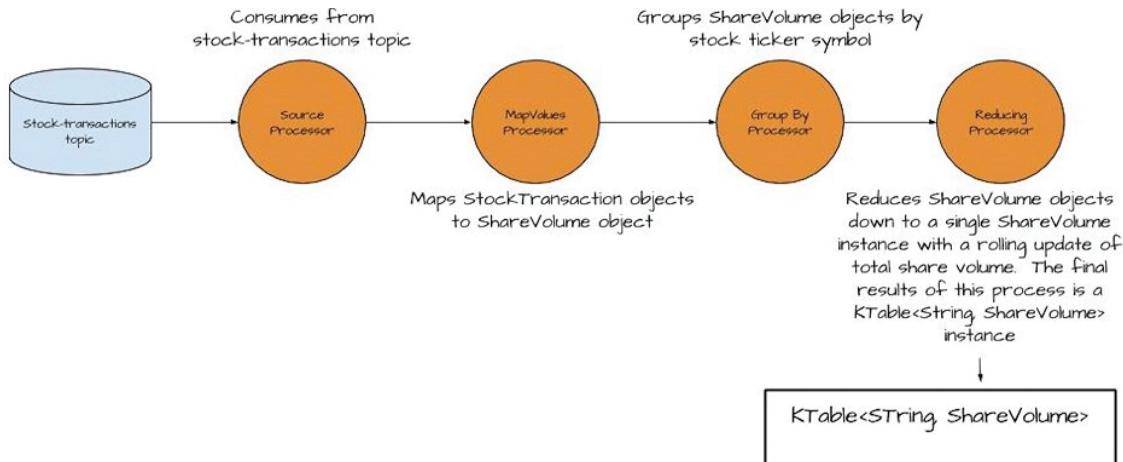
#### NOTE

#### What is the KGroupedStream

When you use `KStream.groupBy` or `KStream.groupByKey` the returned instance is a `KGroupedStream`. The `KGroupedStream` is an intermediate representation of the event stream after grouping by keys, and is never meant for you to work with directly. Instead, the `KGroupedStream` is required to perform aggregation type operations which always result in a `KTable`. Since the result of aggregate operations produce a `KTable` and uses a state store, all updates may not end up being forwarded downstream.

There is an analogous `KGroupedTable` resulting from the `KTable.groupBy` method which is the intermediate representation of the update stream regrouped by key.

Let's pause here for a minute and take a look at an image of what we're building so far:



**Figure 5.9 Mapping and Reducing StockTransaction objects into ShareVolume objects then reducing to rolling total.**

This topology is something we are familiar with by now. Now we'll look at the code behind the topology:

**Listing 5.2 Source for map-reduce of stock transactions found in  
src/main/java/bbejeck/chapter\_5/AggregationsAndReducingExample.java**

```
KTable<String, ShareVolume> shareVolume = builder.stream(STOCK_TRANSACTIONS_TOPIC,
    Consumed.with(stringSerde, stockTransactionSerde))
    .withOffsetResetPolicy(EARLIEST) 1
    .mapValues(st -> ShareVolume.newBuilder(st).build())
    .groupByKey((k, v) -> v.getSymbol(), Serialized.with(stringSerde, shareVolumeSerde)) 2
    .reduce(ShareVolume::reduce); 3
4
```

- 1 Source processor consuming from topic
- 2 Mapping StockTransaction objects to ShareVolume objects
- 3 Grouping the ShareVolume objects by their stock ticker symbol
- 4 Reducing ShareVolume objects to contain a rolling aggregate of share volume

This code is concise, and we are squeezing a lot of power into a few number of lines.

If we look at the first parameter of the `builder.stream` method we see something new. It's the `AutoOffsetReset.EARLIEST` enum (there is also a `LATEST`) that we set with the `Consumed.withOffsetResetPolicy` method. This enum allows us to specify the offset reset strategy for each KStream/KTable. Using the enum takes precedence over the offset reset setting in the streams config.

**IMPORTANT** **GroupByKey vs. GroupBy**

`KStream` has two versions for grouping records `GroupByKey` and `GroupBy` both return a `KGroupedTable`, so you may be asking what's the difference and when to use either one. The `GroupByKey` method is when your `KStream` already has keys, as a result the re-keying step. More importantly, the "needs repartitioning" flag is never set. The `GroupBy` method assumes you have modified the key for the grouping, so the repartition flag is set to `true`, so joins, aggregations etc after calling `GroupBy` will result in an automatic repartitioning. Bottom line, prefer `GroupByKey` over `GroupBy` when ever possible.

While it's clear what `mapValues` and `groupBy` are doing, let's look into our `sum` method:

**Listing 5.3 Source of the ShareVolume sum method found in  
src/main/java/bbejeck/model/ShareVolume.java**

```
public static ShareVolume sum(ShareVolume csv1, ShareVolume csv2) {
    Builder builder = newBuilder(csv1); ①
    builder.shares = csv1.shares + csv2.shares; ②
    return builder.build(); ③
}
```

- ① Using a Builder for a copy constructor.
- ② Setting the number of shares to the total of both ShareVolume objects.
- ③ Calling build and returning a new ShareVolume object.

**NOTE** **Builder Pattern usage**

While we've seen the Builder Pattern in use earlier in the book, we see it here in a somewhat different context. In this example, we are using the Builder to make a copy of an object and update a field without modifying the original object.

The `ShareVolume.sum` method gives us our rolling total of share volume. The outcome of entire processing chain is a `KTable<String, ShareVolume>` object. Now we can see the role of the `KTable`. As `ShareVolume` objects come in, the associated `KTable` keeps the most recent update. It's important to remember every update is reflected in the `shareVolume` `KTable` from above, but each update is not necessarily emitted.

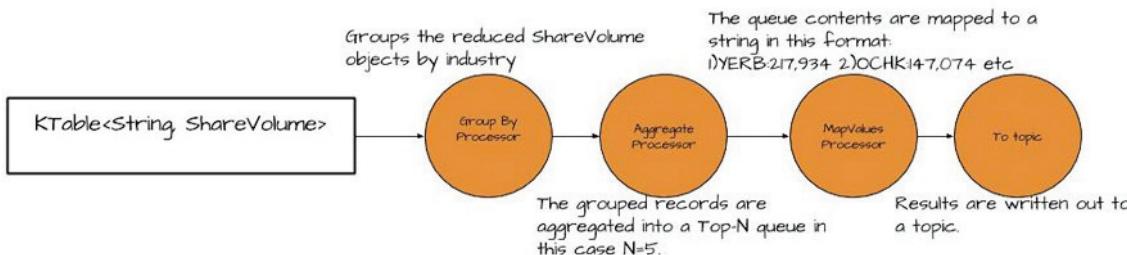
**NOTE****Aggregation vs. Reduce**

Why did we choose to do a reduce over an aggregation? While reducing is a form of aggregation, a reduce operation will yield the same type of object. An aggregation also sums results but could return a different type of object.

Next, we'll take our KTable and use it to perform a top-five aggregation (by share volume) summary. The steps we'll take here are similar to the steps we took for our first operation:

1. We need to perform another groupBy operation. We'll group the individual ShareVolume objects by their industry.
2. With the ShareVolume grouped by industry, we can start to add them together. This time our aggregation object is a fixed size priority queue. The fixed size queue only keeps the top five companies by share volume.
3. Then we'll map the queue into a string reporting the top five stocks per industry by share volume.
4. The string result is then written out to a topic.

Before we look at the code, let's take a look at a topology graph to cement our understanding of the data flow:



**Figure 5.10 Topology for grouping by industry and aggregating by top 5. Then we map the top 5 queue to a string and write the string out to a topic**

We can see from the image of the graph above, this second round of processing is straight forward. With a clear picture of the structure of this second round of processing, it's time to look at the source code:

**Listing 5.4 KTable GroupBy and Aggregation found in  
src/main/java/bbejeck/chapter\_5/AggregationsAndReducingExample.java**

```

Comparator<ShareVolume> comparator = (sv1, sv2) -> sv2.getShares() - sv1.getShares()
FixedSizePriorityQueue<ShareVolume> fixedQueue = new FixedSizePriorityQueue<>(comparator, 5);

shareVolume.groupBy((k, v) -> KeyValue.pair(v.getIndustry(), v),
    Serialized.with(stringSerde, shareVolumeSerde))
    .aggregate(() -> fixedQueue,
        (k, v, agg) -> agg.add(v),
    
```

1  
2  
3

```

    (k, v, agg) -> agg.remove(v),
    Materialized.with(stringSerde, fixedSizePriorityQueueSerde))
    .mapValues(valueMapper)
    .toStream().peek((k, v) -> LOG.info("Stock volume by industry {} {}", k, v))
    .to("stock-volume-by-company", Produced.with(stringSerde, stringSerde));

```

- 4
- 5
- 6
- 7
- 8

- 1 Grouping by industry and provided required Serdes
- 2 The Aggregate initializer an instance of the FixedSizePriorityQueue class (custom for demonstration purposes not for production!)
- 3 Aggregate adder for new updates
- 4 Aggregate remover for removing old updates
- 5 Serde for the Aggregator
- 6 ValueMapper instance to convert aggregator to a string used for reporting.
- 7 Calling toStream() logging the results (to the console) via the peek method
- 8 Writing the results out to the "stock-volume-by-company" topic.

In our initializer there is `fixedQueue` variable. This is a custom object wrapping a `java.util.TreeSet` used to keep track of the top N number of results in the decreasing order of share volume.

In this code example, we've seen the `groupBy`, `mapValues` calls before so we won't go over them again (we call the `KTable#toStream` method as `KTable#print` is deprecated). But we haven't seen the KTable version of aggregate before, so let's take a minute to discuss it.

To review what makes a KTable unique is records with the same key are updates. The KTable replaces the old record with the new one. Aggregation works in the same manner. It's an aggregation of the most recent records with the same key. So as a record arrives we add it to the `FixedSizePriorityQueue` using the `adder` method (second parameter in the `aggregate` call), but if an another record exists with the same key, we remove the old record with the `subtractor` (third parameter in the `aggregate` call).

What this means is our aggregator, `FixedSizePriorityQueue`, does not aggregate *all* values with the same key. Instead, the aggregator keeps a running tally of the top N stocks which have the largest volume. Each record coming has the total volume of shares traded so far. Our `KTable` will show us which companies have the top number of shares traded at the moment; we don't want a running aggregation of each update.

We should pause for a minute and review some key points I want you to remember:

1. We've learned how to group values in a KTable by a common key.
2. Additionally, we've learned how to perform useful operations like reducing and aggregation with those grouped values.

The ability to perform these operations is important when you are trying to make

sense of or determine what your data is telling you as it flows through your Kafka Streams application.

Also, we've brought together some of the key concepts discussed earlier in the book. In chapter 4 we learned about the importance of having fault-tolerant, local state for streaming applications. The first example of this chapter shows why local state is so important. To briefly restate those reasons, state allows you to keep track of what you have seen. Having *local* access does not introduce any network latency and is isolated from other processes.

Anytime we execute a reduction or aggregation operation we provide the name of state store. Reductions and aggregation operations return a KTable instance and The KTable uses the state store to replace older results with the newer ones.

As we've said before, not every update gets forwarded downstream, and that is important because we perform aggregation operations to gather *summary* information. If we didn't use local state, the KTable would forward *every* aggregation or reduction result.

Next, we'll show how to perform aggregation like operations over distinct periods of time called windowing.

### 5.3.2 Windowing Operations

In the first example in this chapter, we provided a "rolling" reduction and aggregation. The application performed a continuous reduction of share volume, then a top 5 aggregation of shares traded in the stock market. Sometimes we'll want an ongoing aggregation/reduction of results. Other times we want only to do operations over a given time range.

How many stock transactions have involved a given company over the last ten minutes? How many users have clicked to view a new advertisement in the last 15 minutes? While the applications perform these operations many times, the results are only for the defined period or "window" of time.

#### COUNTING STOCK TRANSACTIONS BY CUSTOMER

For our next example, we have the following scenario. We'll want to track stock transactions for a handful of traders. These could be large institutional traders or financially savvy individuals. We have two potential reasons for doing this tracking.

One case is we want to see where the market leaders are buying and selling. If these big hitters/savvy investors see an opportunity in the market, we'll follow the same strategy. The other case is we want to track for any indications of insider trading. We'll want to look into the timing of large spikes in trading and correlate them with significant news releases.

Here are the steps we'll take to do this tracking:

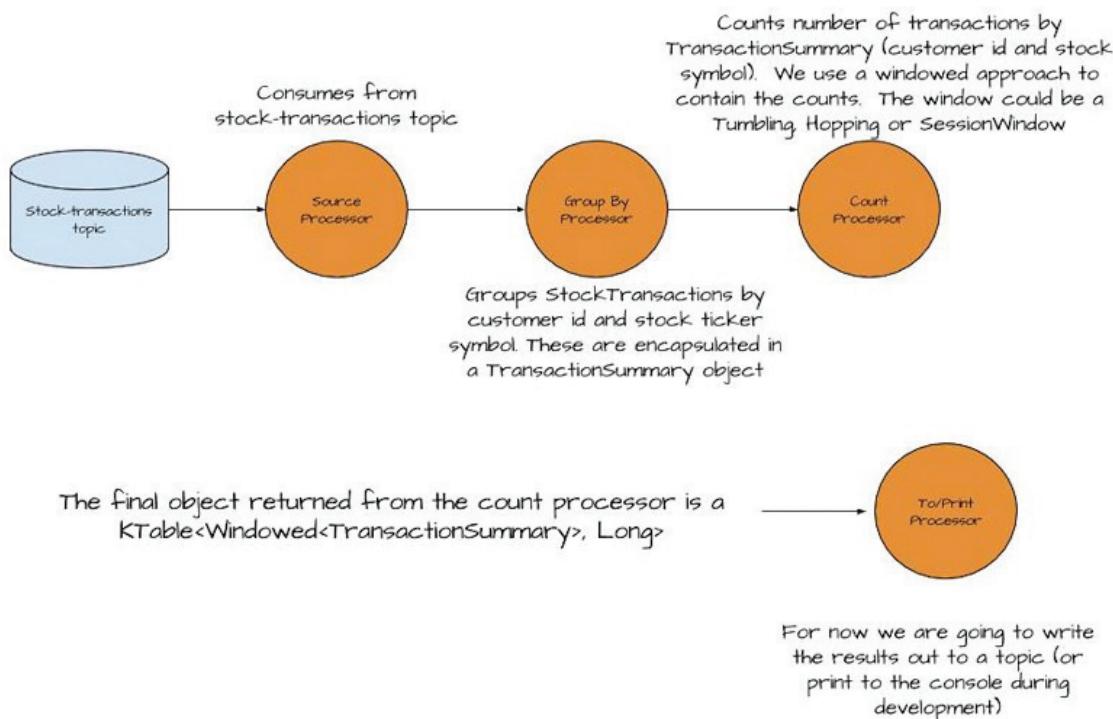
1. Create a stream to read from the "stock-transactions" topic
2. Group incoming records by the customer id and stock ticker symbol. The `groupBy` call returns a `KGroupedStream` instance.
3. After we have grouped the records, we'll use the `KGroupedStream.windowedBy` method

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

to return a windowed stream so we can perform some sort of windowed aggregation. Depending on the window type provided, we'll either get a `TimeWindowedKStream` or `SessionWindowedKStream` in return.

4. Once we have our windowed stream instance, we'll perform a count for the aggregation operation. The windowing stream determines if the record gets included into our count.
5. Write the results to a topic or print results to console during development.

The topology for this application is straightforward, but it's helpful to have a mental picture of the structure. Here's an image of the topology:



We have laid out our topology and have introduced the concept of windows. Our next step is to describe the various windowing functionality and the corresponding code.

## WINDOW TYPES

In Kafka Stream we have three types of windows available to us:

1. Session windows
2. Tumbling windows
3. Sliding/Hopping windows

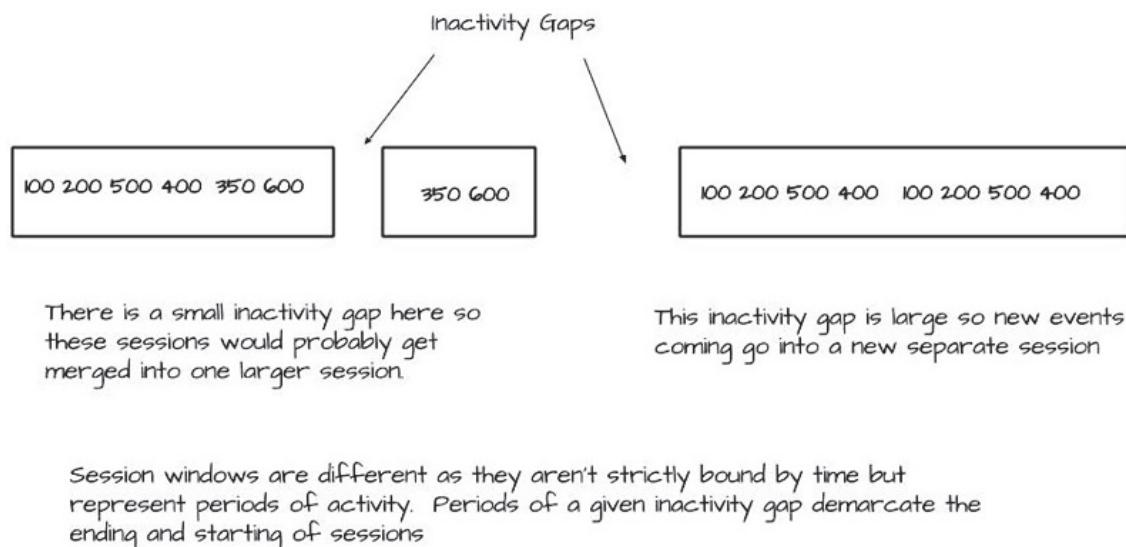
Which one we choose depends on the business requirements. The tumbling and hopping windows are time bound. While session windows are more about user activity the length of the session(s) is driven by the time.

A key point to keep in mind for all windows with respect to time is we are referring to the timestamps in the record and not wallclock time.

Next, we'll implement our topology from above with each of the window types. We'll only show the full code in the first example. Other than changing the type of window operation to use, everything else will remain the same.

## SESSION WINDOWS

Session windows are very different from other windowing operations. Session windows aren't bound strictly by time as much as user activity (or the activity of anything you want to track). We delineate session windows by a given period of inactivity. The image below shows how we could view session windows:



**Figure 5.11 Session windows example. We combine new sessions with a small inactivity gap to an adjacent session to form a new larger session.**

In the image above, you can see how the smaller session gets merged with the session on the left. But the session on the right will mark the start of a new session since the inactivity gap is large.

SessionWindows work on user activity, but use timestamps in the records to determine which "session" the record belongs.

## USING SESSION WINDOWS TO TRACK STOCK TRANSACTIONS

In this section, we'll show how we'll use SessionWindows to capture our stock transactions. The following code example is how we'll implement the session windows image above.

### Listing 5.5 Tracking Stock Transactions with Session Windows found in `src/main/java/bbejeck/chapter_5/CountingWindowingAndKtableJoinExample.java`

```

Serde<String> stringSerde = Serdes.String();
Serde<StockTransaction> transactionSerde = StreamsSerdes.StockTransactionSerde();
Serde<TransactionSummary> transactionKeySerde = StreamsSerdes.TransactionSummarySerde();

long twentySeconds = 1000 * 20;
long fifteenMinutes = 1000 * 60 * 15;

KTable<Windowed<TransactionSummary>, Long> customerTransactionCounts =
    builder.stream(STOCK_TRANSACTIONS_TOPIC, Consumed.with(stringSerde,
        transactionSerde)
        .withOffsetResetPolicy(LATEST))
        .groupByKey((noKey, transaction) -> TransactionSummary.from(transaction),
    
```

1

2

```

    Serialized.with(transactionKeySerde, transactionSerde))
    .windowedBy(SessionWindows.with(twentySeconds).until(fifteenMinutes)).count(); 3  

4  

customerTransactionCounts.toStream().print(Printed.<Windowed<TransactionSummary>,
    Long>toSysOut()
    .withLabel("Customer Transactions Counts")); 5

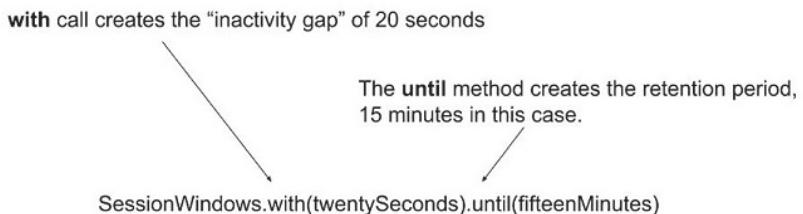
```

- ➊ The KTable resulting from the groupBy and count calls.
- ➋ Creating the stream from the STOCK\_TRANSACTIONS\_TOPIC (a String constant) Here we are using the offset reset strategy enum of LATEST for this stream.
- ➌ Grouping records by customer id and stock symbol; contained in the TransactionSummary object.
- ➍ Windowing the groups with SessionWindow with an inactivity time of twenty seconds and a retention time of fifteen minutes, then performing an aggregation as a count.
- ➎ Converting our KTable output to a KStream and printing result to the console.

We've seen most of the operations specified in this topology before, so we don't need to cover them again. But there are a couple of new items we haven't seen before, and we'll discuss those now.

Any time you do a groupBy operation you'll typically perform some sort of an aggregation operation (aggregate, reduce or count). You can perform a cumulative aggregation where previous results continue to build up or we can perform "windowed" aggregations where records are combined for the given time of the window.

The code on listing 5.5 is demonstration of doing a count over session windows. The following image breaks it down for us:



**Figure 5.12 Creating Session Windows with inactivity periods and retention**

With the call to `windowedBy(SessionWindows.with(twentySeconds).until(fifteenMinutes))` we are creating a SessionWindow with an inactivity gap of twenty seconds and retention period of fifteen minutes. An inactivity time of twenty seconds means the application includes any record arriving within twenty seconds of the current sessions ending or start time in the current (active) session.

We then specify the aggregation operation to perform, a count in this case, on the session windows.

If the record falls outside the inactivity gap (on either side of the timestamp), the application creates a new session. The retention period maintains the session for the specified amount of time. The retained session time is for late-arriving data that is outside the inactivity period of a session but can still be merged. Additionally, as sessions are combined the newly created session has the earliest timestamp and latest timestamp for the start and end of the new session, respectively.

Let's walk through a few records from the `count` method to see sessions in action.

**Table 5.1 Sessioning Table with Twenty Second Inactivity Gap**

Arrival Order	Key	Timestamp
1	{123-345-654,FFBE}	00:00:00
2	{123-345-654,FFBE}	00:00:15
3	{123-345-654,FFBE}	00:00:50
4	{123-345-654,FFBE}	00:00:05

As records come in we look for existing sessions with the same key and ending times less than (current timestamp - inactivity gap) and starting times greater than (current timestamp + inactivity gap). With that in mind here's how the four records in the table end up getting merged into a single session:

1. Record 1 is of course first, so start and end are 00:00:00
2. Record 2 arrives and we look for sessions with the earliest ending of 23:59:55 and the latest start of 00:00:35 and we find record 1. We then merge sessions 1 and 2. We keep the session one start time (earliest) and the session two end time (latest). So now we have one session starting at 00:00:00 and ending at 00:00:15
3. Record 3 arrives, and we look for sessions between 00:00:30 and 00:01:10 and find none. So we add a second session for key 123-345-654, FFBE starting and ending at 00:00:50.
4. Record 4 arrives we search for sessions between 23:59:45 and 00:00:25. This time we find sessions 1 and 2. Now all three are merged into one session with a start time of 00:00:00 and an end time of 00:00:15.

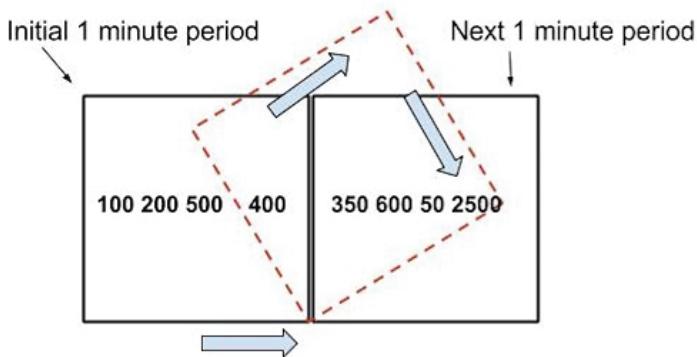
There are a couple of key points to remember from this section. First, sessions are not a fixed size window. Rather the size of a session is driven by the amount of activity within a given timeframe. Secondly, timestamps in the data determine how an event either fits into an existing session or falls into an inactivity gap.

Next, we'll move on to our next windowing option, tumbling windows.

## TUMBLING WINDOWS

Fixed or "Tumbling" windows capture events within a given period. Imagine you want to capture all stock transactions for a company every twenty seconds, so you collect every event for twenty seconds. After that twenty second period is over, your window will "tumble" to a new twenty second observation period. Here's an image to show this situation:

The current time period “tumbles” (represented by the dashed box) into the next time period completely with no overlap.



The box on the left is our first 1 minute window. After one minute it “tumbles” over or updates to capture events in a new 1 minute period.

There is no overlapping of events. The event window one contains [100, 200, 500, 400] and event window two contains [350, 600, 50, 2500].

**Figure 5.13 Tumbling windows reset after a given fixed period**

As you can see, we include each event that has come in for the last minute. However, we wait the length of the specified time to create a new window.

Here’s how we use tumbling windows to capture stock transactions every twenty seconds as defined in our scenario above:

#### **Listing 5.6 Using Tumbling Windows to Count User Transactions found in src/main/java/bbejeck/chapter\_5/CountingWindowingAndKtableJoinExample.java**

```
KTable<Windowed<TransactionSummary>, Long> customerTransactionCounts =
    builder.stream(STOCK_TRANSACTIONS_TOPIC, Consumed.with(stringSerde,
        transactionSerde)
        .withOffsetResetPolicy(LATEST))
    .groupBy((noKey, transaction) -> TransactionSummary.from(transaction),
        Serialized.with(transactionKeySerde, transactionSerde))
    .windowedBy(TimeWindows.of(twentySeconds)).count();
```

①

- Specifying a tumbling window of twenty seconds.

So with a minor change, we can use a tumbling window with the `TimeWindows.of` call. In this example, we did not add the `until` method. By not specifying the duration of the window we’ll get the default retention of twenty-four hours.

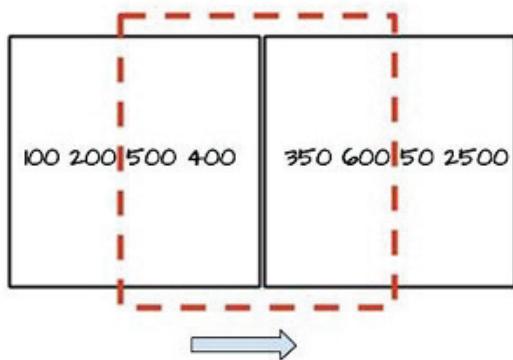
Next, we’ll move on to the last of our windowing options, hopping windows.

## SLIDING OR HOPPING WINDOWS

Sliding or "Hopping" windows are like fixed windows but with a small difference. A sliding window won't wait the entire time before launching another window to process events that have occurred during the last interval. Sliding windows perform a new calculation after waiting for an interval smaller than the duration of the entire window.

To relate how hopping windows differ from tumbling windows, let's recast the stock transaction count example. We still want to count the number of transactions, but we don't wait the entire period before we update the count.

Instead, we update the count at *smaller* intervals. For example, we'll still count the number of transactions every twenty seconds, but we update the count every five seconds. Consider the following example in this graphic:



The box on the left is our first 1 minute window. But the window "slides" over or updates after 30 seconds to start a new window. Now we have overlapping of events. Window one contains [100, 200, 500, 400], window two contains [500, 400, 350, 600] and window three is [350, 600, 50, 2500]

**Figure 5.14 Sliding windows update frequently and may contain overlapping data**

As you can see here, we still have twenty second windows. But we recalculate every five seconds. We now have three result windows with overlapping data.

And here's how to use hopping windows in the code example

### Listing 5.7 Specifying Hopping Windows to Count Transactions found in `src/main/java/bbejeck/chapter_5/CountingWindowingAndKtableJoinExample.java`

```
KTable<Windowed<TransactionSummary>, Long> customerTransactionCounts =
    builder.stream(STOCK_TRANSACTIONS_TOPIC, Consumed.with(stringSerde,
        transactionSerde)
        .withOffsetResetPolicy(LATEST))
    .groupBy((noKey, transaction) -> TransactionSummary.from(transaction),
        Serialized.with(transactionKeySerde, transactionSerde))
    .windowedBy(TimeWindows.of(twentySeconds).advanceBy(fiveSeconds).
        until(fifteenMinutes)).count();
```

①

- Using a hopping window of twenty seconds and advancing every five seconds

With the addition of the `advanceBy` method, we've converted our tumbling window to a hopping window. This time we've specified a retention time of fifteen minutes.

**NOTE**

You'll notice with all the windowing examples presented here, the only code that changes is the `windowedBy` calls. To that end I instead of having 4 nearly identical sample classes, I've added 4 different windowing example lines in the `src/main/java/bbejeck/chapter_5/CountingWindowingAndKtableJoinExample.java` file. To see a different windowing operation in action, simply comment out the current windowing operation and uncomment the one you want to execute.

To sum up, we've covered how to put our aggregation results in time windows. In particular, I want you to remember three things from this section:

1. Session windows aren't fixed by time but are driven by user activity.
2. Tumbling windows give us a *distinct* set picture of events within the specified time frame.
3. Hopping windows are a window of size N, but is frequently updated and can contain overlapping records in each window.

Next, we'll demonstrate how we can convert a `KTable` back into a `KStream` to perform a join.

### 5.3.3 Joining KStreams and KTables

Remember from Chapter 4 we discussed joining two `KStreams`. Now we are going to show how we can join a `KTable` and a `KStream`. The reason why we'll do this is simple. `KStreams` are record streams, and the `KTable` is a stream of record updates. But there may be cases where we need to add some additional context to our record stream with the updates in a `KTable`.

What if we wanted to take our customer stock transaction counts and join them with financial news from relevant industry sectors? It's easy to see stock transactions as a stream of events. While economic news can be frequent, it's more static where today's news can be considered an update or replacement for old news.

To do this with our existing code, we'll need to take a few steps to make this happen:

1. Convert the `KTable` of stock transaction counts into a `KStream` where we change the key to the industry of the count by ticker symbol.
2. Create a `KTable` that reads from a topic of financial news. Luckily the new `KTable` is categorized by industry.
3. Join the news updates with our stock transaction counts by industry.

With these steps laid out for us, let's walk through how we can accomplish these tasks.

## CONVERTING THE KTABLE TO KSTREAM

To do the KTable to KStream conversion, we'll take the following steps:

1. Call the `KTable.toStream()` method.
2. Then use the `KStream.map` call to change the key to the industry name and extract the `TransactionSummary` object from the `Windowed` instance.

These steps are chained together in the following manner:

### **Listing 5.8 Converting A KTable to a KStream found in src/main/java/bbejeck/chapter\_5/CountingWindowingAndKtableJoinExample.java**

```
KStream<String, TransactionSummary> countStream =
    customerTransactionCounts.toStream().map((window, count) -> {
        TransactionSummary transactionSummary = window.key();
        String newKey = transactionSummary.getIndustry();
        transactionSummary.setSummaryCount(count);
        return KeyValue.pair(newKey, transactionSummary);
    });
    
```

1  
2  
3  
4  
5

- 1 Calling `toStream` immediately followed by the `map` call.
- 2 Extracting the `TransactionSummary` object from the `Windowed` instance.
- 3 Setting the key to the industry segment of the stock purchase.
- 4 Taking the count value from the aggregation and placing it in the `TransactionSummary` object.
- 5 Returning the new `KeyValue` pair for the `KStream`.

We should note here that by performing a `KStream.map` operation, repartitioning for the returned `KStream` instance is automatically done when used in a join.

Now we have the conversion process completed our next step is to create the KTable to read the financial news.

## CREATING THE FINANCIAL NEWS KTABLE

Fortunately, to create the KTable is one line of code:

### **Listing 5.9 KTable for Financial News found in src/main/java/bbejeck/chapter\_5/CountingWindowingAndKtableJoinExample.java**

```
KTable<String, String> financialNews = builder.table("financial-news",
    Consumed.with(EARLIEST));
    
```

1

- 1 Creating the KTable with `EARLIEST` enum, topic of "financial-news".

It's worth noting here we didn't provide any serdes as the configuration is using string serdes. Also by using the enum `EARLIEST`, we're populating the table with records

on startup.

Now we'll move on to our last step, setting up the join.

## JOINING NEWS UPDATES WITH TRANSACTION COUNTS

Setting up the join is very simple. In this case, we'll use a left join in case there is no financial news for the industry involved in the transaction.

### **Listing 5.10 Setting up the Join Between the KStream and KTable found in src/main/java/bbejeck/chapter\_5/CountingWindowingAndKtableJoinExample.java**

```
ValueJoiner<TransactionSummary, String, String> valueJoiner = (txnct, news) ->
    String.format("%d shares purchased %s related news [%s]", txnct.getSummaryCount(),
        txnct.getStockTicker(), news); 1

KStream<String, String> joined = countStream.leftJoin(financialNews, valueJoiner,
    Joined.with(stringSerde,
        transactionKeySerde, stringSerde)); 2

joined.print(Printed.<String, String>toSysOut().withLabel("Transactions and News")); 3
```

- ① ValueJoiner used for combining the values from the join result.
- ② The left-join statement for the countStream KStream and the financial news KTable providing serdes with a Joined instance.
- ③ Printing results to console (in production we'd write this to a topic with a "to('topic-name') call)

The left-join statement straight-forward. You'll notice that unlike the joins in Chapter 4, we did not provide a `JoinWindow`. That's because when performing a KStream-KTable join, there is only one record per key in the KTable. So performing the join is not a matter of time, the record is present in the KTable or not.

The key point here is we can use KTables to provide less frequently updated lookup data to enrich our KStream counterparts.

Next, we'll take a look at a more efficient way to enhance our KStream events.

### 5.3.4 GlobalKTables

At this point, we've established the need to enrich or add context to our event streams. We've seen joins between two KStreams in Chapter 4, and the previous example demonstrated a join between a KStream and a KTable. In all of these cases, if we mapped the keys to a new type or value, a repartitioning of the stream needs to happen. Sometimes we do the repartitioning explicitly ourselves, and other times Kafka Streams does it automatically. Repartitioning makes sense, since the keys have been changed and will end up on a new partition or the join doesn't happen (this was discussed in Chapter 4 section 2.4 "Repartitioning the Data").

## REPARTITIONING HAS A COST

But the repartitioning is not free. There is additional overhead in this process with the creation of intermediate topics, storage of duplicate data in another topic, and the increased latency of writing to and reading from another topic. Additionally, if we want to join on more than one facet or dimension, we'll need to chain joins and map the records with a new key and repeat the repartitioning process.

## JOINING WITH SMALLER DATASETS

In some cases, the lookup data you want to join against is relatively small and entire copies of the lookup data could fit locally on each node. In situations where the lookup data is reasonable small enough, Kafka Streams provides the `GlobalKTable`. `GlobalKtables` are unique because the application replicates all of the data to each node. Since the entirety of the data is on each node, we no longer have the requirement that our event stream is partitioned by the key of the lookup data, making it available to all partitions. Another important part is we can now do non-key joins. Let's revisit one of our previous examples to demonstrate this capability.

## JOINING KSTREAMS WITH GLOBALKTABLES

Earlier in the chapter, we performed a windowed aggregation of stock transactions per customer. The output of the aggregation looked like this:

### **Listing 5.11 Windowed Stock Transactions Aggregation Output**

```
{customerId='074-09-3705', stockTicker='GUTM'}, 17
{customerId='037-34-5184', stockTicker='CORK'}, 16
```

While this output accomplished our goal, it would be more impactful if we could see the client's name and the full company name. While we could perform regular joins to fill out the customer and company name, we would need to do two key mappings and repartitioning. But with `GlobalKTables` we can avoid the expense of those operations.

To accomplish this, we'll take our first `countStream` from above and join it against two `GlobalKTables`. For review let's take our code for calculating the number of stock transactions using session windows:

### **Listing 5.12 Aggregation Stock Transactions using Session Windows found in src/main/java/bbejeck/chapter\_5/GlobalKTableExample.java**

```
KStream<String, TransactionSummary> countStream =
    builder.stream( STOCK_TRANSACTIONS_TOPIC, Consumed.with(stringSerde,
        transactionSerde) .withOffsetResetPolicy(LATEST) )
        .groupBy((noKey, transaction) -> TransactionSummary.from(transaction),
            Serialized.with(transactionSummarySerde, transactionSerde))
        .windowedBy(SessionWindows.with(twentySeconds)) .count()
        .toStream() .map(transactionMapper);
```

Since we've covered this previously, we won't review it again here. But I would like to point out we've abstracted the code in the `toStream()` `.map` function into function

object vs. having an in-line lambda for readability purposes.

Our next step is to define two `GlobalKTable` instances:

### **Listing 5.13 Defining the GlobalKTables for Lookup Data found in src/main/java/bbejeck/chapter\_5/GlobalKTableExample.java**

```
GlobalKTable<String, String> publicCompanies = builder.globalTable(COMPANIES.topicName()); ①
GlobalKTable<String, String> clients = builder.globalTable(CLIENTS.topicName()); ②
```

- ① The `publicCompanies` lookup is for finding companies by their stock ticker symbol
- ② The `clients` lookup us for getting customer names by customer id.

Note that in the example above, we've defined the topic names using enums. Now we have our components in place, we need to construct the join:

### **Listing 5.14 Joining a KStream with Two GlobalKTables found in src/main/java/bbejeck/chapter\_5/GlobalKTableExample.java**

```
countStream.leftJoin(publicCompanies, (key, txn) ->
    txn.getStockTicker(), TransactionSummary::withCompanyName) ①
    .leftJoin(clients, (key, txn) -> txn.getCustomerId(),
    TransactionSummary::withCustomerName) ②
    .print(Printed.<String, TransactionSummary>toSysOut().withLabel("Resolved
    Transaction Summaries")); ③
```

- ① Setting up the `leftJoin` with the `publicCompanies` table, keying by stock ticker symbol and returning the `transactionSummary` with company name added.
- ② Setting up the `leftJoin` with the `clients` table, keying by customer id and returning the `transactionSummary` with the customer named added.
- ③ Printing results out to the console.

Although there are two joins, we've chained them together as we don't use any of the results alone. We print the results at the end of the entire operation. Once we run the join operation we now get results like this:

### **Listing 5.15 Resolved Windowed Stock Transaction Aggregates**

```
{customer='Barney, Smith' company="Exxon", transactions= 17}
```

While the facts haven't changed, these results are more clear to read. Including Chapter 4, we've seen several types of joins in action. At this point it will be useful to summarize the join types we have available in a table:

**Table 5.2 Kafka Streams Join Table**

Left Join	Inner Join	Outer Join
KStream-KStream	KStream-KStream	KStream-KStream
KStream-KTable	KStream-KTable	N/A
KTable-KTable	KTable-KTable	KTable-KTable
KStream-GlobalKTable	KStream-GlobaKTable	N/A

This table represents the state of join options as of Kafka Streams 1.0.0, but this may change in future releases.

In conclusion, the key thing to remember is we can combine our event streams (KStream) and update streams (KTable) leveraging local state. Additionally, for those times where the lookup data is of a manageable size, we can use a GlobalKTable. GlobalKTables replicate all partitions to each node in the Kafka Streams application making all data available regardless of which partition the key maps to.

Next, we'll demonstrate a capability of Kafka Streams where we can observe changes to our state, without having to consume data from a Kafka topic.

### 5.3.5 Queryable State

At this point in the book, we've performed several operations involving state. In all cases, we've printed results to the console (development) or written them to a topic (production). The implications of writing the results to a topic are we need to use a Kafka Consumer to view the results. Reading the data from these topic(s) could be considered a form of "materialized views."

We'll define materialized views as "a database object that contains the results of a query. For example, it may be a local copy of data located remotely, or may be a subset of the rows and/or columns of a table or join result, or may be a summary using an aggregate function" <sup>13</sup>.

---

Footnote 13 [en.wikipedia.org/wiki/Materialized\\_view](https://en.wikipedia.org/wiki/Materialized_view)

But Kafka Streams offers "interactive queries" from state stores, giving us the ability to read these materialized views directly. It's important to note querying state stores is a read-only operation. By making the queries read only, we don't have to worry about creating an inconsistent state while our application continues to process data.

The impact of making the state stores directly queriable is significant. This direct querying ability means we can create dashboard applications without having to consume the data from a Kafka Consumer first. There are also some gains in efficiency by not writing the data out again:

1. Since the data is local, you can access it very fast
2. We avoid duplication of data by not copying data to an external store <sup>14</sup>.

---

Footnote 14 Adapted from

[www.confluent.io/blog/unifying-stream-processing-and-interactive-queries-in-apache-kafka/](https://www.confluent.io/blog/unifying-stream-processing-and-interactive-queries-in-apache-kafka/)

The main thing I want to remember from this section is the ability to query state from the application directly. I can't stress enough the power this feature gives you. Instead of consuming from Kafka and storing records in a database to feed your application, you can directly query the state stores for the same results. The impact of direct queries on state stores means less code (no consumer) and less software (no need for database table to store results).

We've covered a lot in this chapter, so we'll stop at this point in our discussion of interactive queries on state stores.

But fear not, in Chapter Nine we'll build a simple dashboard application with interactive queries. This dashboard application will use some of our current examples to demonstrate interactive queries and how you can add it into your Kafka Streams applications.

## 5.4 Summary

Again, we have another chapter where we've covered a lot of ground. This chapter builds on the concepts of the previous one. We can consider this chapter as the capstone chapter of the KStreams API. Although we touched on several topics here are the key takeaways:

- KStreams represent event streams of comparable to inserts into a database. KTables are update streams and are more akin to updates into a database. The size of KTable doesn't continue to grow; it replaces older records with the newer records.
- KTables are essential for performing aggregation operations.
- We can place our aggregated data into "time buckets" with windowing operations.
- GlobalKTables give us lookup data across the entire application, regardless of the partition.
- We can perform joins with KStreams, KTables, and GlobalKTables.

So far we've focused on the "high-level" KStreams DSL to build our Kafka Streams applications. While the high-level approach gives us nice concise programs, everything is a trade-off. By working with the KStreams DSL, we trade a level of control for more concise code. In the next chapter, we'll cover the "low-level" or Processor API. We'll make the different trade-offs using the Processor API; we won't have the conciseness of our applications we've built so far. But we'll gain the ability to perform or create virtually any kind of processor we need.

# The Processor API

In this chapter:

- The Tradeoffs of Higher Abstractions vs. More Control
- Working with Sources, Processors, and Sinks to create a Topology
- Digging Deeper into the Processor API with a Stock Analysis Processor
- Creating a CoGrouping Processor.
- Integrating the Processor API and the Kafka Streams API

Up to this point in the book, we have been working with the "High Level" or Kafka Streams API. The Kafka Streams API is a DSL allowing developers to create robust applications with minimal code. The ability to quickly put together processing topologies is an important feature of the Kafka Streams DSL. The DSL allows you to iterate quickly to flesh out ideas for working on your data without getting bogged down with intricate setup details some other frameworks may need.

But at some point, even when working with the best of tools, you always get to *those* one-off situations. You have a problem to solve requiring you to deviate from the traditional path. Whatever the case may be, we need a way to dig down and write some code that just isn't possible with a higher level abstraction.

## 6.1 The Tradeoffs of Higher Abstractions vs. More Control

To me, the classic example of this trade-off is using ORM (Object Relational Mapping) frameworks. A good ORM framework maps your domain objects to database tables and will create the correct SQL queries for you at runtime. When you have simple to moderate SQL operations (simple select or join statements), using the ORM framework saves you a lot of time. But no matter how good the ORM framework is it seems inevitable that there are those few queries (very complex joins, select statements with nested sub-select statements) that just aren't working the way you want. To get the job done you need to write raw SQL to get the information from the database in the format you need. Many times you'll be able to mix in raw SQL with the mappings provided with the framework.

So we can see the tradeoff of a higher level abstraction vs. more programmatic control now. You can use a higher level abstraction that does much of the low-level work for you. In the case of the previous example the ORM framework, once configured, allows you work almost unaware of the fact you are materializing objects from a database.

To be clear, I'm not trying to draw a direct comparison between the Kafka Streams API and using an ORM framework. I chose the example because I feel it represents the trade-off of higher abstraction vs. more programmatic control many developers experience.

This chapter is about those times you want to do stream processing in a way that the Kafka Streams DSL doesn't seem to handle the way you want. For example, we've seen from working with the `kTable` API the framework controls the timing of forwarding records downstream.

But you may find yourself in a situation where you want explicit control when a record gets sent. For example, you are tracking trades on Wall Street, and you only want to forward records when a stock crosses a price threshold. To gain this type of control, I am talking about using the Processor API. What the Processor API lacks in ease of development, it makes up for in power. You can write your custom processors to do almost anything you want.

In the following sections we'll learn how we can use the Processor API to handle situations like these:

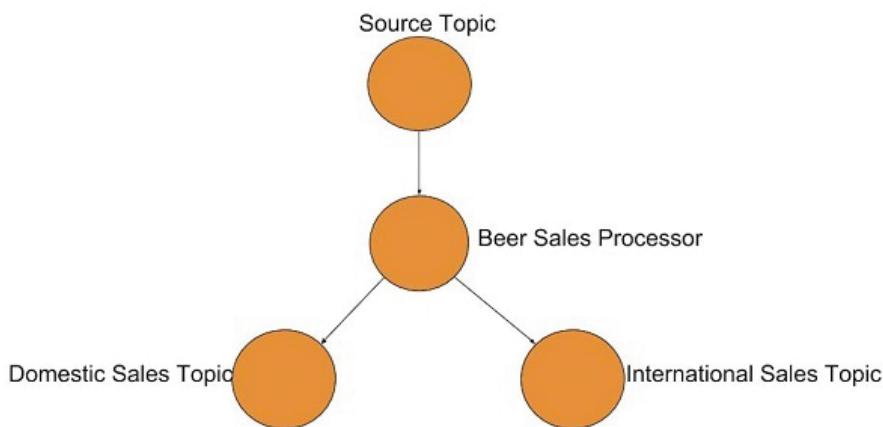
1. Schedule actions to occur at regular intervals (either by timestamps in the records or wall-clock time).
2. Full control over when records sent downstream.
3. Forwarding records to specific child nodes.
4. The ability to create functionality that doesn't exist in the Kafka Streams API. We'll see an example of this later when we build a CoGrouping processor.

First, let's understand how to use the Processor API by developing a topology step by step.

## 6.2 Working with Sources, Processors, and Sinks to create a Topology

Let's say you're the owner of a successful brewery (Pops Hops) with several locations. You've recently expanded your business to accept orders from distributors online, including international sales to Europe. You want to route orders within the company based on whether the order is domestic or international, converting any European sales from either British Pounds or Euros to U.S Dollars. With these details in mind, we are going to use this example to demonstrate how the Processor API gives you the flexibility to select specific child nodes when forwarding records.

If you were to sketch out the flow of operation it would look something like this:



**Figure 6.1 Beer Sales Distribution Pipeline**

The best way to learn about the Processor API is to build a simple application so we can talk about each step. First, we'll start by creating a source node.

### 6.2.1 Adding a Source Node

The first step we take when constructing a topology is establishing the source node(s). Here's the line of code where we set the source of data for our new topology:

**Listing 6.1 Creating the Beer Application source node found in  
src/main/java/bbejeck/chapter\_6/PopsHopsApplication.java**

```

topology.addSource(LATEST,
    purchaseSourceNodeName,
    new UsePreviousTimeOnInvalidTimestamp(),
    stringDeserializer,
    beerPurchaseDeserializer,
    Topics.POPS_HOPS_PURCHASES.topicName() )
  
```

- 1 Specifying the offset reset to use.
- 2 The name of this node.
- 3
- 4
- 5
- 6

- 1 Specifying the offset reset to use.
  - 2 The name of this node.
  - 3
- ©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

- The `TimestampExtractor` to use for this source.
- ④ Key deserializer.
- ⑤ Value deserializer.
- ⑥ The name of the topic to consume data.

In the `Topology#addSource` method, there are some new parameters we didn't use in the DSL. At listing number two, we are naming the source node. When using the Kafka Streams API, we didn't need to pass in a name, because KStreams generates a name for the node. But when using the Processor API, we need to provide the names of the nodes in our topology.

The node name is used to wire up a child node to a parent node. And we've seen that's how data flows in a Kafka Streams application, from parent to child node.

In listing three we specify the timestamp extractor to use with this source. In Chapter 4 section 4.5.1 we discussed the different timestamp extractors available to use per stream source. Here we have an example of using the `UsePreviousTimeOnInvalidTimestamp` class, and all other sources in the application will use the default `FailOnInvalidTimestamp` class.

In Listings four and five we provide a key deserializer and a value deserializer which represents another departure of the Processor API from the Kafka Streams DSL.

In the Kafka Streams DSL, we supplied `Serde` instances when creating a source or sink nodes. The `Serde` itself contains a serializer and deserializer, and the Kafka Streams DSL uses the appropriate one depending on if you are going from object to byte array or byte array to object.

#### NOTE

#### Serde vs. Deserializers

Another difference between the Kafka Streams API and the Processor API is the serialization/deserialization of objects. In the Kafka Streams API, we use a `Serde` which contains a serializer and deserializer for the object in question. In the Processor API, we directly pass in either a serializer or deserializer as required.

With the Processor API being a lower level abstraction, we directly provide a deserializer when creating a source node and a serializer when creating a sink node.

Finally, in listing six, we provide the name of the source topic.

Next, let's look at how we'll work with purchase records coming into our application.

### 6.2.2 Adding a Processor Node

Now we'll add a processor to work with the records coming in from the source node:

#### **Listing 6.2 Adding a processor node found in src/main/java/bbejeck/chapter\_6/PopsHopsApplication.java**

```
BeerPurchaseProcessor beerProcessor = new BeerPurchaseProcessor(domesticSalesSink, internationalSales
©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and
other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.
```

```

topology.addSource(LATEST,
    purchaseSourceNodeName,
    new UsePreviousTimeOnInvalidTimestamp(),
    stringDeserializer,
    beerPurchaseDeserializer,
    Topics.POPS_HOPS_PURCHASES.topicName())
    .addProcessor(purchaseProcessor, ①
        () -> beerProcessor, ②
        purchaseSourceNodeName); ③

```

- ① Naming the processor node.
- ② Adding the processor defined in line one above, using a Java 8 lambda instead of a ProcessorSupplier. A ProcessorSupplier is a single method interface used to provide a Processor instance. Since it's a single method interface, we can substitute a Java 8 lambda.
- ③ Specifying the name of the parent node(s);

You can see we still use the fluent interface pattern for constructing the topology. The difference from the Kafka Streams API lies in the return type. With the Kafka Streams API, every call on a `KStream` operator results in returning a new `KStream/KTable` instance. In the Processor API, each call to the `Topology` returns the same `Topology` instance each time.

In listing number two, we pass in the processor instantiated on the first line of the code example. The `Topology.addProcessor` method takes an instance of a `ProcessorSupplier` interface for the second parameter, but since the `ProcessorSupplier` is a single method interface, we can substitute it with a lambda expression.

The key point in this section is the third parameter is the same value as the second parameter of the `addSource` method. This process establishes the parent-child relationship between nodes. The parent-child relationship determines the flow data in a Kafka Streams application as demonstrated in the image below.

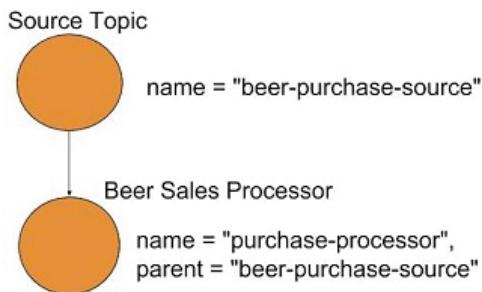
```
builder.addSource(LATEST,
    purchaseSourceNodeName,
    new UsePreviousTimeOnInvalidTimestamp()
    stringDeserializer,
    beerPurchaseDeserializer,
    "pops-hops-purchases");
```

The **name** of the source node is used for the **parent name** of the processing node added here. This establishes the parent-child relationship, which directs data flow in Kafka Streams

```
builder.addProcessor(purchaseProcessor,
    () -> beerProcessor,
    purchaseSourceNodeName);
```

**Figure 6.2 Wiring up nodes in Processor API using a node name for a parent node name.**

We should take this opportunity to review what we've built so far:



**Figure 6.3 Processor API topology built so far including node names and parent names.**

Now let's take a second to discuss the `BeerPurchaseProcessor` functions. The processor has two responsibilities:

1. Convert international sales amounts (in Euros) to US Dollars.
2. Based on the origin of sale, (domestic or international) route the record to the appropriate sink node.

All of this action takes place in the `process()` method. To quickly summarize here's what code does:

1. Check the currency type. If it's not in dollars, then convert it to dollars.
2. If you have a non-domestic sale, after converting the sales amount, forward the updated record to the international sales topic.
3. Otherwise, we can forward the record directly to the domestic sales topic.

The code for this processor is below:

found in `src/main/java/bbejeck/chapter_6/processor/BeerPurchaseProcessor.java`

## .BeerPurchaseProcessor

```

public class BeerPurchaseProcessor extends AbstractProcessor<String, BeerPurchase> {

    private String domesticSalesNode;
    private String internationalSalesNode;

    public BeerPurchaseProcessor(String domesticSalesNode, String internationalSalesNode) {
        this.domesticSalesNode = domesticSalesNode;
        this.internationalSalesNode = internationalSalesNode;
    }

    @Override
    public void process(String key, BeerPurchase beerPurchase) {
        Currency transactionCurrency = beerPurchase.getCurrency();

        if (transactionCurrency != DOLLARS) {
            BeerPurchase dollarBeerPurchase;
            BeerPurchase.Builder builder = BeerPurchase.newBuilder(beerPurchase);
            double internationalSaleAmount = beerPurchase.getTotalSale();
            String pattern = "###.##";
            DecimalFormat decimalFormat = new DecimalFormat(pattern);
            builder.currency(DOLLARS);
            builder.totalSale(Double.parseDouble(decimalFormat.format(transactionCurrency.
                convertToDollars(internationalSaleAmount))));

            dollarBeerPurchase = builder.build();
            context().forward(key, dollarBeerPurchase, internationalSalesNode);
        } else {
            context().forward(key, beerPurchase, domesticSalesNode);
        }
    }
}

```

- ➊ Setting the names for different nodes to forward records.
- ➋ The process method, where the action takes place.
- ➌ Converting international sales to U.S Dollars.
- ➍ Using the ProcessorContext, we forward records to a topic for international transactions.
- ➎ Sending records for domestic sales to the right child node.

Here we are extending the `AbstractProcessor`, a class with overrides for `Processor` interface methods except for the `process()` method. The `Processor.process()` method is where we perform actions for the records flowing through the topology. Listings four and five demonstrate the main point of this example of the Processor API, the ability to select specific child nodes when forwarding records.

**NOTE****The Processor Interface**

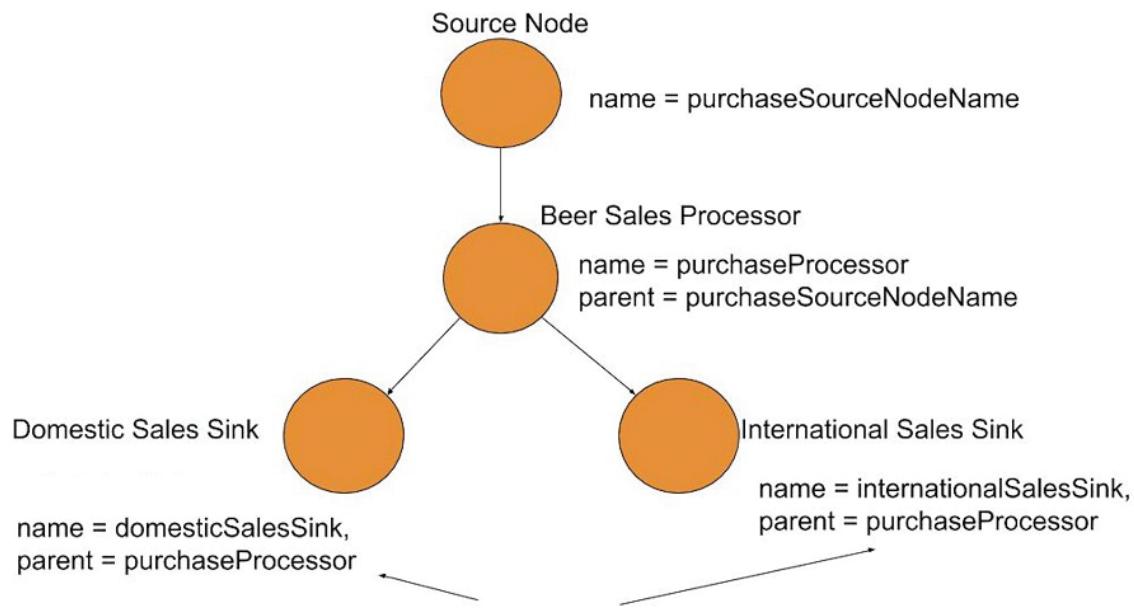
The `Processor` interface provides the `init`, `process`, `punctuate` and `close` methods. The `Processor` is the main driver of your application logic to work with records in your streaming application. In our examples, we'll mostly use the `AbstractProcessor` class, so we only override the methods we want. The `AbstractProcessor` class initializes the `ProcessorContext` for us, so if you don't need to do any set up in your class, you don't need to override the `init` method.

The `context()` method shown here (listings four and five) retrieves the reference to the `ProcessorContext` object for this processor. All processors in a topology receive a reference to the `ProcessorContext` via the `init()` method, which is executed by the `StreamTask` when initializing the topology. Since we don't need to initialize anything in this example, we've relied on the

Now that we've covered how we process our records, the next step is to demonstrate how we connect a sink node (topic) for us to write records back to Kafka.

### 6.2.3 Adding a Sink Node

By now you probably have a good feel for the flow of using the Processor API. To add source we used `addSource` for adding a processor we have an `addProcessor` so as you can imagine we use the `addSink` method to wire up a sink node (topic) to a processor node. Our updated topology will look like this:



**Figure 6.4 Completing the topology by adding sink nodes.**

Let's update the topology we're building by adding a sink nodes now in the code now:

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

### Listing 6.3 Adding a sink node found in src/main/java/bbejeck/chapter\_6/PopsHopsApplication.java

```

topology.addSource(LATEST,
    purchaseSourceNodeName,
    new UsePreviousTimeOnInvalidTimestamp(),
    stringDeserializer,
    beerPurchaseDeserializer,
    Topics.POPS_HOPS_PURCHASES.topicName())
    .addProcessor(purchaseProcessor,
        () -> beerProcessor,
        purchaseSourceNodeName)

    .addSink(internationalSalesSink,          ①
        "international-sales",                ②
        stringSerializer,                   ③
        beerPurchaseSerializer,             ④
        purchaseProcessor)                 ⑤

    .addSink(domesticSalesSink,               ①
        "domestic-sales",                  ②
        stringSerializer,                 ③
        beerPurchaseSerializer,            ④
        purchaseProcessor);              ⑤

```

- ① Name of the sink.
- ② The topic this sink represents.
- ③ Serializer for the key.
- ④ Serializer for the value.
- ⑤ Parent node for this sink.

In this case, we've added two sink nodes, one for Dollars and another for Euros. As depending on the currency of the transaction, we write our records out to two distinct topics. The key point to notice when adding two sink nodes here is both have the same parent name.

By supplying the same parent name to both sink nodes, we have wired both of them to our processor (as shown in the image above). While the Processor API is a little more verbose than the Kafka Streams API, we can see it's still easy to construct topologies.

The key points I want you to remember from our first example is how we wire topologies together and our ability to forward records to specific child nodes.

Our next example we'll explore more of the flexibility the Processor API provides.

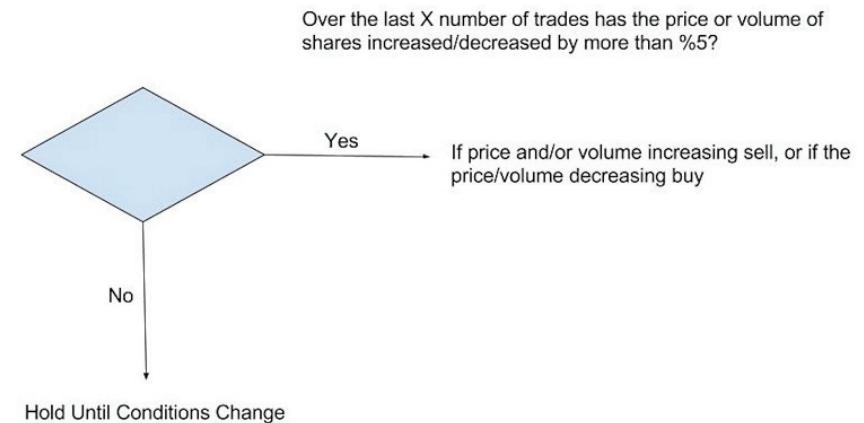
## 6.3 Digging Deeper into the Processor API with a Stock Analysis Processor

Now we are going to return to the world of finance. I want you to put on your day trading hat. As day traders, we want to analyze how stocks are changing in price with the intent of picking the best time to buy and sell. Our goal is to take advantage of market fluctuations and make a quick profit. We have a few key indicators to consider hoping they will tip us off when we want to make a move. Our entire list of requirements are as follows:

1. Show the current value of the stock.
2. Indicate if the price per share is trending up or down.
3. Include the total share volume so far, and if the volume is trending up or down.
4. Restrict sending records downstream to those displaying 2% trending (up or down).

Now let's take a walk through how we might handle this analysis manually when trying to determine when we see a deal to make. In the figure below (Figure 6.5) we show the decision tree of sorts we'll want to create to help us make decisions.

This is the current status of stock trading as XXYY:  
Symbol:XXYY, Share Price: \$10.79, Total Volume: 5,123,987



**Figure 6.5 Stock Trends Updates**

We have a handful calculations we need to perform for our analysis. Additionally, we'll use our calculation results to determine if and when we should forward records downstream.

This restriction for sending records means we can't rely on the standard mechanisms of commit time or cache flushes to handle the flow for us, which rules out using the Kafka Streams API. It goes without saying that we'll also require state so we can keep track of changes over time. What we need here is the ability to write a custom processor.

**IMPORTANT****Demo Purposes Only**

I'm pretty sure it goes without saying, but I'll state the obvious anyway. These stock price evaluations are for demonstration purposes only. Please don't infer any real market forecasting ability from this example. This model bears no similarity to a real-life approach and is presented only as a way of demonstrating more complex processing situations. Remember you're not a day trader, and neither am I!

In the next section, we are going to present the solution to our problem.

### **6.3.1 The Stock Performance Processor Application**

Here's the topology we are going to build:

**Listing 6.4 Stock performance application with custom processor found in  
src/main/java/bbejeck/chapter\_6/StockPerformanceApplication.java**

```
Topology topology = new Topology();
String stocksStateStore = "stock-performance-store";
double differentialThreshold = 0.02; 1

KeyValueBytesStoreSupplier storeSupplier = Stores.inMemoryKeyValueStore(stocksStateStore);
StoreBuilder<KeyValueStore<String, StockPerformance>> storeBuilder = 2
    Stores.keyValueStoreBuilder(storeSupplier, Serdes.String(), stockPerformanceSerde); 3

topology.addSource("stocks-source", stringDeserializer,
    stockTransactionDeserializer, "stock-transactions")
    .addProcessor("stocks-processor", () -> 4
        new StockPerformanceProcessor(stocksStateStore, differentialThreshold),
    "stocks-source") 5
    .addStateStore(storeBuilder, "stocks-processor")
    .addSink("stocks-sink", "stock-performance", stringSerializer, 6
        stockPerformanceSerializer, "stocks-processor");
```

- ① Setting the percentage differential for forwarding stock information.
- ② Creating an in-memory key-value state store.
- ③ Creating the StoreBuilder to place in the topology.
- ④ Adding the processor to the topology.
- ⑤ Adding the state store to the "stocks-processor"
- ⑥ Adding the sink for writing results out, although we could use a printing sink as well.

Building this topology has the same flow as our previous example, so we'll focus our attention on the new features we'll be presenting in our processor.

In the previous example, we didn't have any setup to do, so we relied on the `AbstractProcessor#init` method to initialize our `ProcessorContext` object. In this

example, however, we need to use a state store (listing 3), and we also want to schedule when we emit records versus forwarding records each time we receive records.

The first step then is to look at the `init` method in our processor:

**Listing 6.5 Init Method tasks found in  
src/main/java/bbejeck/chapter\_6/processor/StockPerformanceProcessor.java**

```

@Override
public void init(ProcessorContext processorContext) {
    super.init(processorContext); 1

    keyValueStore = (KeyValueStore) context().getStateStore(stateStoreName);
    StockPerformancePunctuator punctuator = new StockPerformancePunctuator(differentialThreshold,
        context(), 2
        keyValueStore); 3

    context().schedule(10000, PunctuationType.WALL_CLOCK_TIME, punctuator); 4
}
}

```

- ➊ Initializing ProcessorContext via the super class AbstractProcessor.
- ➋ Retrieving state store created when building topology.
- ➌ Initializing the Punctuator to handle the scheduled processing;
- ➍ Scheduling to have the Punctuator#punctuate called every ten seconds.

First, we need to initialize the `AbstractProcessor` with the `ProcessorContext`, so we first call the `init` method on the super class. Next, we grab a reference to the state store we created in the topology. Initially, all we need to do is set the state store to a variable for use later in the processor.

In listing 3, we see a new feature - an instance of a `Punctuator`, an interface which is essentially a callback to handle scheduled execution of processor logic but encapsulated in the `Punctuator#punctuate` method.

**TIP**

**Cancelling a Punctuation Call**

The `ProcessorContext#schedule(long, PunctuationType, Punctuator)` method returns a type of `Cancellable` allowing for canceling a punctuation and managing more advanced scenarios, like those found in the [punctuate use-cases discussion](#). We don't have examples or a discussion here, but I will present some examples in `src/main/java/bbejeck/chapter_6/cancellation`.

Then in listing 4, we use the `ProcessorContext` to schedule the `Punctuator` to execute every 10 seconds. Notice the second parameter `PunctuationType.WALL_CLOCK_TIME`, specifying we want to call `Punctuator.punctuate` every 10 seconds based on *wall-clock time*. Our other option is to specify `PunctuationType.STREAM_TIME`, which means the execution of `Punctuator.punctuate` is still scheduled every 10 seconds but driven by the time

elapsed via timestamps in the data. Let's take a moment to discuss the difference between the two `PunctuationType` settings.

### 6.3.2 Punctuation Semantics

Let's start our conversation on punctuation semantics with `STREAM_TIME` as it requires a little more explanation. The image below walks through the concept stream time:

For the two partitions below, the letter represents the record and the number is the timestamp. For this example we'll assume that `punctuate` is scheduled to run every 5 seconds.



Since Partition A has the smallest timestamp it's chosen first

- 1) Process called with record A
- 2) Process called with record B

Now Partition B has the smallest timestamp

- 3) Process called with record C
- 4) Process called with record D

Switch back to Partition A has smallest time stamp again.

- 5) Process called with record E
- 6) Punctuate called because time elapsed from timestamps is 5 seconds
- 7) Process called with record F

Finally switch back to Partition B

- 8) Process called with record G
- 9) Punctuate called again as 5 more seconds elapsed by timestamps

**Figure 6.6 Punctuation Scheduling**

Using the image above to guide us let's walk through how the schedule is determined:

1. The `StreamTask` extracts the *smallest* timestamp from the `PartitionGroup`. The `PartitionGroup` is a set of partitions for a given `StreamThread` and contains all timestamp information for all partitions in the group.
2. During processing of records, the `StreamThread` iterates over all of its `StreamTask` object and each task will end up calling `punctuate` for each of its processors that are eligible for punctuation. If you recall, we collect a minimum of 20 trades before we examine an individual stock's performance.
3. If the scheduled timestamp from the processor (the last time executing `punctuate` plus the scheduled time) is less than or equal to extracted timestamp from the `PartitionGroup`, then we call that processors `punctuate` method.

The key point here is the application advances timestamps via the `TimestampExtractor`, so `punctuate` calls are consistent only with data arriving at a constant rate.

The bottom line on `PunctuationType.STREAM_TIME` is if your flow of data is sporadic, the `punctuate` method will not get executed at the regularly scheduled intervals.

With `PunctuationType.WALL_CLOCK_TIME` on the other hand, the execution of the `Punctuator#punctuate` is more predictable, as it uses wall-clock time. Note system-time semantics is best-effort, as the wall-clock time is advanced in the polling

interval and the granularity is dependant on how long it takes to complete a polling cycle. So with our example above, we can expect our punctuation activity will get executed closer to every ten seconds, regardless of data activity.

The decision of which approach to use is entirely dependent on your needs. If you need some activity performed on a regular basis regardless of data flow, then system time is probably the best bet. But on the other hand, if you only need calculations performed from incoming data and some lag time between executions is acceptable, then try stream time semantics.

**NOTE      Deprecated Punctuation Methods**

Before Kafka 0.11.0, punctuation involved the `ProcessorContext#schedule(long time)` method which in turn called the `Processor#punctuate` method at the scheduled interval. This approach only worked on stream time semantics, and both methods are now deprecated. I'll try to mention deprecated methods in the book, but we are only going to use the latest punctuation methods in our examples.

Now we've covered scheduling and punctuation, let's move on to how we handle incoming records.

### THE PROCESS METHOD

In the `process()` method is where we perform all of our calculations to evaluate stock performance. We have several steps to take when we receive a record.

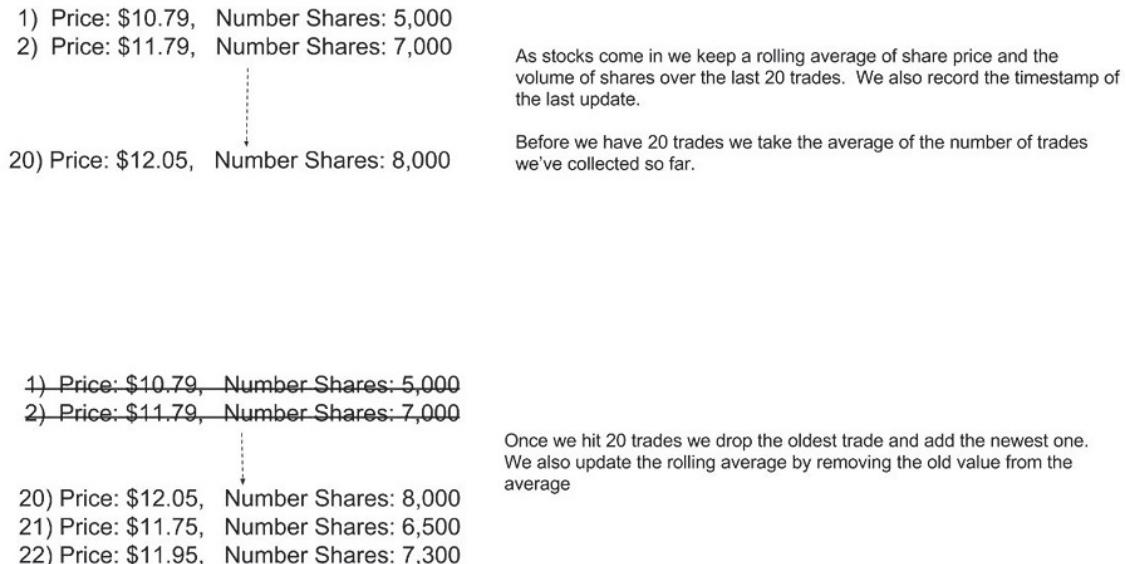
1. First, we check the state store to see if we have a corresponding `StockPerformance` object for the given stock ticker symbol.
2. If the store does not contain the object, we'll record the current share price and volume.
3. Otherwise, we'll add the current share price and share volume to the object and update our calculations.
4. We only start performing calculations once we hit twenty transactions for any given stock.

While financial analysis is beyond the scope of this book, we should take a minute to explain our calculations. For both the share price and volume we are going to perform a simple moving average. In the financial trading world, a simple moving average or SMA is used to calculate an average for data points of size N.

For our purposes, we are going to set N to twenty. Setting a max size means as new trades come in we'll collect the share price and amount of shares traded up to twenty. Once we hit that threshold, we'll remove the oldest value and add the latest one. Using the SMA, we'll get a "rolling average" of stock price and volume over the last 20 trades. It's important to note we won't have to recalculate the entire amount as new values come in.

Now let's take a look at the following image for a high-level walk through of our process method, with the mindset of what we would do if we were to perform these steps

manually:



**Figure 6.7 Stock Analysis Process**

As we can see from our walk through the process method is where we perform all of the calculations. Now let's take a look at the code that makes up our process method:

**Listing 6.6 Process Implementation found in  
src/main/java/bbejeck/chapter\_6/processor/StockPerformanceProcessor.java**

```

@Override
public void process(String symbol, StockTransaction transaction) {
  StockPerformance stockPerformance = keyValueStore.get(symbol); ①

  if (stockPerformance == null) {
    stockPerformance = new StockPerformance(); ②
  }

  stockPerformance.updatePriceStats(transaction.getSharePrice()); ③
  stockPerformance.updateVolumeStats(transaction.getShares()); ④
  stockPerformance.setLastUpdateSent(Instant.now()); ⑤

  keyValueStore.put(symbol, stockPerformance);
}
  
```

- ① Retrieving previous performance stats, possibly null.
- ② Creating a new StockPerformance object when not in the state store.
- ③ Updating the price statistics for this stock.
- ④ Updating the volume statistics for this stock.
- ⑤ Setting the timestamp of the last update.
- ⑥ Placing the updated StockPerformance object into the state store.

In our process method, we take the latest share price and the number of shares involved in the transaction and add them to the `StockPerformance` object. Notice that we've abstracted all details of how we perform the updates inside the `StockPerformance` object. Keeping most of the business logic out of the processor is a good idea, and we'll come back to that point when we cover testing in Chapter eight.

Let's briefly talk about the two key calculations performed, determining the moving average and the differential of stock price/volume from the respective current average. First, we don't want to calculate an average until we've collected data from twenty transactions. So we defer doing anything until the processor receives twenty trades. When getting data from twenty trades for an individual stock, we calculate our first average. Then we'll take the current value of the stock price or a number of shares and divide by our moving average, converting the result to a percentage.

**NOTE****Stock Performance Code**

For those that want to see the calculations, the `StockPerformance` code can be found in `src/main/java/bejeck/model/StockPerformance.java`

In the previous Processor example, once we worked our way through the `process` method, we forwarded our records downstream. But in this case, we store our final results in the state store and leave the forwarding of records to the `Punctuator#punctuate` method.

### 6.3.3 The Punctuator Execution

Having discussed the punctuation semantics and scheduling previously, we'll jump straight into the code for the `Punctuator#punctuate` method:

**Listing 6.7 Punctuation code found in  
`src/main/java/bejeck/chapter_6/processor/punctuator/StockPerformancePunctuator`**

```
@Override
public void punctuate(long timestamp) {
    KeyValueIterator<String, StockPerformance> performanceIterator = keyValueStore.all(); ①

    while (performanceIterator.hasNext()) {
        KeyValue<String, StockPerformance> keyValue = performanceIterator.next();
        String key = keyValue.key;
        StockPerformance stockPerformance = keyValue.value;

        if (stockPerformance != null) {
            if (stockPerformance.priceDifferential() >= differentialThreshold || ②
                stockPerformance.volumeDifferential() >= differentialThreshold) { ③
                context.forward(key, stockPerformance);
            }
        }
    }
}
```

- ① Retrieving the iterator to go over all key values in the state store.

- ② Checking the threshold for the current stock.
- ③ If we've met or exceeded the threshold, forward the record.

The procedure in the `Punctuator#punctuate` method is simple. We iterate over the key-value pairs in the state store, and if the value has crossed over our pre-defined threshold, we forward the record downstream.

An important concept to remember here is before we relied on a combination of committing or cache flushing to forward records, now we defined our terms when records get forwarded. Additionally, even though we expect to execute this code every 10 seconds, it does not guarantee we emit records, they must meet the differential threshold. I should also note at this point the `Processor#process` and `Punctuator#punctuate` methods are not called concurrently.

**NOTE** **Kafka Streams does not share data**

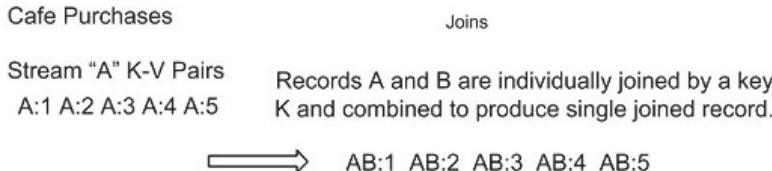
While we're demonstrating access to a state store, it's a good time to review Kafka Streams architecture and go over a few key points. Each `StreamTask` has its *own* copy of a *local* state store and `StreamThread` objects don't share tasks or data. Lastly, as records make their way through the topology, each node is visited in a depth-first manner, meaning there is never concurrent access to state stores from any given processor.

While this example worked to give us an excellent introduction to writing a custom processor, we can take writing custom processors a bit further by adding a new data structure and an entirely new way of aggregating data that does not currently exist in the API. With this in mind, we'll move on to adding a Cogroup processor.

## 6.4 The CoGroup Processor

Back in Chapter four, we discussed joins between two streams; specifically, we joined purchases from different departments within a given timeframe to promote business.

We use joins to bring together records with the same key and arriving in the same time window. With joins there's an implied one-to-one mapping of records from stream A to stream B. The following image depicts this relationship clearly:



Electronics Purchases

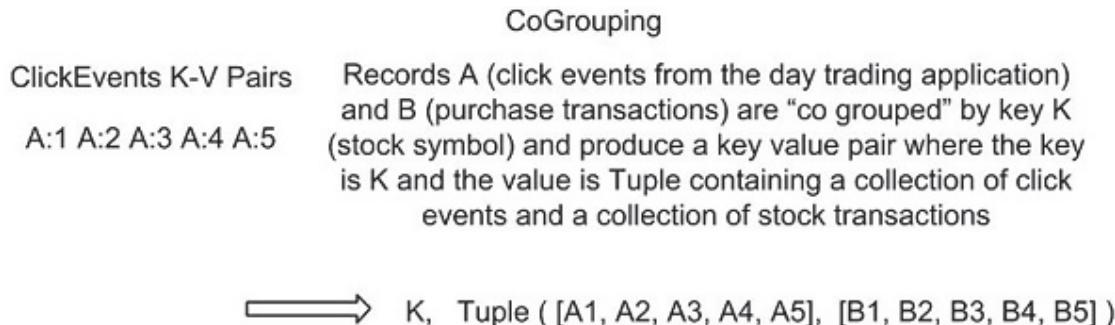
Stream "B" K-V Pairs  
B:1 B:2 B:3 B:4 B:5

For this example we are assuming the window time lines up to give us only one to one joins, so each record uniquely matches up only one other record.

**Figure 6.8 Records A and B joined by a common key**

Now let's imagine you want to do a similar type of analysis, but instead of a one-to-one join by key, you want two collections of data by a common key, or a *co-grouping* of data.

You are the manager of a popular online day-trading application. Day traders are on your application for several hours a day, sometimes the entire time the stock market is open. One of the metrics your application tracks is the notion of "events." You have defined an event as when a user clicks on the ticker symbol to read more about a company and its financial outlook. You want to do some deeper analysis between user clicks on the application and when users purchase stocks. But you want course-grained results, by comparing multiple clicks and purchases to determine some overall pattern. What you want is a tuple with two collections of each event type by the company trading symbol as pictured here:



StockPurchase K-V Pairs

B:1 B:2 B:3 B:4 B:5

For this example we populate each collection with what's available at the time punctuate is called, there might be an empty collection at any point

**Figure 6.9 Output of a Key with a Tuple containing two collections of data, a co-grouped result.**

Your desire is to combine "snapshots" of click events and transactions for a given company every N seconds. But you are not waiting for records from either stream to arrive. When the specified amount of time goes by, you want a "co-grouping" of click events and transactions by company ticker symbol. If either type of event is not present, then one of the collections in the tuple is empty. For those readers familiar with Apache Spark or Apache Flink this functionality is similar to [PairRDDFunctions.cogroup](#) method and [CoGroupDataSet](#) respectively.

Next, let's start walking through the steps we'll take in constructing this processor.

#### 6.4.1 Building the CoGroup Processor

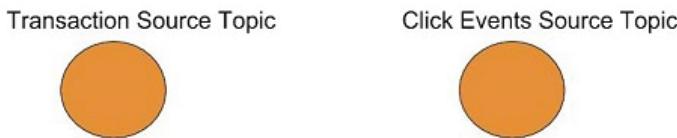
To create the co-grouping processor we have a few pieces we need to tie together:

1. Define sources for two topics (stock-transactions, click-events)
2. Add two processors to consume records from the topics
3. A third processor is acting as an aggregator/co-grouping for the two processors above.
4. A state-store for the aggregating processor to keep the state of both events.
5. A sink node to write our results (and/or a printing processor to print results to console).

Now we'll now walk through the steps to put this processor together.

#### DEFINING THE SOURCE NODES

By this point we are familiar with the first step, creating the source nodes. This time we'll create two source nodes to support reading both the click-event stream and the stock transaction stream. To help us keep track of where we are in the topology, we'll build on the following diagram for a graphical display of the topology:



**Figure 6.10 Cogrouping Source Nodes**

The accompanying code for creating our source nodes looks like this:

**Listing 6.8 Source nodes for CoGrouping Processor found in  
src/main/java/bbejeck/chapter\_6/CoGroupingApplication.java**

```
//Note that we've left out configuration and (de)serializer creation for clarity.
topology.addSource("Txn-Source", stringDeserializer, stockTransactionDeserializer,
    "stock-transactions")
    .addSource("Events-Source", stringDeserializer, clickEventDeserializer, "events")
```

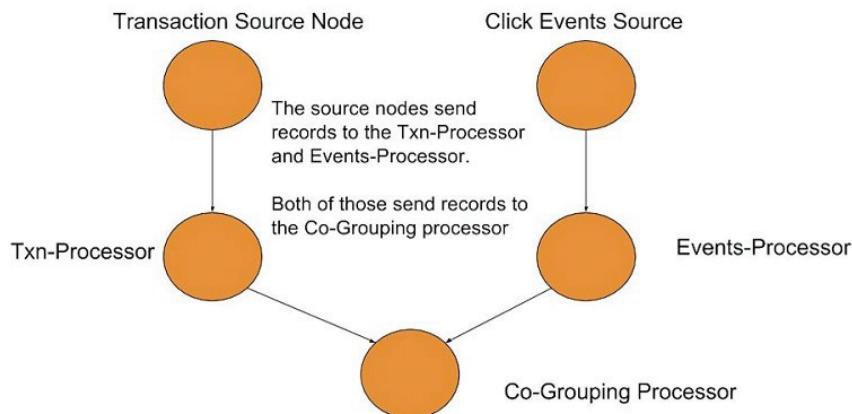
①  
②

- ① Source node for the stock transactions topic.
- ② Source node for the click-events topic.

With the sources for the topology in place, let's move on to the next step.

## ADDING THE PROCESSOR NODES

Now we'll add the workhorses of the topology, the processors. The topology graph is updated accordingly:



**Figure 6.11 Adding Processor Nodes**

And let's look at the code for adding these new processors:

**Listing 6.9 The Processor Nodes found in  
src/main/java/bbejeck/chapter\_6/CoGroupingApplication.java**

```

.addProcessor("Txn-Processor", StockTransactionProcessor::new, "Txn-Source")  

.addProcessor("Events-Processor", ClickEventProcessor::new, "Events-Source")  

.addProcessor("CoGrouping-Processor", CogroupingProcessor::new, "Txn-Processor",  

    "Events-Processor")
    
```

- 1
- 2
- 3

- ① Adding the StockTransactionProcessor
- ② Adding the ClickEventProcessor
- ③ Adding the CogroupingProcessor which is child node of both processors

In listings 1 and two above the parent names are the names of source nodes reading from the "stock-transactions" and "events" topics respectively. But you'll notice the third processor (listing 3) has names of both processors given as the parent nodes. This means both processors will feed the aggregation processor.

For the `ProcessorSupplier` instances we're again taking a Java 8 shortcut. This time we've shortened the form, even more, to use a method handle, in this case, it's a constructor call to create the associated processor.

**TIP****Why use a method handle?**

With single method, no-arg interfaces in Java 8 we can use a lambda in the form of () doSomething. But since the ProcessorSupplier's only role is to return a (possibly new each time) Processor object we can shorten the form, even more, to use a method handle for the constructor of the Processor type. Note this only applies for no-arg constructors.

Now, let's take a second to explain why we have setup of our processors in this manner. This example is an aggregation operation, and the role of the StockTransactionProcessor and ClickEventProcessor is to wrap their respective object into smaller aggregate objects and then forward to another processor for a total aggregation. So both the StockTransactionProcessor and the ClickEventProcessor perform the smaller aggregation and forward their records to the CogroupingProcessor.

Then the CogroupingProcessor performs the co-grouping; then it forwards on regular intervals (intervals driven by timestamps that is) to an output topic. With this description now let's take a look at the code for the processors.

**Listing 6.10 The StockTransactionProcessor found in  
src/main/java/bbejeck/chapter\_6/processor/cogrouping/StockTransactionProcessor.java**

```
public class StockTransactionProcessor extends AbstractProcessor<String, StockTransaction> {
    @Override
    @SuppressWarnings("unchecked")
    public void init(ProcessorContext context) {
        super.init(context);
    }

    @Override
    public void process(String key, StockTransaction value) {
        if (key != null) {
            Tuple<ClickEvent, StockTransaction> tuple = Tuple.of(null, value); ①
            context().forward(key, tuple); ②
        }
    }
}
```

- ① Creating an aggregate object with the StockTransaction
- ② Forwarding the tuple to the CogroupingProcessor

As you can see, the StockTransactionProcessor adds the StockTransaction to the aggregator (the Tuple) and forwards the record.

**NOTE****Tuple is a custom class**

The Tuple<L, R> shown here is a custom object for examples in this book and is found in src/main/java/bbejeck/util/collection/Tuple.java

Now let's take a look at the `ClickEventProcessor` code

**Listing 6.11 The ClickEventProcessor found in  
src/main/java/bbejeck/chapter\_6/processor/cogrouping/ClickEventProcessor.java**

```
public class ClickEventProcessor extends AbstractProcessor<String, ClickEvent> {

    @Override
    @SuppressWarnings("unchecked")
    public void init(ProcessorContext context) {
        super.init(context);
    }

    @Override
    public void process(String key, ClickEvent clickEvent) {
        if (key != null) {
            Tuple<ClickEvent, StockTransaction> tuple = Tuple.of(clickEvent, null); ①
            context().forward(key, tuple); ②
        }
    }
}
```

- ① Adding the ClickEvent to the initial aggregator object
- ② Forwarding the tuple to the CogroupingProcessor

Now, the `ClickEventProcessor` adds the `ClickEvent` to the `Tuple` aggregator, very similar to the previous example. Now to complete the picture of how we perform the aggregation we need to look at the `CogroupingProcessor` code. The code for the `CogroupingProcessor` is a more involved, so we'll break up our review of the code by method, starting with `CogroupingProcessor#init()`:

**Listing 6.12 The CogroupingProcessor init method found in  
src/main/java/bbejeck/chapter\_6/processor/cogrouping/AggregatingProcessor.java**

```
public class CogroupingProcessor extends AbstractProcessor<String,
    Tuple<ClickEvent, StockTransaction>> {

    private KeyValueStore<String, Tuple<List<ClickEvent>, List<StockTransaction>>> tupleStore;
    public static final String TUPLE_STORE_NAME = "tupleCoGroupStore";

    @Override
    @SuppressWarnings("unchecked")
    public void init(ProcessorContext context) {
        super.init(context);
        tupleStore = (KeyValueStore) context().getStateStore(TUPLE_STORE_NAME); ①
        CogroupingPunctuator punctuator = new CogroupingPunctuator(tupleStore, context()); ②
        context().schedule(15000L, STREAM_TIME, punctuator); ③
    }
}
```

- ① Retrieving the configured state store
- ② Creating a Punctuator instance `CogroupingPunctuator`, handling all scheduled calls

### ③ Scheduling to call Punctuator#punctuate method every 15 seconds

As we'd expect, the `init` method handles the details of setting up the class. We grab the state store configured in the main application and save it in a variable for use later on. We create the `CogroupingPunctuator` to handle the scheduled punctuations calls.

#### TIP

#### Method handles for Punctuator

We can also specify a method handle for the `Punctuator` instance. Simply declare a method in the processor accepting a single parameter of type `long` and a `void` return type. You then schedule punctuation like this

```
context().schedule(15000L, STREAM_TIME,
this::myPunctuationMethod); For an example of this look in
src/main/java/bbejeck/chapter_6/processor/cogrouping/
CogroupingMethodHandleProcessor.java
```

We're scheduling `punctuate` for every 15 seconds. Since we are using `PunctuationType.STREAM_TIME` semantics, the timestamps in the *arriving* data drive the calls to punctuate. Remember if our flow of data is not relatively constant, we may have more than 15 seconds between calls to `Punctuator#punctuate`.

#### NOTE

#### Additional examples in the source code

If you recall from our previous discussion on punctuate semantics we have two choices: `PunctuationType.STREAM_TIME` and `PunctuationType.WALL_CLOCK_TIME`. In this example, we are using `STREAM_TIME` semantics. There's an additional processor example showing `WALL_CLOCK_TIME` semantics in `src/main/java/bbejeck/chapter_6/processor/cogrouping/CogroupingSystemTimeProcessor.java` so you can observe any differences in performance/behavior.

Next, let's move on and look at how `CogroupingProcessor` performs one of its main tasks in the `process` method:

#### Listing 6.13 The CogroupingProcessor process method found in src/main/java/bbejeck/chapter\_6/processor/cogrouping/CogroupingProcessor.java

```
@Override
public void process(String key, Tuple<ClickEvent, StockTransaction> value) {

    Tuple<List<ClickEvent>, List<StockTransaction>> cogroupedTuple = tupleStore.get(key);
    if (cogroupedTuple == null) {
        cogroupedTuple = Tuple.of(new ArrayList<>(), new ArrayList<>()); ①
    }

    if (value._1 != null) {
        cogroupedTuple._1.add(value._1); ②
    }
}
```

```

        if (value._2 != null) {
            cogroupedTuple._2.add(value._2);
        }

        tupleStore.put(key, cogroupedTuple);
    }
}

```

③  
④

- ① Initialize the total aggregation if it doesn't exist yet
- ② If the ClickEvent is not null add it to the list of click events
- ③ If the StockTransaction is not null add it to the list of stock transactions
- ④ Place the updated aggregation into the state store

As we process incoming smaller aggregates of our overall co-grouping, our first step checking if we have an instance in our state store already if we don't we create a Tuple with empty collections of ClickEvent and StockTransaction.

Next, we check the incoming, smaller aggregation and if either of a ClickEvent or StockTransaction is present, we add to the overall aggregation.

The last step in the process method is putting the Tuple back into the store, updating our aggregation total.

#### NOTE

#### Non-Concurrent Access

While we have two processors forwarding records to one processor and accessing one state store, we don't have to be concerned with concurrency issue. Remember parent processors forward records to child-processors in a depth-first manner, so in this case, each parent processor serially calls the child processor. Additionally, Kafka Streams only uses one thread per task, so there is never any concurrency access issues.

Our next step is to look at how punctuation is handled (in this case we are using the updated API, so we won't look at the Processor#punctuate call as that is deprecated).

#### **Listing 6.14 The CogroupingPunctuator punctuate method found in src/main/java/bbejeck/chapter\_6/processor/cogrouping/CogroupingPunctuator.java**

```

// leaving out class declaration and constructor for clarity
@Override
public void punctuate(long timestamp) {
    KeyValueIterator<String, Tuple<List<ClickEvent>, List<StockTransaction>>>
        iterator = tupleStore.all();

    while (iterator.hasNext()) {
        KeyValue<String, Tuple<List<ClickEvent>, List<StockTransaction>>>
            cogrouping = iterator.next();

        // if either list contains values forward results
        if (cogrouping.value != null && (!cogrouping.value._1.isEmpty() ||

```

①  
②

```
!cogrouping.value._2.isEmpty())) {  
    List<ClickEvent> clickEvents = new ArrayList<>(cogrouping.value._1);  
    List<StockTransaction> stockTransactions = new ArrayList<>(cogrouping.value._2);  
  
    context.forward(cogrouping.key, Tuple.of(clickEvents, stockTransactions));  
    cogrouped.value._1.clear();  
    cogrouped.value._2.clear();  
    tupleStore.put(cogrouped.key, cogrouped.value);  
}  
}  
}
```

- ① Getting iterator of all cogroupings in the store
  - ② Retrieving the next cogrouping
  - ③ Ensuring the value is not null, and either collection contain data
  - ④ Making defensive copies of cogenerated collections
  - ⑤ Forward the key and aggregated co-grouping
  - ⑥ Putting cleared out tuple back into the store

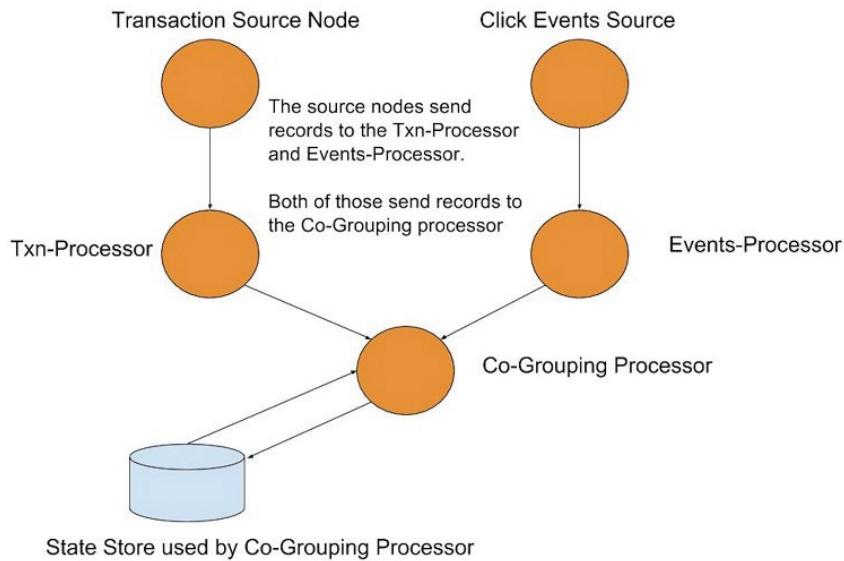
So during each punctuate call we retrieve all of the stored records in a `KeyValueIterator`, and we start to pull out each cogrouped result contained in the iterator.

Then we make defensive copies of the collections and create a new cogrouped Tuple and forward it downstream. In this case, we send the cogrouped results to a sink node. Finally, we remove the current cogrouped results and store the tuple back in the store ready for the next round of records to arrive.

Now that we've covered the co-grouping functionality let's complete building our topology.

## ADDING THE STATE STORE

As we've seen so far, the ability to perform aggregations in a Kafka streaming application requires having state. So we'll need to add a state store to the `CogroupingProcessor` for it to function properly. Let's take a look at our updated topology:



**Figure 6.12 Adding a state store to the co-grouping processor in the topology**

Now let's take a look at the code for adding the state store:

**Listing 6.15 Adding a State Store node found in  
src/main/java/bbejeck/chapter\_6/CoGroupingApplication.java**

```
// this comes earlier in source code, included here for context
Map<String, String> changeLogConfigs = new HashMap<>();
changeLogConfigs.put("retention.ms", "120000" );
changeLogConfigs.put("cleanup.policy", "compact,delete");

KeyValueBytesStoreSupplier storeSupplier =
    Stores.persistentKeyValueStore(TUPLE_STORE_NAME);
    StoreBuilder<KeyValueStore<String, Tuple<List<ClickEvent>,
```

1

```
        List<StockTransaction>>> storeBuilder = ③
            Stores.keyValueStoreBuilder(storeSupplier,
                Serdes.String(),
                eventPerformanceTuple).withLoggingEnabled(changeLogConfigs);
```

2

```
    .addStateStore(storeBuilder, "CoGrouping-Processor")
```

4

5

- ① Specifying how long to keep records, and use compaction and delete for cleanup
- ② Creating the store supplier for a persistent store (RocksDB)
- ③ Creating the store builder
- ④ Adding the changelog configs to the store builder
- ⑤ Adding the store to the topology with the name of processor that will access the store

As we can see here, we're adding a persistent state store. We've decided to use a persistent store, because we might get infrequent updates for some keys. With the

in-memory and LRU based stores, infrequently used keys and values might eventually be removed. We want to ability to retrieve information for any key we've worked with before.

**TIP****Configuring the change log**

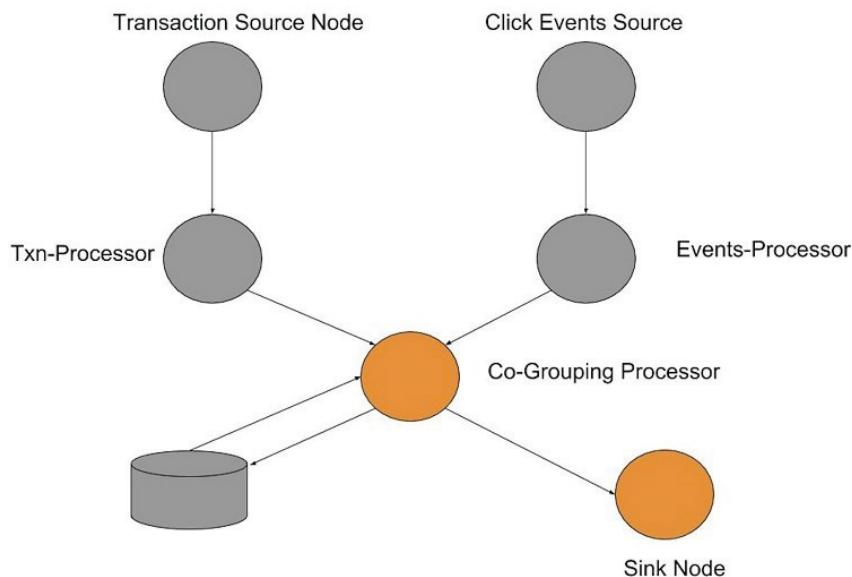
In listing four we are adding specific configurations for the state store, this is done so we can keep the size of the change log at a manageable size. Remember we can configure change log topics with any valid topic configuration.

This code is very straight forward, so we'll keep the comments to minimum. One last point to notice is we've named the "CoGrouping-Processor" as the only processor with the ability to access this store.

At this point, we have one step left to complete our topology, the ability to read the results of our co-grouping.

**ADDING THE SINK NODE**

For our co-grouping topology to be of use to us, we need to write our data out to a topic (or the console). Let's update our topology one more time:



**Figure 6.13 Adding a sink node completes the co-grouping topology.**

Now our co-grouped aggregation results are written out to a topic for use in further analysis. Please note that in several examples I will talk about adding a sink node, but in the source code there is a sink that writes to the console, with the sink writing out to a topic is commented out. For development purposes, I use sink node writing to topic or stdout interchangeably.

### Listing 6.16 The sink node and a printing processor found in src/main/java/bbejeck/chapter\_6/CoGroupingApplication.java

```
.addSink("Tuple-Sink", "cogrouped-results", stringSerializer, tupleSerializer,
    "CoGrouping-Processor"); ①  

topology.addProcessor("Print", new KStreamPrinter("Co-Grouping"),
    "CoGrouping-Processor"); ②
```

- ① The sink node writing our co-grouped tuples out to a topic.
- ② A processor that prints results to stdout for use during development.

In this final piece of our topology we add a sink node as a child of the `CoGrouping-Processor` which writes the co-grouping results out to a topic. In the line below we add an additional processor for printing results to the console for development, which is also a child node of the `CoGrouping-Processor`. Remember with the Processor API the order of defining nodes does not establish the parent-child relationship. The parent-child relationship is determined by providing the name(s) of previously defined processors.

We've reached the end of building our co-grouping processor. The key point I want you to remember from this section is that while you have more code to write the Processor API gives you the flexibility to create virtually any kind of streaming topology you need.

Let's wrap up this chapter with how we can integrate some Processor API functionality into a KStreams application.

## 6.5 Integrating the Processor API and the Kafka Streams API

So far, our coverage on the KStream and the Processor API have been separate, but this is not to say that we can't combine approaches. Why would you want to mix the two approaches?

Let's say you've used both the KStream and Processor API for a while. You've come to prefer the KStream approach, but you want to include some of your previously defined processors in a KStream application because it provides some of the lower lever control you need.

The Kafka Streams API offers three methods `KStream#process`, `KStream#transform`, and `KStream#transformValues` which allows us to plug in functionality built using the Processor API. We already have some experience with this approach, as we worked with the `ValueTransformer` in Chapter 4 section 4.2.3.

The `KStream#process` method creates a terminal node, while the `KStream#transform` (or `KStream#transformValues`) method returns a new `KStream` instance allowing us to continue adding processors to that node. We should also note the

transforms methods are stateful, so we also provide a state store name when using the transform methods. Since the `KStream#process` results in a terminal node, in most cases, you'll want to use either `KStream#transform` or `KStream#transformValues`.

From there you can replace your `Processor` with a `Transformer` instance. The main difference between the two interfaces is the `Processor` main action method is `process` which has a `void` return while the `Transformer` uses `transform` and expects a return type of `R`. Both offer the same punctuation semantics.

In most cases, replacing the `Processor` is a matter of taking the logic from the `Processor#process` method and placing it in the `Transformer#transform` method. Remember you'll need to account for returning a value, but returning null and forwarding results with the `ProcessorContext#forward` is an option.

#### TIP Specifying the return of multiple values

As we stated here the transformer returns a value. In this case, we return a null, which gets us around the requirement to return a value. If we wanted to return multiple values instead, we'd return a `List<KeyValue<K,V>>` and then attach `flatMapValues` afterward to send individual records downstream. An example of this can be found in `src/main/java/bbejeck/chapter_6/StockPerformanceStreamsAndProcessorMultipleValuesApp.java`.

From there you would just plug in the `Transformer` (or `ValueTransformer`) instance using the `KStream#transform` or `KStream#transformValues` method.

A great example of combining the `KStream` and `Processor` API can be found in `src/main/java/bbejeck/chapter_6/StockPerformanceStreamsAndProcessorApplication.java`. I didn't present the example here as the logic is for the most part identical to our first example, `StockPerformanceApplication` so it's left as an exercise for you to look up. Additionally, there is a `Processor` API version of the original `ZMart` application in `src/main/java/bbejeck/chapter_6/ZMartProcessorApp.java`.

## 6.6 Conclusion

We've reached the end of our coverage of the low-level `Processor` API. While we covered a lot of ground here are the key points I want you to take away from this chapter:

- The `Processor` API gives you more flexibility at the cost of more code.
- Although the `Processor` API is more verbose than the Kafka Streams API, it's still easy to use, and the `Processor` API is what the Kafka Streams API uses itself under the covers.
- When faced with a decision on which API to use, consider using the Kafka Streams API and integrate lower-level method (`process`, `transform`, `transformValues`) when needed.

At this point in the book, we've covered how we can *build* applications with Kafka Streams. Our next step is to talk about how we can optimally configure these applications and monitor them for maximum performance and spotting potential issues.

# P3

## *Part 3: Administering Kafka Streams*

Now you'll take all your new skills in building a Kafka Streams application and shift your focus on how you can measure the performance of your application and monitor it. Additionally, you'll learn how to test your Kafka Streams code so that you know it's doing the right thing and will work as expected and gracefully handle errors.



# Monitoring and Performance

In this chapter:

- Basic Kafka Monitoring
- Intercepting Messages
- Measuring Performance
- Observing the State of the application.

So far we've learned how to build a Kafka Streams application from the bottom up. We've heard about the high level Kafka Streams DSL and have seen the power of using a declarative API. We've also learned about the Processor API and have seen how we can give up some flexibility to gain more control in writing our streaming applications.

At this point, it's time for us to change gears a little bit. Now we are going to put on our forensic investigator hats and dig into our application from a different perspective. Our focus is going to shift from *how* we get things to work to more of *what* is going on. In some respects, the building or the original creation of an application is the most comfortable part. Getting the application to run successfully is always the more significant challenge.

There's a practical reason for observing what our application is doing. We need to make sure the answers we are getting are correct. At any point in the application, we need to observe what goes on under the hood either to confirm our processors are performing correctly, to locate performance bottlenecks or to maintain a constant flow of information.

Building an application is one thing, but it's another thing to be able to it scale correctly and ensuring that everything is working as expected. When you are writing an application, despite your best efforts, there is always a situation you didn't account for.

In this chapter, we learn how to check the running state of our Kafka Streams application. We'll also show how we can measure the performance of the application so that we can spot performance bottlenecks. Also, we'll also demonstrate techniques for getting notified of various states of the application and viewing the structure of the topology.

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

We'll find what metrics are available, how we collect metrics, and how we can observe the collected metrics as the application is running.

Let's start with monitoring our Kafka Streaming application.

## 7.1 Basic Kafka Monitoring

Since the Kafka Streams API is a part of Kafka, it goes without saying that to monitor your application will require some monitoring of Kafka as well. Full blown surveillance of a Kafka cluster is a big topic, so we'll keep our discussion on Kafka performance where the two meet; this means we'll talk about monitoring Kafka consumers and producers. [More information on monitoring a Kafka cluster](#) can be found in the documentation.

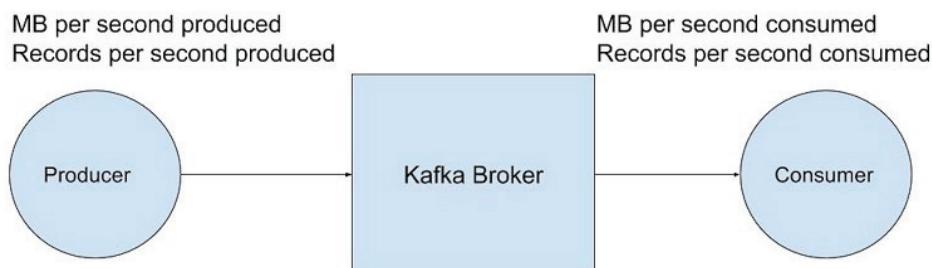
### NOTE

#### Measuring Kafka Streams includes Kafka

One thing we should note here is that to measure Kafka Streams performance also means we need to include measuring Kafka itself. So at times some of our coverage of performance will bleed over to the Kafka side of things. But since this is a book on Kafka Streams we're going to focus our coverage on Kafka Streams.

### 7.1.1 Measuring Consumer and Producer performance

For our discussion on consumer and producer performance let's start by taking a look at the following image depicting one of the fundamental performance concerns for a producer and consumer:



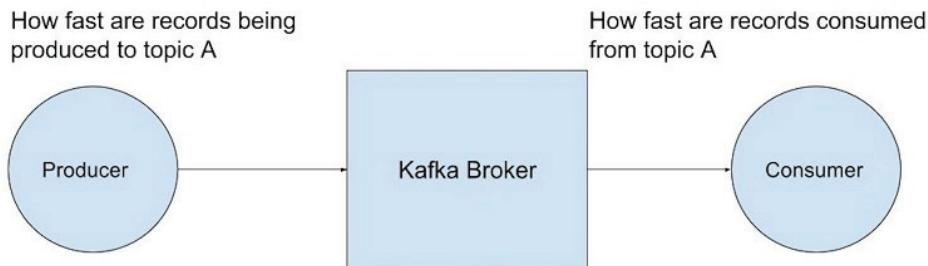
**Figure 7.1 Kafka producer and consumer performance concerns writing to and reading from a broker**

As we can see from the image above, producer and consumer performance are very similar in that both are concerned with throughput. But where we put the emphasis is just different enough to be considered two sides of the same coin. Let me explain a little what I mean here.

For producers, we care mostly about how fast the producer is sending messages to the broker. Obviously the higher the throughput, the better.

Now for consumers, we are also concerned with performance or how fast we can read the messages from a broker. But there is another way to measure consumer performance, consumer lag. To help with our explanation here let's review the image of measuring

producer and consumer throughput with slightly different annotations this time:

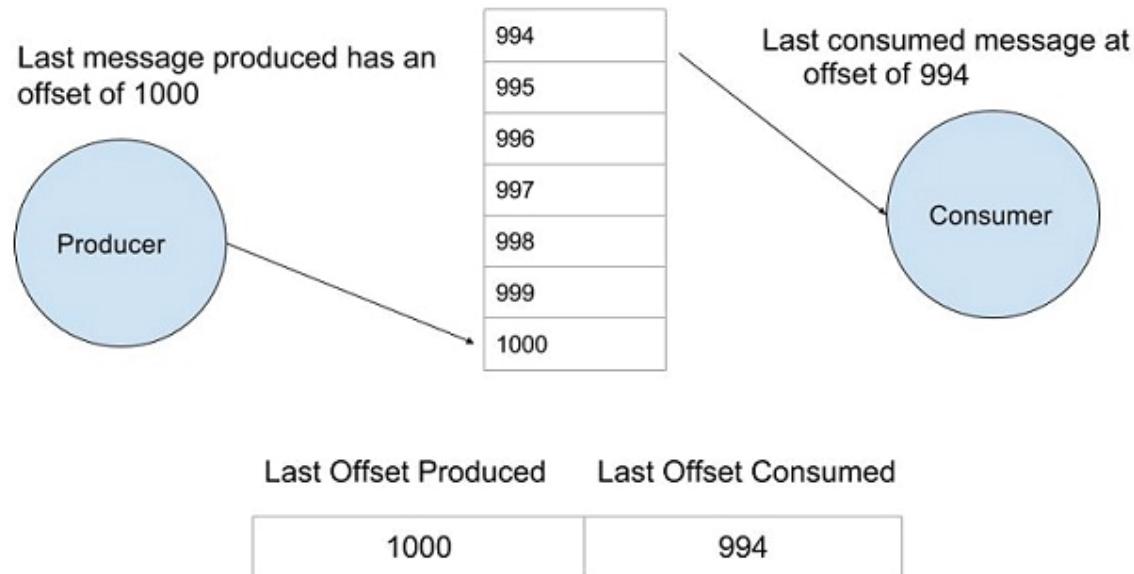


Does the consumer keep up with the rate of records from the producer?

**Figure 7.2 Kafka producer and consumer performance revisited**

So with these different annotations, we can see we care about how much and how fast our producers can publish to a broker while we simultaneously care about how quickly our consumers can read those messages from the broker.

The difference between how fast the producers place records on the broker and when consumers read those messages is called consumer lag. In Figure 7.3 you can see that consumer lag is the difference between the last committed offset from the consumer and the latest offset from a message written to the broker.



The difference between the last offset produced and the last committed consumed offset (from the same topic) is known as consumer lag.

In this case we can see the consumer “lags” behind the producer by 6 records

**Figure 7.3 Consumer lag is the difference in offsets committed by the consumer and offsets written by the producer**

While there is bound to be some lag from the consumer, ideally you'd like to see the consumer catch up or at least have the lag remain constant and not growing over time.

Now that we have defined our performance parameters for producers and consumers let's see how we can go about monitoring those for performance or troubleshooting issues.

### 7.1.2 Checking For Consumer Lag

To check for consumer lag, Kafka provides a convenient command line tool `kafka-consumer-groups.sh` found in the `<kafka-install-dir>/bin` directory. The script has a few options, but here we want to focus on the list and describe options as these two options give us the information we need on consumer group performance. First, let's use the list command to find all active consumer groups using this command:

#### **Listing 7.1 Listing Active Consumer Groups**

```
<kafka-install-dir>/bin/kafka-consumer-groups.sh --bootstrap-server localhost:9092 --list
```

You can view the results of running the command above in this image:

```
oddball:bin bbejeck$ ./kafka-consumer-groups.sh --list --bootstrap-server localhost:9092
Note: This will only show information about consumers that use the Java consumer API (non-ZooKeeper-based consumers).

console-consumer-59026
```

**Figure 7.4 Listing available consumer groups from the command line**

With the information above we can choose a consumer group name and now run this command:

#### **Listing 7.2 Getting Status of a Consumer Group**

```
<kafka-install-dir>/bin/kafka-consumer-groups.sh
--bootstrap-server localhost:9092 --group <GROUP-NAME> --describe
```

Then we'll get a status on how this consumer is performing as shown here:

Number of messages read = 3      Number of messages sent to topic = 10

```
oddball:bin bbejeck$ ./kafka-consumer-groups.sh --bootstrap-server localhost:9092 --group console-consumer-59026 --describe
Note: This will only show information about consumers that use the Java consumer API (non-ZooKeeper-based consumers).

Consumer group 'console-consumer-59026' has no active members.

TOPIC          PARTITION  CURRENT-OFFSET  LOG-END-OFFSET  LAG    CONSUMER-ID  HOST        CLIENT-ID
consumer_log_demo  0           3                  10             7     -           -           -
```

$$10 \text{ (sent messages)} - 3 \text{ (messages read)} = 7 \text{ (Lag or records behind)}$$

**Figure 7.5 Status of a consumer group**

From the results here we can see we have a small consumer lag. Having a consumer lag is not always indicative of a problem, as consumers read messages in batches and will

not retrieve another batch until they are finished processing the current batch. Processing the records takes time, so a little lag is not entirely unsurprising. However, a lag that grows over time is something that we should investigate.

The critical point here is a small lag or one that stays constant is ok, but a lag that continues to grow over time is an indication you'll need to give your consumer more resources.

For example, it could be you need to increase the partition count and hence increase the number of threads consuming from the topic. Another consideration is your processing after reading the message is too heavyweight, and after consuming a message hand it off to an async queue where another thread can pick up the message and do the processing at that point.

In this section, you've learned how to observe how quickly a consumer is reading messages from the broker. Next, we'll dig a little deeper into observing behavior for debugging purposes by showing how to intercept what the producers and consumers are sending and receiving (respectively) *before* your Kafka Streams application sends or consumes records.

### **7.1.3 Intercepting the Producer and Consumer**

Early in 2016 [KIP-42](#) introduced the ability to monitor or "intercept" information on client (consumer and producer) behavior.

The goal of the KIP was to provide "The ability to quickly deploy tools to observe, measure, and monitor Kafka client behavior, down to the message level" <sup>15</sup>.

---

Footnote 15

[cwiki.apache.org/confluence/display/KAFKA/KIP-42:+Add+Producer+and+Consumer+Interceptors](https://cwiki.apache.org/confluence/display/KAFKA/KIP-42:+Add+Producer+and+Consumer+Interceptors)

While the interceptors aren't typically your first line for debugging, they could prove useful in observing the behavior of your Kafka streaming application and are a valuable addition to your toolbox. An excellent example of using an interceptor (producer) is keeping track of the message offsets your Kafka Streams application is producing back to Kafka.

#### **IMPORTANT**

#### **Consumer and Producer Interceptors in Kafka Streams**

Because Kafka Streams can consume or produce any number of key and value types, the internal `Consumer` and `Producer` are configured to work with `byte[]` keys and `byte[]` values, hence never receive unserialized data. Serialized data means you can't inspect messages without an extra deserialization/serialization step.

So let's get started by discussing the Consumer Interceptor first.

## CONSUMER INTERCEPTOR

The consumer interceptor gives you two access points to "intercept." The first is `ConsumerInterceptor#onConsume` which reads `ConsumerRecords` between the point they are retrieved from the broker and before the messages are returned from the `Consumer#poll` method. The following pseudo-code gives you an idea of where the consumer interceptor is doing its work:

### **Listing 7.3 Where the Consumer interceptor is doing its work**

```
ConsumerRecords<String, String> poll(long timeout) {
    ConsumerRecords<String, String> consumerRecords = consume records from broker... ①
    return interceptors.onConsume(consumerRecords); ②
```

- ① Fetching new records from the broker
- ② Running records through the interceptor chain and returning the results.

Although this pseudo code bears no resemblance to actual `KafkaConsumer` code, it illustrates the point. Interceptors accept the `ConsumerRecords` returned from the broker inside the `Consumer#poll` method and have the opportunity to perform any operation including filtering or modification before the `KafkaConsumer` returns the records from the `poll` method.

ConsumerInterceptors are specified via the `ConsumerConfig#INTERCEPTOR_CLASSES_CONFIG` with a collection of one or more `ConsumerInterceptor` implementor classes. Multiple interceptors are chained together and executed in the order specified in the configuration.

A `ConsumerInterceptor` accepts and returns a `ConsumerRecords` instance. If there are multiple interceptors, the returned `ConsumerRecords` from one interceptor serves as the input parameter for the next interceptor in the chain. Thus any modifications made by one interceptor is propagated to the next interceptor in the chain.

Exception handling is an important consideration when chaining multiple interceptors together. If an `Exception` occurs in an interceptor, it logs the error, but it *does not* short-circuit the chain. Thus the `ConsumerRecords` continues to work its way through the remaining interceptors.

For example consider you have three interceptors A, B, and C. All three interceptors modify the records and rely on changes made by the previous interceptor in the chain. However, if interceptor A encounters an error, the `ConsumerRecords` object continues to interceptor B and C, but without the expected modifications, rendering the results from the interceptor chain invalid.

For this reason, it's best not to have an interceptor relying on modified `ConsumerRecords` from a previous interceptor in the chain.

The second interception point is the `ConsumerInterceptor#onCommit` method.

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

After the consumer commits its offsets to the broker, the broker returns a `Map<TopicPartition, OffsetAndMetadata>` containing information with the topic, partition, and the committed offsets along with associated meta-data (time of commit, etc.). The commit information can be useful for tracking purposes. Here's an example of simple `ConsumerInterceptor` used for logging purposes.

**Listing 7.4 Logging Consumer Interceptor found in  
src/main/java/bbejeck/chapter\_7/interceptors/StockTransactionConsumerIntercep**

```
public class StockTransactionConsumerInterceptor implements ConsumerInterceptor<Object, Object> {

    // some details left out for clarity
    private static final Logger LOG =
        LoggerFactory.getLogger(StockTransactionConsumerInterceptor.class);

    public StockTransactionConsumerInterceptor() {
        LOG.info("Built StockTransactionConsumerInterceptor");
    }

    @Override
    public ConsumerRecords<Object, Object> onConsume(ConsumerRecords<Object,
        Object> consumerRecords) {
        LOG.info("Intercepted ConsumerRecords {}", ❶
            buildMessage(consumerRecords.iterator()));
        return consumerRecords;
    }

    @Override
    public void onCommit(Map<TopicPartition, OffsetAndMetadata> map) {
        LOG.info("Commit information {}", map); ❷
    }
}
```

- ❶ Logging the consumer records and metadata before the records are processed
- ❷ Logging the commit information once the Kafka Streams consumer commits offsets to the broker.

Now let's take a minute to cover the producing side of intercepting.

## PRODUCER INTERCEPTOR

The `ProducerInterceptor` works similarly and has two access points `ProducerInterceptor#onSend` and `ProducerInterceptor#onAcknowledgement`.

With the `onSend` method, the interceptor(s) can perform any action including mutating the `ProducerRecord`. Each producer interceptor in the chain receives the returned object from the previous interceptor.

Exception handling is the same as we saw on the consumer side, so the same caveats apply here as well.

The `ProducerInterceptor#onAcknowledgement` method is called when the broker acknowledges the record. If sending the record fails, `onAcknowledgement` is called at that point as well.

Here's a simple logging `ProducerInterceptor` example:

## Listing 7.5 Logging Producer Interceptor found in src/main/java/bbejeck/chapter\_7/interceptors/ZMartProducerInterceptor.java

```
public class ZMartProducerInterceptor implements ProducerInterceptor<Object, Object> {
    // some details left out for clarity
    private static final Logger LOG = LoggerFactory.getLogger(ZMartProducerInterceptor.class);

    @Override
    public ProducerRecord<Object, Object> onSend(ProducerRecord<Object, Object> record) {
        1
        LOG.info("ProducerRecord being sent out {}", record);
        return record;
    }

    @Override
    public void onAcknowledgement(RecordMetadata metadata, Exception exception) {
        if (exception != null) {
            2
            LOG.warn("Exception encountered producing record {}", exception);
        } else {
            LOG.info("record has been acknowledged {} ", metadata);
        }
    }
}
```

- 1 Logging right before the message is sent to the broker
- 2 Logging broker acknowledgement or if error occurred (broker side) during the produce

The ProducerInterceptor is specified with the ProducerConfig#INTERCEPTOR\_CLASSES\_CONFIG and takes a Collection of one or more ProducerInterceptor classes.

### TIP Configuring Interceptors

When configuring interceptors in a Kafka Streams application you need to prefix the consumer and producer interceptors property names with props.put(StreamsConfig.consumerPrefix(ConsumerConfig.INTERCEPTOR\_CLASSES\_CONFIG) and StreamsConfig.producerPrefix(ProducerConfig.INTERCEPTOR\_CLASSES\_CONFIG) respectively.

To see the interceptors in action there is `src/main/java/bbejeck/chapter_7/StockPerformanceStreamsAndProcessorMetricsApp` which uses a consumer interceptor, and `src/main/java/bbejeck/chapter_7/ZMartKafkaStreamsAdvancedReqsMetricsApp`. uses a producer interceptor, so both classes the configuration required for using interceptors.

As a side note, since interceptors work on every record in the Kafka Streams application the output of the logging interceptors is significant. As a result, the interceptor results are output to the `consumer_interceptor.log` and `producer_interceptor.log` found in the `logs` directory in the base of the source code installation.

We've spent some time looking at metrics on consumer performance how we can ©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

intercept records coming into and out of a Kafka Streams application. But this information is coarse-grained and *outside* of a Kafka Streams application, so let's go inside a Kafka Streams application and see what's going on under the hood.

Our next step is going to dig a little deeper and take a look at how we can measure performance inside our topology, which we'll cover in the next section, gathering metrics.

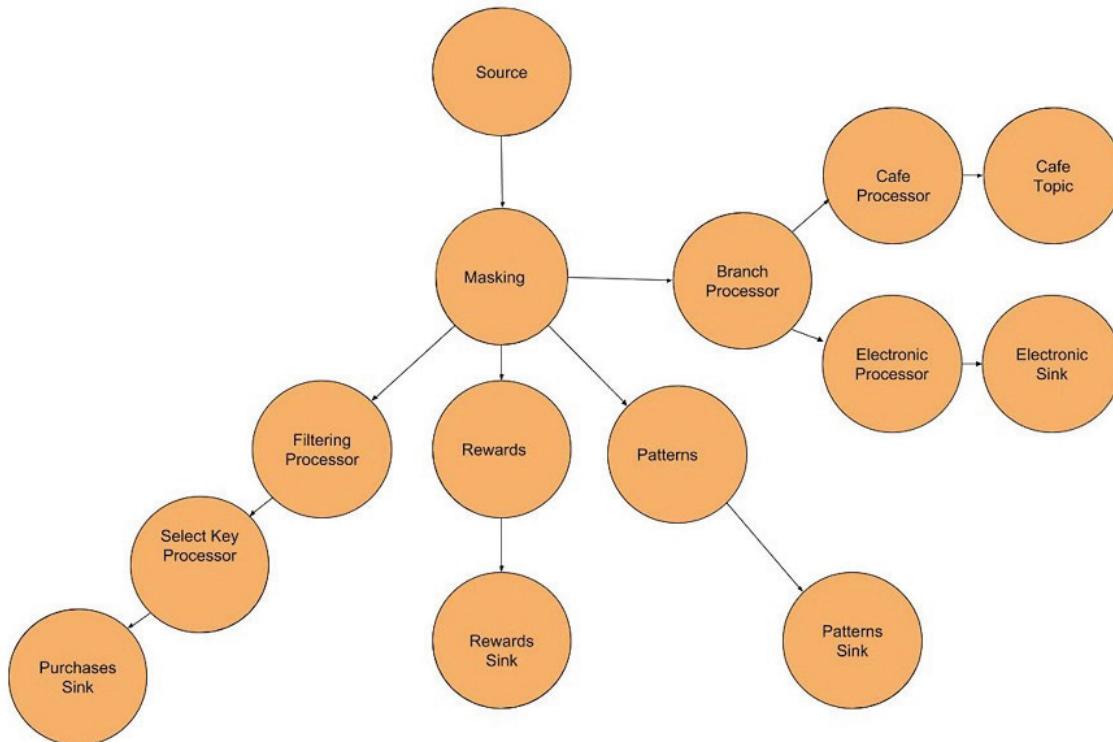
## 7.2 Application Metrics

When it comes to measuring the performance of an application, we can get a sense of how long it takes us to process one record and measuring end-to-end latency is undoubtedly a good indicator of overall performance. But if we want to improve performance we'll need to know the exact spots where things are slowing down.

Measuring performance is essential for streaming applications, for the mere fact that you are using a streaming application implies you want to process data or information as it's available.

So it stands to reason if your business needs dictate a streaming solution then you'll want the most *efficient* and *correct* streaming process you can.

Before we discuss the actual metrics we'll focus on, let's revisit one of the applications we built in Chapter 3, the advanced ZMart application. The ZMart app is a good candidate for metrics tracking as there are several processing nodes, so we'll use this topology for our concrete example. Figure 7.6 shows the topology we created:



**Figure 7.6 ZMart Advanced Application Topology with a lot of nodes an excellent candidate for metrics observation**

Keeping the ZMart topology in mind, let's take a look at the categories of metrics we

have.

1. Thread Metrics
  1. Average time for commits, poll, process operations
  2. Tasks created per second, tasks closed per second,
2. Task Metrics
  1. Average number of commits per second
  2. Average commit time.
3. Processor Node Metrics
  1. Average and max processing time
  2. Average number of process operations per second.
  3. Forward Rate
4. State Store Metrics
  1. Average execution time for put, get and flush operations.
  2. Average number put, get and flush operations per second.

Note that this is not an exhaustive list of [the possible metrics](#), but I've chosen the ones here as in my opinion they offer excellent coverage of the most common performance scenarios.

Now we have what we're going to measure, let's talk about how we capture the information.

### **7.2.1 Metrics Configuration**

Kafka Streams already provides the mechanism for collecting performance metrics. For the most part, we just need to give some configurations values (although in a later section we'll cover how to provide a custom metric instance). Since the collection of metrics does incur a performance cost, we have two levels INFO and DEBUG.

While an individual metric itself may not be that expensive, when one considers that some metrics may involve every record flowing through the Kafka Streams application, we can see how the impact to performance can add up.

We use the metrics levels like we use logging levels. When troubleshooting an issue or observing how the application behaves we want more information, so we'll use the DEBUG level. Other times we don't need all the information, we'll use the INFO level.

Typically we don't want to use DEBUG in production as the cost of performance would be too high. Each of the metrics above are available at different levels let's look at the table to display which metrics we are available for INFO and DEBUG.

**Table 7.1 Metrics Availability by Levels**

Metrics Category	DEBUG	INFO
Thread	x	x
Task	x	
Processor Node	x	
State Store	x	
Record Cache	x	

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

So from the table above we can see Thread metrics are available at any level, while the rest of the metric categories are only collected when the using the DEBUG level. Setting the level is done when setting the configurations of your Kafka Streams application.

It's been there all along with our other application configurations; it's just we've accepted the default settings up to this point. The default level of metrics collection is INFO. Let's update the configs from our advanced ZMart application and but to turn on the collection of all metrics we use this configuration setting:

#### **Listing 7.6 Updating the configs for DEBUG metrics**

**src/main/java/bbejeck/chapter\_7/ZMartKafkaStreamsAdvancedReqsMetricsApp.java**

```
private static Properties getProperties() {
    Properties props = new Properties();
    props.put(StreamsConfig.CLIENT_ID_CONFIG, "metrics-client-id");
    props.put(ConsumerConfig.GROUP_ID_CONFIG, "metrics-group-id");
    props.put(StreamsConfig.APPLICATION_ID_CONFIG, "metrics-app-id");
    props.put(StreamsConfig.METRICS_RECORDING_LEVEL_CONFIG, "DEBUG");
    props.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
    return props;
}
```

- 1
- 2
- 3
- 4
- 5

- 1 Client ID
- 2 Group ID
- 3 Application ID
- 4 Setting the metrics recording level to DEBUG
- 5 Setting the connection for the broker(s)

So now we've enabled the collection/recording of DEBUG level metrics. The key points I want you to remember from this section is there are metrics built in to measure the full scope of a Kafka Streams application and to carefully consider before turning on metrics collection at the DEBUG level.

Now that we've discussed what metrics are available and how we collect metrics the next step is how we can observe the collected metrics.

#### **7.2.2 How to Hook Into the Collected Metrics**

The metrics in a Kafka Streams application are collected and distributed to metrics reporters. As you might have guessed, Kafka Streams provides a default reporter via JMX(Java Management Extensions).

So at this point, once we've enabled collecting metrics at the DEBUG level we have nothing left to do but observe the metrics. One thing to keep in mind is JMX only works with live running applications. So the metrics we are going to look at will be when the application is running.

**TIP Programmatic Metrics Access**

You can also access metrics programmatically. For an example of programmatic metrics a  
y o u c a n l o o k  
src/main/java/bbejeck/chapter\_7/StockPerformanceStreamsAndProcessorMetricsApplicatio

While I imagine most readers are familiar with using JMX or have at least heard of it, in the next section, I'm going to give a brief overview on how to get started using JMX for those readers who have less experience with it. If you are an experienced JMX user, feel free to skip this next section.

**7.2.3 Using JMX**

JMX or Java Management Extensions is a standard way of viewing the behavior of programs running on the JVM. We can also use JMX to see how the JVM (Java Virtual Machine) is performing itself. In a nutshell, JMX gives us the infrastructure to expose parts of our running program.

Fortunately, we won't need to write any code to do this monitoring. We'll simply need to connect either [JVisualVM](#), [JConsole](#) or [Java Mission Control \(JMC\)](#).

**TIP JMC Usage Notes**

Java Mission Control or JMC is very powerful and can be a great tool for monitoring, but it requires a commercial license for use in production. Since JMC ships with the JDK, you can start JMC directly from the command-line with the command `jmc` (assuming JDK bin directory is on your path). Additionally, you'll need to add these flags when starting your Kafka streaming application

```
-XX:+UnlockCommercialFeatures  
-XX:+FlightRecorder
```

As JConsole is the most straightforward approach, we'll start with it for now.

**NOTE****What is JMX?**

The Java Management Extensions (JMX) technology is a standard part of the Java Platform, Standard Edition (Java SE platform), added to the platform in the 5.0 release.

The JMX technology provides a simple, standard way of managing resources such as applications, devices, and services. Because the JMX technology is dynamic, you can use it to monitor and control resources as they are created, installed and implemented. You can also use the JMX technology to monitor and manage the Java Virtual Machine (Java VM).

The JMX specification defines the architecture, design patterns, APIs, and services in the Java programming language for management and monitoring of applications and networks<sup>16</sup>

---

Footnote 16 [docs.oracle.com/javase/tutorial/jmx/intro/index.html](https://docs.oracle.com/javase/tutorial/jmx/intro/index.html)

**STARTING JCONSOLE**

JConsole ships with the JDK, so by having Java installed, we already have JConsole. Starting JConsole is as simple as running `jconsole` from a command prompt (assuming Java is on your path). Once started you will see a GUI come up that looks like this:

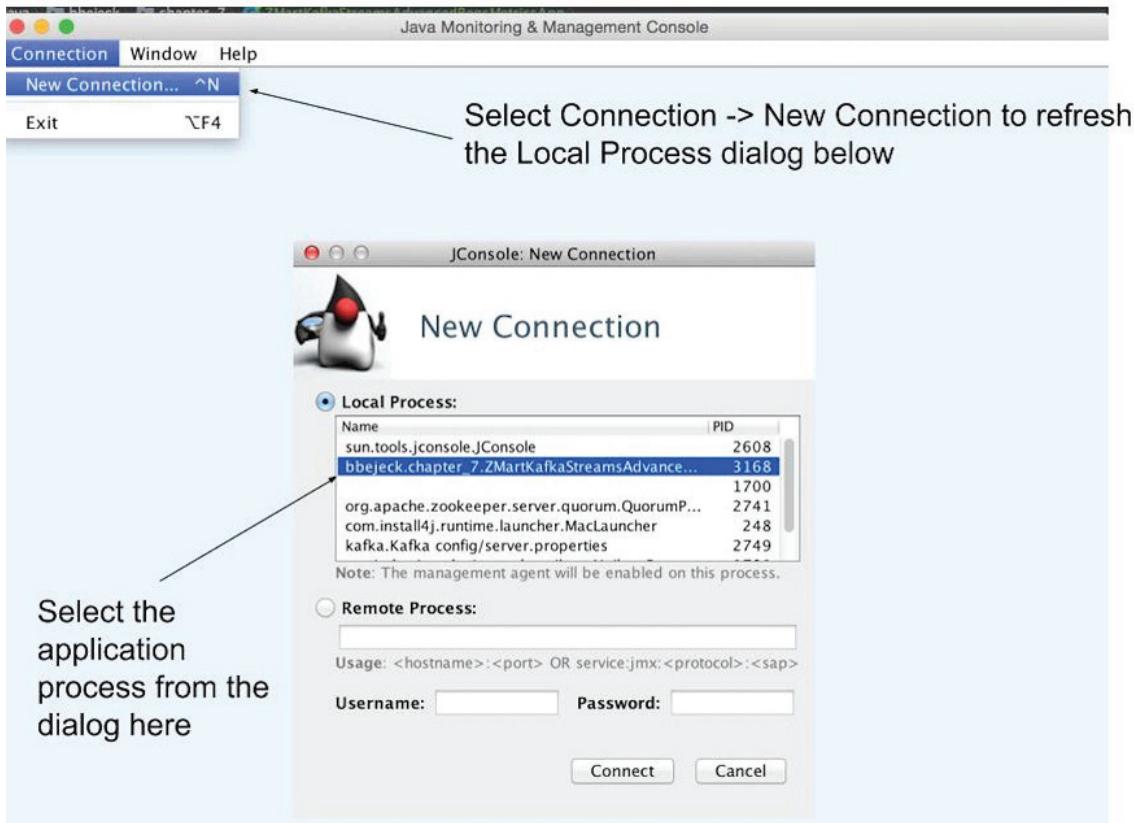


**Figure 7.7 JConsole Start Menu**

Now we have jconsole up; our next step is to use it to look at some metric data!

## START MONITORING A RUNNING PROGRAM

If we look at the center of the jconsole GUI, we'll see a "New Connection" dialog box right in the middle. The image below shows us the starting point for jconsole:



**Figure 7.8 JConsole Connect to Program**

For now, we are only concerned with the java processes listed under the "Local Process" section.

### NOTE

#### Remote JMX and Protecting Access

You can use jconsole to monitor remote applications, and you can secure access to JMX access. You can see the "Remote Process" and "Username" and "Password" sections in the image above. For the scope of this book, however, we are going to limit our discussion here to local access only during development. However, the internet is full instructions on remote and secure jconsole access and Oracle's documentation [is a great starting point](#).

If you haven't already started your Kafka Streams application, do so now. Then to get your application to show up in the "Local Process" window, click on "Connection" then

"New Connection" (as pictured in the image above). The processes listed under "Local Process" should refresh and you'll see your Kafka Streams application (again as shown in the picture above). Double-click on the Kafka Streams application process.

Now chances are after you double-click on the program you want to connect to you'll be greeted with a warning similar to this:



**Figure 7.9 JConsole Connect Warning No SSL**

Since we are on our local machine, we will select the "Insecure connection" button. Now we are all set to look at the metrics being collected by our Kafka Streams application.

#### CAUTION

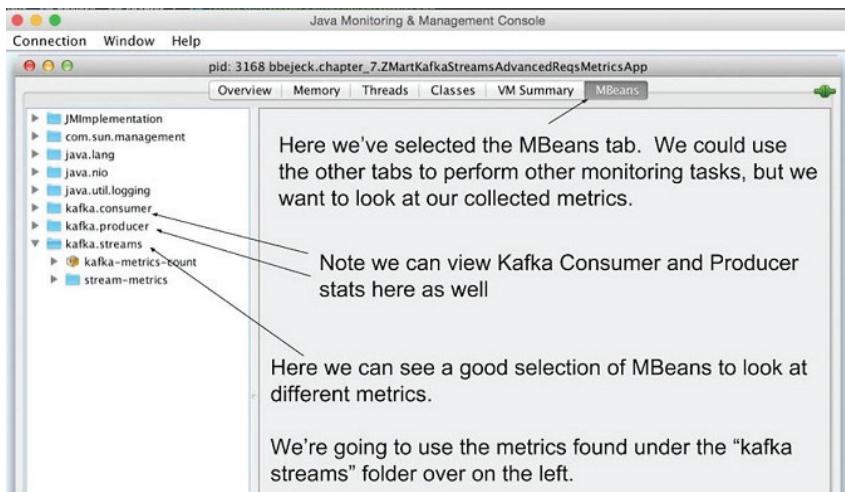
#### Insecure connections are for local development only

We're using an insecure connection for development on our local machine. But in practice, you should always secure access to remote services accessing the internal state of your application.

Now you are connected to your running program, the next step to look at the information available.

#### VIEWING THE INFORMATION

Once connected you'll see a GUI screen looking something like this image:



**Figure 7.10 JConsole Started**

Viewing the image above we can see jconsole gives us several handy options to peek inside the internals of our running application. But out of the Overview, Memory, Thread, Classes, VM Summary and MBeans tabs, we are only going to use the MBeans tab, as MBeans contain the collected statistics about the performance of our Kafka Streams program.

The other tabs provide relevant information, but that info relates more to overall application health and how our program is utilizing resources. The metrics collected in the MBeans contain information regarding the internal performance of the topology.

So we've concluded our introduction to using jconsole, our next step is to start viewing the recorded metrics on the topology.

#### 7.2.4 Viewing Metrics

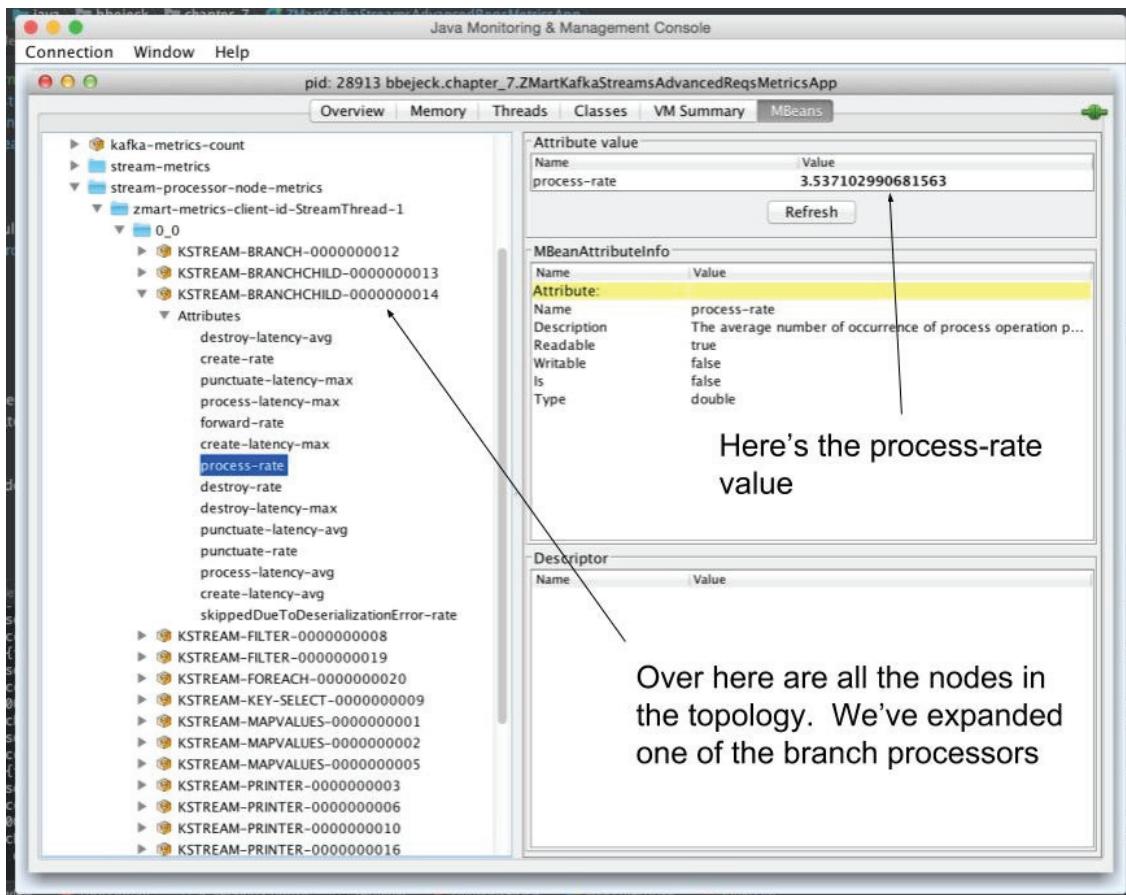
Now we'll start looking at the metrics collected from our running Kafka Streams application.

##### TIP

##### Trying out the metrics

Since JMX only works with running applications, some of the example applications found in `src/main/java/bbejeck/chapter_7` will run continuously so that you can play around with the metrics. As a result, you will need to explicitly stop either in the IDE or by a `Ctrl+C` from the command line

Now let's take a look at some metrics via jconsole while running the ZMart application:



**Figure 7.11 JConsole metrics for  
src/main/java/bbejeck/chapter\_7/ZMartKafkaStreamsAdvancedReqsMetricsApp.java**

As you can see from the image above, we can drill down to all processors and nodes in the topology to view performance (either throughput or latency). Here we've selected the `process-rate` metric which tells us the average number of records processed per millisecond. From looking at the upper right under `Attribute Value` you can see the `process-rate` averages to 3.537 records per millisecond (3,537 records per second).

Additionally, if you remember as we saw earlier, we can also see the producer and consumer metrics from jconsole as well.

#### TIP Providing Custom metrics

While the provided metrics are comprehensive, there may be cases where we want to provide custom metrics. Since this is a low-level detail and probably not a very common use case, we won't walk through an example in detail. But I'll point you to the `StockPerformanceMetricsTransformer#init` method for an example of how you can add a custom metric. The `StockPerformanceMetricsTransformer#transform` method for an example of how you utilize it. The `StockPerformanceMetricsTransformer` is found in `src/main/java/bbejeck/chapter_7/transformer/StockPerformanceMetricsTransformer.java`

Now that we've covered viewing Kafka Streams metrics let's move on to other useful

techniques for observing what's going on in our application.

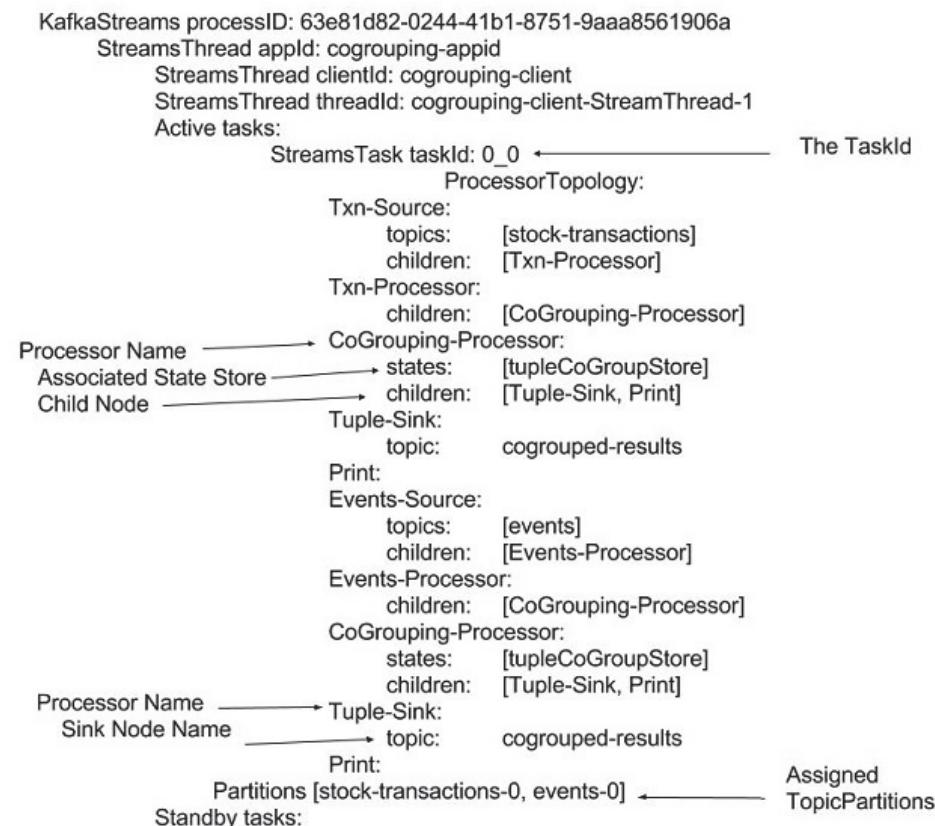
### 7.3 More Kafka Streams Debugging Techniques

In this section, we'll see more ways we can observe and debug our Kafka streaming application. While the previous section was more about performance, the techniques we'll present in this section are more about getting notified of various states of the application and viewing the structure of the topology.

#### 7.3.1 Viewing a representation of the application

After you're up and running, you might run into situations where you need to debug the application. You'd like to get a second pair of eyes on the job, but for whatever reason, you can't share the code. Or you'd like to see the `TopicPartition` assigned to the tasks of the application.

The option we have for getting general information on the structure of the application is the `Topology.describe()` method, which prints out beneficial information regarding the structure of the program, including any internal topics created to support repartitioning. The following image displays the results of calling `describe` on the `CoGroupingListeningExampleApplication` from chapter 7:



**Figure 7.12 Showing Node names and associated Child nodes Processor API Describe Example for  
src/main/java/bbejeck/chapter\_7/CoGroupingListeningExampleApplication.java**

As you can see, the `Topology.describe()` method prints out a nice concise view of

the application structure. There are two things we should note about this example. First, the `CoGroupingListeningExampleApplication` used the Processor API, so all the nodes in the topology have the names we've chosen. With the Kafka Streams API the names of the nodes are a little more generic i.e

### **Listing 7.7 To String sample on Kafka Streams API application**

```
KSTREAM-SOURCE-0000000000:
topics: [transactions]
children: [KSTREAM-MAPVALUES-0000000001]
```

#### **TIP**

#### **Accessing the Topology from the DSL**

When you are using the Kafka Stream DSL API, you don't directly use the `Topology`, but it is easily accessed. If you want to print the topology from the DSL use the `StreamsBuilder.build()` method which returns a `Topology` object, and then you can call `Topology.describe()` just like the example above.

Secondly, to get the full picture of what's going on, you'll want to wait until the application is fully initialized and in the `RUNNING` state before printing the topology. In the example above we used `StateListener` to trigger printing the topology at the correct time. We'll get into states and listeners in our next section.

#### **TIP**

#### **StreamThread Information**

Getting information about the `StreamThread` object(s) in your application can be useful as well. To access the `StreamThread` info use the `KafkaStreams.localThreadsMetadata()` method

### **7.3.2 Getting notification on various states of the application**

When we start our Kafka Streams application, it does not automatically begin processing data as there is some coordination that has to happen first. The consumer needs to fetch metadata and subscription information; the application needs to start `StreamThread` instance(s) and assign `TopicPartition`(s) to `StreamTask`(s).

This process of assigning or redistributing tasks (workload) is called rebalancing. Rebalancing means Kafka Streams can "auto-scale" either up or down automatically. This ability is a crucial strength; you can add a new application instance(s) *while an existing application is already running*, and the rebalance process will redistribute the workload.

For example assume you have a Kafka Streams application with two source topics and each topic has two partitions, resulting in 4 `TopicPartition` objects needing assignment. Initially, you start the application with one thread. Kafka Streams determines

the number of tasks to create by taking the max partition size among all input topics. In this case each topic has two partitions so the max is two and you end up with two tasks. The rebalance process then assigns 2 tasks each two `TopicPartition` objects.

But after running for a little while you decide you want to process records faster.

In this case, all you need to do is start another version of the application *with the same application-id* and the rebalance process will distribute the load across the new application thread resulting in assigning the 2 tasks across both threads. You have just doubled the scale of your application, while the original version is still running, no need to shut the initial application down.

Other causes of rebalancing include another Kafka Streams instance (with the same application id) starting or stopping, adding partitions to a topic, or in the case of a regex defined source node, adding or removing topics matching the regex pattern.

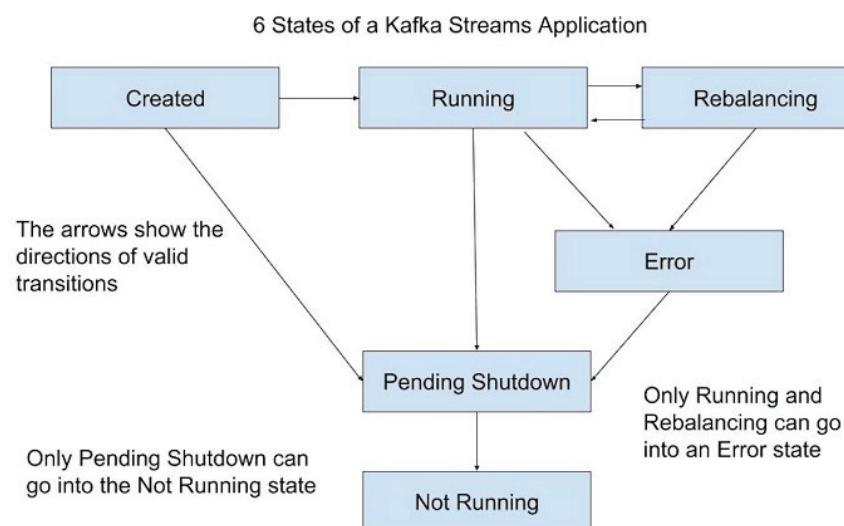
But during the rebalance phase external interaction temporarily paused until the application has completed assignments of topic-partitions to stream tasks, so we'd like to be aware of this point in the lifecycle of the application.

For example, the queryable state stores (we'll build such an application in Chapter 9) are unavailable, so we'd like to have the ability to restrict requests to view the contents of the store until the stores are available again.

But how can we check if our other applications are going through a rebalance? Fortunately, Kafka Streams provides just such a mechanism, the state listener, which we'll talk about in the next section.

### 7.3.3 Using the State Listener

A Kafka Streams application at any point in time can be in one of six states. The following diagram shows the possible valid states a Kafka Streams application can be in:



**Figure 7.13 States of the Kafka Streams Application**

While we can see from the image above there are few state change scenarios we could discuss, we are going to focus on the transitions between `RUNNING` and `REBALANCING`.

The transition between these two states is the most frequent and has the most impact on performance, as during the REBALANCING phase no processing occurs.

To capture these state changes you'll use the `KafkaStreams#setStateListener` method.

The method takes an instance of the `StateListener` interface. It's a single method interface, so we'll be able to use Java 8 lambda syntax, as demonstrated in this code sample:

**Listing 7.8 Adding a State Listener found in  
src/main/java/bbejeck/chapter\_7/ZMartKafkaStreamsAdvancedReqsMetricsApp.java**

```
KafkaStreams.StateListener stateListener = (newState, oldState) -> {
    if (newState == KafkaStreams.State.RUNNING && oldState ==
        KafkaStreams.State.REBALANCING) {1
        LOG.info("Application has gone from REBALANCING to RUNNING ");
        LOG.info("Topology Layout {}", streamsBuilder.build().describe());2
    }
};
```

- ① Checking that we are transitioning from REBALANCING to RUNNING
- ② Printing out the structure of the Topology

**TIP**

**Trying out the listener**

Since this example running `ZMartKafkaStreamsAdvancedReqsMetricsApp.java` involve viewing at JMX metrics and the state transition notification, we've turned off printing the streaming results to the console and are writing solely out to Kafka topics. When running the app, you should see the listener output in the console.

For our first `StateListener` implementation, we are just going to log the state change out to the console. If you remember from the section above where we discussed printing the topology structure, we spoke of needing to wait until the application has completed rebalancing. That is what we've done here as we print out the structure once all tasks and assignments are complete.

Let's take this example a little further and show notification of when the application is going into rebalancing state:

We would update our code to handle this additional state transition to this:

**Listing 7.9 Updating the state listener when entering REBALANCING found in  
src/main/java/bbejeck/chapter\_7/ZMartKafkaStreamsAdvancedReqsMetricsApp.java**

```
KafkaStreams.StateListener stateListener = (newState, oldState) -> {
    if (newState == KafkaStreams.State.RUNNING && oldState == KafkaStreams.State.REBALANCING) {
        LOG.info("Application has gone from REBALANCING to RUNNING ");
        LOG.info("Topology Layout {}", streamsBuilder.build().describe());
    }
};
```

```

if (newState == KafkaStreams.State.REBALANCING) {
    LOG.info("Application is entering REBALANCING phase");
}

```

1

### ① Adding an action when entering the rebalance phase

Even though we are using simple logging statements, it should be evident as to how we can add more sophisticated logic to handle the state changes within our application.

#### **NOTE**

Since Kafka Streams is a library and not a framework, you can run a single instance on one server. If you do run multiple instances on different machines, you'll only see results from state changes on that local machine.

The critical point of this section is we can hook into the current state of our Kafka Streams application making it less of a black box operation. Next, since we've mentioned rebalancing in our example of monitoring state, it's a great segue to discuss rebalancing a little more in depth.

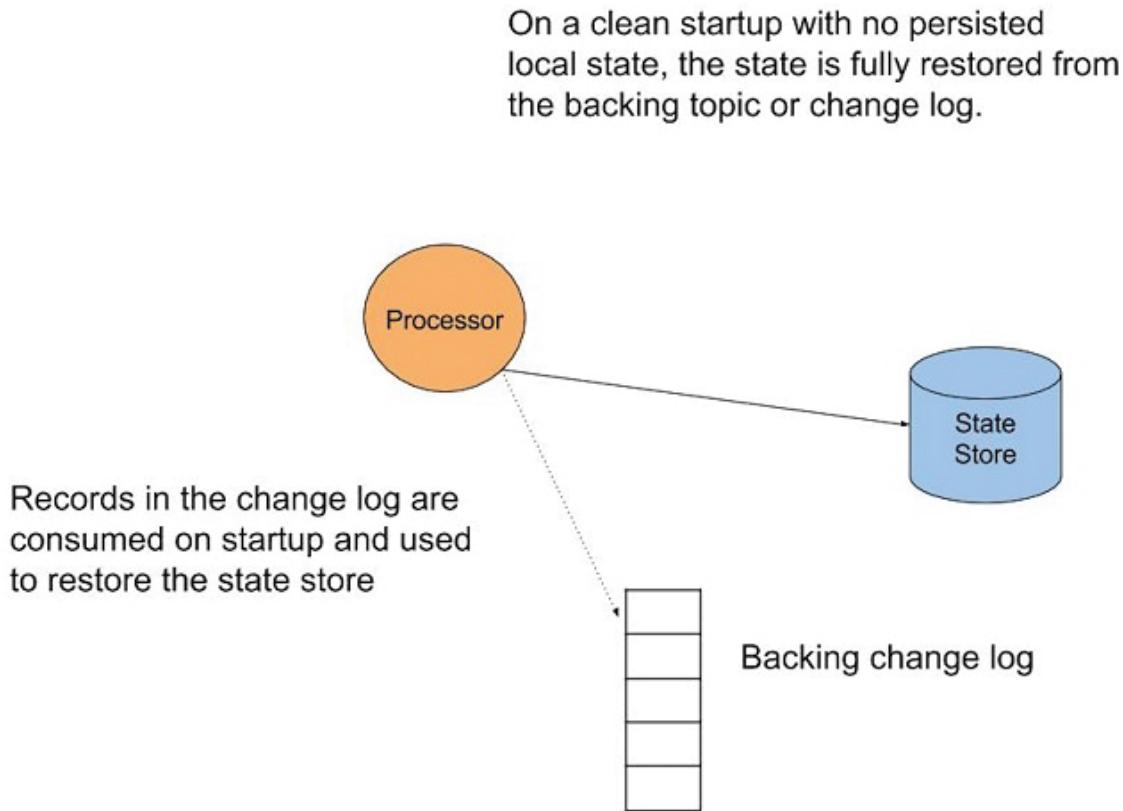
While the ability to automatically rebalance the workload is a strength of Kafka Streams, we want to keep the number of rebalances to a minimum because when a rebalance occurs, we are not processing data, and we'd like to have our application processing data as much as possible.

#### **7.3.4 State Restore listener**

We learned in chapter 4 about state stores and the importance of having the state stores backed up in case of a failure. In Kafka Streams, we use topics frequently called change-logs as the backup for the state stores.

The change-log is recording the updates to the state store as the update occurs. When a Kafka Streams application fails, or you restart it, there are the local state files for the state store to recover from, as depicted in the following image:

## State Store Restoration



**Figure 7.14 Restoring a state store from clean start/recovery**

However, in some circumstances, we may need to a full recovery of state store from the change-log. For example, if you are running your Kafka Streams application in a stateless environment like Mesos or you encounter a severe failure and the files on local disk are wiped out.

Depending on the amount of data you have to restore, this restoration process could take a non-negligible amount of time.

During this recovery period, any state stores you have exposed for querying are unavailable, so it'd be nice if we can get an idea of how long this restoration process is going to take and how progress is coming along. Additionally, if we have a custom state store, we'd like notification when the restore is starting and ending so we can do any set-up or tear-down tasks we need to.

To accomplish this task `StateRestoreListener` interface, much like the `StateListener` allows us to get notification of what is going on inside the application.

The `StateRestoreListener` has three methods `onRestoreStart`, `onBatchRestored`, and `onRestoreEnd`. We use the `KafkaStreams#setGlobalRestoreListener` method to specify the global restore listener to use.

**IMPORTANT****Restore Listener is shared**

The provided `StateRestoreListener` is shared application wide and is expected to be stateless. If you need to keep track of any state in the listener, you'll need to provide the synchronization.

Let's take a walk through our listener code to get an idea of how this notification process can work. We'll start with the declaration of the variable and the `onRestoreStart` method:

**Listing 7.10 Here is a logging restore listener found in  
src/main/java/bbejeck/chapter\_7/restore/LoggingStateRestoreListener.java**

```
public class LoggingStateRestoreListener implements StateRestoreListener {

    private static final Logger LOG = LoggerFactory.getLogger(LoggingStateRestoreListener.class);
    private final Map<TopicPartition, Long> totalToRestore = new ConcurrentHashMap<>();
    private final Map<TopicPartition, Long> restoredSoFar = new ConcurrentHashMap<>();

    @Override
    public void onRestoreStart(TopicPartition topicPartition, String store, long start, long end) {
        long toRestore = end - start;
        totalToRestore.put(topicPartition, toRestore);
        LOG.info("Starting restoration for {} on topic-partition {} total to restore {}",
            store, topicPartition, toRestore);
    }

    // other methods left out for clarity covered below
}
```

- ➊ Creating `ConcurrentHashMap` instances for keeping track of restore progress
- ➋ Storing the total amount to restore for the given `TopicPartition`

So our first steps are to create two `ConcurrentHashMap` instances for keeping track of restoration progress. In the `onRestoreStart` method we store the total number of records we need to restore and log out the fact we are starting.

Next, let's move on to the code that handles each batch restored.

**Listing 7.11 Handling onBatchRestored found in  
src/main/java/bbejeck/chapter\_7/restore/LoggingStateRestoreListener.java**

```
@Override
public void onBatchRestored(TopicPartition topicPartition,
    String store, long start, long batchCompleted) {
    NumberFormat formatter = new DecimalFormat("#.##");

    long currentProgress = batchCompleted + restoredSoFar.getOrDefault(topicPartition, 0L);
    double percentComplete = (double) currentProgress / totalToRestore.get(topicPartition);

    LOG.info("Completed {} for {}% of total restoration for {} on {}",
        batchCompleted, formatter.format(percentComplete * 100.00),
```

```

        store, topicPartition);
restoredSoFar.put(topicPartition, currentProgress);
}

```

3

4

- ➊ Calculating total number of records restored
- ➋ Determining percentage of restoration completion
- ➌ Logging out the percent restored
- ➍ Storing number of records restored so far

The restoration process uses an internal consumer to read the changelog topic, so it follows the application restores records in batches from each `consumer.poll()` method call. As a consequence, the max size of any batch will equal to the `max.poll.records` setting.

The `onBatchRestored` method is called after the restore process has loaded the latest "batch" into the state store. In listing one, we add the size of the current batch to the accumulated restore count. In listing 2, we calculate the percentage of restoration completion and log out the results in listing 3. Finally, in listing 4, we store the new total number of records computed earlier.

The last step we will cover is when the restoration process completes.

#### **Listing 7.12 Method called when restoration completed**

**src/main/java/bbejeck/chapter\_7/restore/LoggingStateRestoreListener.java**

```

@Override
public void onRestoreEnd(TopicPartition topicPartition, String store, long totalRestored) {
    LOG.info("Restoration completed for {} on topic-partition {}", store, topicPartition);
    restoredSoFar.put(topicPartition, 0L); ①
}

```

- ➊ Keeping track of restore progress for a TopicPartition

Once the application completes the recovery process, we make one final call to the listener with the total number of records restored. In our example here we log out the finished state and update the full restoration count map to zero.

Finally, to use the `LoggingStateRestoreListener` in our application we use the following:

#### **Listing 7.13 Specifying the global restore listener**

**src/main/java/bbejeck/chapter\_7/CoGroupingListeningExampleApplication.java**

```
kafkaStreams.setGlobalStateRestoreListener(new LoggingStateRestoreListener());
```

This is a primary example of using a `StateRestoreListener`. In chapter nine we'll show an example where you'll get a graphical representation of the restore progress when

we present our queryable state application.

**TIP****Viewing the log files**

To view the log file generated by running the CoGroupingListeningExampleApplication example look for a log file named logs/state\_restore\_listener.log in the root directory where you installed the source code.

### 7.3.5 Uncaught Exception Handler

From time to time I think it's fair to say every developer has encountered an unaccounted for `Exception` and the big stack trace in the console/log as your program suddenly quits. Although this situation does not quite fit into a "monitoring" example, the ability to get a notification and handle any cleanup with the unexpected occurs is good practice. Kafka Streams provides the `KafkaStreams#setUncaughtExceptionHandler` for dealing with the unexpected error.

**Listing 7.14 Basic Uncaught Exception handler found in  
src/main/java/bbejeck/chapter\_7/CoGroupingListeningExampleApplication.java**

```
kafkaStreams.setUncaughtExceptionHandler((thread, exception) -> {
    CONSOLE_LOG.info("Thread [" + thread + "] encountered [" + exception.getMessage() + "]");
});
```

While this is undoubtedly a bare-bones implementation, it serves to demonstrate where you can set a hook for dealing with unexpected errors either by logging the error as shown here or performing any required clean-up and shutting down the streams application.

At this point, we've wrapped up the coverage of how we can monitor our Kafka Streams application.

## 7.4 Conclusion

In this chapter, we went from building applications to peeking under the hood of what we've built. While sometimes monitoring is not as exciting as developing new features, it's just as important.

The key points I want you to remember from this chapter are:

1. To monitor Kafka Streams at some point, you'll need to look at the Kafka brokers as well.
2. Enabling metrics reporting from time to time should be done from time to time to see the how the performance of the application is doing.
3. Peeking under the hood is required, and sometimes you need to go to lower levels and use command line tools included with Java such as `jstack` (thread dumps) and `jmap/jhat`(for heap dumps) to understand what your application is doing.

While in this chapter you focused on *observing* behavior, in our next chapter, we'll

shift our focus the making sure your application handles errors consistently, adequately and provides *expected* behavior by doing regular testing.

# 8

## *Testing a Kafka Streams Application*

In this chapter:

- Testing a Topology
- Testing Individual Processors and Transformers
- Integration Testing with an Embedded Kafka

Up to this point in the book, we've covered the essential building blocks for creating a Kafka Streams application. But there is one crucial part of our application development we've left out until now, that is how to test our application. While I'll assume most readers of this book are aware of the importance of testing, we'll review the top two reasons (at least to me) for why testing is just as necessary as the development process itself.

One of the critical concepts we'll focus on is placing your business logic in stand-alone classes that are entirely independent of a Kafka Streams application as that makes your logic much more accessible to test.

First, as we develop our code, we are creating an implicit contract of what we and others can expect how the code will execute. The only way to prove that the code works is by testing thoroughly. So we'll use testing to provide a good breadth of possible inputs and scenarios to make sure our code works appropriately under reasonable circumstances.

The second reason we need an excellent suite of tests is dealing with the inevitable change of software. Having a good set of tests gives an immediate feedback when our new code has broken the expected set of behaviors. Additionally, when we do a major refactor or add new functionality, passing tests gives us a level of confidence about the impact of the changes (provided we have good tests).

While we know the importance of testing, it's not always straightforward how we test our Kafka Streams applications. We still have the option to run a simple topology and

observe the results, but there is one drawback with that approach. We'd like a suite of *repeatable* tests that we can run at any time, and as part of a build, so we'd like the ability to test our applications without a Kafka cluster and Zookeeper ensemble.

That's what we're going to cover in this chapter. First, we'll show how we can test a topology *without* Kafka running so we can see the entire topology working from a unit test. We'll also cover testing a processor or transformer independently and mock any required dependencies.

**NOTE**
**What is a Mock Object?**

While I'm assuming most of the reading audience has experience testing with mock objects, here's more information for those who are interested [Mock\\_object](#)

While unit testing is critical for repeatability and quick feedback, integration testing is important as well, as sometimes we need to see the moving parts of our application in action. For example consider the case of rebalancing, an essential part of a Kafka Streams application, to get rebalancing to work in a unit test is nearly impossible. Here is a table summarizing the differences:

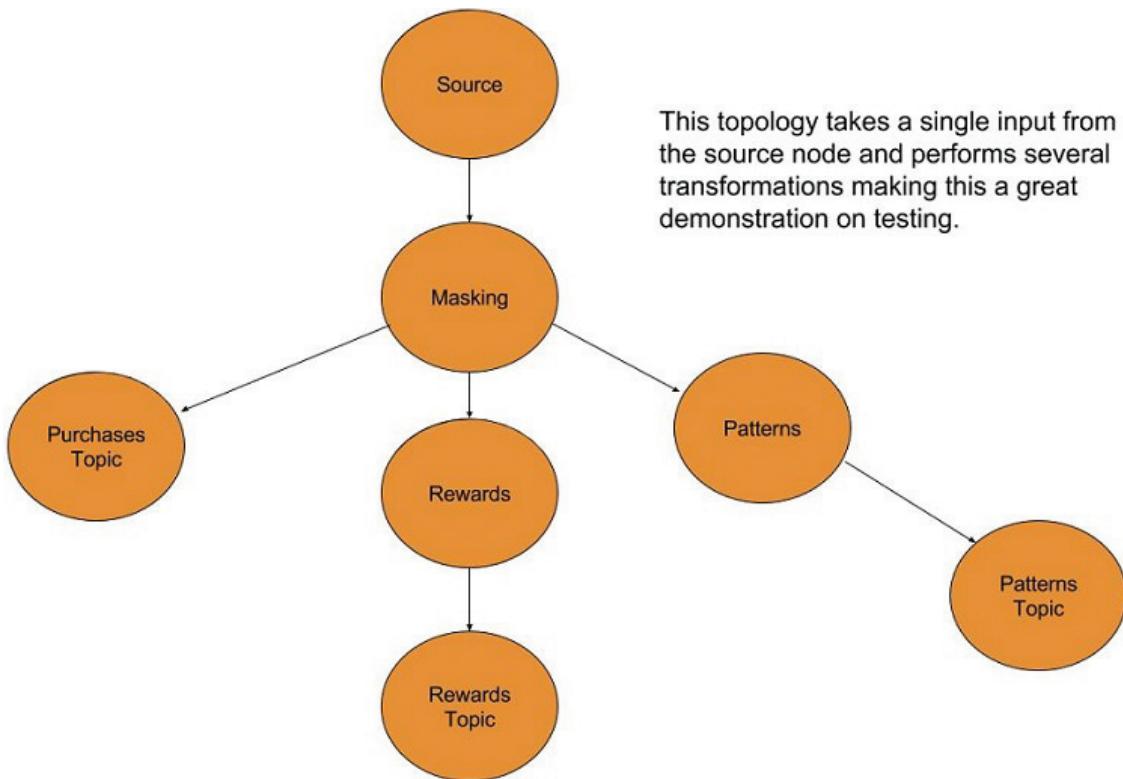
**Table 8.1 Testing Approaches**

Test Type	Purpose	Testing Speed	Level of Use
Unit	Testing individual parts of functionality in isolation	Fast	Predominate should comprise bulk of tests
Integration	Testing integration points between whole systems	Longer to run	Small percentage of tests

You need to trigger an actual rebalance under realistic conditions. For those situations, we need to run with a live instance of a Kafka cluster, but we don't want to rely on having an external cluster setup, so we'll cover how to use an embedded Kafka and Zookeeper for the integration testing.

## 8.1 Testing a Topology

If we go back to Chapter 3 and the first topology we built, we can see that it is relatively complicated. Here's an image just to refresh our memory of what we built



**Figure 8.1 Initial Complete Topology for ZMart Kafka Streams program**

While the processing logic is pretty straightforward, as we can see from the structure, it has several nodes in the topology. But it has one key thing that will help us to demonstrate testing, it takes one input, an initial purchase, and performs several transformations.

**TIP**

**Most Business Logic Should be outside the topology**

In most cases, you'll want to have the logic in separate classes so we can unit test the business logic separately from the topology. In the case of the ZMart topology most of our logic was very simple and represented as Java 8 lambda expressions, so in this case, we'll test the topology flow.

This will make our testing somewhat easy in that we only need to supply a single purchase value and we can confirm all the appropriate transformations take place.

But we want a repeatable stand-alone test, so we'll use the `ProcessorTopologyTestDriver` that will allow us to write just such a test, *without* needing Kafka to run the test. Remember, the ability to test your topology without a live Kafka instance makes testing faster and more lightweight hence leading to a quicker development cycle. Also take note that the `ProcessorTopologyTestDriver` is a generic testing framework that tests the individual Kafka Streams `Topology` object(s) you have created.

**TIP****Using Kafka Streams Testing Utilities**

To use testing utilities from Kafka Streams, we'll need to update our build.gradle file with

```
testCompile group:'org.apache.kafka', name:'kafka-streams',
           version:'1.0.0', classifier:'test'
testCompile group:'org.apache.kafka', name:'kafka-clients',
           version:'1.0.0', classifier:'test'
```

If you are using maven you will use:

```
<dependency>
    <groupId>org.apache.kafka</groupId>
    <artifactId>kafka-streams</artifactId>
    <version>1.0.0</version>
    <scope>test</scope>
    <classifier>test</classifier>
</dependency>

<dependency>
    <groupId>org.apache.kafka</groupId>
    <artifactId>kafka-clients</artifactId>
    <version>1.0.0</version>
    <scope>test</scope>
    <classifier>test</classifier>
</dependency>
```

When we initially built this topology, we did it entirely within the `ZMartKafkaStreamsApp#main` method, which at the time was fine for quick development, but does not lend itself to being testable. What we are going to is refactor the topology into a stand-alone class which will enable us to test our topology. Note that since the logic isn't changing and we'll move the code as is, I won't show the conversion here, but point you to `src/main/java/bbejeck/chapter_8/ZMartTopology.java` for viewing later.

With the code transitioned, let's get on with constructing a test.

### **8.1.1 Building the Test**

Now let's move on to building our unit test for the `ZMartTopology`. For this test, we can use a standard JUnit test, and we'll have some expected set-up work before running the test as shown here:

#### **Listing 8.1 Set up method for topology test found in src/main/java/bbejeck/chapter\_8/ZMartTopologyTest.java**

```
@Before
public void setUp() {
    // properties construction left out for clarity
    StreamsConfig streamsConfig = new StreamsConfig(props);
    Topology topology = ZMartTopology.build();
```

1

```
topologyTestDriver = new ProcessorTopologyTestDriver(streamsConfig, topology);
}
```

2

- ➊ Refactored ZMart topology, now we can get the topology from method call
- ➋ Creating the ProcessorTopologyTestDriver

The critical point of this code section is the creation of the `ProcessorTopologyTestDriver` which we'll use in the following code section when we run our test.

### **Listing 8.2 Testing the topology**

**src/main/java/bbejeck/chapter\_8/ZMartTopologyTest.java**

```
@Test
public void testZMartTopology() {
    // serde creation left out for clarity

    Purchase purchase = DataGenerator.generatePurchase(); ➌

    topologyTestDriver.process("transactions",
        null,
        purchase,
        stringSerde.serializer(),
        purchaseSerde.serializer()); ➍

    ProducerRecord<String, Purchase> record = topologyTestDriver.readOutput("purchases",
        stringSerde.deserializer(),
        purchaseSerde.deserializer()); ➎

    Purchase expectedPurchase = Purchase.builder(purchase).maskCreditCard().build();
    assertEquals(record.value(), equalTo(expectedPurchase));
}
```

➌

➍

➎

➏

➐

- ➊ Creating a test object, reusing the generation code from running the topology
- ➋ Sending an initial record into the topology
- ➌ Reading a record from the "purchases" topic
- ➍ Converting our test object to the expected format
- ➎ Verifying the record from the topology matches expected record

There are two critical sections from this code sample that we'll discuss here, specifically listings 2 and 3. In listing 2, we feed a record into the `transactions` topic as it's the source for our entire topology.

With the loading of the topology completed, now we can verify the correct actions have taken place. In listing 3, we read the record out from the `purchases` topic one of the sink nodes defined in our topology. In listing 4, we create our expected output record, and in listing five, we assert the results are what we expect.

We also have two other sink nodes in the topology, so let's complete the test by

verifying we get the correct output from the rest of the topology:

**Listing 8.3 Testing the rest of the topology found  
src/test/java/bbejeck/chapter\_8/ZMartTopologyTest.java**

```
@Test
public void testZMartTopology() {
    // continuing test from the previous section

    RewardAccumulator expectedRewardAccumulator =
        RewardAccumulator.builder(expectedPurchase).build();

    ProducerRecord<String, RewardAccumulator> accumulatorProducerRecord =
        topologyTestDriver.readOutput("rewards",
            stringSerde.deserializer(),
            rewardAccumulatorSerde.deserializer()); 1

    assertThat(accumulatorProducerRecord.value(), equalTo(expectedRewardAccumulator)); 2

    PurchasePattern expectedPurchasePattern = PurchasePattern.builder(expectedPurchase).build();

    ProducerRecord<String, PurchasePattern> purchasePatternProducerRecord =
        topologyTestDriver.readOutput("patterns",
            stringSerde.deserializer(),
            purchasePatternSerde.deserializer()); 3

    assertThat(purchasePatternProducerRecord.value(), equalTo(expectedPurchasePattern)); 4
}
```

- ① Reading a record from the "rewards" topic
- ② Verifying our output for "rewards" topic output
- ③ Reading a record from the "patterns" topic
- ④ Verifying "patterns" topic output matches expectations

As you add another processing node to the test, you see the same pattern from the previous example. You read records from each topic and verify our expectations with an assert statement. The critical point to keep in mind with this test is that we are running records through our topology, but we are doing so *without* Kafka.

So now you have a repeatable test running a record through our entire topology, without the overhead of running Kafka.

The ProcessorTopologyTestDriver also supports testing topologies with a state store, so let's take a look at how we'd accomplish that in our next section.

### 8.1.2 Testing a State Store in the Topology

To demonstrate testing a state store we'll refactor another one of our applications, StockPerformanceStreamsAndProcessorApplication to have the Topology returned from a method call. You'll find the class in src/main/java/bbejeck/chapter\_8/StockPerformanceStreamsProcessorTopology and since we've haven't made any changes to the logic, we won't review it here.

The test set up is the same as the previous test, so we'll keep the description to the

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

parts that are new:

#### **Listing 8.4 Testing the state store**

**src/test/java/bbejeck/chapter\_8/StockPerformanceStreamsProcessorTopologyTes**

```
StockTransaction stockTransaction = DataGenerator.generateStockTransaction(); ①

topologyTestDriver.process("stock-transactions",
    stockTransaction.getSymbol(),
    stockTransaction,
    stringSerde.serializer(),
    stockTransactionSerde.serializer()); ②

KeyValueStore<String, StockPerformance> store =
    topologyTestDriver.getKeyValueStore("stock-performance-store"); ③

assertThat(store.get(stockTransaction.getSymbol()), notNullValue()); ④
```

- ① Generating a test record
- ② Processing the record with the test driver
- ③ Retrieving the state store from the test topology
- ④ Assert the store contains the expected value

As you can see here in listing four, we can quickly verify that our code is using the state store as expected.

We've seen the `ProcessorTopologyTestDriver` in action and how we can achieve end to end testing of a topology. The topologies we test can be very simple with one processing node to the very complex, consisting of several sub-topologies.

Even though we are doing this testing without a Kafka broker, make no mistake this is a full test of the topology and will exercise all parts including serialization/deserializing of records.

So far we've seen how we can do end to end testing of a topology. But we also would like to test the internal logic of our `Processor` and `Transformer` objects. Testing an entire topology is great, but verifying the behavior inside each class requires a more fine-grained approach which we'll cover next.

### **8.1.3 Testing Processors and Transformers**

To verify the behavior inside a single class requires a true unit test, where there is only *one* class under test. Writing a unit test for a `Processor` or `Transformer` should not be very challenging but remember both classes have a dependency on the `ProcessorContext` for obtaining any state stores and for scheduling punctuation actions.

Now we don't want to create a "real" `ProcessorContext` object, but rather a *stand-in* we can use for testing purposes, in other words, a mock object. Now when it comes to using a mock object, we have two paths we can follow.

One option is to use a mock object framework such as [mockito](#) to generate mock objects in our test. Another option is to use the `MockProcessorContext` object found in the same test library as the `ProcessorTopologyTestDriver`. The choice of which one to use depends on how we need to use them.

If you need the mock object strictly as placeholder for a real dependency, then concrete (mocks not created from a framework) are a good choice.

But if you want to verify either the parameters passed to a mock, the value returned or any other behavior, then using a mock object generated from a framework is a good choice. This is because mock object frameworks (like mockito mentioned above) come with a rich API for setting expectation and verifying behavior, saving you development time and speeding up your testing process.

In our next example, we are going to use both types of mock objects. We will use the Mockito framework for creating the `ProcessorContext` mock, as we want to verify parameters during the `init` call as well as validate we forward the expected values from the `punctuate` method. We'll also use a custom mock object for the key-value store, which we'll see in action as we step through the code example.

Now, let's move on to testing a `Processor` using mock objects to help us. We'll start with a test for the `AggregatingMethodHandleProcessor` named `AggregatingMethodHandleProcessorTest` located in `src/test/java/bbejeck_chapter6/processor/cogrouping/`.

First, we want to verify the parameters used in the `init` method

**Listing 8.5 Testing the init method found in  
src/test/java/bbejeck/chapter\_6/AggregatingMethodHandleProcessorTest.java**

```
// some details left out for clarity
private ProcessorContext processorContext = mock(ProcessorContext.class);
private MockKeyValueStore<String, Tuple<List<ClickEvent>,
    List<StockTransaction>>> keyValueStore = new MockKeyValueStore<>();
private AggregatingMethodHandleProcessor processor =
    new AggregatingMethodHandleProcessor();
@Test
@DisplayName("Processor should initialize correctly")
public void testInitializeCorrectly() {
    processor.init(processorContext);
    verify(processorContext).schedule(eq(15000L), eq(STREAM_TIME), isA(Punctuator.class));
    verify(processorContext).getStateStore(TUPLE_STORE_NAME);
}
```

- ➊ Mocking the `ProcessorContext` with Mockito
- ➋ A mock `KeyValueStore` object
- ➌ The class under test
- ➍ Calling the `init` method on the processor triggering method calls on

## ProcessorContext

- ⑤ Verifying the parameters for ProcessorContext.schedule method
- ⑥ Verifying retrieving the state store

Our first test is a simple one; we call the `init` method on the processor under test with the mocked `ProcessorContext`. We then validate the parameters used to schedule the `punctuate` method and that the state store is retrieved.

Next, let's test the `punctuate` method to verify we forward the records as expected.

### **Listing 8.6 Testing the punctuate method found in src/test/java/bbejeck/chapter\_6/AggregatingMethodHandleProcessorTest.java**

```
@Test
@DisplayName("Punctuate should forward records")
public void testPunctuateProcess() {
    when(processorContext.getStateStore(TUPLE_STORE_NAME)).thenReturn(keyValueStore); ①

    processor.init(processorContext);
    processor.process("ABC", Tuple.of(clickEvent, null)); ②
    processor.process("ABC", Tuple.of(null, transaction)); ③

    Tuple<List<ClickEvent>,List<StockTransaction>> tuple =
        keyValueStore.innerStore().get("ABC"); ④
    List<ClickEvent> clickEvents = new ArrayList<>(tuple._1);
    List<StockTransaction> stockTransactions = new ArrayList<>(tuple._2);

    processor.cogroup(124722348947L); ⑤

    verify(processorContext).forward("ABC", Tuple.of(clickEvents, stockTransactions)); ⑥

    assertThat(tuple._1.size(), equalTo(0));
    assertThat(tuple._2.size(), equalTo(0));
}
```

- ① Setting mock behavior to return a `KeyValueStore` when called
- ② Calling `init` method on processor
- ③ Processing a `ClickEvent` and a `StockTransaction`
- ④ Extracting the entries put into the state store in the `process` method
- ⑤ Calling the `cogroup` method, which is the method we used to schedule for `punctuate`
- ⑥ Validating the `ProcessorContext` forwards the expected records
- ⑦ Validating the collections within the tuple are cleared out

This test is a little more involved, and it utilizes a mix of mock and real behavior to run the test. Let's take a brief walk through the test to talk about what we are doing here.

Our first step is to specify the behavior for our mock `ProcessorContext` to return our stubbed out `KeyValueStore` when the `ProcessorContext#getStateStore` method

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

is called. This line alone is an interesting mix of generated mock vs. a stubbed out mock object.

While we could easily use Mockito to generate a mock `KeyValueStore`, we chose not to do so for two reasons. First, a generated mock returning another generated mock seems a bit unnatural (in my opinion). Second, we want to verify and use the stored values in the `KeyValueStore` during the test vs. setting expectations with a canned response.

The second and third steps are running the actual processor through its usual steps; first initializing and then processor records.

The fourth step is where to have a real working `KeyValueStore` is important. Since the `KeyValueStore` is a simple stub, we use a `java.util.HashMap` underneath for the actual storage. So in listing 4, we retrieve the contents from the store placed there by the process method calls in step 3. We create new `ArrayList` objects with the contents of the `Tuple` (again this is a custom class developed for the sample code in this book) pulled from the state store by the provided key.

In step five we drive the `punctuate` method of the processor. Since this is a unit test, we don't care about testing how time is advanced, that would constitute testing the Kafka Streams API itself, which we don't want here. Our goal is to verify the behavior of the method we defined as our `Punctuator` (via a method reference) when called.

In step six we verify the main point of the test, take the expected key and value are forwarded downstream via the `ProcessorContext#forward` method. Now, this is portion of the test is a good example of the usefulness of a generated mock object. Using the Mockito framework we just need to tell the mock to expect a `forward` call with the given key and value and verify the test executed the code precisely in this manner.

We wrap up the test with step seven, where we verify our processor cleared out the collections of `ClickEvent` and `StockTransaction` objects after forwarding downstream.

As you can see from this test, we've been able to isolate the class under test with a mix of generated and stubbed out mock objects. Although it was stated earlier in this chapter, it bears repeating here. The bulk of testing in your Kafka Streams API application should be unit tests on your business logic and on any individual `Processor` or `Transformer` objects you have, as Kafka Streams is thoroughly tested so you'll want to focus your efforts on new, untested code.

We don't always want to wait to deploy our application to see how it interacts with a Kafka cluster. We'd like to sanity check our code especially when the only way to test the real code (without any mocking or stubs) is to run against a live Kafka broker, which will require us to move on to integration testing.

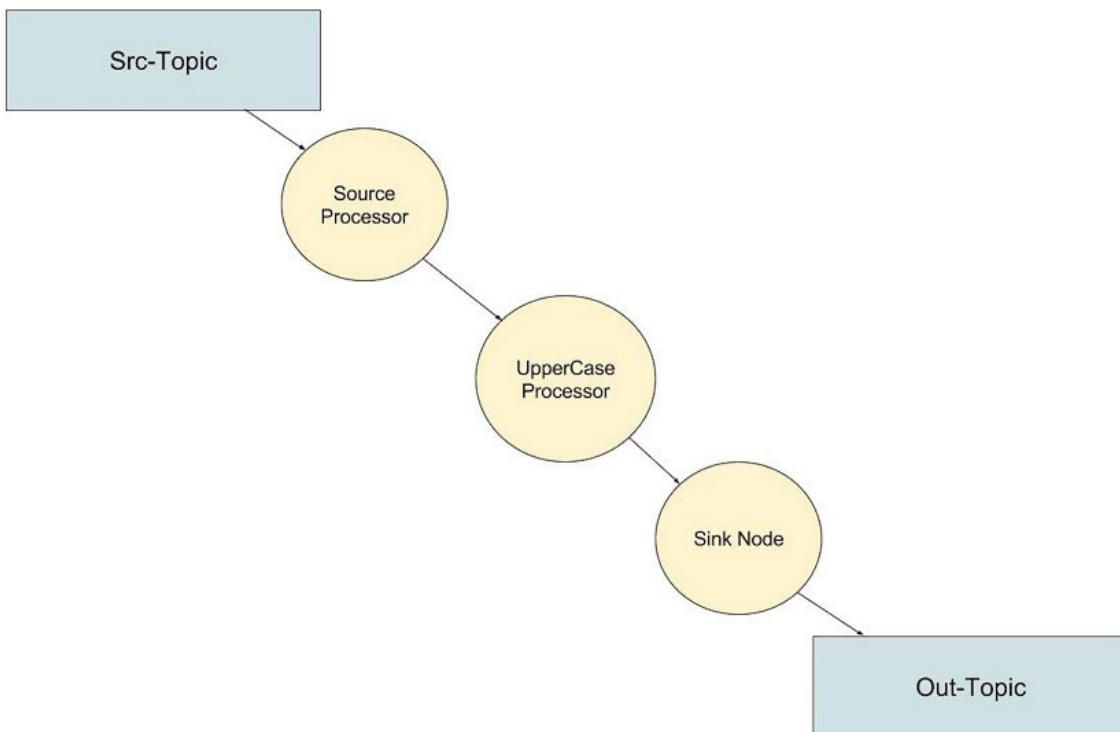
Next, we'll see how we can *locally* test against a real Kafka broker.

## 8.2 Integration Testing

So far we've seen how we can test an entire topology or an individual component in a unit test. For the most part, we prefer these types of tests as they are quick to run and they validate specific parts of our code.

But there are times where we need to test all working parts together end to end, in other words, an integration test. Usually an integration test is required when you have some functionality that can't be covered in a unit test.

For example of when you might need an integration test let's go back to our very first application the Yelling App. Since we created the topology so long ago, let's have another look at the following image:



**Figure 8.2 Another Look at the Yelling App Topology**

You've decided to change the source from a single named topic to any topic matching a regex, `yell-[A-Za-zA-Z0-9-]`. You want to confirm that if the topic matching the pattern, `yell-at-everyone` for example, is added while your application is deployed and running, you'll start reading information from that newly added topic.

Now we aren't going to update the original Yelling Application since it's so small, so we'll use a modified version directly in the test and here's how it will look:

### Listing 8.7 Updating the Yelling App found in src/java/bbejeck/chapter\_3/KafkaStreamsYellingIntegrationTest.java

```

streamsBuilder.<String, String>stream(Pattern.compile("yell.*"))
    .mapValues(String::toUpperCase)
  
```

①

②

```
.to(OUT_TOPIC);
```

- ① Subscribing to any topic starting with "yell"
- ② Converting all text to uppercase
- ③ Writing out to topic, or yelling at people!

Since you add topics at the Kafka Broker level, the only real way to test if your application picks up a newly created topic is to add one while your application is running. To run this scenario is all but impossible with a unit test. But does this mean you need to deploy your updated app to test it?

Fortunately for you, the answer is no.

You can use the *embedded* Kafka cluster available with Kafka test libraries.

By using the embedded Kafka cluster, we can run integration test requiring a Kafka cluster at any point on your machine either individually or as part of your entire battery of tests, which speeds up your development cycle.

#### NOTE

#### What's embeddable mean?

We are using the term *embeddable* here to mean taking a large application like Kafka or Zookeeper and have the ability to run in a local standalone mode, or "embed" them in an existing application.

Now let's move on to building an integration test.

### **8.2.1 Building an Integration Test**

So our first step to using the embedded Kafka server requires us to add three more *\*testing\** dependencies, the `scala-library-2.12.4.jar`, `kafka_2.12-1.0.0-test.jar`, and `kafka_2.12-1.0.0.jar` to our `build.gradle` or `pom.xml` file. We've already covered the syntax for providing a test jar in the "Testing a Topology" section above so that we won't repeat it here.

While it may seem like the number dependencies is starting to increase, remember anything we add in the chapter is a *testing* dependency. Testing dependencies are not packaged up and deployed with your application code, so in effect adding these testing dependencies should be somewhat transparent to you.

Now that we've added the required dependencies let's start defining our integration test with an embedded Kafka broker. We'll use a standard JUnit approach for creating our integration test.

#### **ADDING THE EMBEDDEDKAFKA CLUSTER**

Adding the embedded Kafka broker to the test is as simple as one line as the following code listing shows:

### Listing 8.8 Adding the Embedded Kafka Broker found in src/java/bbejeck/chapter\_3/KafkaStreamsYellingIntegrationTest.java

```

private static final int NUM_BROKERS = 1; 1

@ClassRule
public static final EmbeddedKafkaCluster EMBEDDED_KAFKA =
    new EmbeddedKafkaCluster(NUM_BROKERS); 2
3

```

- 1 Defining the number of brokers
- 2 The JUnit ClassRule annotation
- 3 Creating an instance of the EmbeddedKafkaCluster

In listing two, we create the `EmbeddedKafkaCluster` instance that serves as the cluster for running the tests in the class. The key point in this example is the `@ClassRule` annotation. While a full description of testing frameworks and JUnit is beyond the scope of this book, we'll take a minute here to explain the importance of the `@ClassRule` and how it drives the test.

#### JUNIT RULES

The release of JUnit 4 (or one of the JUnit 4 releases) introduced [the concept of Rules](#). Here's a brief definition of what a JUnit Rule:

Rules allow very flexible addition or redefinition of the behavior of each test method in a test class. Testers can reuse or extend one of the provided Rules below, or write their own.

JUnit provides three types of rules, and the `EmbeddedKafkaCluster` class uses the [ExternalResource rules](#). You use external resource rules for setting up and tearing down external resources, such as our `EmbeddedKafkaCluster` needed for a test.

JUnit provides class `ExternalResource` which has two no-op methods `before()` and `after()`. Any class extending the `ExternalResource` override the `before()` and `after()` methods setting up and tearing down the external resource needed for testing.

Rules provide an excellent abstraction for using external resources in your tests. After you create your class extending `ExternalResource` all you need to do is create a variable in your test and use the `@Rule` or `@ClassRule` annotation and all the setup and tear down methods are executed automatically.

The difference between a `@Rule` and `@ClassRule` is how often the `before()` and `after()` are called. The `@Rule` annotation executes the `before()` and `after()` methods for *each individual* test in the class. The `@ClassRule` executes the `before()` and `after()` method *once*; `before()` is executed prior to any test execution and `after()` is called when the last test in the class completes.

Setting up the `EmbeddedKafkaCluster` is relatively resource intensive, so it makes sense we only want to set it up once per test class. We'll stop here with our detour in describing JUnit rules and get back to building our integration test.

Now we have created our `EmbeddedKafkaCluster` our next step is to create any topics we'll need initially.

## CREATING TOPICS

Now that we have our embedded Kafka cluster available we can use it to create topics, as shown in the following code listing:

**Listing 8.9 Creating the topics for testing found in  
src/java/bbejeck/chapter\_3/KafkaStreamsYellingIntegrationTest.java**

```
@BeforeClass
public static void setUpAll() throws Exception {
    EMBEDDED_KAFKA.createTopic(YELL_A_TOPIC); ②
    EMBEDDED_KAFKA.createTopic(OUT_TOPIC); ③
}
```

- ① `@BeforeClass` annotation
- ② Creating the first source topic
- ③ Creating the output topic

Creating topics for the test is also something we'll want to do only once for all tests, so we use a `@BeforeClass` annotation that creates the required topics before the execution of any tests.

For this test we only need topics with a single partition and a replication factor of one, so we use the convenience method `EmbeddedKafkaCluster#createTopic(String name)`. If we need to have more than one partition, a replication factor greater than one require configurations different from the defaults we can use one of the following overloaded `createTopic` methods:

1. `EmbeddedKafkaCluster#createTopic(String topic, int partitions, int replication)`
2. `EmbeddedKafkaCluster#createTopic(String topic, int partitions, int replication, Properties topicConfig)`

Now we have all the pieces in place for our embedded Kafka cluster to run, let's move on to testing the topology with the embedded broker.

## TESTING THE TOPOLOGY

At this point, we have all our pieces in place, and now we're going to do the following steps to execute our integration test:

1. Start the streams application

2. Write some records to the source topic, assert the correct results
3. Create a new topic matching our pattern
4. Write some additional records to the newly created topic, assert the accurate results.

Let's start with the first two parts of our test:

**Listing 8.10 Starting the Streams application and asserting first set of values found in `src/java/bbejeck/chapter_3/KafkaStreamsYellingIntegrationTest.java`**

```
// some setup code left out for clarity

kafkaStreams = new KafkaStreams(streamsBuilder.build(), streamsConfig);
kafkaStreams.start(); 1

List<String> valuesToSendList = Arrays.asList("this", "should", "yell", "at", "you");
List<String> expectedValuesList = valuesToSendList.stream()
    .map(String::toUpperCase) 2
    .collect(Collectors.toList()); 3

IntegrationTestUtils.produceValuesSynchronously(YELL_A_TOPIC,
    valuesToSendList,
    producerConfig,
    mockTime); 4

int expectedNumberOfRecords = 5;
List<String> actualValues =
    IntegrationTestUtils.waitUntilMinValuesRecordsReceived(consumerConfig,
        OUT_TOPIC,
        expectedNumberOfRecords); 5

assertThat(actualValues, equalTo(expectedValuesList)); 6
```

- 1 Starting the Kafka Streams application
- 2 The list of values to send
- 3 Creating the list of expected values
- 4 Producing the values to embedded Kafka
- 5 Consuming records from Kafka
- 6 Asserting the values read equal the expected values

This portion of our test is pretty standard testing code. We "seed" our streaming application by writing records to the source topic. Our streaming application is already running, so it consumes, processes and writes out records as part of its standard processing. To verify the application is performing as we expect, the test consumes records from the sink node topic and compares the expected values to the actual values.

Now looking at listings 4 and five we can see some static utility methods making the construction of this integration test much more manageable. These producing and consuming utility methods are part of the `kafka-streams-test.jar`. Let's discuss these methods briefly.

## PRODUCING AND CONSUMING RECORDS IN A TEST

The `IntegrationTestUtils#produceValuesSynchronously` creates a `ProducerRecord` for each item in the given collection with a null key. This method is synchronous as it takes the resulting `Future<RecordMetadata>` from the `Producer#send` call and immediately calls `Future#get()` which blocks until the the produce request returns.

Since this method is sending records synchronously, we know the records are available for consuming once the method returns.

There is another method `IntegrationTestUtils#produceKeyValuesSynchronously` which takes a collection of `KeyValue<K, V>` if you want to specify a value for the key.

For consuming records in our test above we are using the `IntegrationTestUtils#waitForMinValuesReceived` method. As we can guess from the name, this method will attempt to consume the expected number of records from the given topic. By default this method will wait up to thirty seconds and if the expected number of records has not been consumed an `AssertionError` is thrown, failing the test.

If you need to work with the consumed `KeyValue` vs. just the value, there is the `IntegrationTestUtils#waitForMinKeyValueRecordsReceived`, which works in the same manner, but returns a Collection of `KeyValue` results. Additionally, there are overloaded versions of the consuming utility where you can specify a custom amount of time to wait via a parameter of type `long`.

With our explanation of the test utility methods complete, let's get back to finish describing our test.

## DYNAMICALLY ADDING A TOPIC

Now we are at the point in our test where we want to test the dynamic behavior we need a live Kafka broker for. The previous portion of the test was done to verify our starting point. Now we are going to use the `EmbeddedKafkaCluster` to create a new topic, and we'll test the application consumes from the new topic and process records as expected.

### **Listing 8.11 Starting the Streams application and asserting values from dynamically added topic found in src/java/bbejeck/chapter\_3/KafkaStreamsYellingIntegrationTest.java**

```

1 EMBEDDED_KAFKA.createTopic(YELL_B_TOPIC);

2 valuesToSendList = Arrays.asList("yell", "at", "you", "too");

3 expectedValuesList = valuesToSendList.stream()
    .map(String::toUpperCase)
    .collect(Collectors.toList());

4 IntegrationTestUtils.produceValuesSynchronously(YELL_B_TOPIC,
    valuesToSendList,

```

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

```

    producerConfig,
    mockTime);

expectedNumberOfRecords = 4;
List<String> actualValues =
    IntegrationTestUtils.waitUntilMinValuesRecordsReceived(consumerConfig,
        OUT_TOPIC,
        expectedNumberOfRecords);

assertThat(actualValues, equalTo(expectedValuesList));

```

5

6

- 1 Creating the new topic
- 2 A new list of values to send
- 3 Creating the expected values
- 4 Producing the values to source topic of the streaming application
- 5 Consuming the results of the streaming application
- 6 Asserting expected results match the actual results

In listing one, we create a new topic matching the pattern for the source node of the streaming application. After that point, we go through the same steps of populating the new topic with data, then consuming records from the topic backing the sink node of the streaming application. At the last part of the test in listing six, we verify the consumed results match the expected results.

Now you can run this test from inside your IDE, and you should see a successful result. Now you have completed your first integration test!

While we don't want to use an integration test for everything as unit tests are easier to write and maintain, they can be indispensable in times when the only way to verify the behavior of your code is working with a live Kafka broker.

#### **IMPORTANT**

#### **Don't Use the Embedded Broker For All Tests**

While it may be tempting to make all your tests using the `EmbeddedKafkaCluster`, it's best if you don't. If you run the sample integration test we just built, you'll notice it takes much longer to run than our unit tests. While the few extra seconds taken from one test might not seem like much, when you multiply the time by several hundred to one thousand plus tests, the time it takes to run your test suite is substantial. Additionally, we always want to keep our tests small and focused on one specific piece of functionality vs. exercising all parts of the application chain.

### **8.3 Conclusion**

We've wrapped up how to test our Kafka Streams application. Testing is an essential part of the development cycle, and in this chapter, we learned how we could ensure our streaming application performs as expected. The key takeaways from this chapter I want you to remember are:

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

1. We should strive to keep business logic extracted to stand-alone classes that are entirely independent of a Kafka Streams application and that makes them very easy to unit test.
2. It's useful to have at least one test using the `ProcessorTopologyTestDriver` to test your topology from end to end. This type of test does not use a Kafka broker so it's fast and you can see end-to-end results.
3. For testing individual `Processor` or `Transformer` instances try to use a mock object framework only when you need to verify the behavior of some class in the Kafka Streams API.
4. Integration tests with the `EmbeddedKafkaCluster` should be used sparingly and only when you have interactive behavior that can only be verified with a live running Kafka broker

It's been a fun journey, and we've learned quite a lot about the Kafka Streams API, and how we can use to handle our data processing needs.

So to conclude our learning path, we'll switch gears from student to master.

The next and final chapter of this book is going to be a capstone project based on everything we've learned so far and in some cases extending out to write custom code that isn't in the Kafka Streams API. The result will be a live end-to-end to end application using the core functionality presented in this book.

Here's a bulleted list highlighting what we'll build:

1. Integration with Kafka Connect to bring in outside data for our streaming application to work on.
2. A web application leveraging Interactive Queries so we can observe the state of our streaming application, *without* the need for a database to hold state.
3. An introduction to how to run continuous queries against data in Kafka - KSQL

## *Part 4: Advanced Concepts with Kafka Streams*



In this final section, you'll take everything you've learned and applied it to building advanced applications by integrating Kafka Streams with Kafka Connect so that you can stream data even if it's being written to a relational database. Then you'll learn how to leverage the power of Interactive Queries so you can display them in information your Kafka Streams application is building in real-time, directly from Kafka Streams, no need for an external tool. You'll wrap up this final section learning about KSQL, a new tool introduced by Confluent (the company founded by the original developers of Kafka at LinkedIn) and how you can write SQL statements and run continuous queries on data coming into Kafka.

# *Advanced Applications with Kafka Streams*



In this chapter:

- Integrating Outside Data into Kafka Streams with Kafka Connect
- Kicking Your Database to the Curb with Interactive Queries
- KSQL Continuous Queries on Kafka

We've come along way in our journey to learn how to use Kafka Streams. We've covered a lot of ground, and now we should know how to build streaming applications. Up to this point, we've included the core functionality of Kafka Streams, but there is so much more we can do. In this chapter, we'll take what we've learned and build two advanced applications that will allow us to work in more real-world situations.

For example, in many organizations, as you bring in new technology it must mesh with legacy technology or processes. It's not uncommon to see database tables as the main sink for incoming data. We learned back in Chapter 5 that tables are streams, so we know we should be able to treat these tables as streams of data.

So the first advanced application will "convert" the physical database into a streaming application by integrating Kafka Connect with Kafka Streams.

Kafka Connect will listen for new inserts into the table(s) and write those records to a topic in Kafka. This same topic serves as the source for the Kafka Streams application so that we can turn our database table into a streaming application.

While we are speaking of legacy applications, it is also typical to dump data even if it is captured in real time into a database to serve as the data source for dashboard applications.

For our second, advanced application, we'll see how we can use Interactive Queries where we expose our Kafka Streams state stores for direct queries. Meaning our dashboard applications can pull directly from the state stores and show data as it is flowing through the streaming application, eliminating the need for a database.

Then we'll wrap up our coverage of advanced features by looking at a powerful new feature, KSQL. KSQL gives you the ability to write long running SQL queries on data

coming into Kafka. KSQL gives you all the power of Kafka Streams, with ease of writing a SQL query. We'll include KSQL because it uses Kafka Streams under the covers to get the job done.

## **9.1 Integrating with Other Data Sources**

For our first advanced application let's consider you work at an established financial services firm, Big Short Equity (also known as BS Equity or BSE). BSE has a desire to migrate their legacy data operations into the modern era, and that plan includes using Kafka. BSE is part way through the migration and you are tasked with updating their analytics. The desire is to display the latest equities transactions and associated information in real time, and Kafka Streams is the perfect fit.

More specifically, BSE offers funds covering different areas of the financial market. BSE records fund transactions in real time in a relational database. Eventually, BSE plans to write transactions directly into Kafka, but in the short term, the database is the system of record.

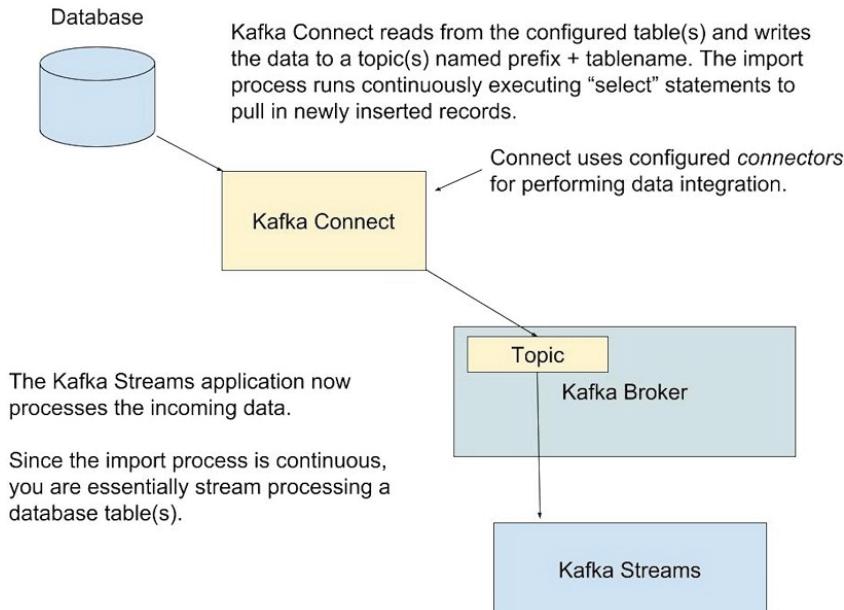
So the question is with the incoming data is fed into a relational database so how can you "bridge the gap" between the database and your emerging Kafka Streams application?

The answer is to use [Kafka Connect](#), which is a framework (included with Apache Kafka) that integrates Kafka with other systems. Once Kafka has the data you no longer are concerned about the location of the source data; you just point your Kafka Streams application at the source topic as you've done with any other streams application.

**NOTE**
**Kafka Connect Integrates with Kafka**

Just to be clear when using Kafka Connect to bring in data from other sources, the integration point is a Kafka topic. This means *any* application using a KafkaConsumer can make use of the imported data. Since this is a book on Kafka Streams, I'm going to emphasize how we'll integrate with a Kafka Streams application.

Figure 9.1 below depicts how this integration between the database and Kafka works:



**Figure 9.1 Kafka Connect Integrating Database Table and Kafka Streams**

In this case, you'll use Kafka Connect to monitor a database table and stream updates into a Kafka topic which is the source of your financial analysis application.

**TIP**

**Kafka Connect Information**

Since this is a book on Kafka Streams, this section is a whirlwind tour of Kafka Connect. For more information readers are encouraged to go to [Apache Kafka Documentation](#) and [Connect Quickstart](#) for more in-depth information.

### 9.1.1 Using Kafka Connect to Integrate Data

Kafka Connect is explicitly designed for streaming data from other systems into Kafka, or stream data from Kafka into another data storage application such as [MongoDB](#) [Elasticsearch](#). With Kafka Connect it is possible to import entire databases into Kafka, or other data (such as performance metrics) that is available to other applications for later analysis<sup>17</sup>.

---

Footnote 17 [kafka.apache.org/documentation/#connect\\_overview](https://kafka.apache.org/documentation/#connect_overview)

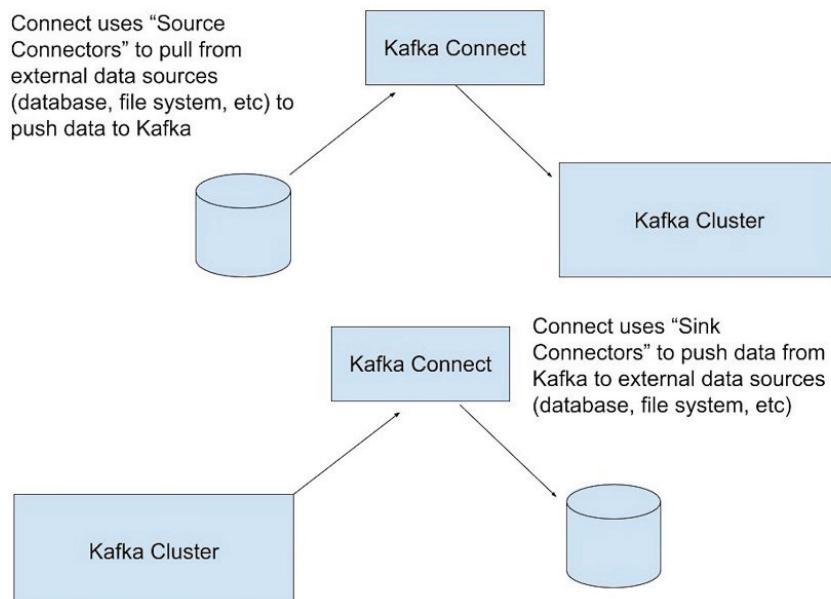
Kafka Connect uses specific *Connectors* to interact with various outside data sources. There are [several connectors](#) available (many of them developed by the connector community) making integration between Kafka and almost any other storage system possible. If there is not a connector available for your purposes, you can [implement a connector of your own](#).

### 9.1.2 Setting Up Kafka Connect

Kafka Connect runs in two flavors, distributed and standalone mode. For most production environments running in distributed mode makes sense as you can exploit the parallelism and fault tolerance from running multiple Connect instances.

I'm assuming you will run the examples on your local machine, so everything is configured in standalone mode.

As I mentioned before, Kafka Connect uses specific *connectors* to interact with various outside data sources. Connectors are one of two types, either a source connector or sink connector. How Kafka Connect uses source and sink connectors is depicted in the following image:



**Figure 9.2 Kafka Connect Source and Sink Connectors**

As you can see from the image above, source connectors bring data into Kafka, and sink connectors receive data from Kafka for use by another system.

For this example here, we'll use the [Kafka JDBC Connector](#). While the JDBC connector is available on GitHub ([github.com/confluentinc/kafka-connect-jdbc](https://github.com/confluentinc/kafka-connect-jdbc)) it's packaged with the source code distribution of the book as a convenience.

When using Kafka Connect, you'll need to do some minor configuration. There are some configurations for Kafka Connect itself and another set of configs the individual connector you are using to import or export data.

First, let's take a look at the configuration parameters you'll work with for Kafka Connect:

1. `bootstrap.servers` - comma-separated list of the Kafka broker(s) used by Connect.
2. `key.converter` - the class of the converter which controls the serialization of the key from Connect format to the format written to Kafka.
3. `value.converter` - the class of the converter controlling the serialization of the value from Connect format to the format written to Kafka. For the first example, you will use

the built in `org.apache.kafka.connect.json.JsonConverter`.

4. `value.converter.schemas.enable` - should Connect include the Schema for the value indicated by a `true` or `false`. In our first integration example, we'll have this value set to `false`, and we'll explain why in the next section on using Transformations.
5. `plugin.path` - We need to tell Kafka Connect the location of any connector and its dependencies that you're using. This location can be a single top-level directory containing uber jar file or multiple jar files. It's also possible to provide several paths by listing comma-separated list of locations.
6. `offset.storage.file.filename` - This is the file containing the offsets stored by the Connect consumer.

We'll also need to provide some configuration for our JDBC connector. Let's review those some of the settings now:

1. `name` - the name of the connector
2. `connector.class` - the class of the connector
3. `tasks.max` - total number of tasks,
4. `connection.url` - the URL used for connecting to the database
5. `mode` - the method the JDBC source connector uses for detecting changes
6. `incrementing.column.name` - name of column tracked for changes
7. `topic.prefix` - Connect write each table to a topic where the topic is named `topic.prefix+Table name`

While most of these configurations are straightforward, we should discuss two of them (`mode` and `incrementing.column.name`) in a little more detail as they have an active role in how the connector runs. The JDBC source connector uses the `mode` configuration for detecting which rows it needs to load from the last time it pulled data.

For our example, you will use the `incrementing` setting, which relies on an auto-incrementing column where each insert increments the column value by one. By tracking an incrementing column, you will only pull in *new* records; updates go unnoticed. But since your Kafka Streams application is concerned with only pulling in the latest equities product purchase, this setting is ideal. The `incrementing.column.name` is the name of the column of the auto-incrementing key.

#### TIP

#### Examples Provide Configuration

The source code for the book contains nearly completed configuration for both connect and the JDBC connector. The config files are located in the `src/main/resources` directory of the source code distribution. You'll still need to provide some information about the path where you have extracted the source code repository. Be sure to look at the `README.md` file for full instructions.

This brings us to the end of our overview of Kafka Connect and the JDBC Source Connector, but we have one more integration point to cover, which we'll discuss in the next section.

**NOTE****More Information the JDBC Connector**

While this has been a quick overview of the JDBC Source Connector, you can find more information from the Confluent documentation [jdbc-source-connector](#). Additionally, there are other incremental query modes you should look over as well [features](#).

### **9.1.3 Transforming Data**

Before getting this new assignment, you had already developed a Kafka Streams application using very similar data. As a result, you already have existing model and Serde objects (using Gson underneath for the JSON serialization and deserialization). So to keep your development velocity high, you'd prefer not have to write any new code to support working with Connect. As you'll see in the next section, you'll be able to do just that, import data from Connect seamlessly.

**TIP****What's GSON?**

GSON is an Apache Licensed library developed by Google for the serialization and deserialization of Java objects into and from JSON. You can find more information on the GitHub page: [github.com/google/gson](https://github.com/google/gson) and the user guide: [github.com/google/gson/blob/master/UserGuide.md#overview](https://github.com/google/gson/blob/master/UserGuide.md#overview)

To enable this seamless integration, you'll have some minor additional *configuration* changes to your JDBC-connector properties. Before we do that, let's revisit section 9.1.2 where discussed some of the configuration settings. Specifically I said we'd use the `org.apache.kafka.connect.json.JsonConverter` with schemas disabled, hence the value is converted into a simple JSON format.

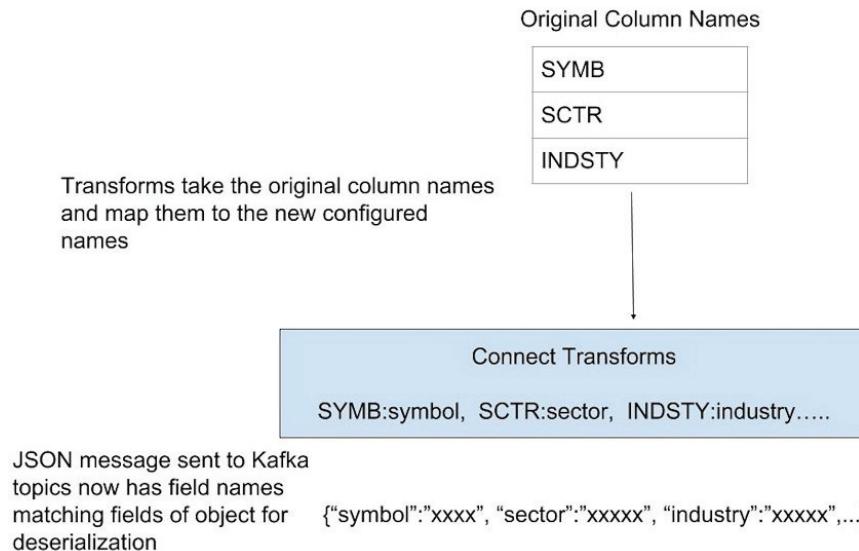
While the JSON is what you want to consume in your Kafka Streams application, there are two issues.

First, when converting the data into JSON format, the column names are used for names of the fields in the converted JSON. The names are all in the BSE abbreviation format that has no meaning outside the organization, so when your Gson Serde converts from JSON to the expected model object all the fields are null since the names don't match.

Second, the date and time is stored in the database as a timestamp, which is entirely expected, but the provided Gson Serde has not defined a custom TypeAdapter ([google.github.io/gson/apidocs/com/google/gson/TypeAdapter.html](https://google.github.io/gson/apidocs/com/google/gson/TypeAdapter.html)) for the `Date` type, thus all dates will need to be a `String` formatted like so:  
`yyyy-MM-dd'T'HH:mm:ss.SSS-0400`.

Luckily, Kafka Connect provides a mechanism which allows you to handle these two issues with ease.

Kafka Connect has the concept of *Transformations* that provide the ability to perform *lightweight* transformations before Connect writes data to Kafka. Where this transformation process takes place is depicted in figure 9.3:



**Figure 9.3 Transforming the names of the columns to match expected field names.**

To tackle these issues you are going to use two built-in transformations, the `TimestampConverter` and `ReplaceField`. As I mentioned previously, to use these transformations you'll just need to add a few lines of configuration to the `connector-jdbc.properties` file. Let's take a look at those configuration additions now:

#### **Listing 9.1 Updated JDBC Connector Properties files src/main/resources/conf/connector-jdbc.properties:**

```

1 transforms=ConvertDate,Rename
2 transforms.ConvertDate.type=org.apache.kafka.connect.transforms.TimestampConverter$Value
3 transforms.ConvertDate.field=TXNTS
4 transforms.ConvertDate.target.type=string
5 transforms.ConvertDate.format=yyyy-MM-dd'T'HH:mm:ss.SSS-0400
6 transforms.Rename.type=org.apache.kafka.connect.transforms.ReplaceField$Value
7 transforms.Rename.renames=SMBL:symbol, SCTR:sector,....

```

- ➊ Defining the aliases for the transformers
- ➋ Specifying the type for the `ConvertDate` alias
- ➌ The date field to convert
- ➍ The output type of the converted date field
- ➎ Format of date
- ➏ Specifying the type for the `Rename` alias

- ⑦ List of column names (truncated for clarity) to replace format is  
Original:Replacement

These properties are relatively self-descriptive so we won't spend too much time on them, but as you can see these properties provide you with exactly what you need for your Kafka Streams application to successfully deserialize messages imported into Kafka by Connect and the JDBC-Connector.

Now with all the Connect pieces in place, to complete your integration between the database table and your Kafka Streams application is just a matter of using a topic with the specified prefix in the `connector-jdbc.properties` file.

### **Listing 9.2 Kafka Streams**

`src/main/java/bbejeck/chapter_9/StockCountsStreamsConnectIntegrationApplicat`

```
Serde<StockTransaction> stockTransactionSerde = StreamsSerdes.StockTransactionSerde(); ①
StreamsBuilder builder = new StreamsBuilder();
builder.stream("dbTxnTRANSACTIONS", Consumed.with(stringSerde, stockTransactionSerde)) ②
    .peek((k, v) -> LOG.info("transactions from database key {} value {}", k, v)); ③
```

- ① Serde for StockTransaction object
- ② Using the topic Connect writes to as the source for the stream
- ③ Printing the messages out to the console

So at this point, we are stream processing records from a database table in Kafka Streams, but there is more for you left to do. You are streaming stock transaction data, and to do any analysis, you're going to need to group the transactions by the stock ticker symbol.

Now, we've seen already how to select a key and repartition the records, but it's more efficient if the records come into Kafka keyed, that way your Kafka Streams application can save the repartitioning step which saves processing time and disk space. Let's revisit our configuration for Kafka Connect.

First you can add a `ValueToKey` transformer that takes a list of field names in the value to extract and use for the key. So let's update the `connector-jdbc.properties` file :

### **Listing 9.3 Updated JDBC Connector Properties files**

`src/main/resources/conf/connector-jdbc.properties:`

```
transforms=ConvertDate,Rename,ExtractKey ①
transforms.ConvertDate.type=org.apache.kafka.connect.transforms.TimestampConverter$Value
transforms.ConvertDate.field=TXNTS
transforms.ConvertDate.target.type=string
transforms.ConvertDate.format=yyyy-MM-dd'T'HH:mm:ss.SSS-0400
```

```

transforms.Rename.type=org.apache.kafka.connect.transforms.ReplaceField$Value
transforms.Rename.renames=SMBL:symbol, SCTR:sector,....
transforms.ExtractKey.type=org.apache.kafka.connect.transforms.ValueToKey ②
transforms.ExtractKey.fields=symbol ③

```

- ① Adding the ExtractKey transform
- ② Specifying the class name of the ExtractKey transform
- ③ Listing the field(s) to extract for the key.

In listing 1, you added the `ExtractKey` transform, and in listing two, we provide Connect the name of the transformer class the `ValueToKey` class. In listing three we provide the name of the field `symbol` to use for the key, but the provided value could consist of multiple comma separated values, but in this case, we only need one value. Note you used the renamed version for the field as this transformer is executed *after* the `Rename` transformer.

The result of the `ExtractKey` field is a `Struct` of one value, but you only want the value contained in the `Struct` for the key, the stock ticker symbol. For this operation, we'll add one more transform `FlattenStruct` to pull the ticker symbol out by itself for the key.

#### **Listing 9.4 Updated JDBC Connector Properties files src/main/resources/conf/connector-jdbc.properties:**

```

transforms=ConvertDate,Rename,ExtractKey,FlattenStruct ①
transforms.ConvertDate.type=org.apache.kafka.connect.transforms.TimestampConverter$Value
transforms.ConvertDate.field=TXNTS
transforms.ConvertDate.target.type=string
transforms.ConvertDate.format=yyyy-MM-dd'T'HH:mm:ss.SSS-0400
transforms.Rename.type=org.apache.kafka.connect.transforms.ReplaceField$Value
transforms.Rename.renames=SMBL:symbol, SCTR:sector,....
transforms.ExtractKey.type=org.apache.kafka.connect.transforms.ValueToKey
transforms.ExtractKey.fields=symbol

transforms.FlattenStruct.type=org.apache.kafka.connect.transforms.ExtractField$Key ②
transforms.FlattenStruct.field=symbol ③

```

- ① Adding the last transform
- ② Specifying the class for the transform `ExtractField$Key`
- ③ The name of the field to pull out.

In listing 1, we've added the final transform named `FlattenStruct` and in listing two, we specify the class which is the `ExtractField$Key` which is used by Connect to extract the named field and only include the field in the results, which in this case is the key. Finally, in listing three, we provide the name of the field, which is `symbol` the same name from the previous transform, which makes sense since that is the field used to create the key struct.

Now with just a few extra lines of configuration we can expand the previous Kafka streams application to do more advanced operations without the need to select a key and the follow on repartitioning step:

### **Listing 9.5 Kafka Streams**

**src/main/java/bbejeck/chapter\_9/StockCountsStreamsConnectIntegrationApplicati**

```
Serde<StockTransaction> stockTransactionSerde = StreamsSerdes.StockTransactionSerde();
StreamsBuilder builder = new StreamsBuilder();
builder.stream("dbTxnTRANSACTIONS", Consumed.with(stringSerde, stockTransactionSerde))
    .peek((k, v) -> LOG.info("transactions from database key {} value {}", k, v))
    .groupByKey(Serialized.with(stringSerde, stockTransactionSerde)) 1
    .aggregate(() -> 0L, (symb, stockTxn, numShares) -> numShares + stockTxn.getShares(),
        Materialized.with(stringSerde, longSerde)).toStream() 2
    .peek((k, v) -> LOG.info("Aggregated stock sales for {} {}", k, v))
    .to("stock-counts", Produced.with(stringSerde, longSerde));
```

- ① Grouping By key
- ② Performing an aggregation of total number of shares sold

In listing one, since the data is coming in keyed, you can use a `GroupByKey` which is important because this method won't set the automatic repartitioning flag. From the group-by operation, we can directly go into an aggregation without performing a repartitioning step which is important for performance reasons. The README.md file included with the source code contains instructions for running an embedded database, H2 and Kafka Connect to produce data to the `dbTxnTRANSACTIONS` topic to run the streaming application.

#### **TIP**

#### **Transformation should be lightweight**

While it might seem tempting to use `Transformations` to perform all of the work when importing data into Kafka via Connect, remember these transformations are meant to be *lightweight*. For any transformations beyond the simple ones I've shown here in the examples, it's better to pull the data into Kafka then use Kafka Streams to do the heavy transformational work.

Now that we've seen how to use Kafka Connect for getting data into Kafka so we can process the data with Kafka Streams, we now turn our attention to how we can visualize the state of our data in real time.

## **9.2 Kicking Your Database to the Curb**

In Chapter 4 you learned how to add local state to a Kafka Streams application. Streaming applications need to use state to perform operations like aggregations, reduce and joins. Unless your streaming application works exclusively with individual records, having local state is required.

Going back to your requirements from BSE, you have developed a Kafka Streams

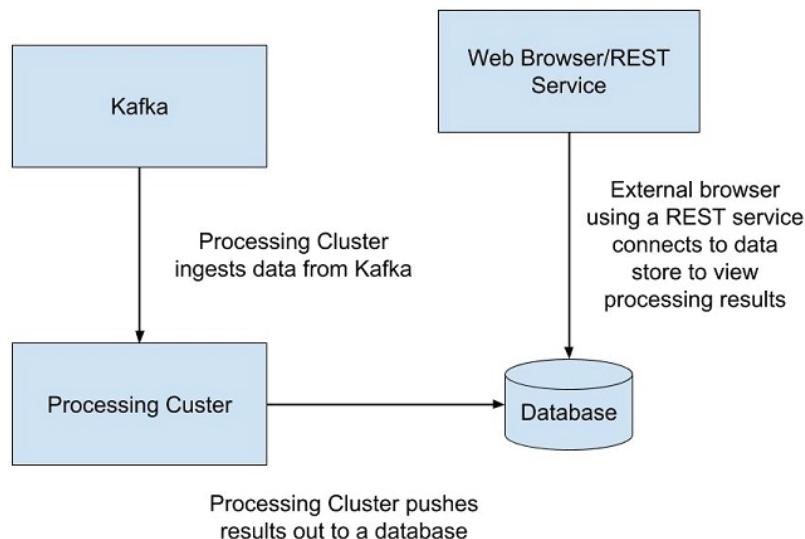
application your application captures three different categories of equity transactions:

1. The total deals by market sector
2. Customer purchases of shares per session
3. The total number of shares traded by stock symbol, in tumbling windows of 10 seconds

In all of the examples, you've reviewed the results either in the console, or read the results from a sink topic. While viewing data in console is suitable for development, it is not the best medium for displaying results. But for doing any analytic work or to quickly understand what is going on, a dashboard application is the better medium to use.

So in this section, you'll see how you can use Interactive Queries in Kafka Streams to develop a dashboard application for viewing the analytics outlined above, but *without* the need for a relational database to hold the state. You will populate the dashboard application directly from Kafka Streams, as the data is streaming. Hence the app continually updates naturally.

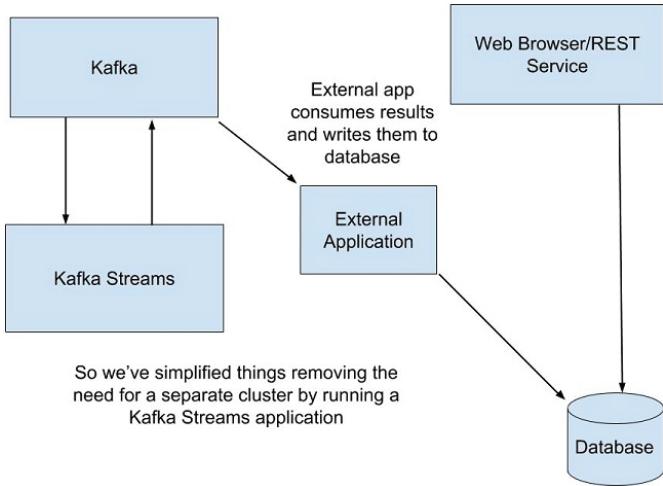
In typical architecture, data captured and operated on is pushed out to a relational database for viewing. The following image depicts this typical architecture setup:



**Figure 9.4 Architecture of Applications Viewing Processed Data.**

Pre Kafka Streams you would ingest data with Kafka then feed some analysis engine then the data is pushed to a relational database table used by a dashboard application.

Now if we add in Kafka Streams, using local state, our architecture changes slightly, as shown below in figure 9.5:

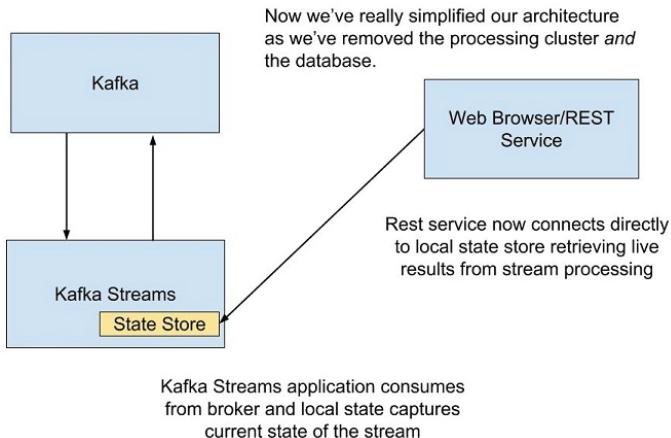


**Figure 9.5 Architecture with Kafka Streams and State Added.**

From looking at the image above we've simplified the structure significantly by removing an entire cluster (not to mention the deployment is much more manageable as well), but Kafka Streams still writes back to Kafka, and the database is still the primary consumer of the transformed data.

In final part of Chapter 5, I talked about interactive queries, so let's revisit the definition briefly. Interactive queries allow for directly viewing the data in the state store *without having to consume the data from Kafka*. In other words, the stream has become the database as well.

Since a picture is worth a thousand words, let's take another look at the image from above, but making adjustments for using interactive queries:



**Figure 9.6 Interactive Queries**

The idea demonstrated above is simple but very powerful. While state stores hold the state of the stream, Kafka Streams provides *read only* access from outside the streaming application via a RESTful interface. It's worth stating again how powerful this concept is; you can view the running state of the stream, without the need for an external database.

Now that you have an understanding of the impact of interactive queries let's walk

through how they work.

### 9.2.1 How Interactive Queries Work

For interactive queries to work, Kafka Streams expose state stores in a *read only* wrapper. It's important to understand that even Kafka Streams makes the stores available for queries, there is no updating for modifying the state store in any way. Kafka Streams exposes state stores from the `KafkaStreams#store` method.

Listing 9.6 demonstrates how this method works:

#### **Listing 9.6 Example of Retrieving a State Store**

```
ReadOnlyWindowStore readOnlyStore = kafkaStreams.store(storeName,
    QueryableStoreTypes.windowStore());
```

The example above is retrieving a `WindowStore`, but `QueryableStoreTypes` provides two other types:

1. `QueryableStoreTypes.sessionStore()`
2. `QueryableStoreTypes.keyValueStore()`

Once you have a reference to the read-only store, it's just a matter of exposing the store to some service (RESTful service for example) for users to query the state of the streaming data. However, retrieving the state store is only part of the picture. The state store extracted in the example above will only contain keys contained in the local store.

**IMPORTANT**

Remember, Kafka Streams assigns a state store per task, and as long as you use the same application id, a Kafka Streams application could consist of multiple instances. Additionally, those instances may not all be located on the same host. So it is possible that the state store you query only contains a subset of all the keys, and other state stores (with the same name, just located on other machine(s)) contain another sub-set of the keys.

Let's use one of the analytics listed above to make this concept clear.

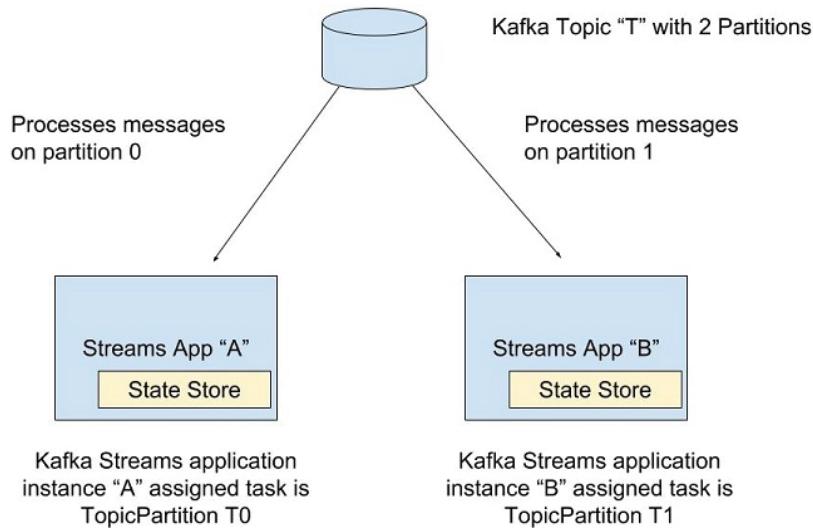
### 9.2.2 State Store Distribution

Consider the first analytic, aggregating stock trades per market sector. Since you are doing an aggregation, state stores come into play. You want to expose the state stores to provide a real-time view into the number of trades per sector to gain insight into which segment of the market is seeing the most action at the moment.

Even though stock market activity generates a significant data volume, I'll only discuss using two partitions to keep the details of the example clear. Additionally, we'll say you are running two single-threaded instances on two separate machines, although

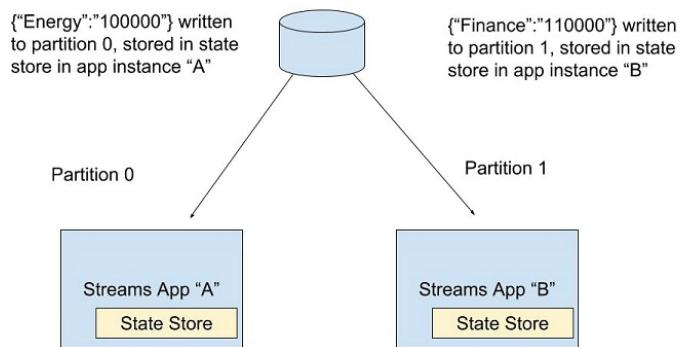
located in the same datacenter. Because of Kafka Streams automatic load balancing, each application will have one task each processing the data from each partition of your input topic.

The image below in figure 9.7 shows the distribution of tasks and state stores.



**Figure 9.7 Task and State Store Distribution**

As you can see instance, "A" handles all records on partition 0 and instance "B" handles all records on partition 1. So let's take a look at what happens when we have two records with keys "Energy" and "Finance" as depicted in the following image:



**Figure 9.8 Key and Value Distribution in State Stores**

So as you can see above, "Energy": "100000" lands in the state store on instance A, and "Finance": "110000" ends up in the state store on instance B. So, going back to our example of exposing the state store for queries, you can clearly see now if we expose the state store on instance A to a web service or any external querying, you can only retrieve the value for the key "Energy".

So, what is the solution to this issue? We certainly don't want to set up an individual web service to query each instance as that approach just won't scale. Fortunately, you won't have to as the solution is as simple as setting a configuration.

### 9.2.3 Setting Up and Discovering Distributed State Store

To enable interactive queries, you need to set the `StreamsConfig.APPLICATION_SERVER_CONFIG` parameter, which consists of the hostname of the Kafka Streams application and the port that a query service will listen on, in `hostname:port` format.

When a Kafka Streams instance receives a query for a given key, you'll need to find out if the key is contained in the local store, and more importantly, if it's not local you'll want to find out which instance contains the key and query against that store.

There are several methods on the `KafkaStreams` object that allow for retrieving information for all *running* instances with the *same* application id and defining the `APPLICATION_SERVER_CONFIG`.

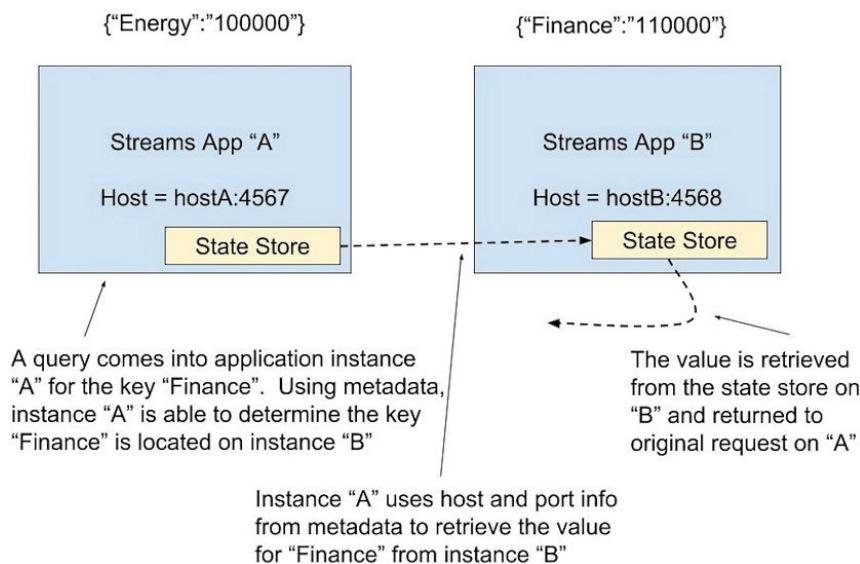
Here's a table listing the method names and the information available:

**Table 9.1 Methods For Retrieving Store Metadata**

Name	Parameter(s)	Usage
<code>allMetadata</code>	n/a	All instances with some possibly remote.
<code>allMetadataForStore</code>	store name	All instances (some remote) containing the named store
<code>allMetadataForKey</code>	key, Serializer	All instances (some remote) with the store containing the key
<code>allMetadataForKey</code>	key, StreamPartitioner	All instances (some remote) with the store containing the key

From the methods above you can use the `KafkaStreams.allMetadata` to obtain information for all instances that are eligible for interactive queries. From what I've found the `KafkaStreams.allMetadataForKey` method is the method I used the most when writing the interactive queries.

Next, let's take another look at the key-value distribution across Kafka Streams instances, but I'll add in the sequence of checking for the "Finance" key which is found and returned from another instance as demonstrated in figure 9.9:



**Figure 9.9 Key and Value Query and Discover Process**

From the image above each Kafka Streams instance has a lightweight embedded server listing to the port specified in the APPLICATION\_SERVER\_CONFIG.

It's important to point out here that you'll only need to query *one* of the Kafka Streams instances, and which one you choose doesn't matter (assuming you've configured it correctly!). By using an RPC mechanism and the metadata retrieval methods, if the instance you've queried does not contain the data you are looking for, the queried Kafka Streams instance will find where it is located, pull the results from that and return the results to the original query.

You can see this in action by tracing the flow of the calls in the image above. Instance "A" does not contain the key "Finance," but discovers that instance "B" does contain the key, so it issues a call to the embedded server on B retrieves the data and returns the result to the original caller.

#### **NOTE      RPC Mechanism Is Not Supplied**

Interactive queries will work on a single node out of the box, but an RPC mechanism is not provided, you have to implement your own. In this section, I've written one *possible* solution that we'll walk through, and you can run in the sample code, but you are free to implement your process, and I'm sure many of you will come up with something better. A great example of another RPC implementation is located in the Confluent kafka-streams-examples GitHub repo:  
[github.com/confluentinc./..MusicPlaysRestService.java](https://github.com/confluentinc./..MusicPlaysRestService.java)

Now that you've learned how to enable interactive queries and how to discover which instances are available it's time to move on to demonstrating interactive queries in action.

#### **9.2.4 Code For Interactive Queries**

The application you'll write for interactive queries will look very similar to the other apps you've written so far, just a couple of small changes. The first difference is you'll need to pass in two arguments when launching the Streams application, the hostname and the port the embedded service will listen to.

##### **Listing 9.7 Setting the hostname and port**

**src/main/java/bbejeck/chapter\_9/StockPerformanceInteractiveQueryApplication.java**

```
public static void main(String[] args) throws Exception {
    if(args.length < 2){
        LOG.error("Need to specify host and port");
        System.exit(1);
    }

    String host = args[0];
    int port = Integer.parseInt(args[1]);
    final HostInfo hostInfo = new HostInfo(host, port);

    Properties properties = getProperties();
}
```

1

```
properties.put(StreamsConfig.APPLICATION_SERVER_CONFIG, host+":"+port);
// other details left out for clarity
```

2

- ➊ Creating a HostInfo object, for later use in the application
- ➋ Setting the config for enabling interactive queries

Up until this point, you've just had to fire up the application without a second thought, but now we need to provide two arguments (the host and the port as shown in the code listing), but this is a change with minimal impact. Now we need to embed the local server for performing the actual queries. For this implementation I've chosen to use the [Spark webserver](#) (not THAT Spark, this is a book on Kafka Streams after all!)

The motivation for going with the Spark webserver is due to small footprint, its convention over configuration approach and the fact that it is purpose-built for microservices. And a microservice is what we can provide with interactive queries. If the Spark web server is not your liking, feel free to replace it with another web server.

**NOTE****What's a Microservice?**

While I think many readers are familiar with the term microservice here's the best definition I've seen that defines the term:

"Microservices - also known as the microservice architecture - is an architectural style that structures an application as a collection of loosely coupled services, which implement business capabilities. The microservice architecture enables the continuous delivery/deployment of large, complex applications. It also enables an organization to evolve its technology stack." - [Microservices.io](#)

Now let's take a look at the point of the code where we embed the Spark server, and some of the supporting code used to manage it.

**Listing 9.8 Setting the hostname and port**

[src/main/java/bbejeck/chapter\\_9/StockPerformanceInteractiveQueryApplication.java](#)

```
// details left out for clarity

KafkaStreams kafkaStreams = new KafkaStreams(builder.build(), streamsConfig);
InteractiveQueryServer queryServer = new InteractiveQueryServer(kafkaStreams, hostInfo);
queryServer.init();

kafkaStreams.setStateListener(((newState, oldState) -> {
    if (newState == KafkaStreams.State.RUNNING && oldState == KafkaStreams.State.REBALANCING) {
        LOG.info("Setting the query server to ready");
        queryServer.setReady(true);
    } else if (newState != KafkaStreams.State.RUNNING) {
        LOG.info("State not RUNNING, disabling the query server");
        queryServer.setReady(false);
    }
}));

kafkaStreams.setUncaughtExceptionHandler((t, e) -> {
```

1

2

3

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

```

LOG.error("Thread {} had a fatal error {}", t, e, e);
shutdown(kafkaStreams, queryServer);
});

Runtime.getRuntime().addShutdownHook(new Thread(() -> {
    shutdown(kafkaStreams, queryServer);
}));

```

4

5

- 1 Creating the embedded web server (actually a wrapper class)
- 2 Adding a StateListener to only enable queries to state stores until ready
- 3 Enabling queries to state stores once Kafka Streams application is in RUNNING state, in condition below disable if state is not RUNNING.
- 4 Setting an UncaughtExceptionHandler to log unexpected errors and close everything down
- 5 Adding a shutdown hook to close everything down when the application exits normally

In listing one, we are creating an instance of the `InteractiveQueryServer` which is a wrapper class containing the Spark web server and the code to manage the web service calls and starting and stopping the web server.

In chapter 7 we covered using a `StateListener` for notification of various states of a Kafka Streams application. Here in listing 2, you can see an efficient use of the listener. If you remember from the section above when running an interactive query, you need to use an instance of `StreamsMetadata` to determine if the data for the given key is local to the instance processing the query. So in listing 3, we set the state of the query server to `true`, allowing access to the metadata needed for queries only if the application is in the `RUNNING` state.

A key point to keep in mind is the metadata returned is a snapshot of the makeup of the Kafka Streams application now. At any point in time, you can scale up (or down) your application. When this occurs (or any other qualifying event like adding topics with a regex source node etc) the Kafka Streams application goes through a rebalancing phase and could change partition assignments. So in this case, I'm choosing only to allow queries when in a running state, but feel free to use whatever strategy you think is appropriate.

Listing 4 is another example of one of the concepts we covered in chapter 7, setting an `UncaughtExceptionHandler`, in this case, we log the error and shut down the application and the query server. Since this application runs indefinitely, I've added a shutdown hook to close everything down once you stop the demo.

Now that you've seen how we instantiate and start the service, let's move on to the code for running the query server.

## 9.2.5 Inside the Query Server

When implementing your RESTful service, the first step is to map URL paths to the correct methods to execute:

### Listing 9.9 Mapping URL Paths to Methods src/main/java/bbejeck/webserver/InteractiveQueryServer.java

```
public void init() {
    LOG.info("Started the Interactive Query Web server");

    get("/kv/:store", (req, res) -> ready ? fetchAllFromKeyValueStore(req.params()) :
        STORES_NOT_ACCESSIBLE);
    get("/session/:store/:key", (req, res) -> ready ? fetchFromSessionStore(req.params()) :
        STORES_NOT_ACCESSIBLE);
    get("/window/:store/:key", (req, res) -> ready ? fetchFromWindowStore(req.params()) :
        STORES_NOT_ACCESSIBLE);
    get("/window/:store/:key/:from/:to", (req, res) -> ready ? fetchFromWindowStore(req.params()):
        STORES_NOT_ACCESSIBLE);
}
```

- ➊ A mapping to retrieve all values from a plain key-value store
- ➋ Mapping for returning all sessions (from a session store) for a given key
- ➌ Mapping for window store no times specified, I'll explain below
- ➍ Mapping for a window store with a from and to time

This code listing highlights my decision to go with the Spark web server, as you can concisely map URLs to a Java 8 lambda expression to handle the request. While these mappings are straightforward, you'll notice we map the retrieval from the window store twice. To retrieve values from a window store, you need to provide a "from" and a "to" time.

One key thing to notice in the url mappings here is the check for the `ready` boolean value. This is value set in the `StateListener` so if `ready` evaluate to false; we don't attempt to process the request and return a message that the stores are not currently accessible.

This makes sense as a window store is segmented by time, with the segment size established when you create the store. We covered windowing in chapter 5, section 5.3.2 and examples are demonstrating creating windows there.

But I'm going to "cheat" here and offer a method accepting only a key and store and provide a default from and to time that we'll explore in the next example.

**NOTE****Changes For Window Store Retrieval**

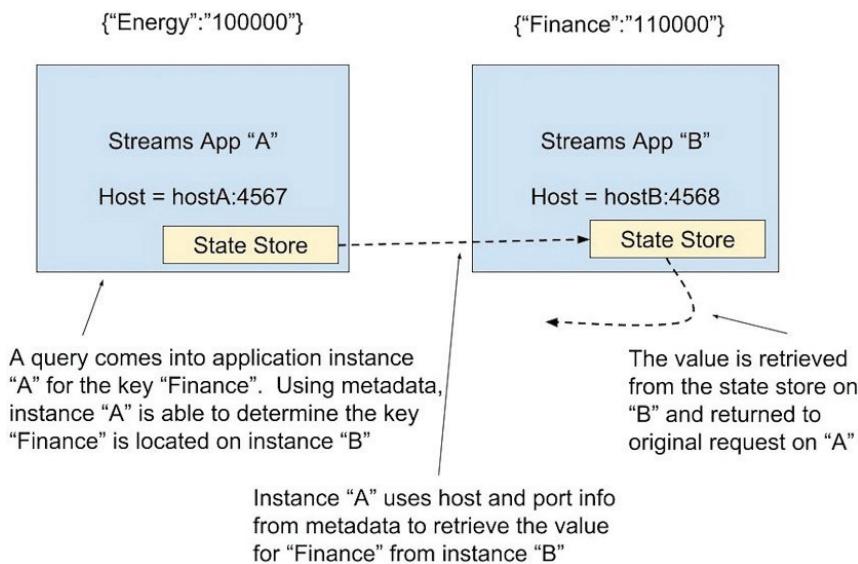
There is a proposal [KIP-205](#) to extend the `ReadOnlyWindowStore` to provide an "all()" method for retrieval of all time segments by key, alleviating the need to specify a "from" and "to" time. This functionality has not been implemented yet but should be included in a future release.

For our example of how the interactive query service works, we'll walk through an example of retrieving from a windowed store. Even though we'll only walk through this one example, the source code will contain instructions to run all types of queries demonstrated above.

### CHECKING FOR STATE STORE LOCATION

To refresh your memory, the requirements you have state that you need to collect different metrics on BSE's sales of securities. You first decide to track sales of individual stocks keeping running total over 10-second windows to identify different stocks that could trend up or down.

We are going to use the mapping `get("/window/:store/:key", (req, res) ready ? fetchFromWindowStore(req.params()) : STORES_NOT_ACCESSIBLE);` for the example to walk through from reviewing the request to returning the response. To help you keep your place where we are in the query process let's re-use figure 9.9 here as a roadmap to follow.



**Figure 9.10 Key and Value Query and Discover Process**

So we'll start with sending an HTTP get request `localhost:4567/window/NumberSharesPerPeriod/xxxx`. Here the xxxx represents the ticker symbol for a given stock.

## Listing 9.10 Mapping the Request and Check for Key Location src/main/java/bbejeck/webserver/InteractiveQueryServer.java

```

private String fetchFromWindowStore(Map<String, String> params) {
    String store = params.get(STORE_PARAM);
    String key = params.get(KEY_PARAM);
    String fromStr = params.get(FROM_PARAM);

    String toStr = params.get(TO_PARAM); 1

    HostInfo storeHostInfo = getHostInfo(store, key); 2

    if(storeHostInfo.host() .equals("unknown")){
        return STORES_NOT_ACCESSIBLE;
    } 3

    if(dataNotLocal(storeHostInfo)){
        LOG.info("{} located in state store on another instance", key);
        return fetchRemote(storeHostInfo,"window", params);
    } 4
}

```

- ① Extracting the request parameters.
- ② Getting the HostInfo for the key
- ③ If the hostname is "unknown" return appropriate message
- ④ Check if the returned hostname matches host of this instance

So the request is mapped to the `fetchFromWindowStore` method, the first step is to pull out the store name and key (stock symbol) from the request parameters map. In listing two, we fetch the `HostInfo` object for the key in the request, and you'll use the hostname to determine if the key is located on this instance or a remote one.

At listing three, we are checking if the Kafka Streams instance is (re)initializing as indicated by the `host()` method returning "unknown"; if so then we stop processing the request here and return the "not accessible" message.

Listing four is where you check if the hostname matches the hostname for the current instance. If the hostname does not match, we get the data from the instance containing the key and return the results.

Next, let's look at how we retrieve and format the results.

## Listing 9.11 Retrieving and Formatting the Results src/main/java/bbejeck/webserver/InteractiveQueryServer.java

```

Instant instant = Instant.now();
long now = instant.toEpochMilli(); 1

long from =  fromStr != null ? Long.parseLong(fromStr) : now - 60000; 2
long to =  toStr != null ? Long.parseLong(toStr) : now; 3

List<Integer> results = new ArrayList<>();

ReadOnlyWindowStore<String, Integer> readOnlyWindowStore = kafkaStreams.store(store,
    QueryableStoreTypes.windowStore()); 4

```

```

try(WindowStoreIterator<Integer> iterator = readOnlyWindowStore.fetch(key, from , to)) {
    while (iterator.hasNext()) {
        results.add(iterator.next() .value);
    }
}
return gson.toJson(results);

```

- ① Getting the current time in milliseconds
- ② Setting the window segment start time, if not provided, the time as of one minute ago
- ③ Setting the window segment ending time, if not provided, the current time
- ④ Retrieving the read ReadOnlyWindowStore
- ⑤ Fetching the window segments
- ⑥ Building up the response
- ⑦ Converting the results to JSON and return to the requestor

I had mentioned earlier, that I would "cheat" on the window store query if the "from" and "to" parameters we not provided in the query. If the user does not specify a range, by default we will return the last minute of results from the window store. Since we have defined a window of 10 seconds, we'll return six-windowed results.

After we fetch the window segments from the store, we iterate over them building a response that indicates the number of shares purchased for each ten-second interval over the last minute.

## RUNNING THE INTERACTIVE QUERY EXAMPLE

To observe the results of this example, you'll need to run three commands:

1. ./gradlew runProducerInteractiveQueries - produces the data needed for the examples
2. ./gradlew runInteractiveQueryApplicationOne - starts a Kafka Streams application with HostInfo using port 4567
3. ./gradlew runInteractiveQueryApplicationTwo - starts a Kafka Streams application with HostInfo using port 4568

Then point your browser to [localhost:4568/window/NumberSharesPerPeriod/AEBB](http://localhost:4568/window/NumberSharesPerPeriod/AEBB). Hit refresh a few time to see different results. Here is a static list of company symbols just for this example (AEBB, VABC, ALBC, EABC, BWBC, BNBC, MASH, BARX, WNBC, WKRP).

A better example is a mini-dashboard web application that updates automatically (via ajax) and displays the results from 4 different Kafka Streams aggregation operations. By running the commands listed above, you have everything setup, so point your browser to localhost:4568/iq or localhost:4567/iq to run the dashboard application. By going to either instance, you'll see how Kafka Stream's interactive queries handle getting results from all instances with the same application id. Look in the README file in the source code for full instructions on how to set up and start the dashboard application.

As you can see from observing the web application, you can view live results of the stream in a dashboard like application. Previously this type of application required a relational database, but not in this case, Kafka Streams can provide the information as needed.

Now we've wrapped up our coverage of Interactive Queries; we move on to an exciting new tool that Confluent<sup>18</sup> has recently released which allows you specify long-running queries against the stream without code, but using SQL.

---

Footnote 18 The company founded by the original developers of Kafka at LinkedIn

---

### 9.3 KSQL

Imagine you are working with business analysts as BSE. The analysts are very interested in your ability to quickly write applications in Kafka Streams to perform real-time data analysis. This interest puts you in sort of a bind.

You want to work with the analysts and write applications for their requests, but you have your workload, so the additional work makes it hard to keep up with everything. The analysts understand the added work they are creating, but they can't write code, so they are dependent on you for getting their analytics written.

However, the analysts are experts on working with relational databases, thus are quite comfortable with SQL queries. If there were just some way to give the analysts a SQL layer over Kafka Streams, then everyone's productivity would increase. Well, now there is.

In August of 2017, Confluent unveiled a powerful new tool for stream processing, [KSQL](#). KSQL is a streaming SQL engine for Apache Kafka, providing an interactive SQL interface allowing for writing power stream processing queries *without* the need for writing code. KSQL is especially adept at fraud detection and real-time applications.

#### NOTE

#### KSQL Coverage

KSQL is a big topic and could take a chapter or two if not an entire book on its own. So the coverage here will be very concise. Fortunately, you've already learned the core concepts underpinning KSQL as it uses Kafka Streams under the covers. For more information, you can get to the [KSQL Documentation](#).

KSQL provides scalable, distributed stream processing including aggregations, joins, windowing and more. Additionally, unlike SQL run against a database or a batch processing system, the results of a KSQL query are *continuous*. Before we dive into writing streaming queries, let's take a minute to review some fundamental concepts of KSQL.

### 9.3.1 KSQL Streams and Tables

In chapter 5 section 5.1.3, we discussed the concept of an event stream vs. an update stream. The event stream was an unbounded stream of individual *independent* events, while the update or record stream is a stream of updates to previous records with the same key.

KSQL has a similar concept of querying from a Stream or a Table. Where the Stream is an infinite series of events or facts, but are immutable, but with a query on a Table the facts are updatable or can even be deleted.

So while some of the terminologies might be different, the concepts are pretty much the same, so if you are comfortable with Kafka Streams, you will feel right at home with KSQL. For more details on these concepts, you can view the KSQL Terminology section at [github.com/confluentinc/ksql/blob/master/docs/concepts.md#terminology](https://github.com/confluentinc/ksql/blob/master/docs/concepts.md#terminology).

### 9.3.2 KSQL Architecture

KSQL uses Kafka Streams under the covers to build and fetch the results of the query. KSQL is made up of two components, the KSQL CLI and the KSQL server. Users of standard SQL tools such as MySql, Oracle or even Hive will feel right at home with CLI when writing queries in KSQL. Best of all KSQL is open-source (Apache 2.0 licensed).

The CLI is also the client connecting to the KSQL Server. The KSQL server is responsible for processing the queries and retrieving data from Kafka as well as writing results into Kafka.

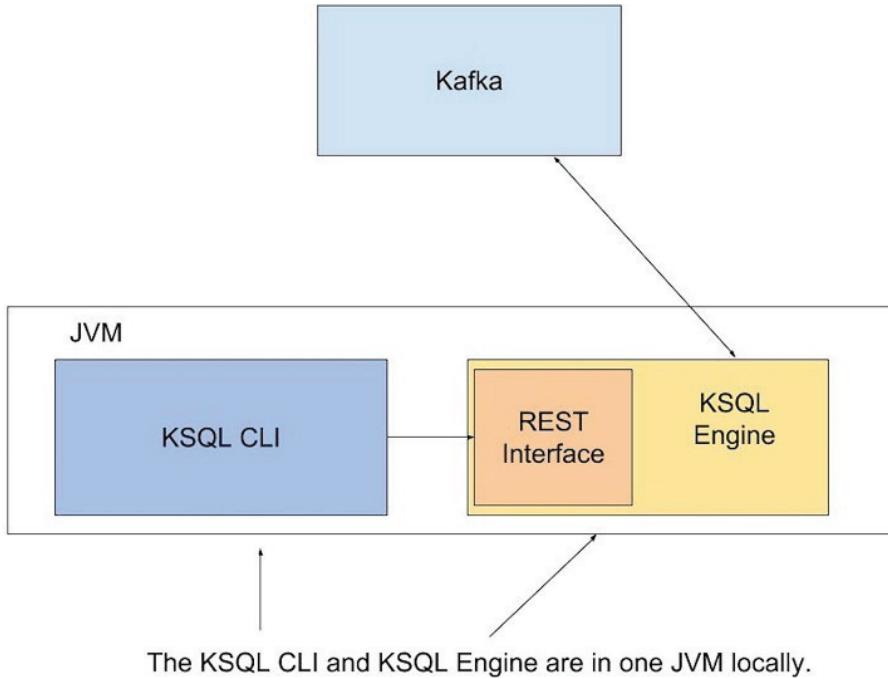
KSQL runs in two modes, standalone which is useful for prototyping and development or in distributed mode which of course is how you would use KSQL when working in a more realistic sized data environment.

**NOTE**

**KSQL is currently in Developer Preview**

As exciting as KSQL is and what it promises to deliver for SQL over streaming data, at the time of this writing, KSQL is considered a developer preview and it's not suggested to run against production clusters. A GA release of KSQL is expected soon, probably around the time this book is in print.

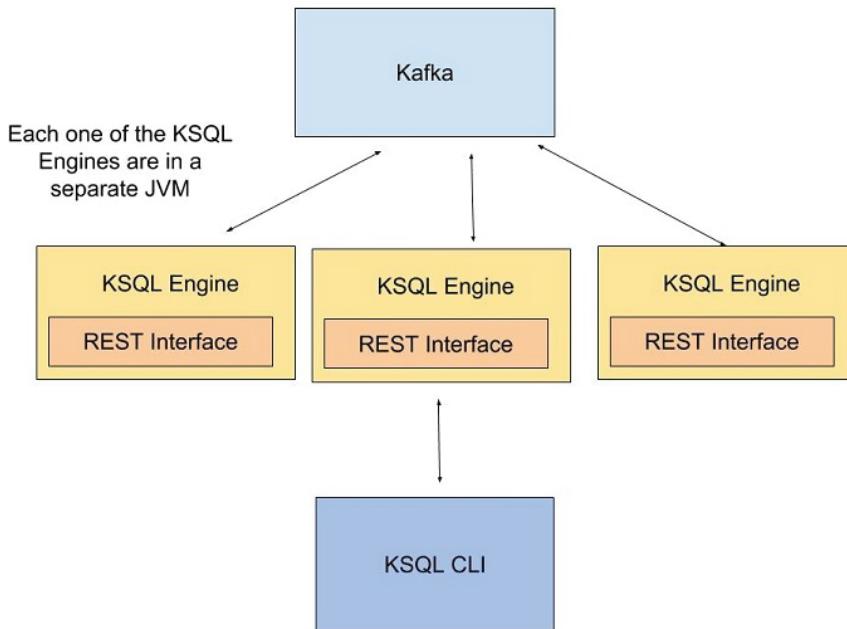
Let's take a look at how KSQL works in both modes starting with local mode demonstrated in figure 9.11



**Figure 9.11 KSQL in Local Mode**

As you can see from the image above the KSQL CLI, REST server and the KSQL engine are all located on the same JVM, which is ideal when running on your laptop.

Now let's take a look at how KSQL looks in the distributed mode.



**Figure 9.12 KSQL in Distributed Mode**

With KSQL in distributed mode, the KSQL CLI is now by itself, and it will connect to one of the remote KSQL servers (we'll cover starting and connections in the next section). A key point here is that although you only explicitly connect to one of the remote KSQL servers, all servers pointing to the same Kafka cluster will share in the workload of the submitted query.

A key point to note here is the KSQL Servers are using Kafka Streams to execute the queries. This means if you need more processing power you can stand up another KSQL server even during live operations (just like we can spin up another Kafka Streams application). The opposite case works just as well, if you have excess capacity, you can stop any number of KSQL servers, with the assumption that you'll leave at least one server operational. Otherwise, your queries will stop running!

Now that you have an idea of how KSQL is set up, let's cover how you get KSQL installed and running.

### 9.3.3 Installing and Running KSQL

For installing KSQL, you'll want to clone the KSQL repo with this command `git clone git@github.com:confluentinc/ksql.git` then cd into the `ksql` directory and execute this command `mvn clean package`, this will build the entire KSQL project. If you don't have git installed or don't want to build from source, you can download the release of the KSQL from [github.com/confluentinc/ksql/releases/download/v0.4/ksql-0.4.tgz](https://github.com/confluentinc/ksql/releases/download/v0.4/ksql-0.4.tgz)

**TIP**

#### Maven Needed for Building KSQL

KSQL is an Apache Maven based project, so you'll need Maven installed to build KSQL. If you don't have maven installed and you are on a Mac and have Homebrew installed simply run `brew install maven`, otherwise you can head over to [maven.apache.org/download.cgi](https://maven.apache.org/download.cgi) and download maven directly, and install instructions can be found at [maven.apache.org/install.html](https://maven.apache.org/install.html).

Now that you have the KSQL project downloaded and build our next step is to start KSQL. Note we'll be using KSQL in local mode for all of the examples in the book, but we'll still cover how to run KSQL in distributed mode. Make sure you are in the base directory of the KSQL project before going any further.

#### **Listing 9.12 Starting KSQL in local mode**

```
./bin/ksql-cli local
```

After running the command above you should see something like this in your console:

```
oddball:bin bbejeck$ ./ksql-cli local
Initializing KSQL...
=====
=   [ \ // __/] / _ \ | |
=   | '/ ( _ | | | | |
=   | < \_ \| | | | |
=   | . \_ ) | | | | |
=   |_\ \_ / \_ \| |
=
=   Streaming SQL Engine for Kafka =
Copyright 2017 Confluent Inc.

CLI v0.4, Server v0.4 located at http://localhost:9098

Having trouble? Type 'help' (case-insensitive) for a rundown of how things work!

ksql>
```

**Figure 9.13 KSQL successful launch result**

Congratulations you have successfully installed and launched KSQL! Next, let's start writing some queries.

### **9.3.4 Creating a KSQL Stream**

Getting back to your work at BSE, you've been approached by one of the analysts who is very interested in one of the applications you've written before and would like to make some tweaks to the application. But now, instead of this request resulting in more work, you spin up a KSQL console and turn the analyst loose to reconstruct your application as a SQL statement!

The example you are going to convert is the last windowed stream from the interactive queries example found in `src/main/java/bbejeck/chapter_9/StockPerformanceInteractiveQueryApplication` from lines 96-103. In that application, you are tracking the number shares sold every ten seconds by company ticker symbol.

You already have the topic defined (the topic maps to a database table) and a model object `StockTransaction` where the fields on the object map to columns in a table. Even though the topic is defined we need to register this information with KSQL by using a `CREATE STREAM` statement:

### **Listing 9.13 Creating a Stream found in `src/main/resources/ksql/create_stream.txt`**

```
CREATE STREAM stock_txn_stream (symbol VARCHAR, sector VARCHAR, \
    industry VARCHAR, shares BIGINT, sharePrice DOUBLE, \
    customerId VARCHAR, transactionTimestamp STRING, purchase BOOLEAN) \
    WITH (VALUE FORMAT = 'JSON', KAFKA TOPIC = 'stock-transactions');
```

- ① The CREATE STREAM statement named stock\_txn\_stream
  - ② Registering the fields of the StockTransaction object as columns
  - ③ Specifying the data format and the Kafka topic serving as the source of the stream (both required parameters)

With this one statement you are creating a KSQL Stream instance that you can now issue queries against. In the `WITH` clause you'll notice two required parameters `VALUE_FORMAT` telling KSQL the format of the data and the `KAFKA_TOPIC` parameter, telling KSQL where to pull the data from.

There are two additional parameters you can use in the `WITH` clause when creating a stream. One is `TIMESTAMP` which associates the message timestamp with a column in the KSQL Stream. Operations requiring a timestamp such as windowing will use this column to process the record.

The other is `KEY` which associates the key of the message with a column on the defined stream. In our case the message key for the `stock-transactions` topic matches the `symbol` field in the JSON value, so we didn't need to specify the key.

But had this not been the case then you would have needed to map the key to named column because you'll always need a key to perform grouping operations, which you'll see when we execute the stream SQL in an upcoming section.

**TIP**
**Listing Topics To Register**

With KSQL the command `list topics;` you will see list of topics on the broker the KSQL CLI is pointing to and whether the topics are "registered" or not.

After you have created your new stream you can view all streams and verify KSQL created the new stream as expected with the following commands:

**Listing 9.14 Listing all Streams and describing the stream you just created**

```
show streams;
describe stock_txn_stream;
```

The results of issuing these commands will give you results as demonstrated in figure 9.14:

```
ksql> show streams;

Stream Name      | Kafka Topic      | Format
-----
STOCK_TXN_STREAM | stock-transactions | JSON
ksql> describe stock_txn_stream;

Field           | Type
-----
ROWTIME         | BIGINT
ROWKEY          | VARCHAR(STRING)
SYMBOL          | VARCHAR(STRING)
SECTOR          | VARCHAR(STRING)
INDUSTRY        | VARCHAR(STRING)
SHARES          | BIGINT
SHAREPRICE      | DOUBLE
CUSTOMERID      | VARCHAR(STRING)
TRANSACTIONTIMESTAMP | VARCHAR(STRING)
PURCHASE        | BOOLEAN
```

**Figure 9.14 Listing Streams and describing your newly created stream**

You'll notice two extra columns `ROWTIME` and `ROWKEY` that KSQL has inserted. The `ROWTIME` column is the timestamp placed on the message (either from the producer or by the broker), and the `ROWKEY` is the key (if any) of the message. Now that you have created the stream, let's run our query on this stream.

### 9.3.5 Writing a KSQL Query

#### **Listing 9.15 SQL for performing stock analysis**

```
SELECT symbol, sum(shares) FROM stock_txn_stream WINDOW TUMBLING (SIZE 10 SECONDS)
    GROUP BY symbol;
```

Once you run this query, you'll results similar to what displayed here in figure 9.15:

ITZL	44694
KPAU	52858
NSTR	74110
ZERA	97959
MONA	29507
MESG	43474

**Figure 9.15 Results of tumbling window query**

#### **IMPORTANT Populating Data for the KSQL examples**

You'll need to run `./gradlew runProducerInteractiveQueries` to provide data for the KSQL examples

The column on the left is the ticker symbol, and the number is the number of shares traded for that symbol over the last ten seconds. With this query, you have specified a tumbling window of 10 seconds, but KSQL supports session and hopping windows as

well. If you'd like to review, we covered windowing in chapter 5, section 5.3.2.

Now you have built a streaming application without writing any code at all, quite an achievement. For a comparison let's take a look at the corresponding application written in the Kafka Streams API:

#### **Listing 9.16 Stock analysis application written in Kafka Streams**

```
KStream<String, StockTransaction> stockTransactionKStream =
    builder.stream(MockDataProducer STOCK_TRANSACTIONS_TOPIC,
        Consumed.with(stringSerde, stockTransactionSerde)
        .withOffsetResetPolicy(Topology.AutoOffsetReset.EARLIEST));

Aggregator<String, StockTransaction, Integer> sharesAggregator = (k, v, i) -> v.getShares() + i;

stockTransactionKStream.groupByKey()
    .windowedBy(TimeWindows.of(10000))
    .aggregate(() -> 0, sharesAggregator,
        Materialized.<String, Integer, WindowStore<Bytes, byte[]>>as("NumberSharesPerPeriod")
            .withKeySerde(stringSerde)
            .withValueSerde(Serdes.Integer())
    )
    .toStream().peek((k,v)->LOG.info("key is {} value is {}", k, v));
```

Even though the Kafka Streams API is concise, the equivalent you wrote in KSQL is a *one* one-liner query.

Before we wrap up our coverage of KSQL, we should discuss some additional features of KSQL.

#### **9.3.6 Creating A KSQL Table**

So far we've demonstrated creating a KSQL Stream, so let's show how to create a KSQL table, we'll use the `stock-transactions` topic as the source for familiarity:

#### **Listing 9.17 Creating a Table found in `src/main/resources/ksql/create_table.txt`**

```
CREATE TABLE stock_txn_table (symbol VARCHAR, sector VARCHAR, \
    industry VARCHAR, shares BIGINT, sharePrice DOUBLE, \
    customerId VARCHAR, transactionTimestamp STRING, purchase BOOLEAN) \
    WITH (KEY='symbol', VALUE_FORMAT = 'JSON', \
    KAFKA_TOPIC = 'stock-transactions');
```

Once you have the table created, you can now execute queries against the table. Now keep in mind, the table will contain updates for each transaction by `symbol` as the `stock-transactions` topic is keyed by the ticker symbol.

A useful experiment is to pick a ticker symbol from the streaming stock performance query above then run the following queries in the KSQL console and notice the difference in output.

#### **Listing 9.18 Comparing Stream and Table queries**

```
select * from stock_txn_stream where symbol='CCLU';
select * from stock_txn_table where symbol='CCLU';
```

The first query will produce several results, as it is a stream of individual events, but the table query returns far fewer results (one record when I ran the experiment). These results are the expected behavior because a table represents updates to facts, while the stream represents a series of unbounded events.

### 9.3.7 Configuring KSQL

Aside from the familiar SQL syntax and the ability write powerful streaming applications quickly, one thing you may have noticed is the lack of configuration.

This is not to say you can't configure KSQL as you are free to override any settings to your needs and any of the Streams, Consumer and Producer configs you can set for a Kafka Streams application are available. To view the properties that are currently set you can run this command `show properties;`

As an example of setting a property, let's see how we can specify the `auto.offset.reset` setting to `earliest`:

#### **Listing 9.19 Setting Auto Offset Reset**

```
SET 'auto.offset.reset'='earliest';
```

This is the approach you take for setting any property in the KSQL shell. However, if you have several configurations to set, typing each one into the console not very convenient, so you can specify a configuration file on start-up:

#### **Listing 9.20 Specifying Configs on Startup**

```
./bin/ksql-cli local --properties-file /path/to/configs.properties
```

This has been a quick tour of KSQL, but hopefully you can see the power and flexibility it gives you for creating streaming applications on Kafka.

## 9.4 Conclusion

We've covered new territory in this last chapter. We changed direction from learning how to write Kafka Streams applications to some key integrations and extensions to Kafka Streams. The key points to take away from this chapter are:

1. How to bring in data from outside sources into Kafka for a Kafka Streams application via Kafka Connect. By using Kafka Connect you can incorporate other data sources into your Kafka Streaming applications.
2. Using interactive queries to view the data in the stream without the need for a relational database. Interactive queries are a potent tool because they allow you to see the data as it's flowing through your Kafka Streams application.
3. The KSQL language which allows for quickly building powerful streaming applications without code. KSQL promises to deliver the power and flexibility of Kafka Streams to groups of workers that aren't developers.

# Additional Configuration Information



In this appendix we'll cover common and not-so common configuration options for a Kafka Streams application. During the course of the book you had a chance to see several examples of configuring a Kafka Streams application. But the configurations usually only included required (application-id, bootstrap-servers) and a handful of other configs (key and value serdes). In this appendix I'll show you some other settings, while not required, will help you keep Kafka Streams application running smoothly.

The appendix will depart some from the format of the previous chapters in that we'll present things in somewhat of a cookbook fashion.

## A.1 Limiting the Number of Rebalances on Startup

When starting up a Kafka Streams application, if you have multiple instances, the first instance gets all the topic partitions assigned from the GroupCoordinator on the broker. If you start another instance, a rebalance occurs taking away current TopicPartition assignments and then re-assigning all TopicPartitions across both Streams instances. This process is repeated until you have started all Streams applications (with the same application id).

While this is normal operation for a Kafka Streams application, during a rebalance processing of records is paused until the re-balance is completed, so we'd like to limit the number of rebalances when starting up if possible.

With the release of Kafka 0.11.0 a new *broker* configuration, `group.initial.rebalance.delay.ms` was introduced. This configuration delays the initial consumer rebalance from the GroupCoordinator when a new consumer joins the group by the amount specified in the `group.initial.rebalance.delay.ms` configuration. The default setting is three seconds. As other consumers join the group the rebalance is continually delayed by the configured amount (up to a limit of `max.poll.interval.ms`). This benefits Kafka Streams as you start new instances, the rebalance is delayed until all instances have come online (assuming you are starting them up one after another). For example if starting four instances with the appropriate rebalance delay setting, you should only have one rebalance after all four instances come online, meaning you'll start processing data faster.

## A.2 Resilience To Broker Outages

To keep your Kafka Streams application resilient in the face of broker failures here are some recommended settings:

1. Set the Producer.NUM\_RETRIES config to Integer.MAX\_VALUE
2. Set the Producer.REQUEST\_TIMEOUT config to 305000 (5 minutes)
3. Set the Producer.BLOCK\_MS\_CONFIG to Integer.MAX\_VALUE
4. Set the Consumer.MAX\_POLL\_CONFIG to Integer.MAX\_VALUE

### **Listing A.1 Setting Properties for Resilience to Broker Outages**

```
Properties props = new Properties();
props.put(StreamsConfig.producerPrefix(ProducerConfig.RETRIES_CONFIG), Integer.MAX_VALUE);
props.put(StreamsConfig.producerPrefix(ProducerConfig.MAX_BLOCK_MS_CONFIG), Integer.MAX_VALUE);
props.put(StreamsConfig.REQUEST_TIMEOUT_MS_CONFIG, 305000);
props.put(StreamsConfig.consumerPrefix(ConsumerConfig.MAX_POLL_INTERVAL_MS_CONFIG),
    Integer.MAX_VALUE);
```

Setting these values should help ensure the all brokers in the Kafka cluster go down, your streams application should stay up and be ready to resume working once the brokers are back online.

## A.3 Handling Deserialization Errors

Since Kafka works with byte arrays for the keys and values, you need to deserialize the keys and values to work with them. This is why you need to provide Serdes for all source and sink processors. It would not be unexpected to have some mal-formed data in the processing of records. Kafka Streams provides a configuration default.serialization.exception.handler or StreamsConfig.DEFAULT\_DESERIALIZATION\_EXCEPTION\_HANDLER\_CLASS\_CONFIG to specify how you want to handle these deserialization errors.

The default setting is the org.apache.kafka.streams.errors.LogAndFailExceptionHandler, which as the name implies, will log the error and your streams application instance will fail (shutdown) from this deserialization exception. There is another provided class, org.apache.kafka.streams.errors.LogAndContinueExceptionHandler which logs the error, but your streams application will continue to run.

You can implement your own deserialization exception handler by creating a class implementing the DeserializationExceptionHandler interface.

### **Listing A.2 Setting Deserialization Handler**

```
Properties props = new Properties();
props.put(StreamsConfig.DEFAULT_DESERIALIZATION_EXCEPTION_HANDLER_CLASS_CONFIG,
    LogAndContinueExceptionHandler.class);
```

I'm only showing how to set the LogAndContinueExceptionHandler handler as the

"log and fail" version is the default setting.

## A.4 Scaling Up Your Application

For all of the examples in the book, the streams applications ran with one stream thread. While that is fine for development, in practice you'll most likely need to run with more than one stream thread. But the question is how many threads to use and how many streams instances. While there are no concrete answers I can give, as only you know your circumstances well enough to answer those questions, we can go over some basic calculations to give a good idea.

If you remember from Chapter 3, Kafka Streams creates a `StreamsTask` per partition of the input topic(s). For our first example we'll only consider a single input topic with 12 partitions to keep the discussion straight forward.

With 12 input partitions, Kafka Streams will create 12 tasks. For the moment let's assume you want to have one task per thread. While you could have one instance with 12 threads, there is one drawback to that approach. If the machine hosting your streams application were to go down, you're entire stream processing will stop.

But if you start instances with 4 threads each, then each instance will process 4 input partitions. The benefit to this approach is that if one of the streams instances were to go down, a rebalance is triggered and the 4 tasks from the non-running instance are assigned to the other two instances, so the remaining applications will now process 6 tasks each. Additionally, when the stopped instance resumes running another rebalance will occur and all three instances will go back to processing 4 tasks.

One important consideration is that when determining the number of tasks to create, Kafka Streams takes the max number of partitions from *all* input topics. The first example we had one topic with 12 partitions, so we ended up with 12 tasks, but if the number of source topics was 4 with 3 partitions each, then you would have 3 tasks, with each task responsible for processing 4 partitions.

A key point to keep in mind is that the number of stream threads beyond the number of tasks will be idle. Going back to our example of 3 streams instances if you were to stand up a fourth instance of 4 threads, after rebalancing you would have 4 idle stream threads among your applications (16 threads only 12 tasks).

This is a key component of Kafka Streams I spoke earlier in the book, this dynamic scaling up or down does not involve taking your application offline, it happens automatically. This feature of Kafka Streams is very helpful if you have an uneven flow of data into the application you can spin up additional instances to handle the load, then take some offline when the volume drops off.

So do you always want a single thread per task? Maybe, but that is hard to say as this completely depends on demands of your application.

## A.5 RocksDB Configuration

For stateful operations, Kafka Streams uses [RocksDB](#) under the covers as the persistence mechanism. RocksDB is a fast and highly configurable key-value store. While there are too many options to make specific recommendations here, Kafka Streams provides a way to override the default settings with the `RocksDBConfigSetter` interface.

To set custom RocksDB settings create a class implementing the `RocksDBConfigSetter` interface, then provide the class name when configuring your Kafka Streams application via the `StreamsConfig.ROCKSDB_CONFIG_SETTER_CLASS_CONFIG` setting. To get an idea of what you can adjust with RocksDB I encourage you to read the [RocksDB Tuning Guide](#).

## A.6 Maybe Create Repartitioning Topics Ahead of Time

In Kafka Streams any time you perform an operation where you can potentially change the `key map`, `transform` or `groupBy` for example, an internal flag is set within the `StreamsBuilder` class indicating a repartitioning will be required. Now performing a `map` or `transform` does not automatically force the creation of a repartitioning topic and the repartition operation, but as soon as you add an operation using the updated key, a repartitioning operation is triggered.

While this is a required step (covered in Chapter 4) there are some cases where it's better to repartition the data yourself ahead of time. Consider the following (abbreviated) example:

```
KStream<String, String> mappedStream = streamsBuilder.stream("inputTopic").map(...); ①

// windowed count option 1
KTable<Windowed<String>, Long> ktable1 = mappedStream.groupByKey().windowedBy...count()
// windowed count option 2
KTable<Windowed<String>, Long> ktable2 = mappedStream.groupByKey().windowedBy...count() ②
// windowed count option 3
KTable<Windowed<String>, Long> ktable3 = mappedStream.groupByKey().windowedBy...count() ③
④
```

- ① Mapping the original input stream to create a new key
- ② Windowed count option 1
- ③ Windowed count option 2
- ④ Windowed count option 3

Here you are mapping the original stream to create a new key to group by. You want to perform three counts with three different windowing options, a perfectly legitimate use case. But since you mapped to a new key in listing 1, *each* windowed count operation in listing 2 through 4 will create a new repartition topic. Again the need for a repartition topic makes sense due to the changed key, but having 3 repartition topics duplicates the data when you only need one repartition topic.

The solution to this issue is simple, after your `map` call you immediately use a `through` operation to partition the data. Then the subsequent `groupByKey` calls will not trigger any repartitioning since the `groupByKey` operator by itself *does not* set the repartition needed flag. Here's what the code will look like now:

```
KStream<String, String> mappedStream =
    streamsBuilder.stream("inputTopic").map(...).through(...);
```

1

- 1 Mapping the original input stream to create a new key and repartitioning

By adding the `through` processor and repartitioning manually you will only have one repartitioning topic instead of three.

## A.7 Internal Topic Configuration

When building a topology, depending on the processors you add, Kafka Streams may create several internal topics. These internal topics can be changelogs for backing up state stores or repartitioning topics. Depending on your data volume these internal topics can start to consume a large amount of space.

Additionally, even though changelog topics are by default created with a cleanup policy of "compact" if you have a lot of unique keys these compacted topics can grow in size.

So with this in mind, it's a good idea to configure your internal topics to keep the size manageable. You have two options for managing the maintenance of internal topics.

One option is provide configs directly when creating state stores either with `StoreBuilder.withLoggingEnabled` or `Materialized.withLoggingEnabled` which of these two methods you'll use depends on how you are creating the state store. Both methods take a `Map<String, String>` containing the topic properties. One example exists in  
`src/main/java/bbejeck/chapter_7/CoGroupingListeningExampleApplication`.

The other option is when configuring your Kafka Streams application you can provide configurations for internal topics like so:

```
Properties props = new Properties();
// other properties set here
props.put(StreamsConfig.topicPrefix("retention.bytes"), 1024 * 1024);
props.put(StreamsConfig.topicPrefix("retention.ms"), 3600000);
```

When using the `StreamsConfig.topicPrefix` approach the provided settings are applied globally to all internal topics. Any topic settings provided when creating a state store will take precedence over the settings provided with `StreamsConfig`.

I can't give any direct advice here for what settings to use, as that depends on your particular use case, but keep in mind the default size of a topic is unlimited and the default retention time is one week, so those two settings (`retention.bytes` and

retention.ms) should be adjusted. Finally for changelogs backing state stores with a lot of unique keys you can set the `cleanup.policy` to `compact,delete` to ensure the topic size stays manageable.

## A.8 Reseting Your Streams Application

At some point you may need to start your Kafka Streams application over and reprocess data, either in development or after a code update. Kafka Streams provides a script `kafka-streams-application-reset.sh` found in the `bin` directory of a Kafka installation.

The script has one *required* parameter, the application id of the Kafka Streams application. The script offers several options for what you can do, but in a nutshell it can reset input topics to the earliest available offset, reset intermediate topics to the latest offset and delete any internal topics. Please note that you will need to call `KafkaStreams.cleanUp` the next time you start your application to delete any local state from previous runs.

## A.9 Cleaning Up Local State

In Chapter 4 I discussed that Kafka Streams stores local state per task on the local file system. During development, testing or when migrating to new instance, it may be desirable to clean out all previous local state.

To clean up any previous state, you can use `KafkaStreams.cleanUp` either before you call `KafkaStreams.start` or after `KafkaStreams.stop`. Using the `cleanUp` method at any other time will result in an error.

# Exactly Once Semantics

Kafka achieved a major milestone with the release of the 0.11.0 version of Kafka, exactly-once semantics. Prior to this release of Kafka, the delivery semantics of Kafka could have been described as "at least once" or "at most once" depending on the producer.

In the case of "at least once" a broker can persist a message, but experience an error before sending the acknowledgement back to the producer, assuming the producer is configured with `acks="all"` and times out waiting for the acknowledgement. If the producer is configured with `retries` greater than zero, it will resend the message, unaware of the fact the previous message was successfully persisted. In this scenario (although rare), a duplicate message is delivered to consumers, hence the term "at least once".

For the "at most once" condition, consider the case when a producer is configured with `retries` set to zero. If we consider the previous example, the message in question would be delivered only once, as there are no retries. But in the case of errors where the broker experiences an error before it can persist the message results in the message not being sent. In this case we've traded receiving all messages for not receiving any duplicate messages.

Now with exactly once semantics, even under situations where a producer resends a message that was previously persisted to a topic, consumers will receive the message exactly once.

To enable transactions or exactly once processing with a `KafkaProducer` you simply add a configuration `transactional.id` and add a couple of method calls. Otherwise producing messages with transactions looks familiar. Note that the following excerpt is not to stand alone and is just provided to highlight what's needed for sending messages with the transactional API.

## **Listing .1 Enabling Transactions in a KafkaProducer**

```
Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092");
props.put("transactional.id", "transactional-id");

Producer<String, String> producer = new KafkaProducer<>(props, new StringSerializer(),
```

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.

<https://forums.manning.com/forums/kafka-streams-in-action>  
**Licensed to Asif Qamar <[asifqamar.com](mailto:asifqamar.com)>**

```

new StringSerializer());  
  

// when setting transactional.id, you need to call this method before any others  

producer.initTransactions();  
  

try {  

    // called right before sending any records  

    producer.beginTransaction();  
  

    ...sending some messages  
  

    // when done sending, commit the transaction  

    producer.commitTransaction();  
  

} catch (ProducerFencedException | OutOfOrderSequenceException | AuthorizationException e) {  

    // The only option with any of above non-recoverable exceptions is to close the producer.  

    producer.close();  

} catch (KafkaException e) {  

    // for any other exception abort and retry  

    producer.abortTransaction();  

}

```

To use `KafkaConsumer` with transactions, you only need to add one configuration

## **Listing .2 Enabling Consumers for Reading Transactional Messages**

```
props.put("isolation.level", "read_committed");
```

In `read_committed` mode, the `KafkaConsumer` will only read successfully committed transactional messages. The default setting is `read_uncommitted` which returns all messages. Non transactional messages are always retrieved in either configuration setting.

The impact of exactly once semantics is a big win for Kafka Streams. With exactly once, or transactions, you are guaranteed to process records through a topology exactly once.

To enable exactly once processing with Kafka Streams, you set the `StreamsConfig.PROCESSING_GUARANTEE_CONFIG` to `exactly_once`. The default setting for `PROCESSING_GUARANTEE_CONFIG` is `at_least_once` or non-transactional processing.

With that simple configuration setting, Kafka Streams handles all required steps for performing transactional processing.

This has been a very quick overview of Kafka's transactional API. For more information the reader is encouraged to view the following resources:

1. [www.manning.com/books/kafka-in-action](http://www.manning.com/books/kafka-in-action) - A book on Kafka from Manning.
2. [www.confluent.io/blog/exactly-once-semantics-are-possible-heres-how-apache-kafka-does-it](http://www.confluent.io/blog/exactly-once-semantics-are-possible-heres-how-apache-kafka-does-it)
3. [www.confluent.io/blog/transactions-apache-kafka/](http://www.confluent.io/blog/transactions-apache-kafka/)
4. [www.confluent.io/blog/enabling-exactly-kafka-streams/](http://www.confluent.io/blog/enabling-exactly-kafka-streams/)